

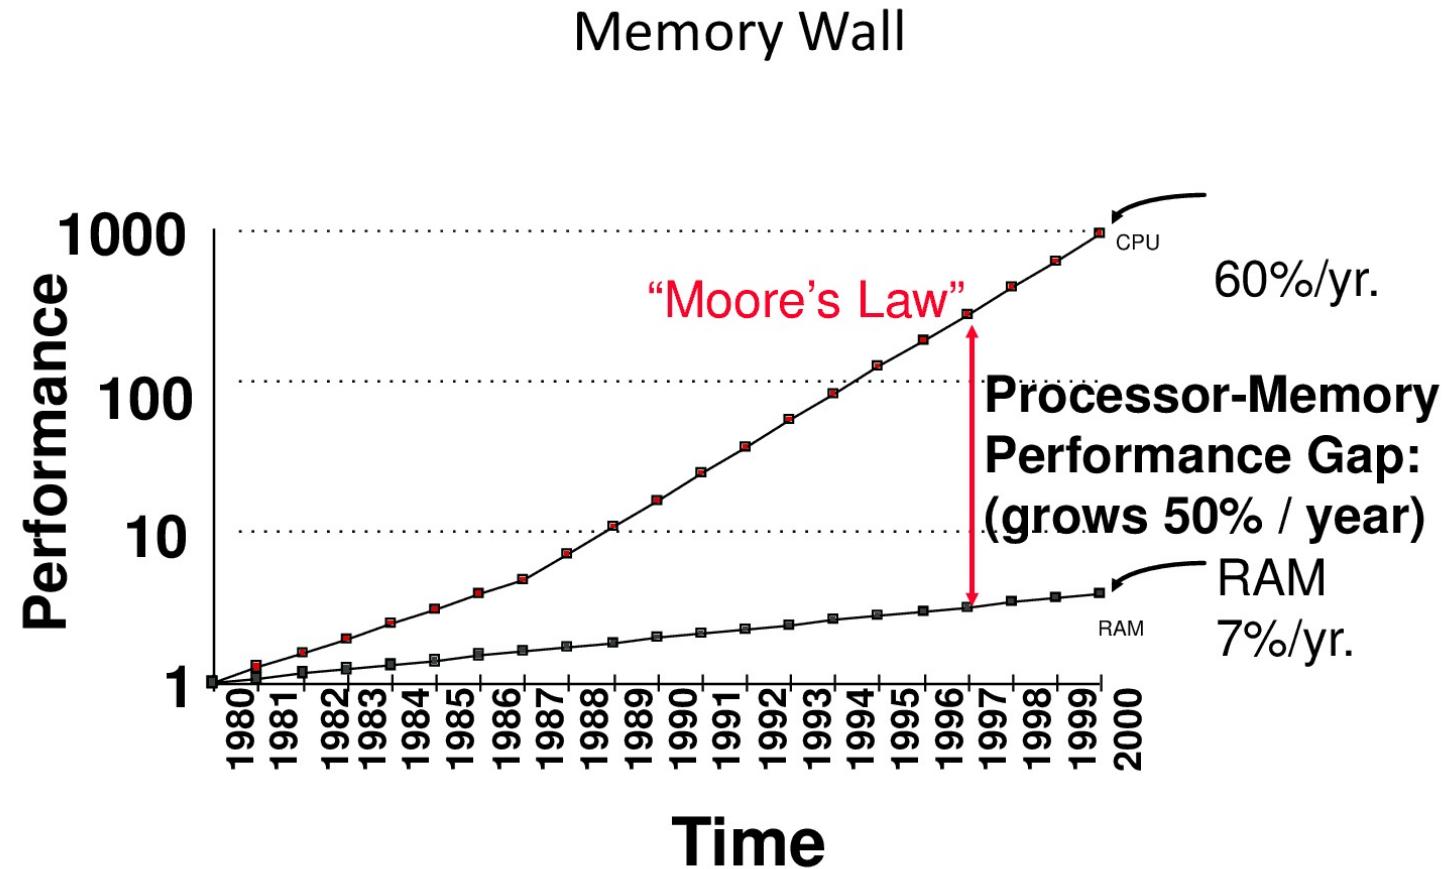
Remedies Towards Breaching The Memory Wall for Sparse Computations

Mert Hidayetoglu¹, Leopold Grinberg², Constantinos Evangelinos²,
Jinjun Xiong², and Wen-mei W. Hwu¹

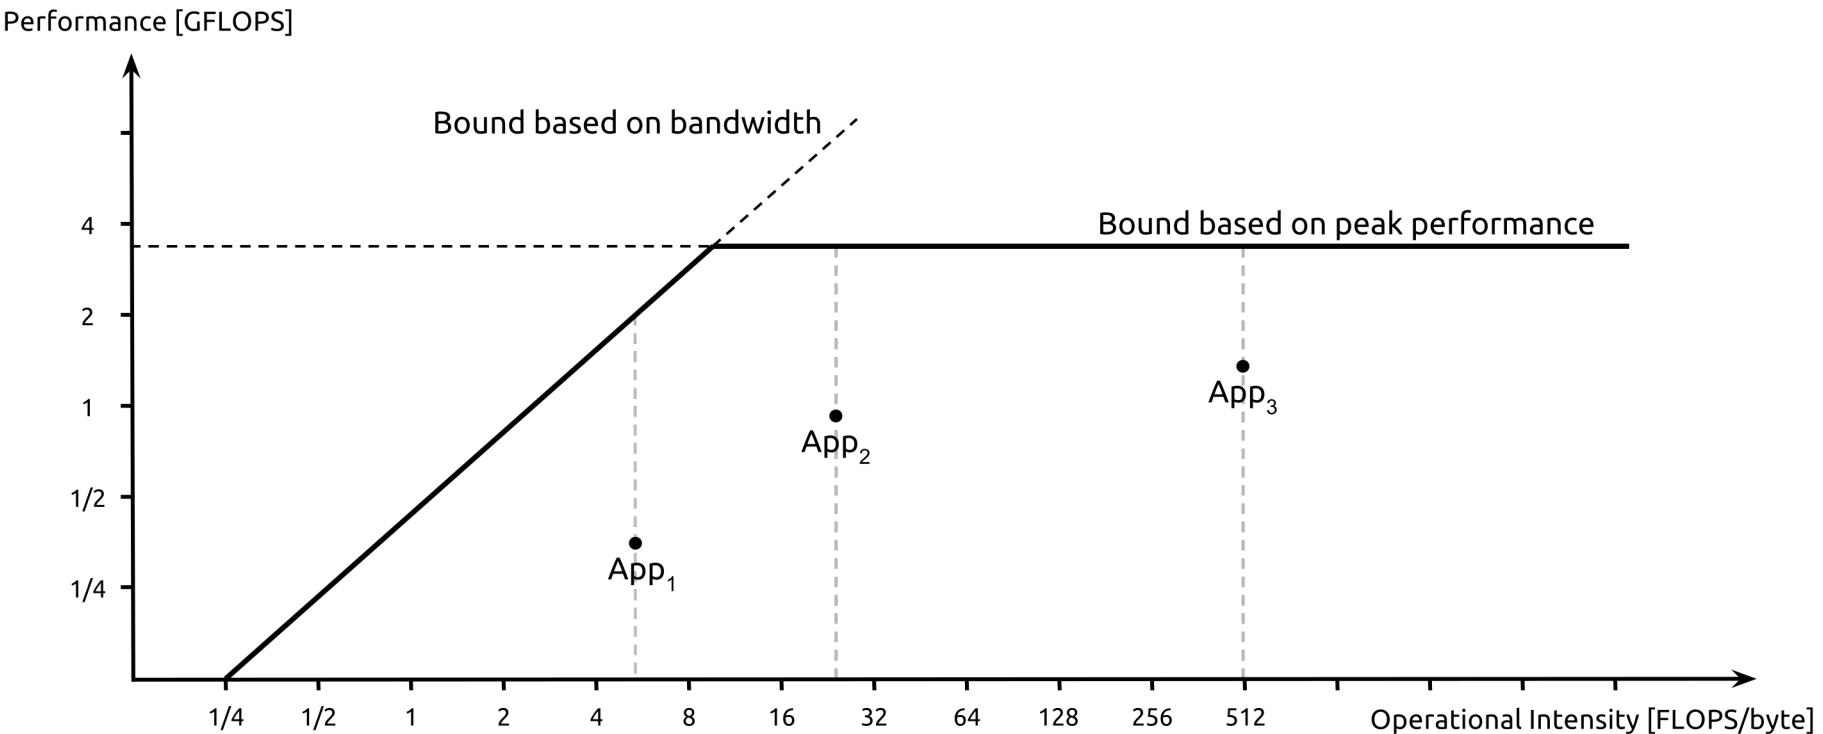
¹University of Illinois at Urbana-Champaign

²IBM T. J. Watson Research Center

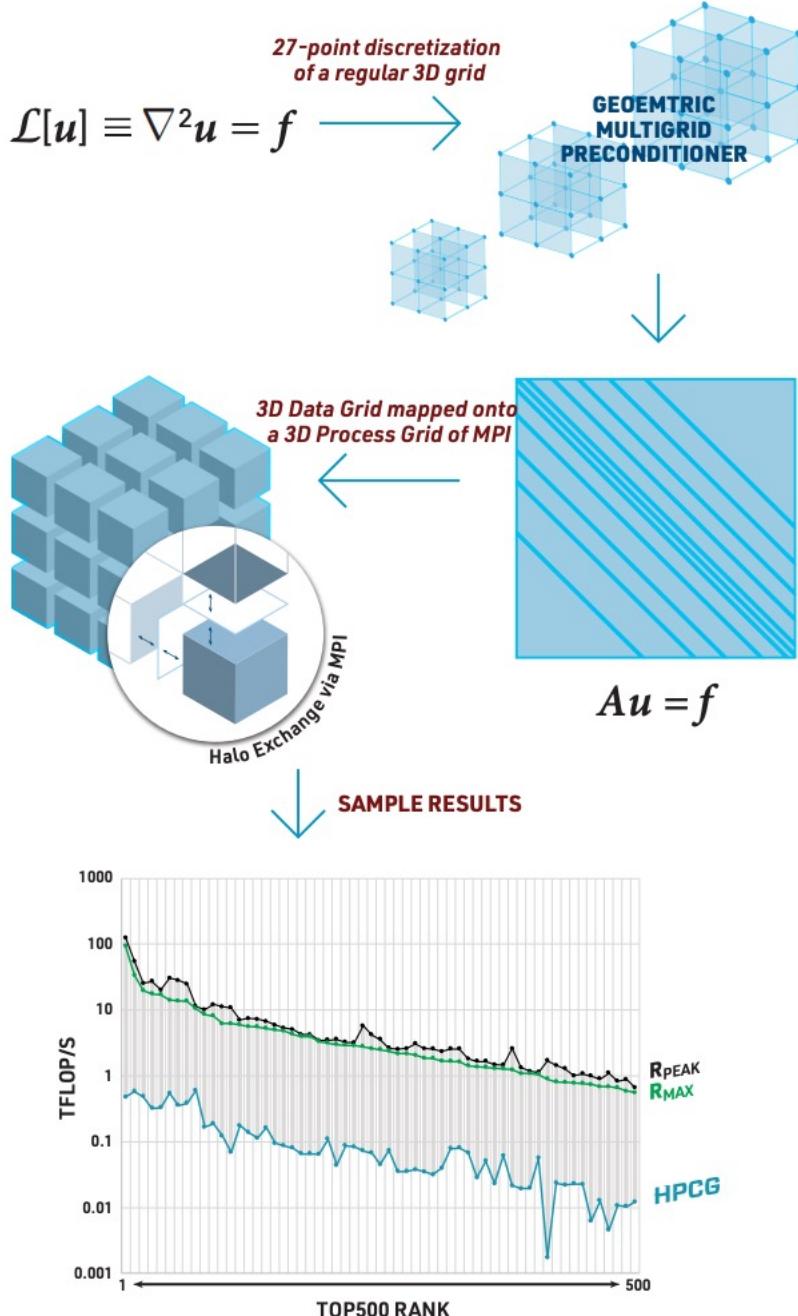
Memory Wall



Roofline Model



```
//SPMV
for(int m = 0; m < nump; m++){
    double reduce = 0;
    for(int n = displ[m]; n < displ[m+1]; n++)
        reduce += value[n]*x[index[n]];
    y[m] = reduce;
}
```



HPCG

HPCG List for June 2019

TOP500			Cores	Rmax (TFlop/s)	HPCG (TFlop/s)
Rank	Rank	System			
1	1	Summit - IBM Power System AC922, IBM POWER9 22C 3.07GHz, NVIDIA Volta GV100, Dual-rail Mellanox EDR Infiniband , IBM DOE/SC/Oak Ridge National Laboratory United States	2,414,592	148,600.0	2925.75
2	2	Sierra - IBM Power System S922LC, IBM POWER9 22C 3.1GHz, NVIDIA Volta GV100, Dual-rail Mellanox EDR Infiniband , IBM / NVIDIA / Mellanox DOE/NNSA/LLNL United States	1,572,480	94,640.0	1795.67
3	20	K computer, SPARC64 VIIIIfx 2.0GHz, Tofu interconnect , Fujitsu RIKEN Advanced Institute for Computational Science (AICS) Japan	705,024	10,510.0	602.74
4	7	Trinity - Cray XC40, Xeon E5-2698v3 16C 2.3GHz, Intel Xeon Phi 7250 68C 1.4GHz, Aries interconnect , Cray Inc. DOE/NNSA/LANL/SNL United States	979,072	20,158.7	546.12
5	8	AI Bridging Cloud Infrastructure (ABCi) - PRIMERGY CX2570 M4, Xeon Gold 6148 20C 2.4GHz, NVIDIA Tesla V100 SXM2, Infiniband EDR , Fujitsu National Institute of Advanced Industrial Science and Technology (AIST) Japan	391,680	19,880.0	508.85

www.top500.org

LINPACK

Rank	System	Cores	Rmax (TFlop/s)	Rpeak (TFlop/s)	Power (kW)
1	Summit - IBM Power System AC922, IBM POWER9 22C 3.07GHz, NVIDIA Volta GV100, Dual-rail Mellanox EDR Infiniband , IBM DOE/SC/Oak Ridge National Laboratory United States	2,414,592	148,600.0	200,794.9	10,096
2	Sierra - IBM Power System S922LC, IBM POWER9 22C 3.1GHz, NVIDIA Volta GV100, Dual-rail Mellanox EDR Infiniband , IBM / NVIDIA / Mellanox DOE/NNSA/LLNL United States	1,572,480	94,640.0	125,712.0	7,438
3	Sunway TaihuLight - Sunway MPP, Sunway SW26010 260C 1.45GHz, Sunway , NRCPC National Supercomputing Center in Wuxi China	10,649,600	93,014.6	125,435.9	15,371
4	Tianhe-2A - TH-IVB-FEP Cluster, Intel Xeon E5-2692v2 12C 2.2GHz, TH Express-2, Matrix-2000 , NUDT National Super Computer Center in Guangzhou China	4,981,760	61,444.5	100,678.7	18,482
5	Frontera - Dell C6420, Xeon Platinum 8280 28C 2.7GHz, Mellanox InfiniBand HDR , Dell EMC Texas Advanced Computing Center/Univ. of Texas United States	448,448	23,516.4	38,745.9	

www.top500.org

Sparse data
Low FLOPS
Stresses memory syst.

$\mathcal{O}(N)$ Memory
 $\mathcal{O}(N)$ Comput.

Dense data
High FLOPS
Stresses processors

$\mathcal{O}(N^2)$ Memory
 $\mathcal{O}(N^3)$ Comput.

Computational Kernels in Each Iteration

Other HPCG Operations

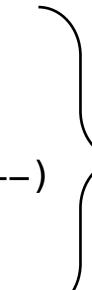
```
//GAUSS-SEIDEL FORWARD SWEEP
for(int m = 0; m < nump; m++){
    double reduce = b[m];
    for(int n = displ[m]; n < displ[m+1]; n++)
        reduce -= value[n]*x[index[n]];
    x[m] += reduce/diag[m];
}
```



Inherently Sequential

Other HPCG Operations

```
//GAUSS-SEIDEL BACKWARD SWEEP
for(int m = nump-1; m > -1; m--){
    double reduce = b[m];
    for(int n = displ[m+1]-1; n > displ[m]-1; n--)
        reduce -= value[n]*x[index[n]];
    x[m] += reduce/diag[m];
}
//SPMV
```



Inherently Sequential

Other HPCG Operations

```
for(int m = 0; m < nump; m++){
    double reduce = 0;
    for(int n = displ[m]; n < displ[m+1]; n++)
        reduce += value[n]*x[index[n]];
    y[m] = reduce;
}
```

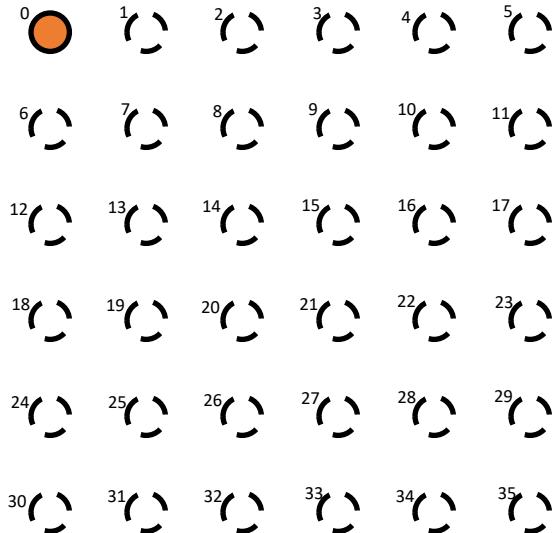


Easily Parallelized

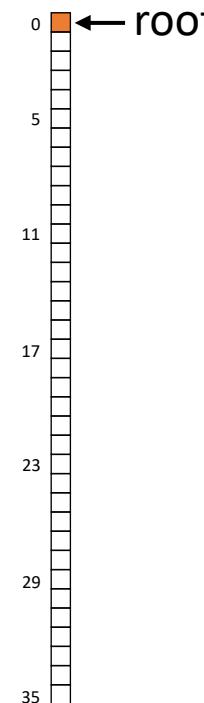
Batching Algorithm

Row-Major Order

Batch 0 (1)



Batch Size: 10

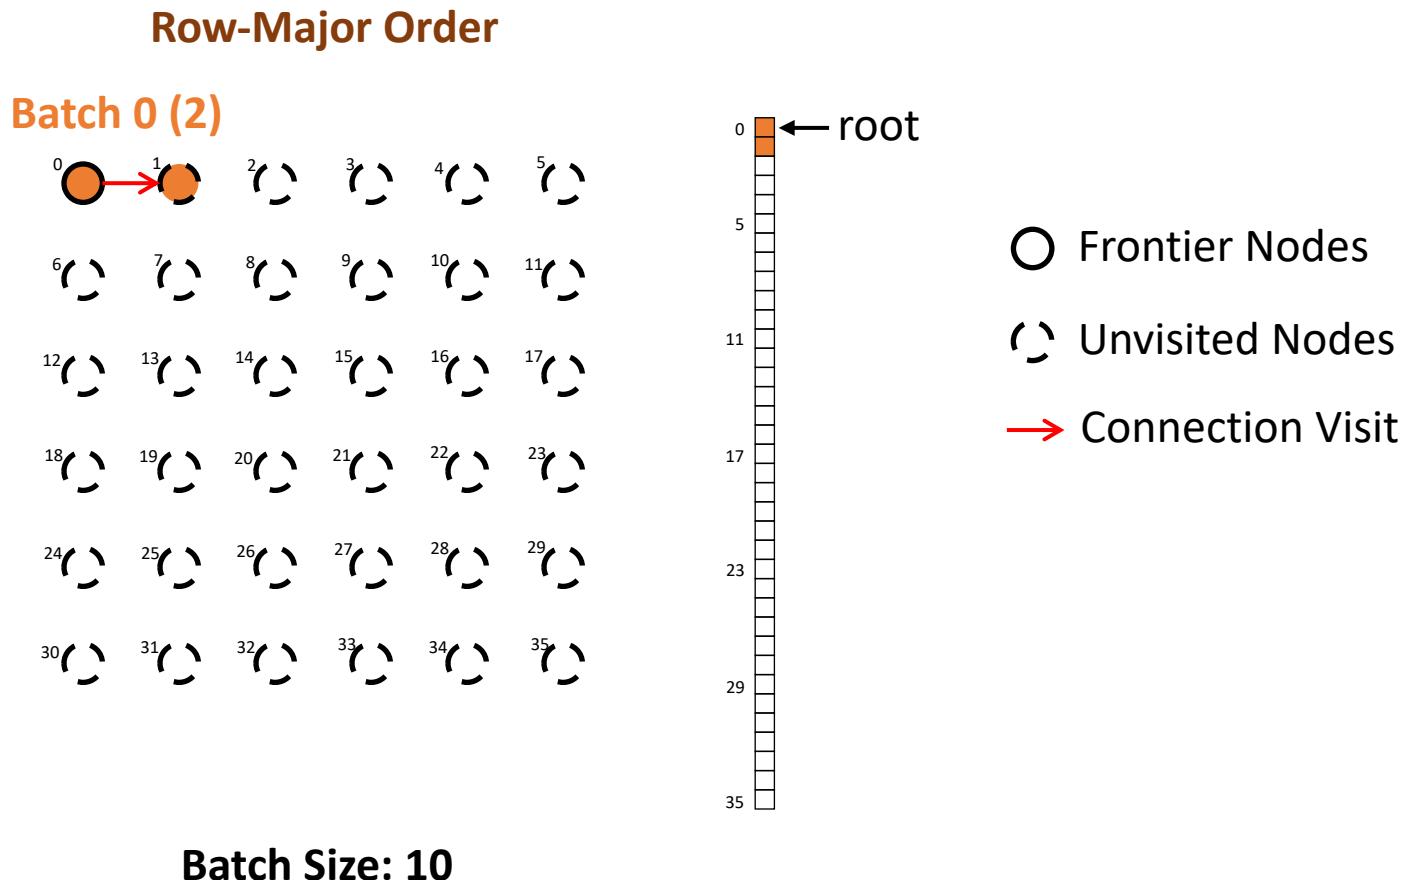


○ Frontier Nodes

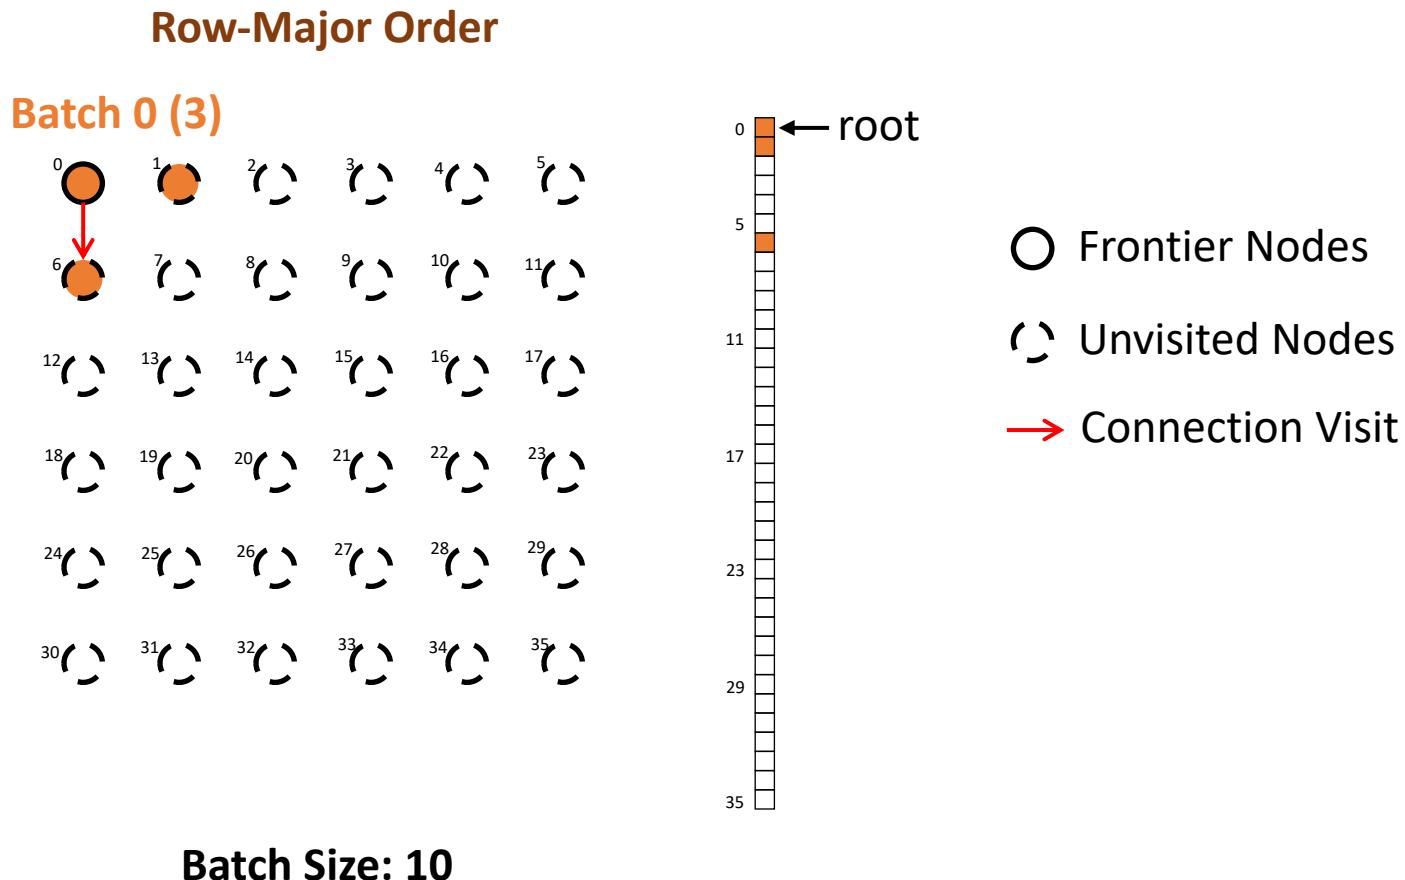
○ Unvisited Nodes

→ Connection Visit

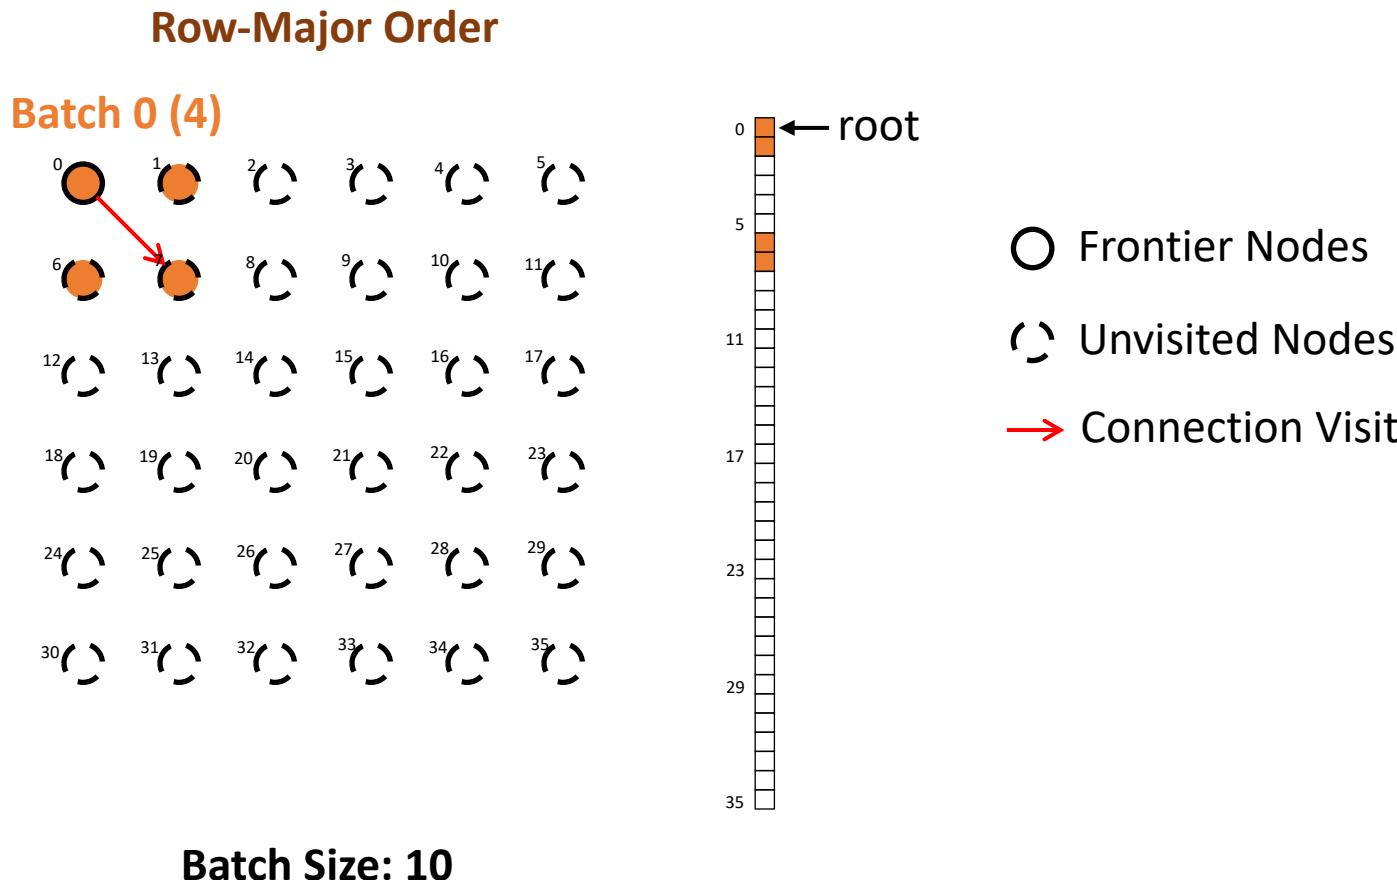
Batching Algorithm



Batching Algorithm



Batching Algorithm

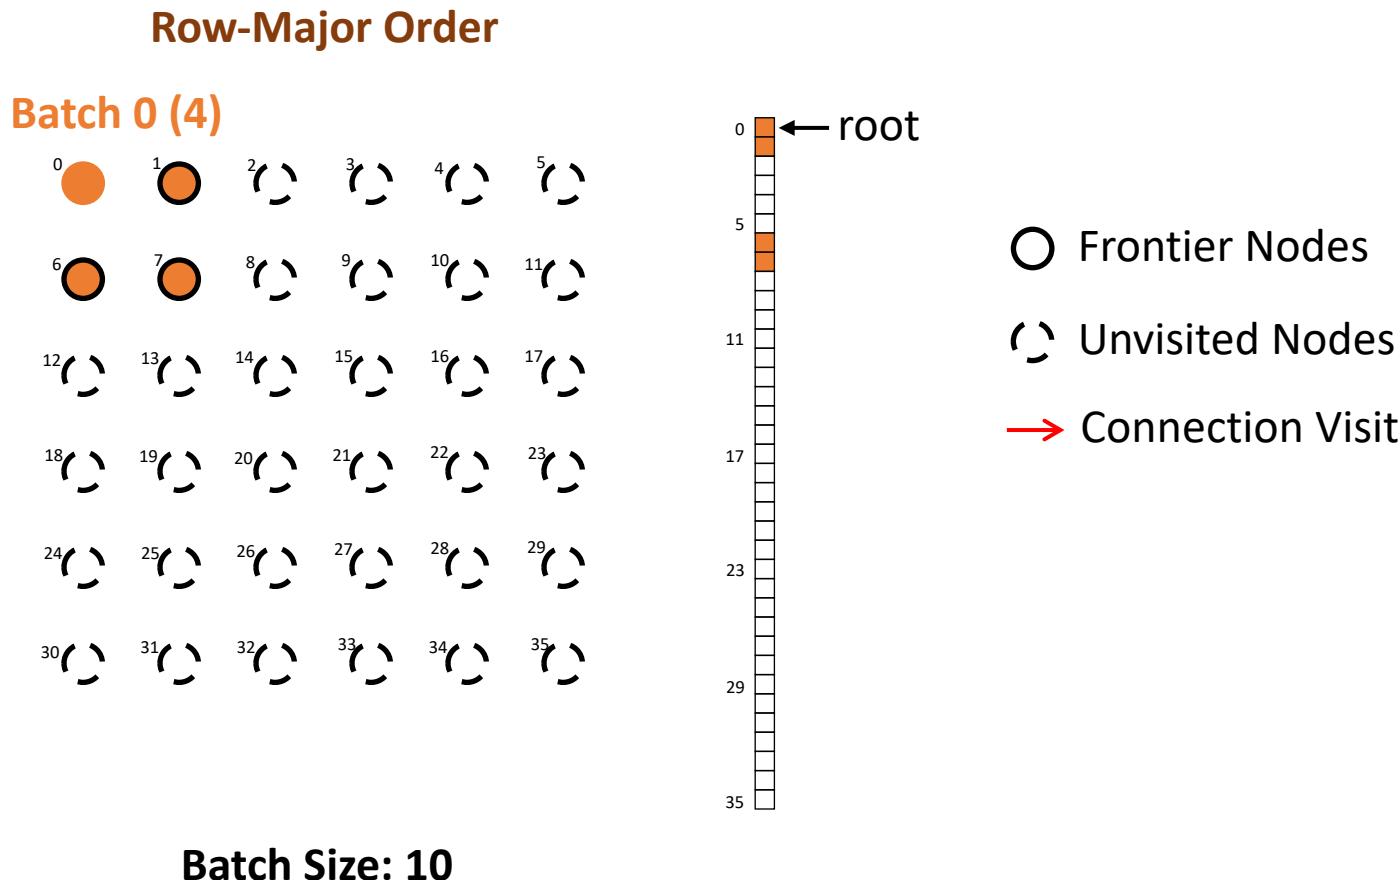


IBM Research

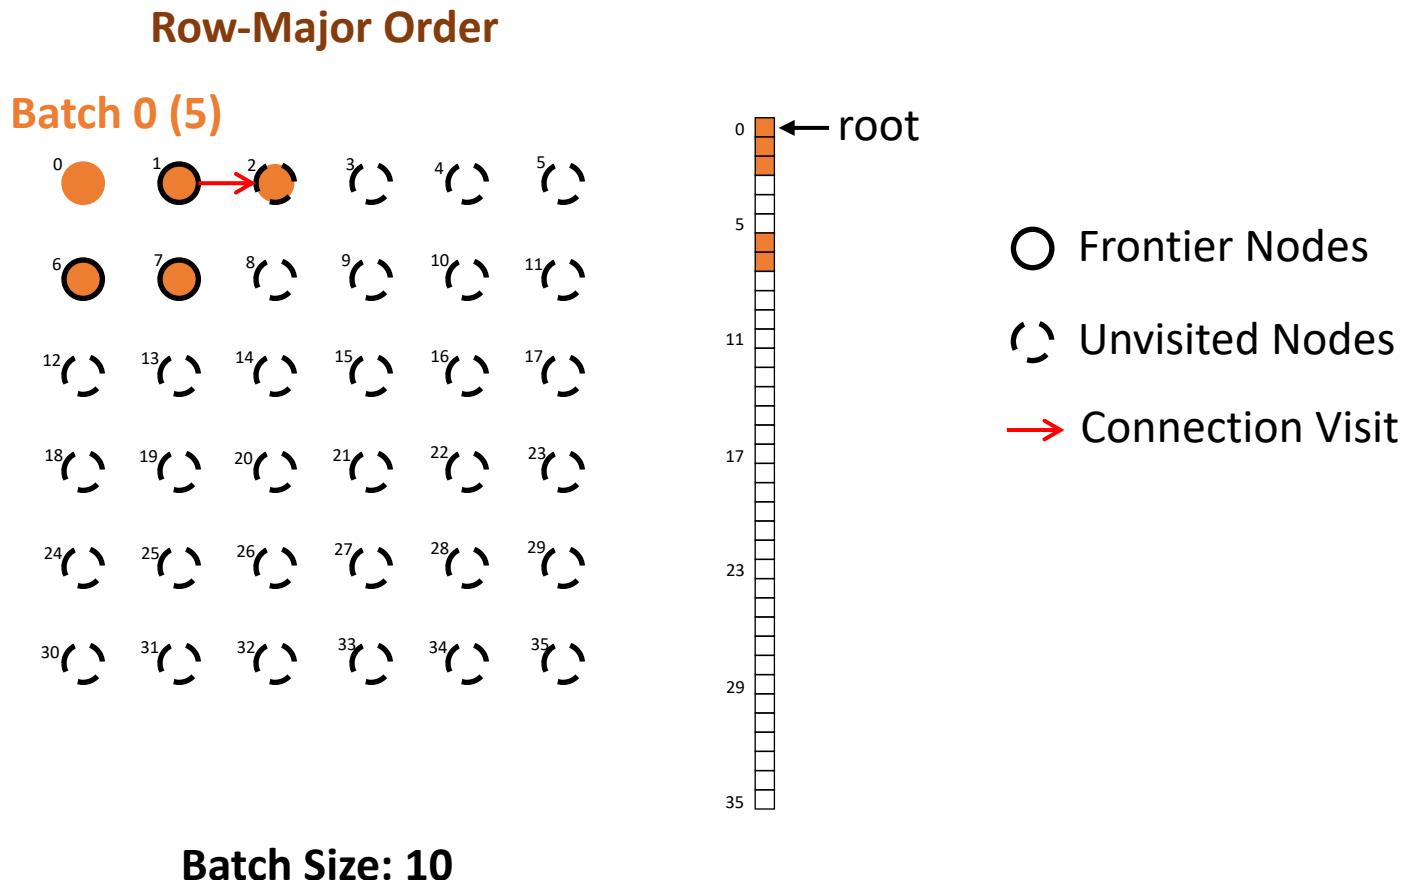


center for
cognitive computing
systems research

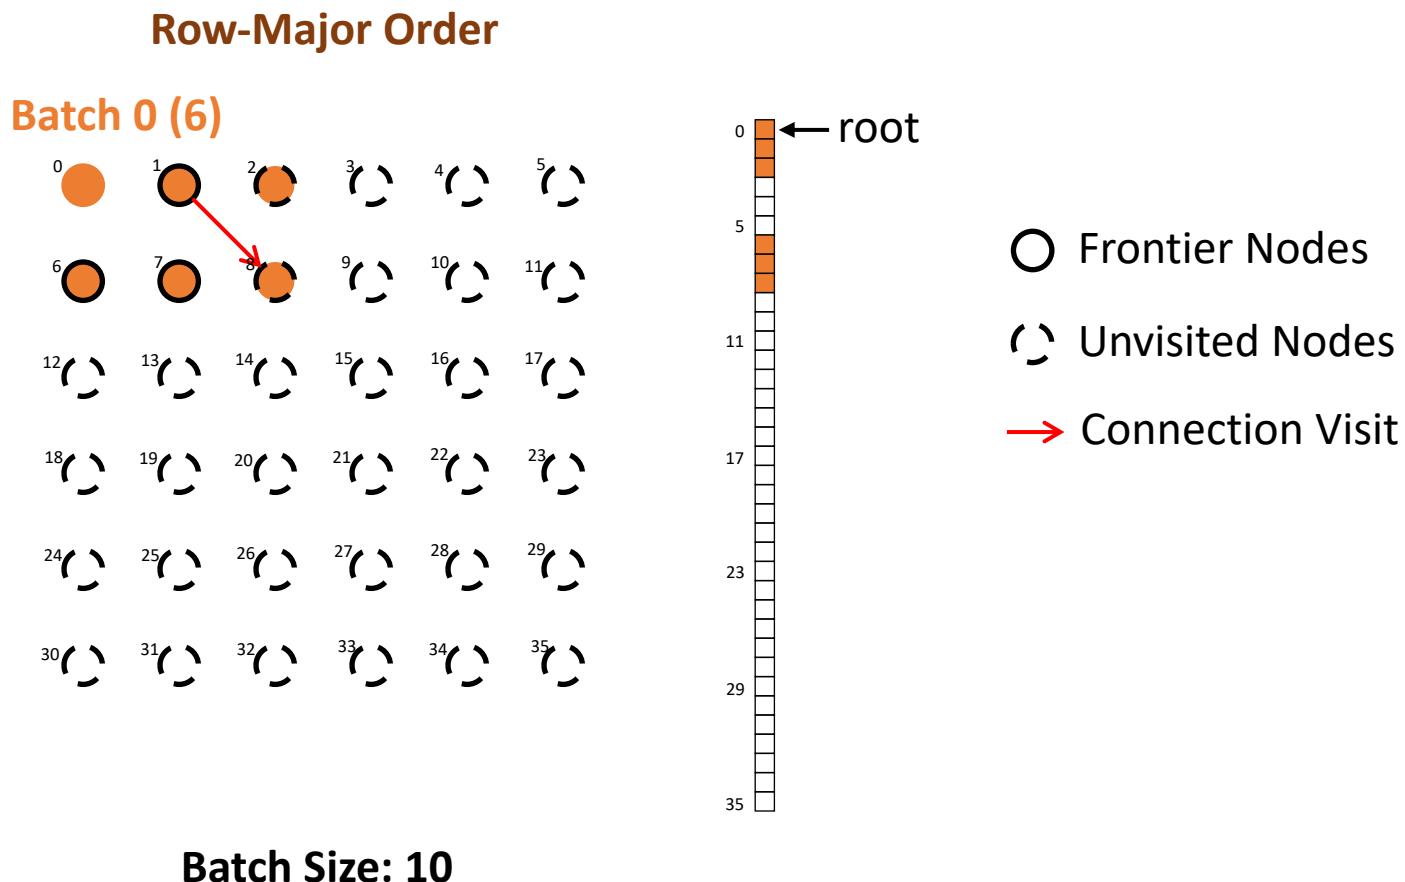
Batching Algorithm



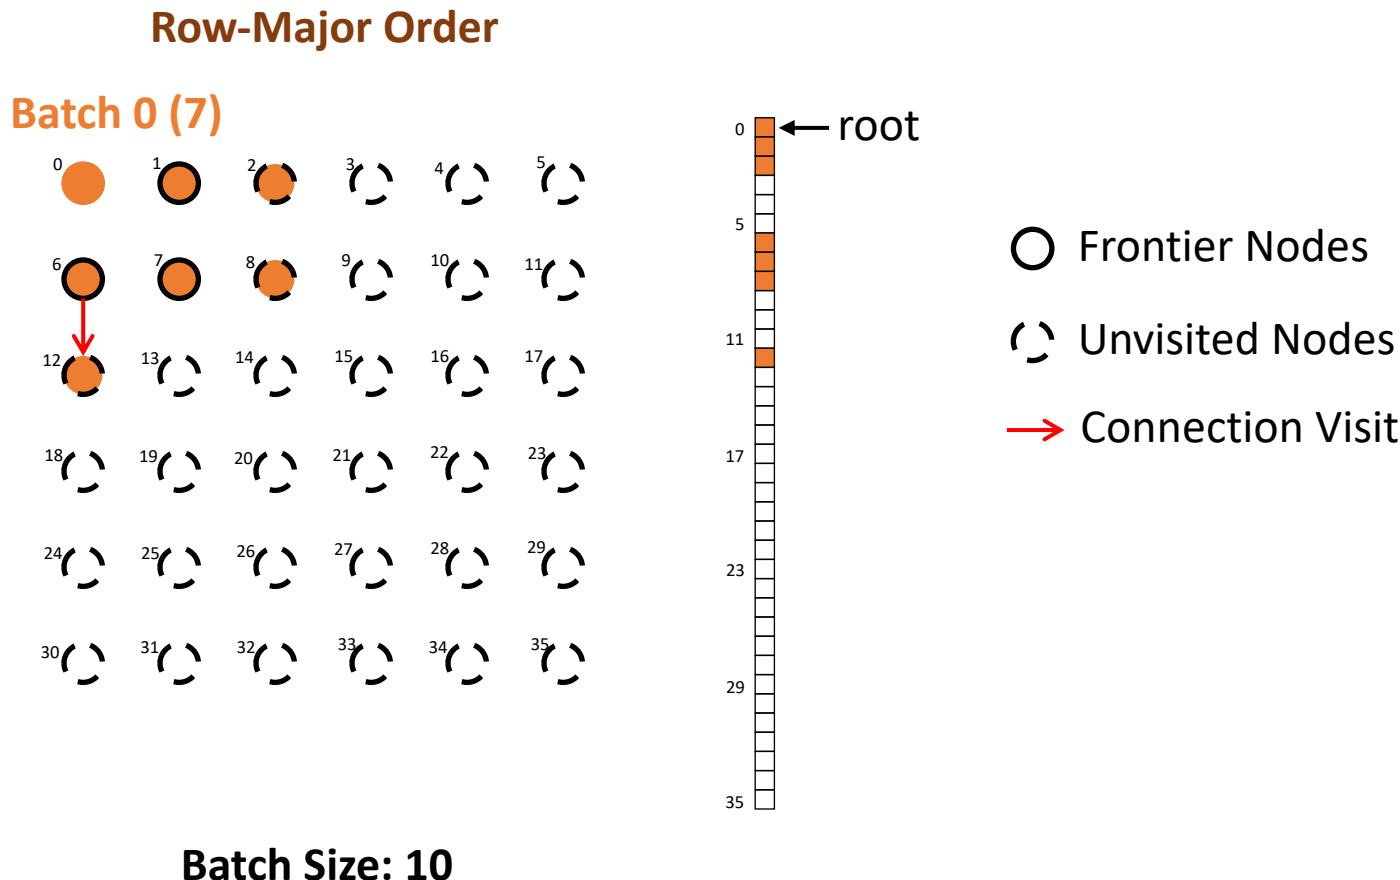
Batching Algorithm



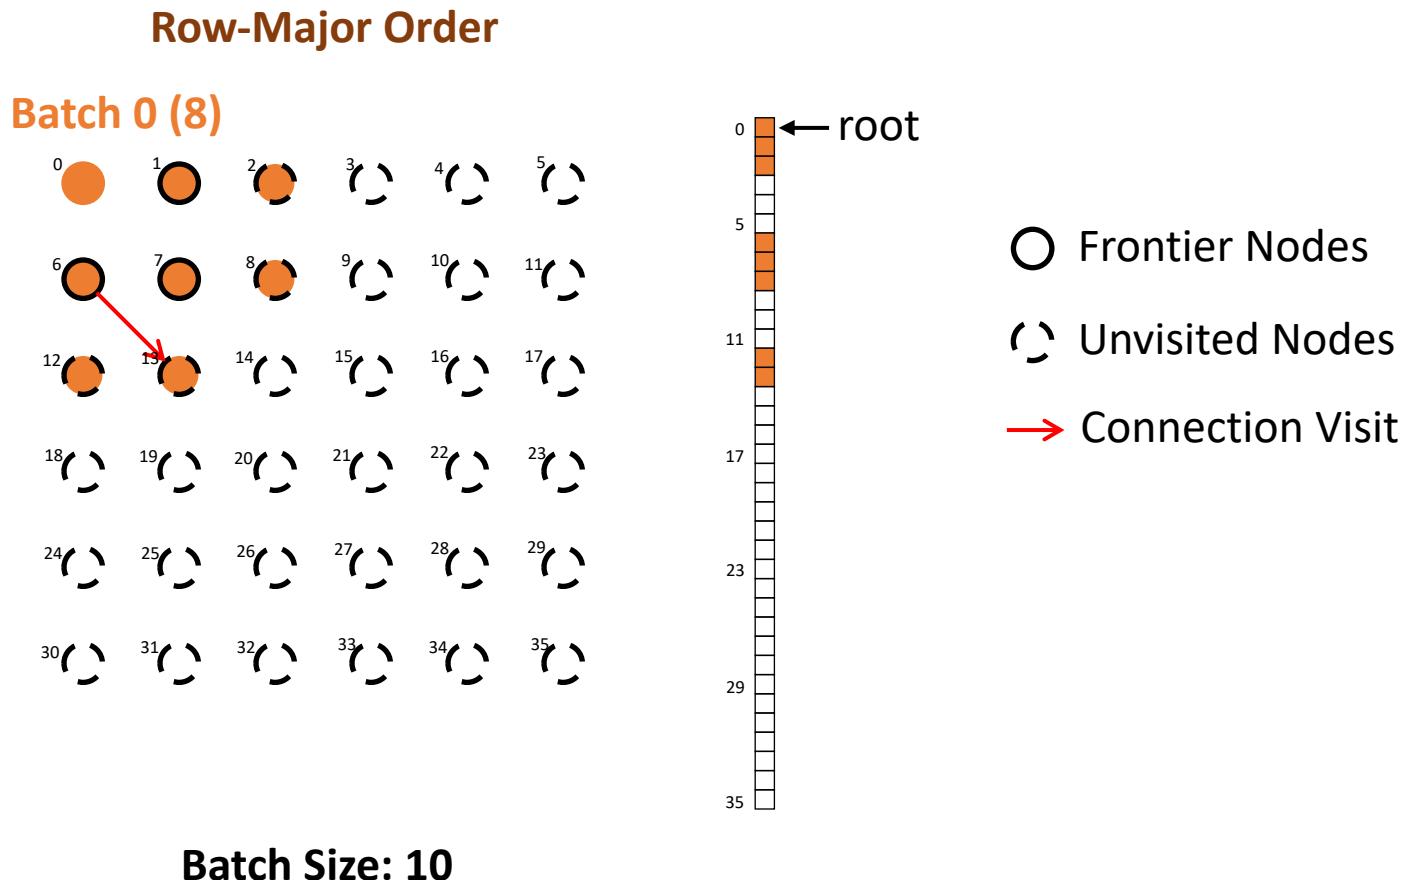
Batching Algorithm



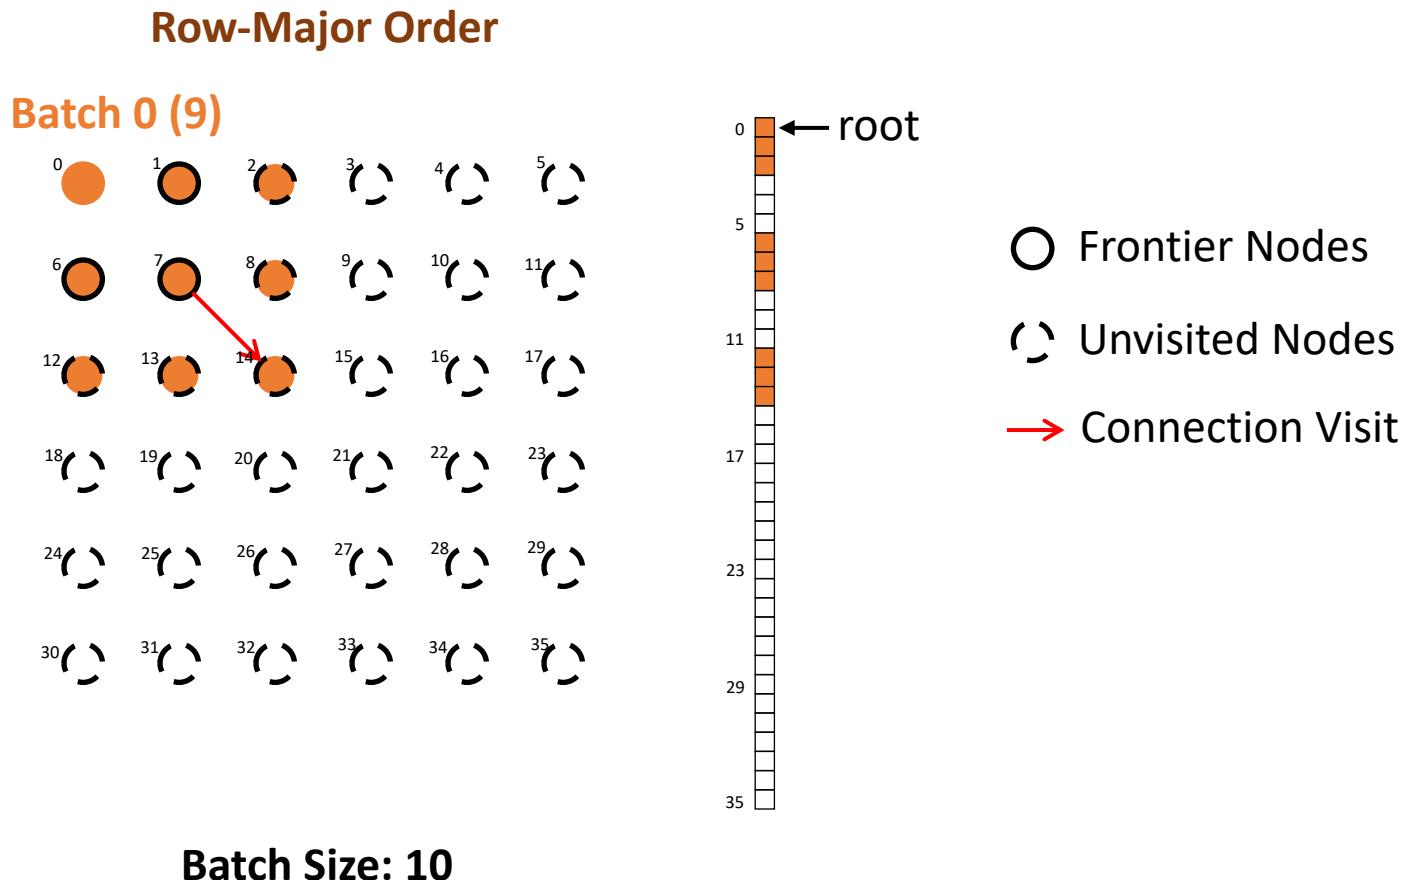
Batching Algorithm



Batching Algorithm



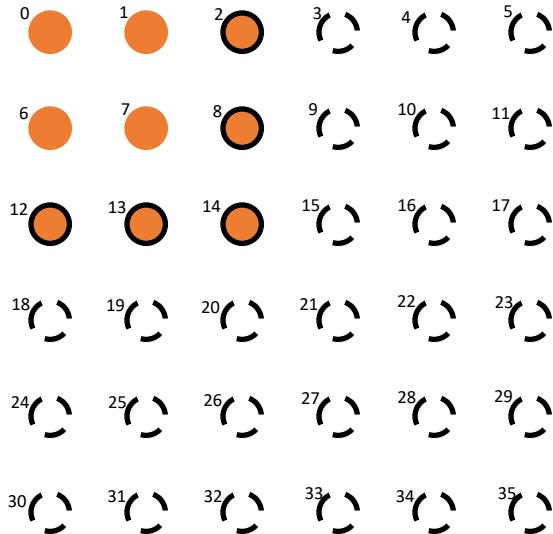
Batching Algorithm



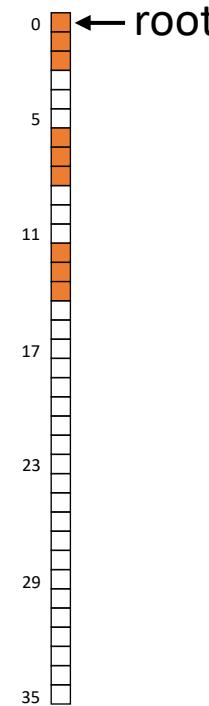
Batching Algorithm

Row-Major Order

Batch 0 (9)



Batch Size: 10

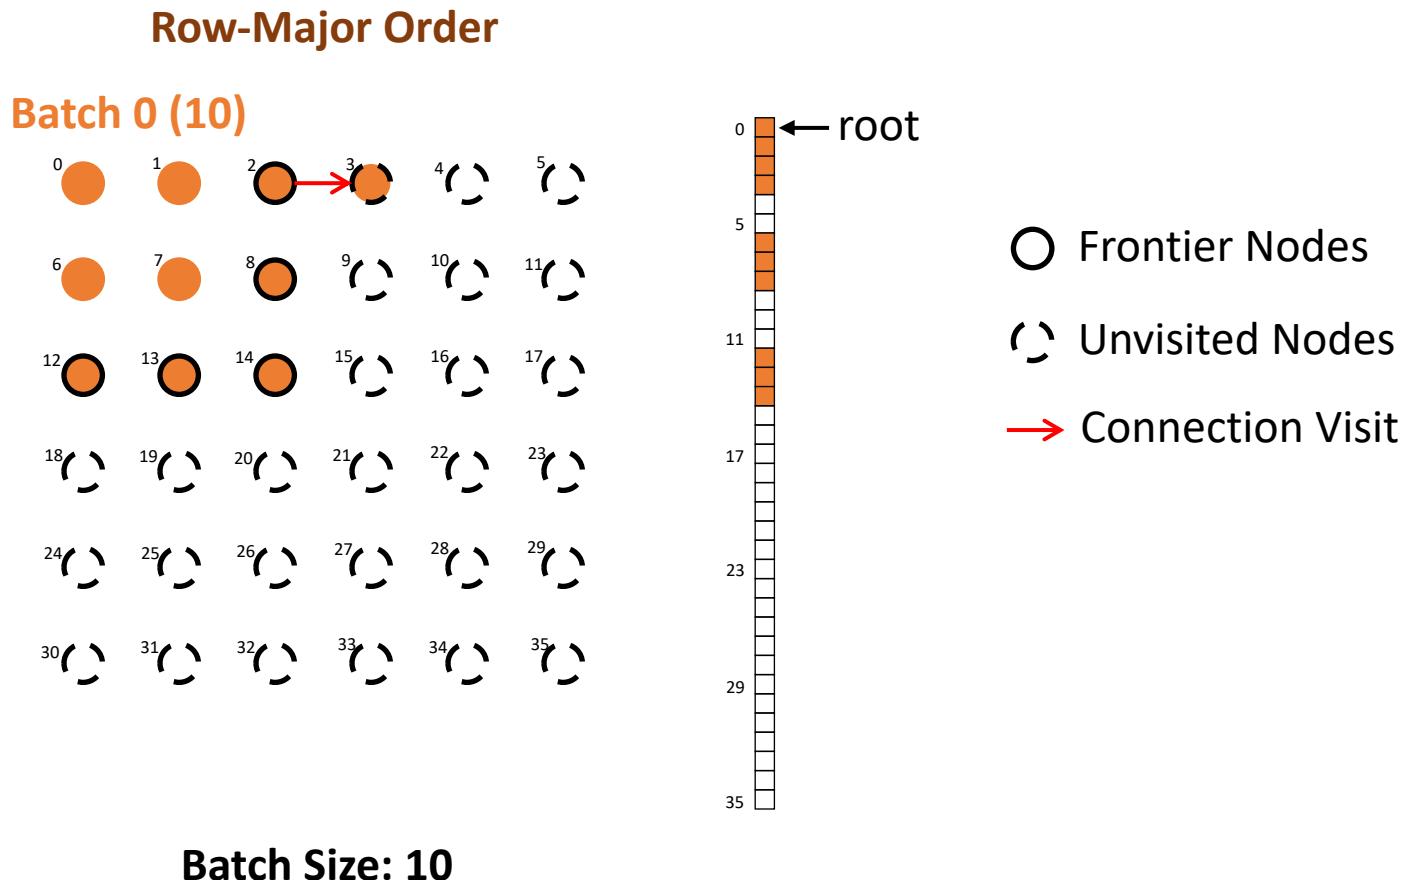


○ Frontier Nodes

○ Unvisited Nodes

→ Connection Visit

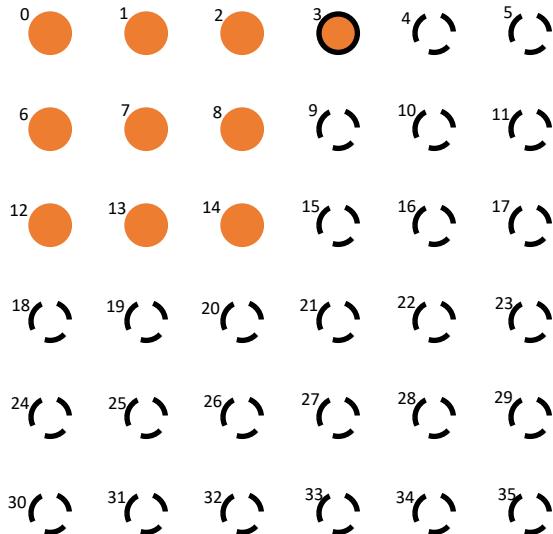
Batching Algorithm



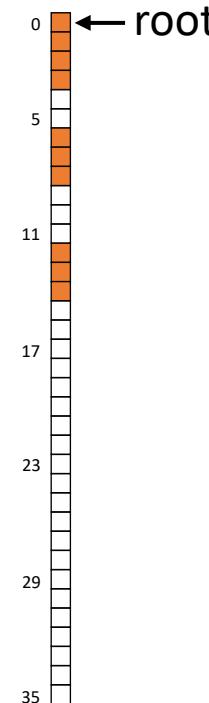
Batching Algorithm

Row-Major Order

Batch 0 (10)



Batch Size: 10

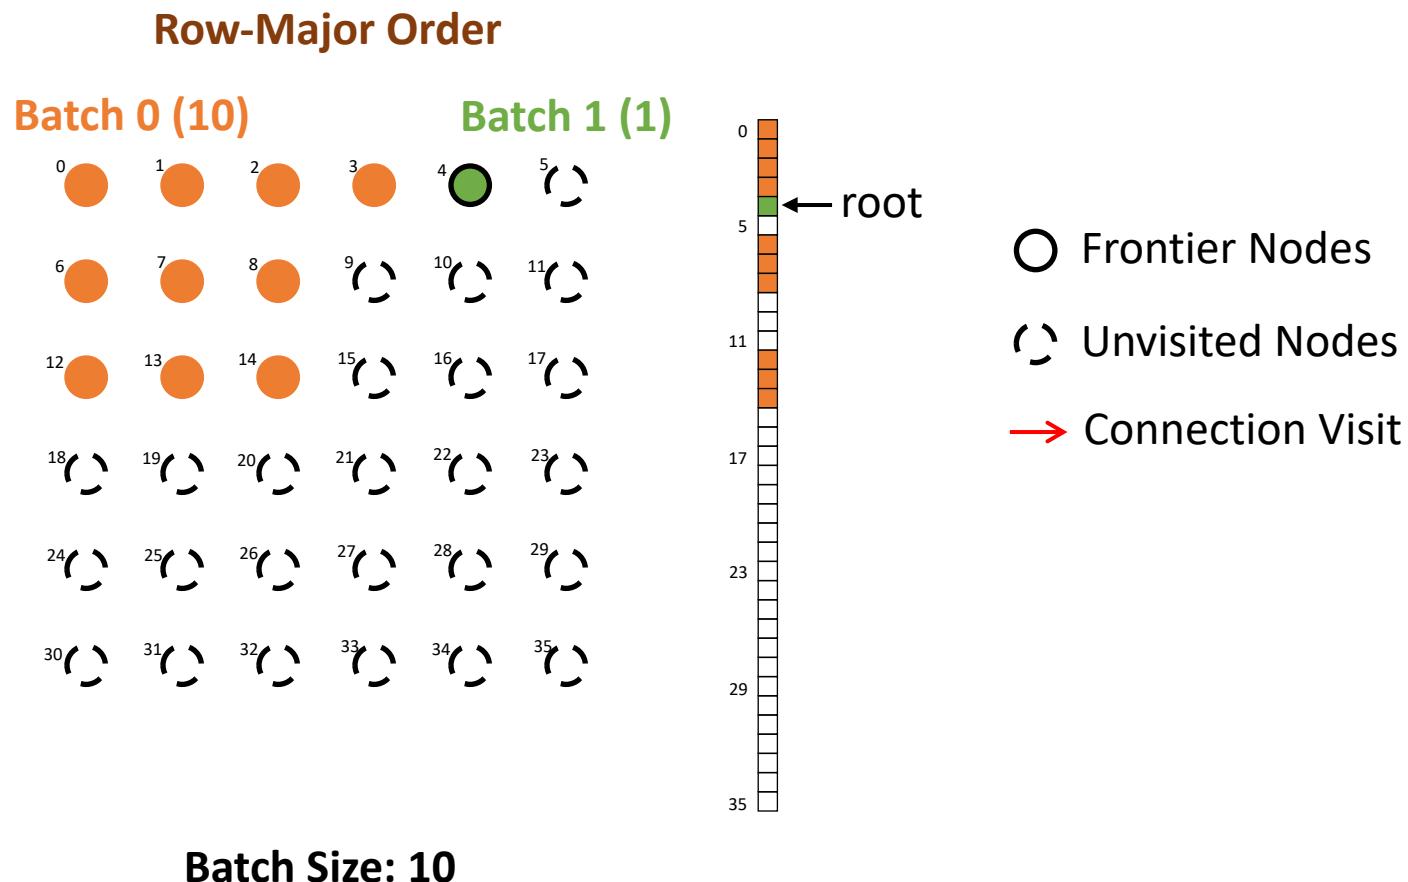


○ Frontier Nodes

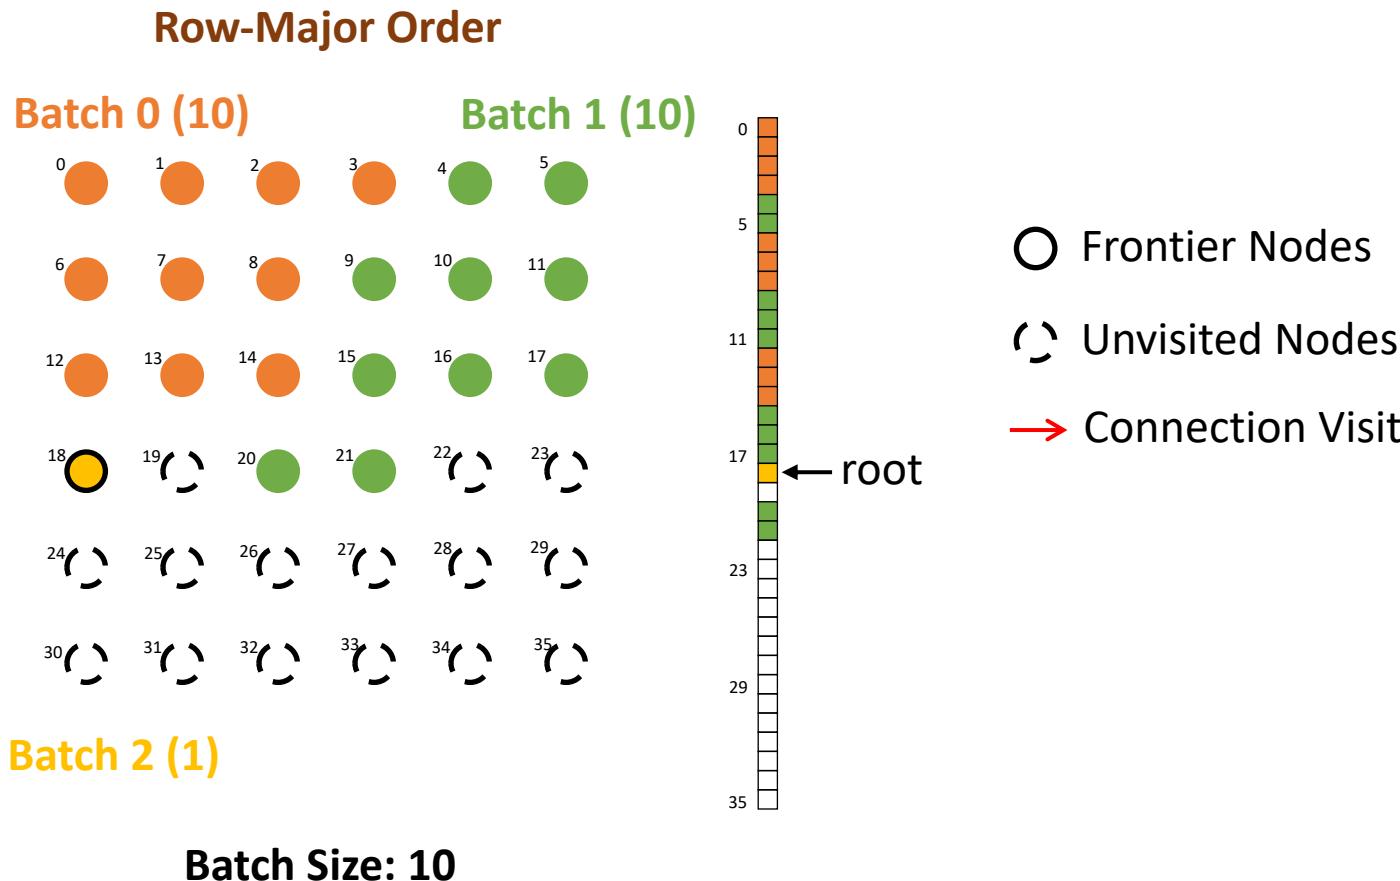
○ Unvisited Nodes

→ Connection Visit

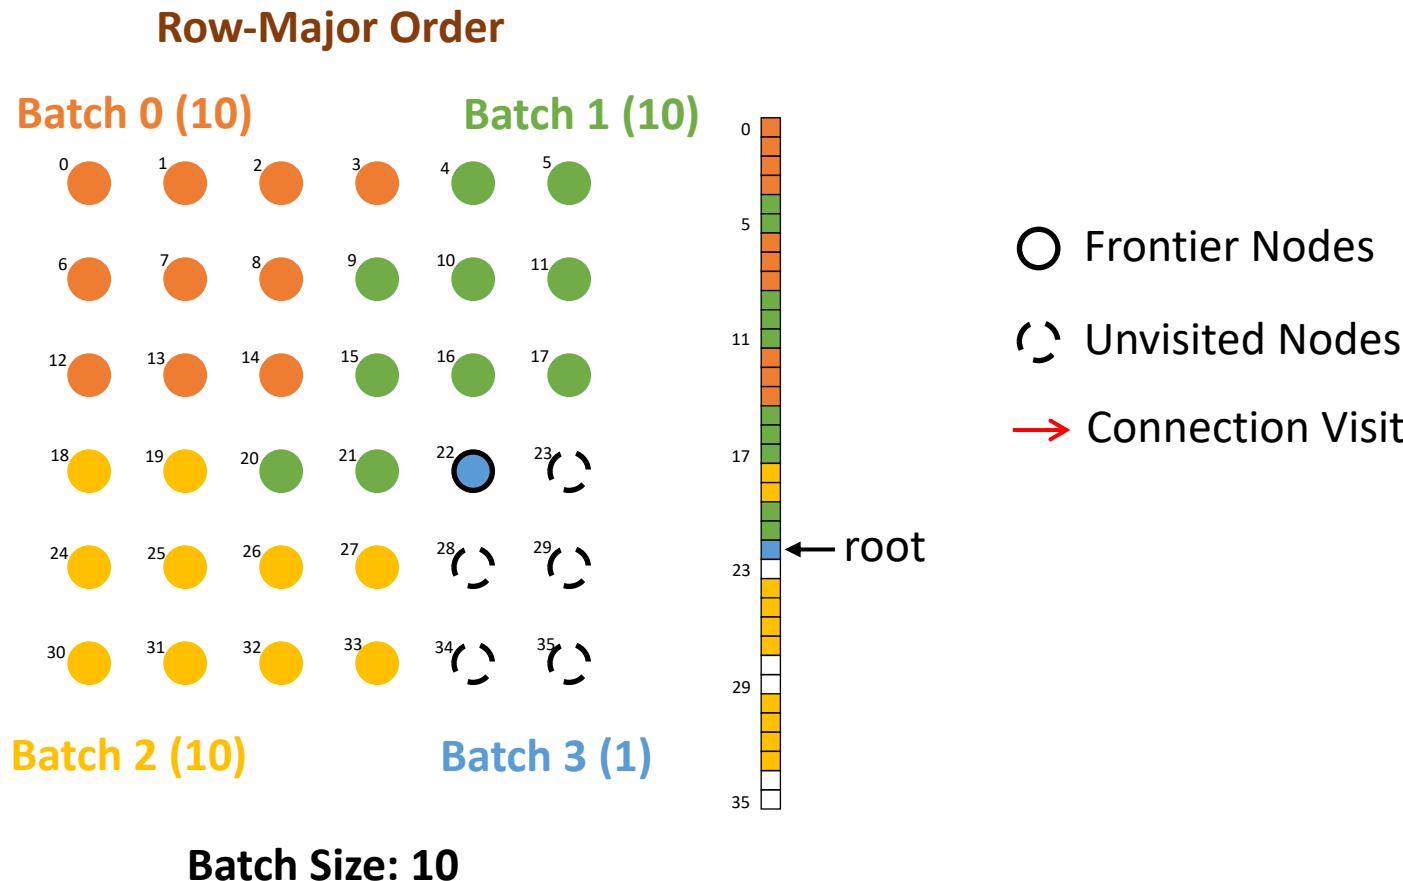
Batching Algorithm



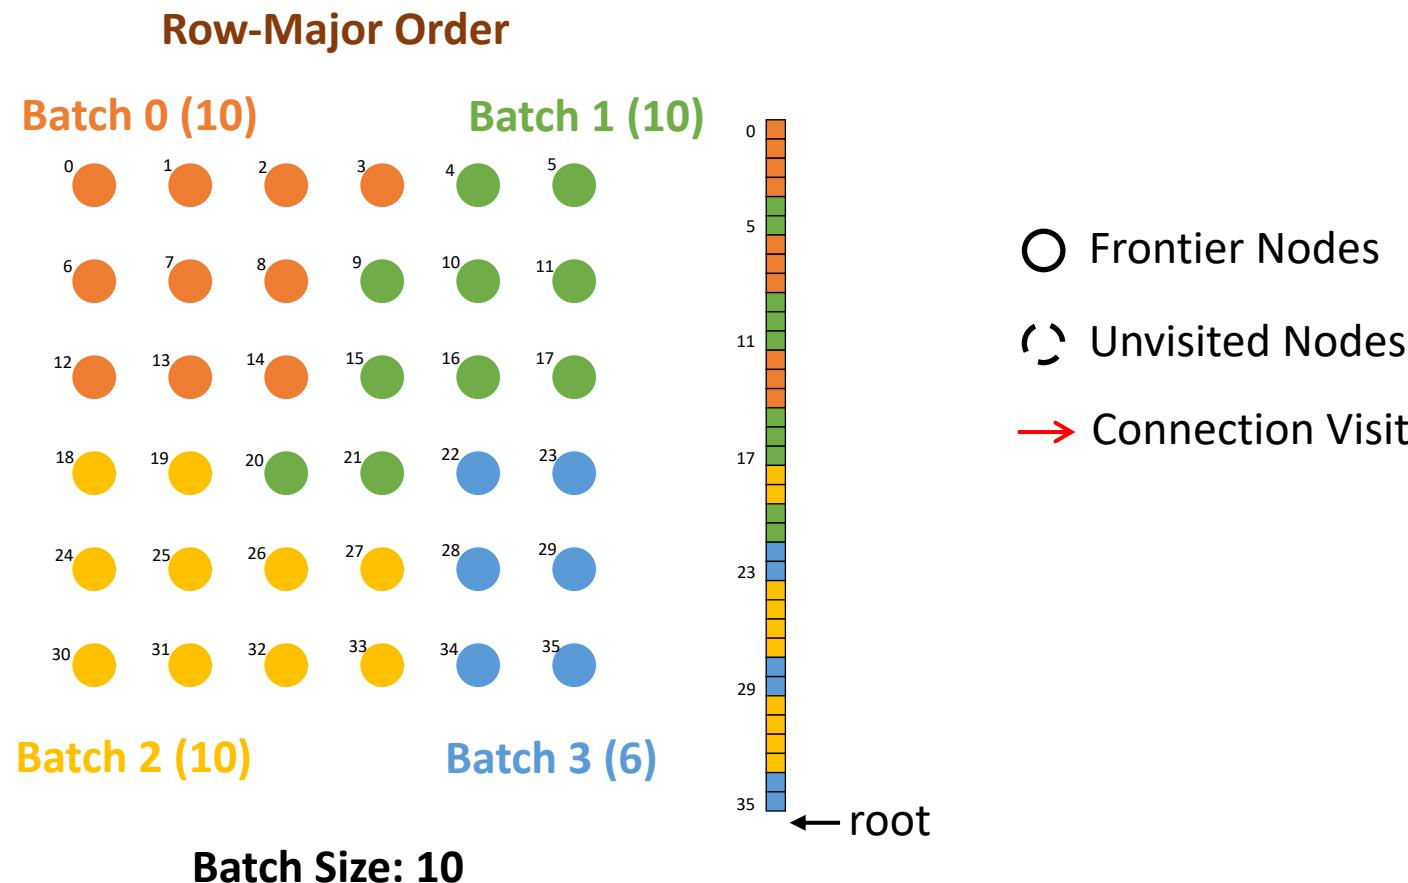
Batching Algorithm



Batching Algorithm



Batching Algorithm

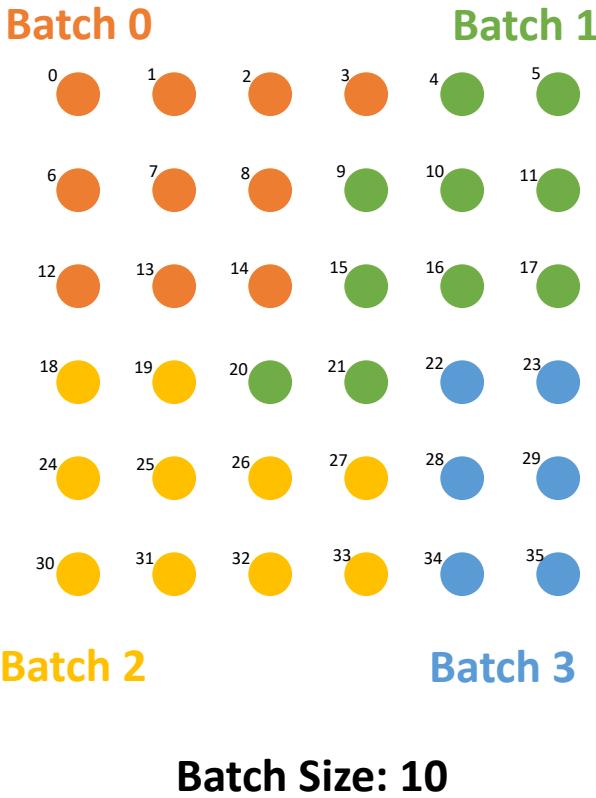




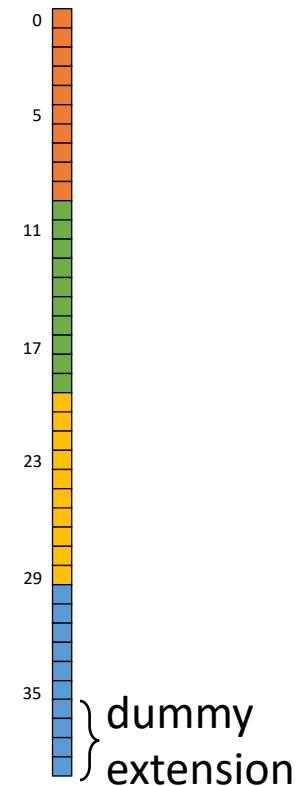
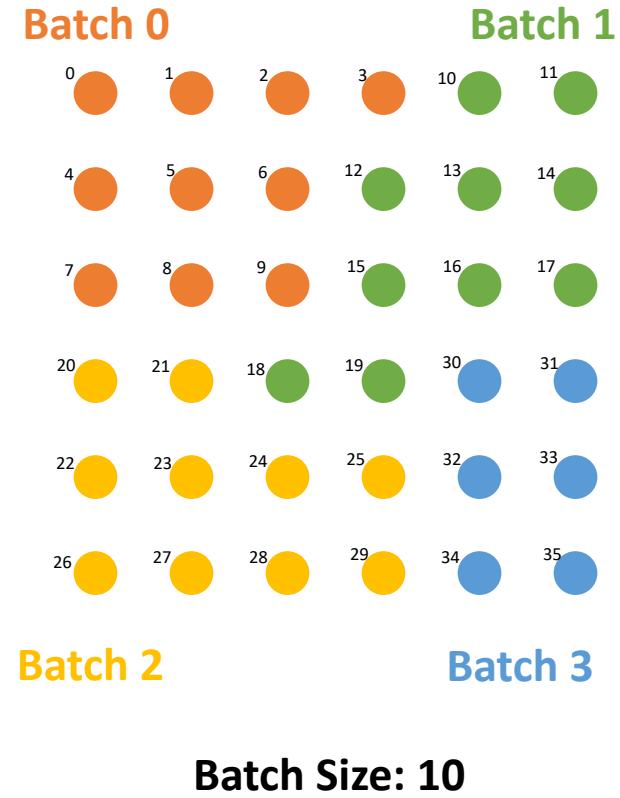
IDEA: Reorder to increase cache reuse

Comput. & Memory Complexity: $\mathcal{O}(N)$

Row-Major Order



Batch Reordering



Original HPCG Gauss-Seidel

```
//GAUSS-SEIDEL FORWARD SWEEP
for(int m = 0; m < nump; m++){
    double reduce = b[m];
    for(int n = displ[m]; n < displ[m+1]; n++)
        reduce -= value[n]*x[index[n]];
    x[m] += reduce/diag[m];
}
```

Memory Transactions

Auxiliary Data: $3N \times 8$ Bytes Read

Matrix Data: $27N \times (8 + 4)$ Bytes Read

Vector Data: $27N \times 8$ Bytes Read *Irregular Access*

$N \times 8$ Bytes Read

$N \times 8$ Bytes Write

Gauss-Seidel with Reordering

```
//GAUSS-SEIDEL FORWARD SWEEP
for(int m = 0; m < nump; m++){
    double reduce = b[m];
    for(int n = displ[m]; n < displ[m+1]; n++)
        reduce -= value[n]*x[index[n]];
    x[m] += reduce/diag[m];
}
```

No kernel code change!

Memory Transactions

Auxiliary Data: $3N \times 8$ Bytes Read

Matrix Data: $27N \times (8 + 4)$ Bytes Read

Vector Data: $27N \times 8$ Bytes Read Maximum cache-reuse: 27

$N \times 8$ Bytes Read

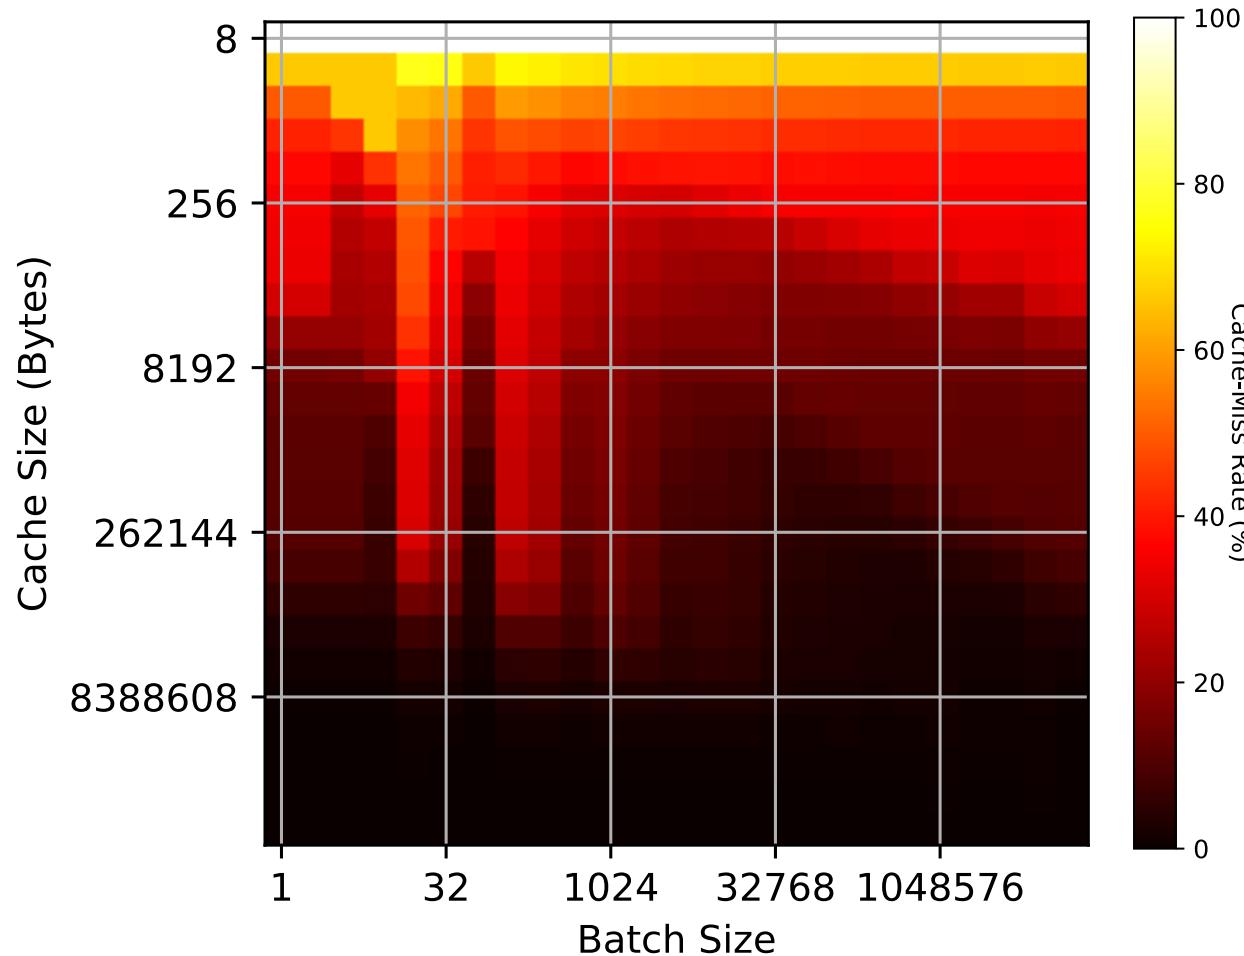
$N \times 8$ Bytes Write



IDEA: Use a buffer for explicit vector data reuse

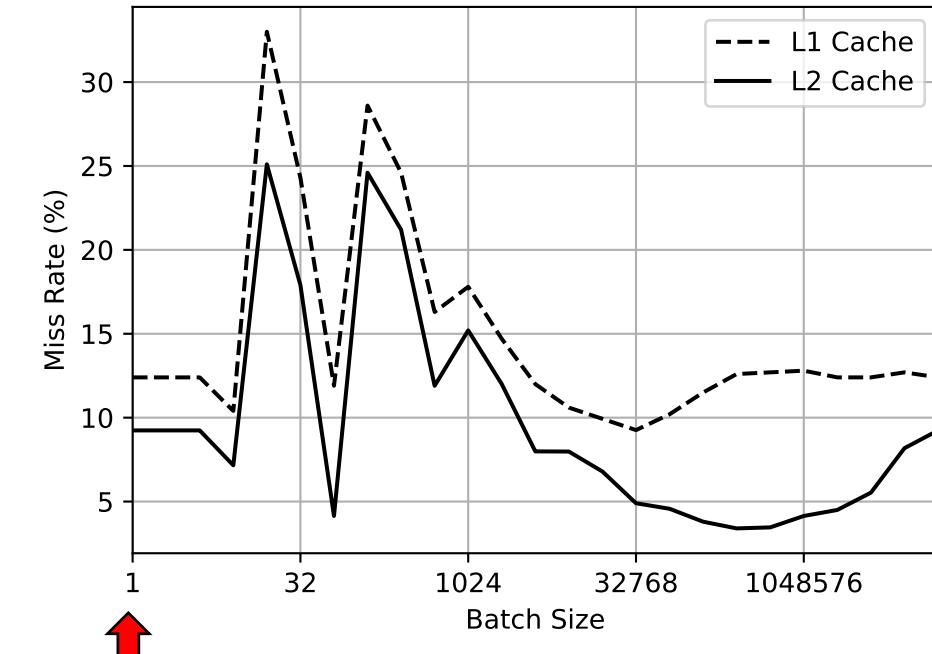
IDEA: Use short index to save memory B/W

Cache-Miss Rate Simulation (for single-level and single-line cache)



POWER9 Specs

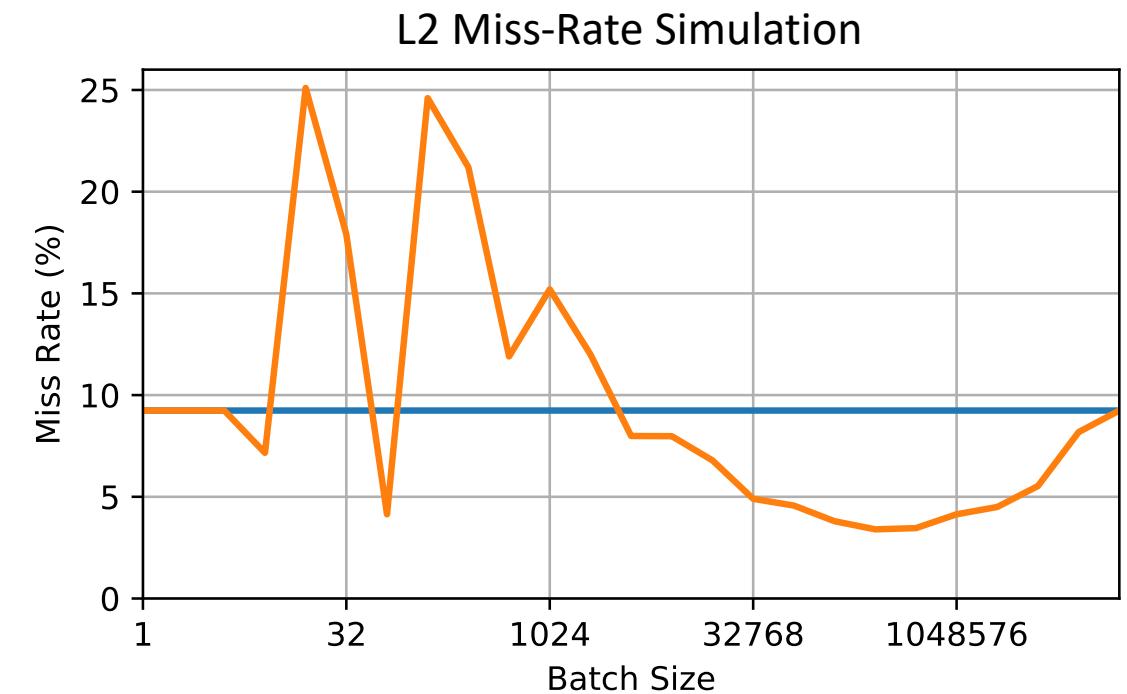
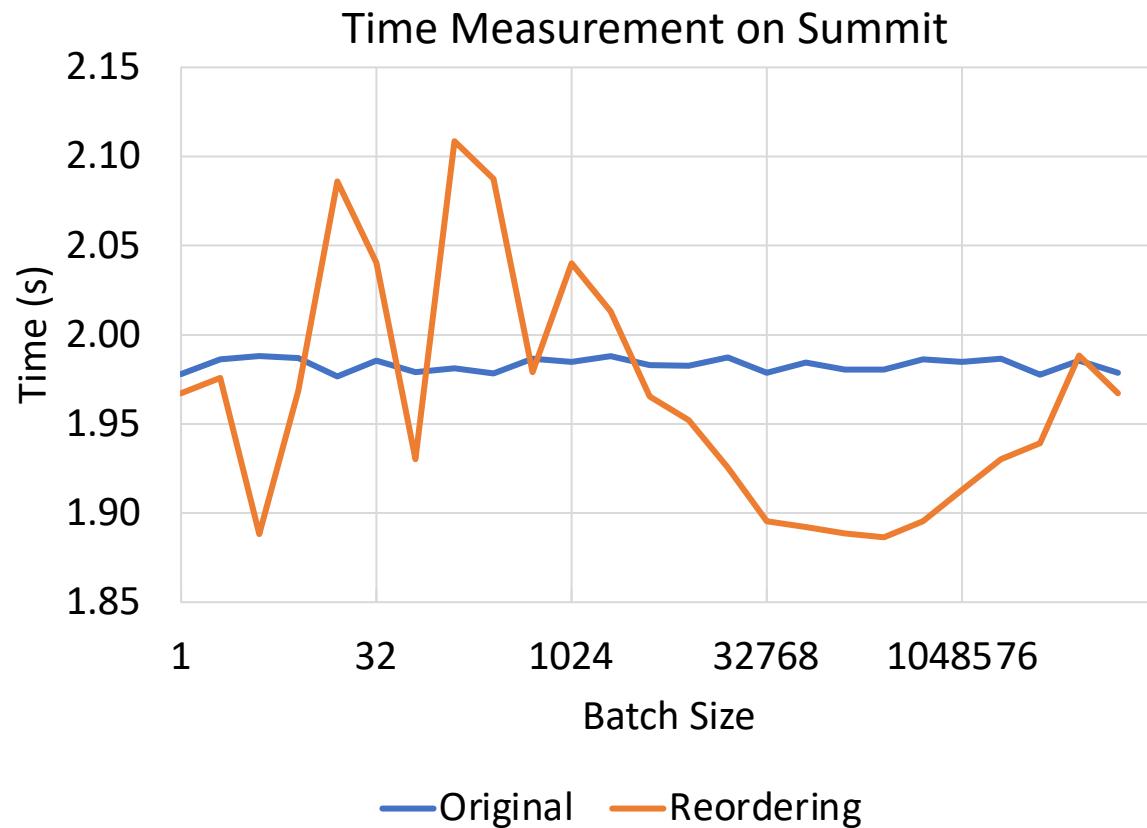
L1: 32 KB per Core
L2: 512 KB per Core
L3: 120 MB per Chip



Batch Size = 1 means no reordering

Cache-Miss Rate Simulation vs. Time Measurements

220x220x220 Grid with 27-Point Stencil

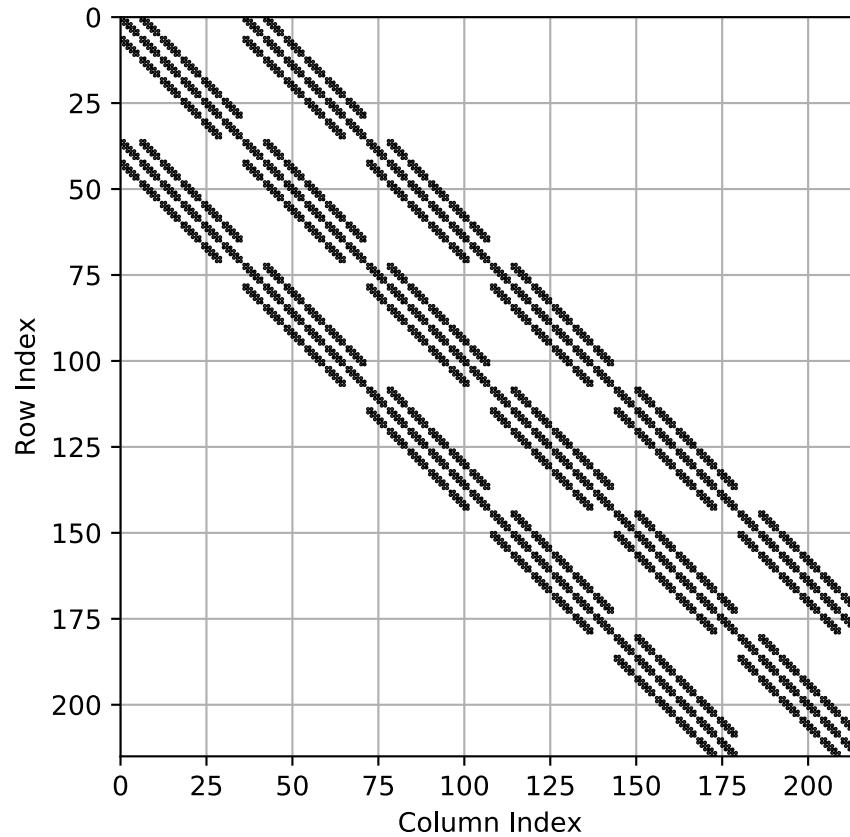


Very Good Match!!!

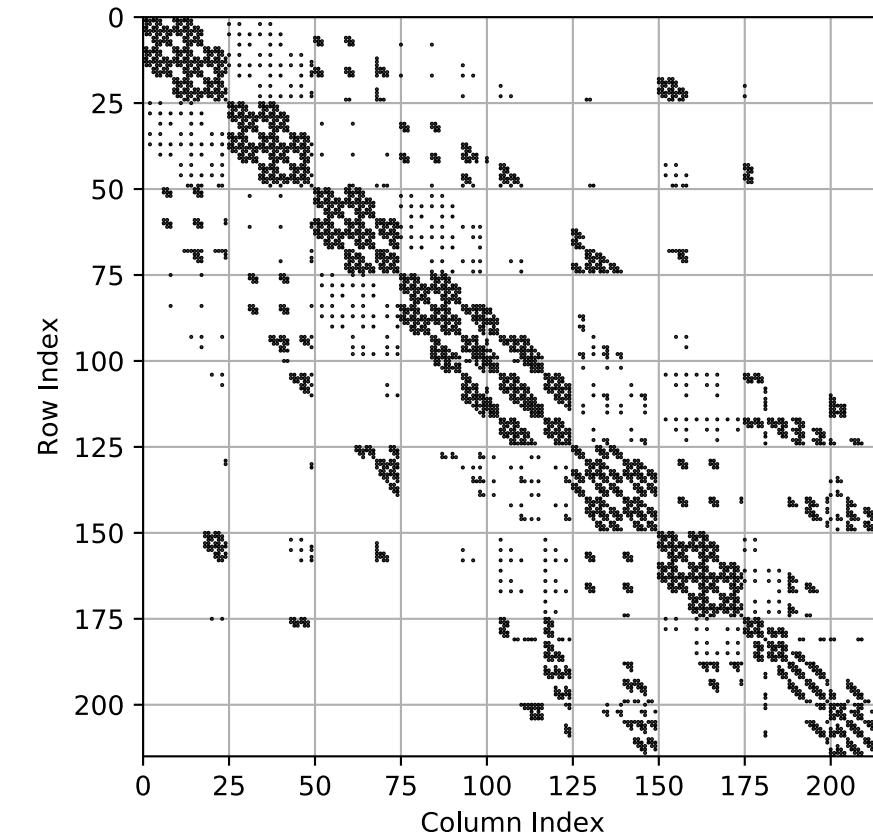
Matrix Sparsity Pattern

6x6x6 27-point stencil
Batch Size: 25

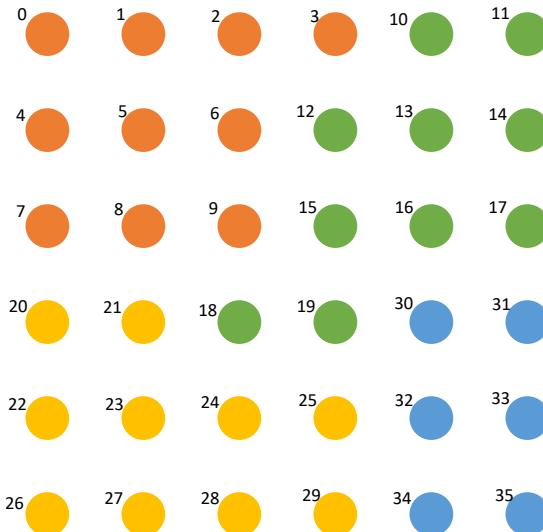
Original



Batch Reordering



Buffering Algorithm

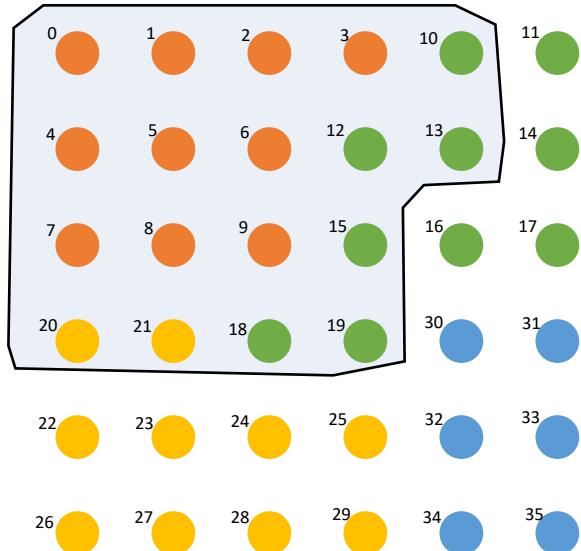


```
//GAUSS-SEIDEL FORWARD SWEEP  
for(int batch = 0; batch < numbatch; batch++){
```

}

Buffering Algorithm

Data Accessed by Batch 0

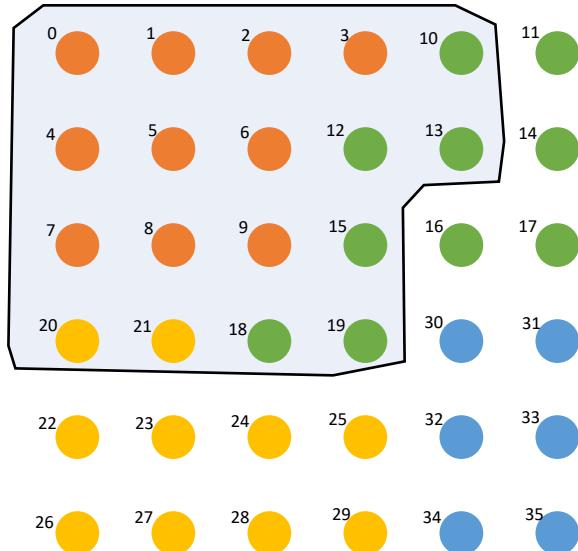


```
//GAUSS-SEIDEL FORWARD SWEEP
for(int batch = 0; batch < numbatch; batch++){
    //READ MAPPING DATA
    int start = mapdispl[batch];
    int mapnz = mapdispl[batch+1]-start;
```

}

Buffering Algorithm

Data Accessed by Batch 0



0
5
10
11
16
17
20
23
24
25
29
30
31
32
33
34
35

Temporary Buffer

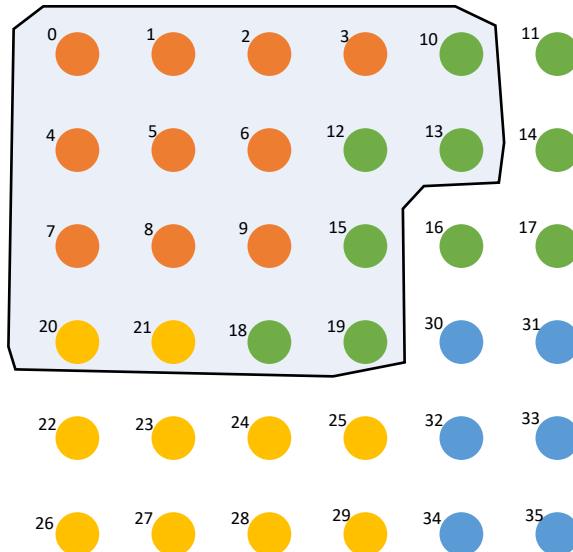
0
5
10
11
16
17
20
23
24
25
29
30
31
32
33
34
35

```
//GAUSS-SEIDEL FORWARD SWEEP
for(int batch = 0; batch < numbatch; batch++){
    //READ MAPPING DATA
    int start = mapdispl[batch];
    int mapnz = mapdispl[batch+1]-start;
    //ALLOCATE BUFFER
    double buffer[batchsize+mapnz];
```

}

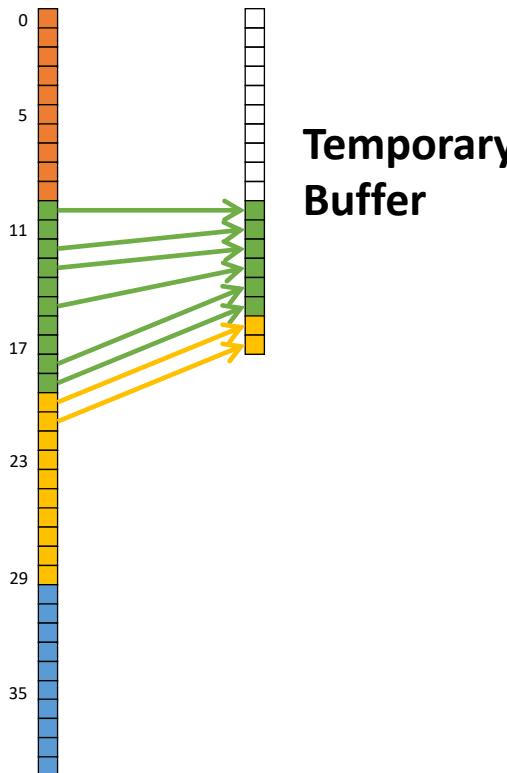
Buffering Algorithm

Data Accessed by Batch 0



Load inter-batch data

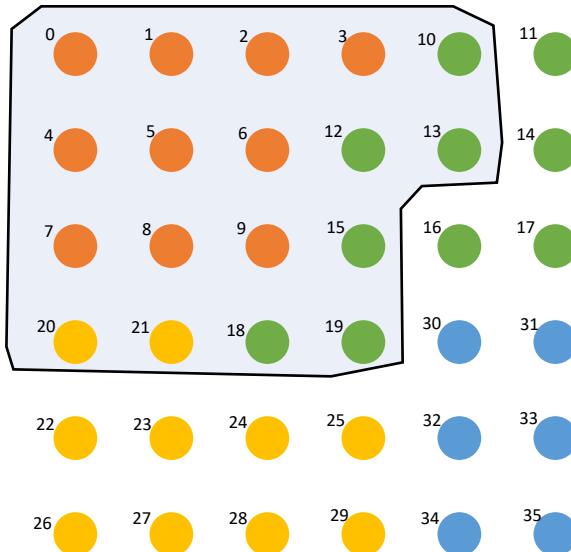
- Irregular Access
- Hardly Vectorizable
- Requires a mapping (additional overhead)



```
//GAUSS-SEIDEL FORWARD SWEEP
for(int batch = 0; batch < numbatch; batch++){
    //READ MAPPING DATA
    int start = mapdispl[batch];
    int mapnz = mapdispl[batch+1]-start;
    //ALLOCATE BUFFER
    double buffer[batchsize+mapnz];
    //LOAD INTER-BATCH DATA
    for(int n = 0; n < mapnz; n++)
        buffer[batchsize+n] = x[map[start+n]];
}
```

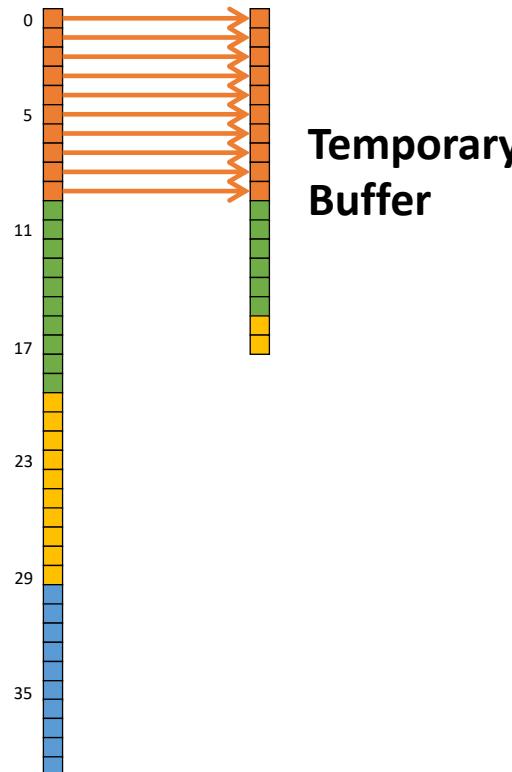
Buffering Algorithm

Data Accessed by Batch 0



Load intra-batch data

- Regular Access
- Vectorizable



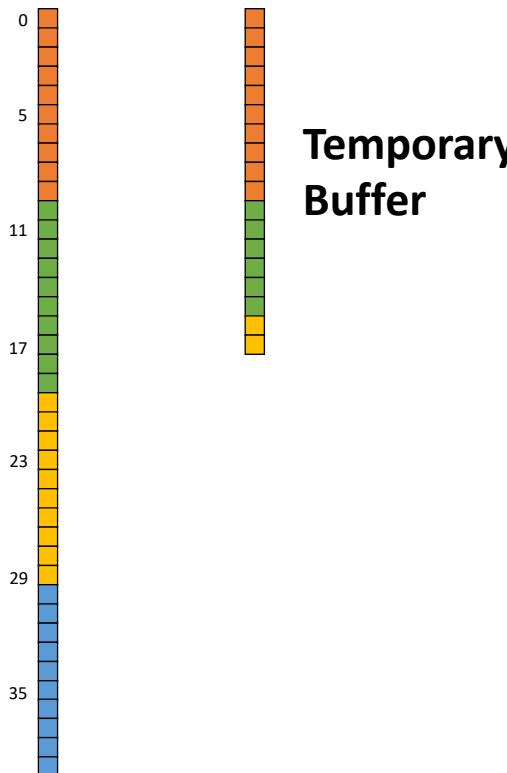
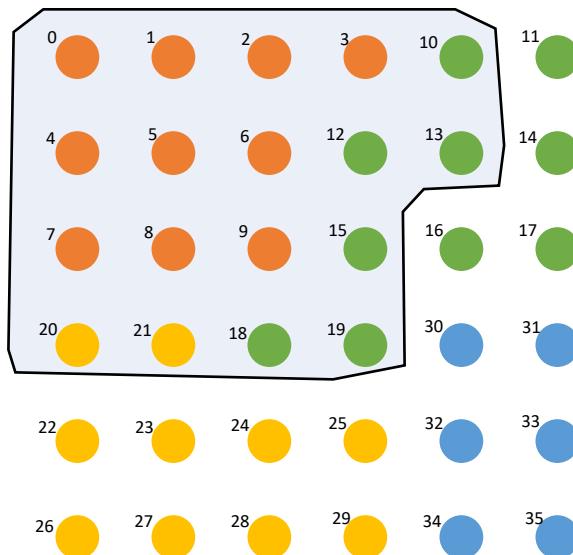
```
//GAUSS-SEIDEL FORWARD SWEEP
for(int batch = 0; batch < numbatch; batch++){
    //READ MAPPING DATA
    int start = mapdispl[batch];
    int mapnz = mapdispl[batch+1]-start;
    //ALLOCATE BUFFER
    double buffer[batchsize+mapnz];
    //LOAD INTER-BATCH DATA
    for(int n = 0; n < mapnz; n++)
        buffer[batchsize+n] = x[map[start+n]];
    //LOAD INTRA-BATCH DATA
    start = batch*batchsize;
    for(int n = 0; n < batchsize; n++)
        buffer[n] = x[start+n];
}
```

Buffering Algorithm

Perform Gauss-Seidel from buffer

- Intra-Batch Data is Updated

Data Accessed by Batch 0

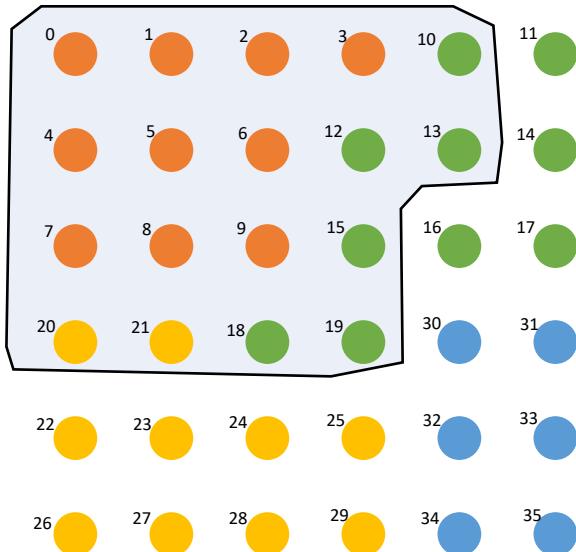


```
//GAUSS-SEIDEL FORWARD SWEEP
for(int batch = 0; batch < numbatch; batch++){
    //READ MAPPING DATA
    int start = mapdispl[batch];
    int mapnz = mapdispl[batch+1]-start;
    //ALLOCATE BUFFER
    double buffer[batchsize+mapnz];
    //LOAD INTER-BATCH DATA
    for(int n = 0; n < mapnz; n++)
        buffer[batchsize+n] = x[map[start+n]];
    //LOAD INTRA-BATCH DATA
    start = batch*batchsize;
    for(int n = 0; n < batchsize; n++)
        buffer[n] = x[start+n];
    //PERFORM GAUSS-SEIDEL PER BATCH
    for(int m = start; m < start+batchsize; m++){
        double reduce = b[m];
        for(int n = displ[m]; n < displ[m+1]; n++)
            reduce -= value[n]*buffer[index[n]];
        buffer[m-start] += reduce/diag[m];
    }
}
```

}

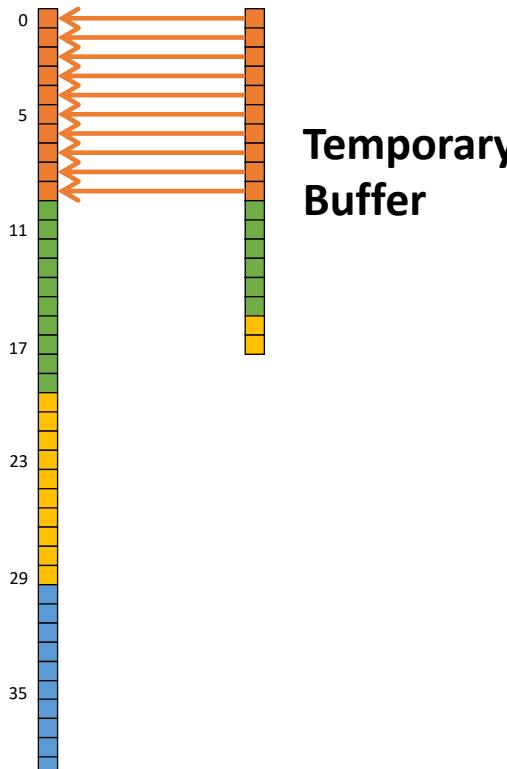
Buffering Algorithm

Data Accessed by Batch 0



Write-back intra-batch data

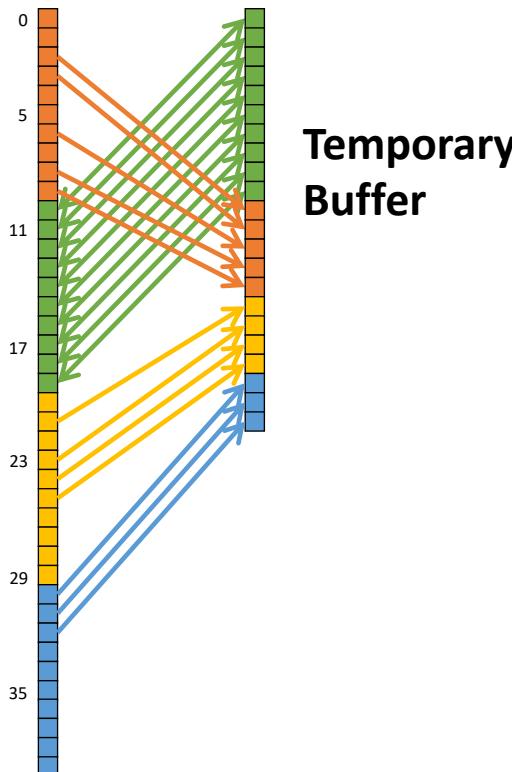
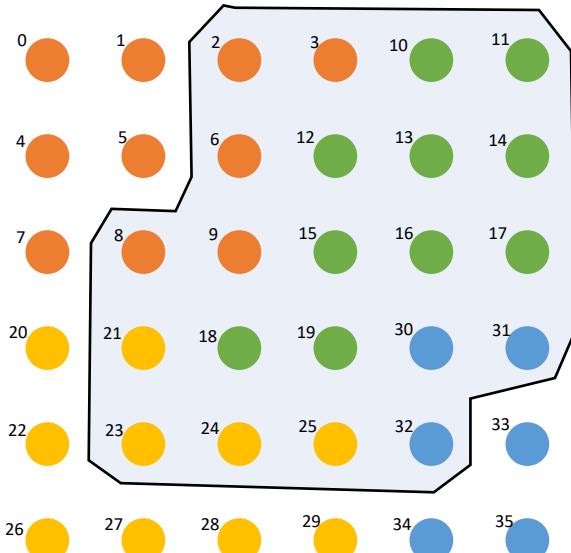
- Regular Access
- Vectorizable



```
//GAUSS-SEIDEL FORWARD SWEEP
for(int batch = 0; batch < numbatch; batch++){
    //READ MAPPING DATA
    int start = mapdispl[batch];
    int mapnz = mapdispl[batch+1]-start;
    //ALLOCATE BUFFER
    double buffer[batchsize+mapnz];
    //LOAD INTER-BATCH DATA
    for(int n = 0; n < mapnz; n++)
        buffer[batchsize+n] = x[map[start+n]];
    //LOAD INTRA-BATCH DATA
    start = batch*batchsize;
    for(int n = 0; n < batchsize; n++)
        buffer[n] = x[start+n];
    //PERFORM GAUSS-SEIDEL PER BATCH
    for(int m = start; m < start+batchsize; m++){
        double reduce = b[m];
        for(int n = displ[m]; n < displ[m+1]; n++)
            reduce -= value[n]*buffer[index[n]];
        buffer[m-start] += reduce/diag[m];
    }
    //WRITE-BACK INTRA-BATCH DATA
    for(int m = 0; m < batchsize; m++)
        x[start+m] = buffer[m];
}
```

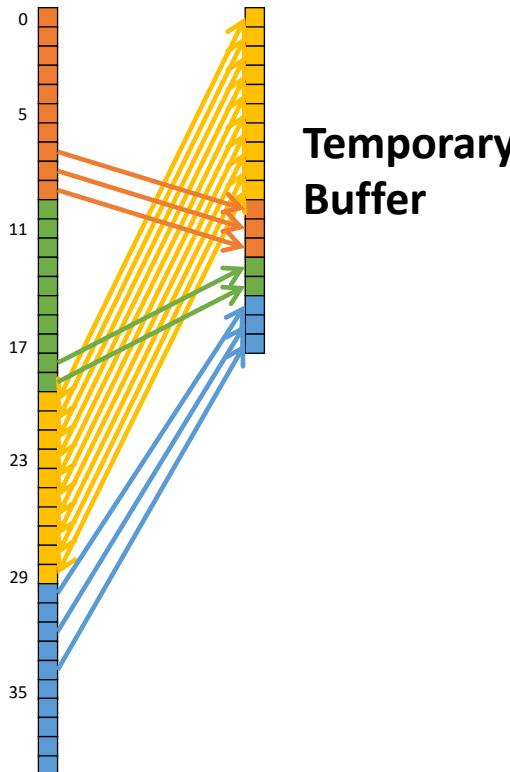
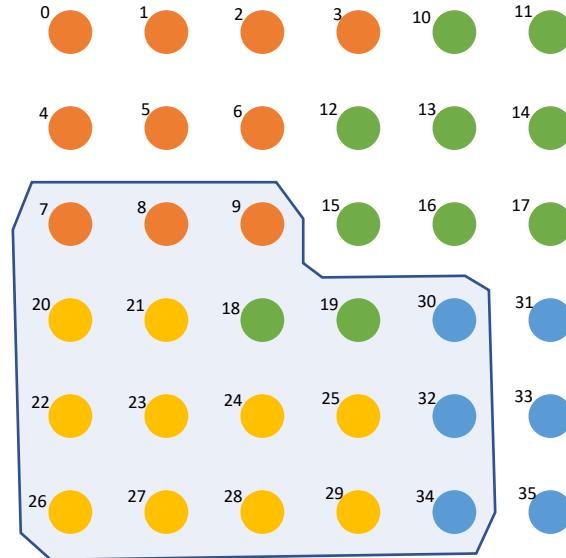
Buffering Algorithm

Data Accessed by Batch 1



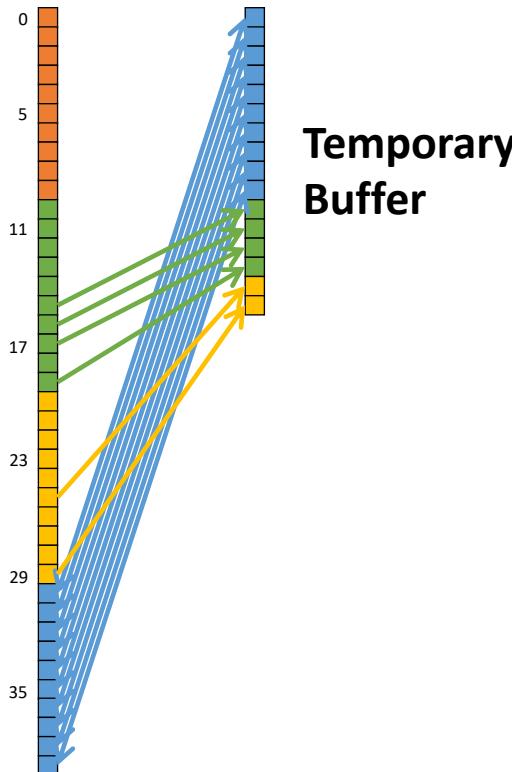
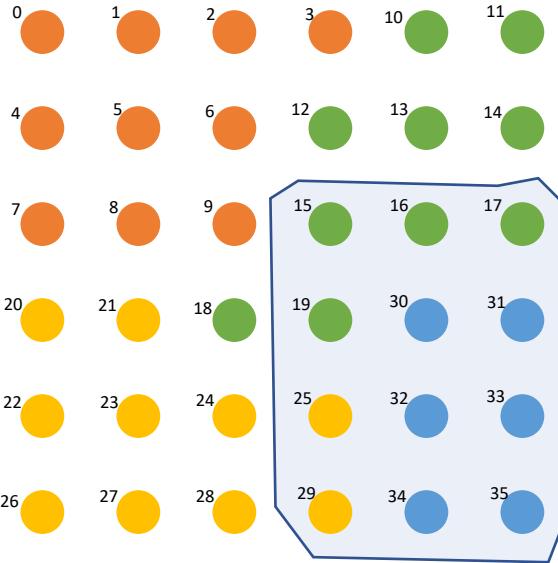
```
//GAUSS-SEIDEL FORWARD SWEEP
for(int batch = 0; batch < numbatch; batch++){
    //READ MAPPING DATA
    int start = mapdispl[batch];
    int mapnz = mapdispl[batch+1]-start;
    //ALLOCATE BUFFER
    double buffer[batchsize+mapnz];
    //LOAD INTER-BATCH DATA
    for(int n = 0; n < mapnz; n++)
        buffer[batchsize+n] = x[map[start+n]];
    //LOAD INTRA-BATCH DATA
    start = batch*batchsize;
    for(int n = 0; n < batchsize; n++)
        buffer[n] = x[start+n];
    //PERFORM GAUSS-SEIDEL PER BATCH
    for(int m = start; m < start+batchsize; m++){
        double reduce = b[m];
        for(int n = displ[m]; n < displ[m+1]; n++)
            reduce -= value[n]*buffer[index[n]];
        buffer[m-start] += reduce/diag[m];
    }
    //WRITE-BACK INTRA-BATCH DATA
    for(int m = 0; m < batchsize; m++)
        x[start+m] = buffer[m];
}
```

Buffering Algorithm



```
//GAUSS-SEIDEL FORWARD SWEEP
for(int batch = 0; batch < numbatch; batch++){
    //READ MAPPING DATA
    int start = mapdispl[batch];
    int mapnz = mapdispl[batch+1]-start;
    //ALLOCATE BUFFER
    double buffer[batchsize+mapnz];
    //LOAD INTER-BATCH DATA
    for(int n = 0; n < mapnz; n++)
        buffer[batchsize+n] = x[map[start+n]];
    //LOAD INTRA-BATCH DATA
    start = batch*batchsize;
    for(int n = 0; n < batchsize; n++)
        buffer[n] = x[start+n];
    //PERFORM GAUSS-SEIDEL PER BATCH
    for(int m = start; m < start+batchsize; m++){
        double reduce = b[m];
        for(int n = displ[m]; n < displ[m+1]; n++)
            reduce -= value[n]*buffer[index[n]];
        buffer[m-start] += reduce/diag[m];
    }
    //WRITE-BACK INTRA-BATCH DATA
    for(int m = 0; m < batchsize; m++)
        x[start+m] = buffer[m];
}
```

Buffering Algorithm



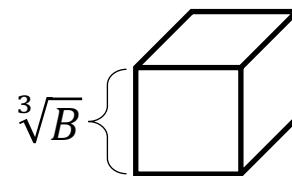
```
//GAUSS-SEIDEL FORWARD SWEEP
for(int batch = 0; batch < numbatch; batch++){
    //READ MAPPING DATA
    int start = mapdispl[batch];
    int mapnz = mapdispl[batch+1]-start;
    //ALLOCATE BUFFER
    double buffer[batchsize+mapnz];
    //LOAD INTER-BATCH DATA
    for(int n = 0; n < mapnz; n++)
        buffer[batchsize+n] = x[map[start+n]];
    //LOAD INTRA-BATCH DATA
    start = batch*batchsize;
    for(int n = 0; n < batchsize; n++)
        buffer[n] = x[start+n];
    //PERFORM GAUSS-SEIDEL PER BATCH
    for(int m = start; m < start+batchsize; m++){
        double reduce = b[m];
        for(int n = displ[m]; n < displ[m+1]; n++)
            reduce -= value[n]*buffer[index[n]];
        buffer[m-start] += reduce/diag[m];
    }
    //WRITE-BACK INTRA-BATCH DATA
    for(int m = 0; m < batchsize; m++)
        x[start+m] = buffer[m];
}
```

```

//GAUSS-SEIDEL FORWARD SWEEP
for(int batch = 0; batch < numbatch; batch++){
    //READ MAPPING DATA
    int start = mapdispl[batch];
    int mapnz = mapdispl[batch+1]-start;
    //ALLOCATE BUFFER
    double buffer[batchsize+mapnz];
    //LOAD INTER-BATCH DATA
    for(int n = 0; n < mapnz; n++)
        buffer[batchsize+n] = x[map[start+n]];
    //LOAD INTRA-BATCH DATA
    start = batch*batchsize;
    for(int n = 0; n < batchsize; n++)
        buffer[n] = x[start+n];
    //PERFORM GAUSS-SEIDEL PER BATCH
    for(int m = start; m < start+batchsize; m++){
        double reduce = b[m];
        for(int n = displ[m]; n < displ[m+1]; n++)
            reduce -= value[n]*buffer[index[n]];
        buffer[m-start] += reduce/diag[m];
    }
    //WRITE-BACK INTRA-BATCH DATA
    for(int m = 0; m < batchsize; m++)
        x[start+m] = buffer[m];
}

```

Overhead Analysis



B : Batch Size

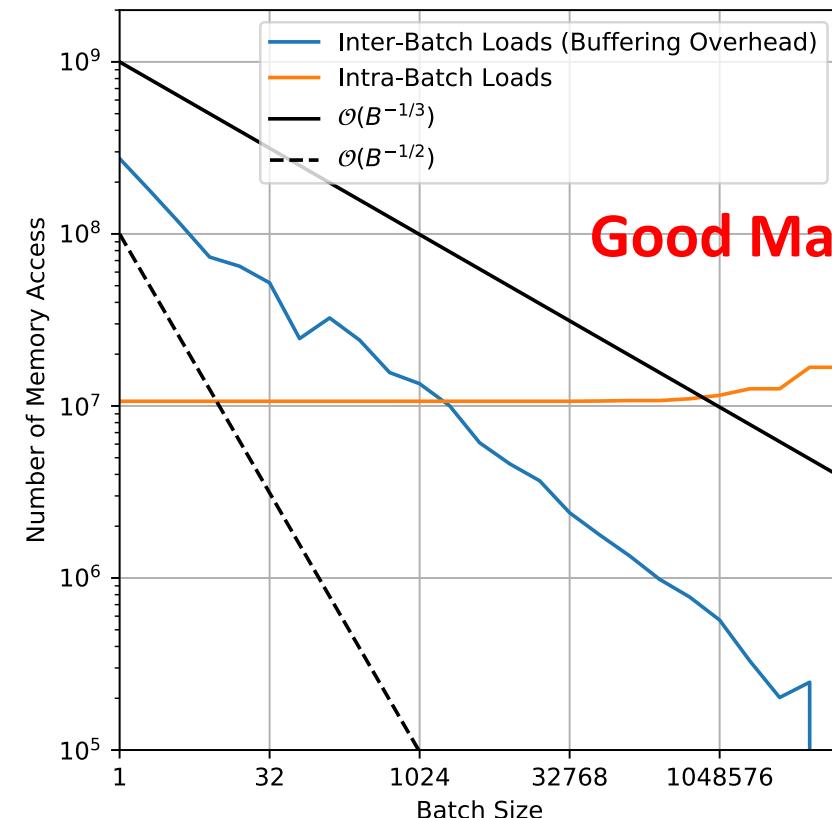
Inter-Batch Loads per Batch: $\mathcal{O}(B^{2/3})$

Number of Batches: N/B

Total Inter-Batch Loads: $\mathcal{O}(NB^{-1/3})$

These are Irregular Access!

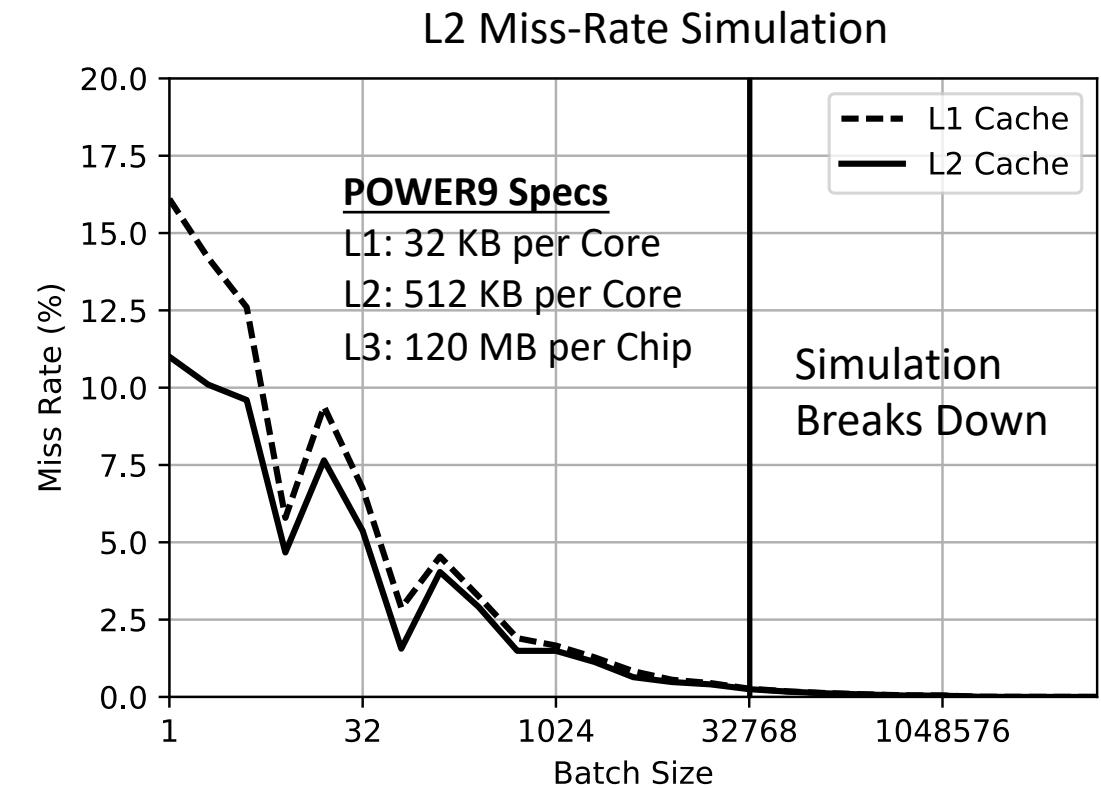
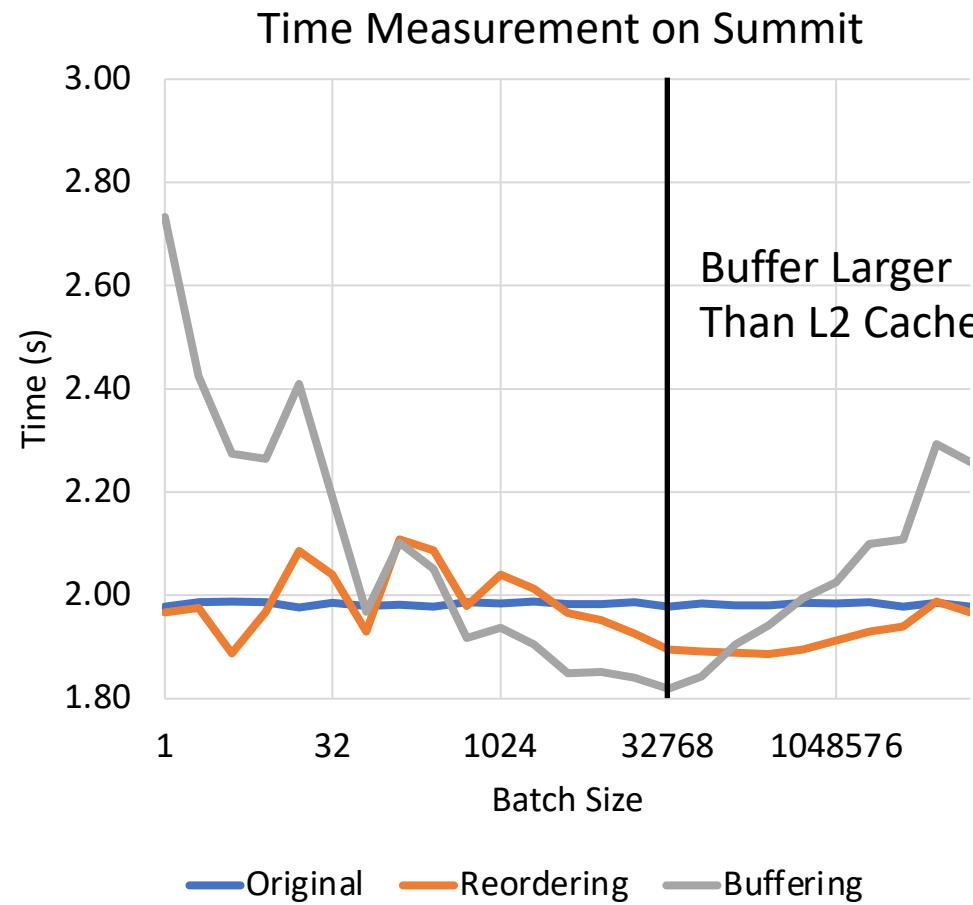
Measured # of Buffer Loads



Good Match!!!

Cache-Miss Rate Simulation (for single-level and single-line cache)

220x220x220 Grid with 27-Point Stencil



Using Small Buffer

```
//GAUSS-SEIDEL FORWARD SWEEP
for(int batch = 0; batch < numbatch; batch++){
    //READ MAPPING DATA
    int start = mapdispl[batch];
    int mapnz = mapdispl[batch+1]-start;
    //ALLOCATE BUFFER
    double buffer[batchsize+mapnz];
    //LOAD INTER-BATCH DATA
    for(int n = 0; n < mapnz; n++)
        buffer[batchsize+n] = x[map[start+n]];
    //LOAD INTRA-BATCH DATA
    start = batch*batchsize;
    for(int n = 0; n < batchsize; n++)
        buffer[n] = x[start+n];
    //PERFORM GAUSS-SEIDEL PER BATCH
    for(int m = start; m < start+batchsize; m++){
        double reduce = b[m];
        for(int n = displ[m]; n < displ[m+1]; n++)
            reduce -= value[n]*buffer[sndex[n]];
        buffer[m-start] += reduce/diag[m];
    }
    //WRITE-BACK INTRA-BATCH DATA
    for(int m = 0; m < batchsize; m++)
        x[start+m] = buffer[m];
}
```

Use unsigned short
instead of int

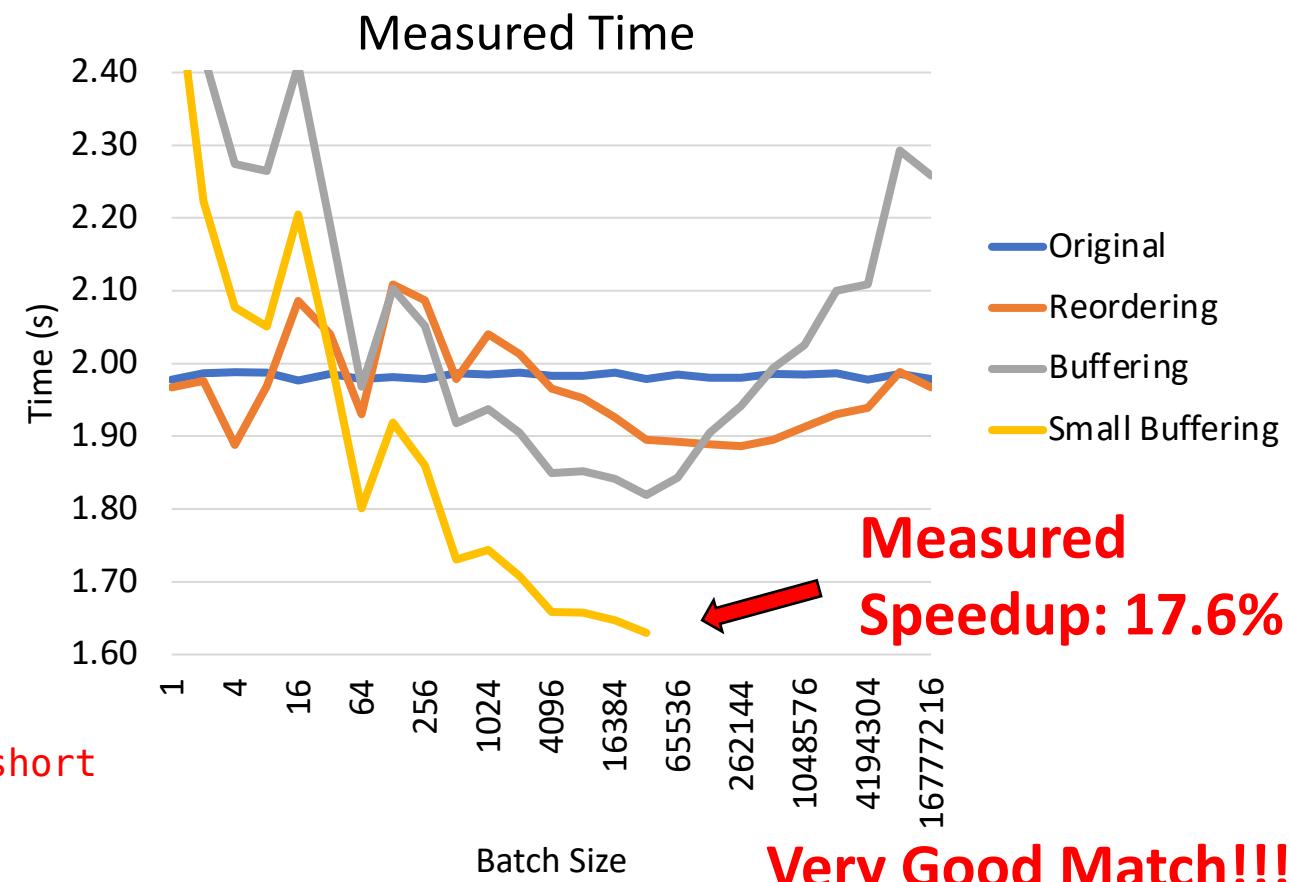
Memory Transactions

Auxiliary Data: $3N \times 8$ Bytes Read

Matrix Data: $27N \times (8 + 2)$ Bytes Read → 17% B/W Saving

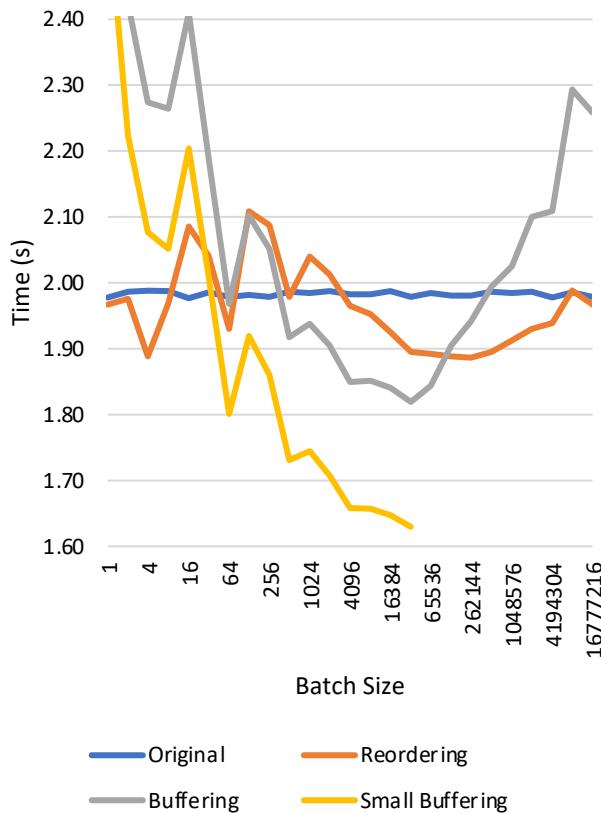
Vector Data: $N \times 8$ Bytes Read → Irregular Access Eliminated
 $N \times 8$ Bytes Write

Theoretical Speedup: 17%



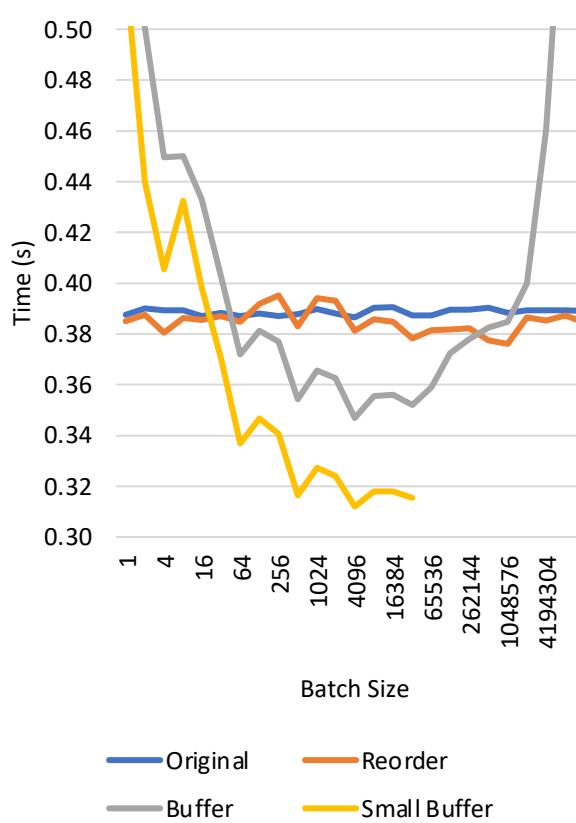
Preliminary Results on Summit Node

Grid Size: 220x220x220



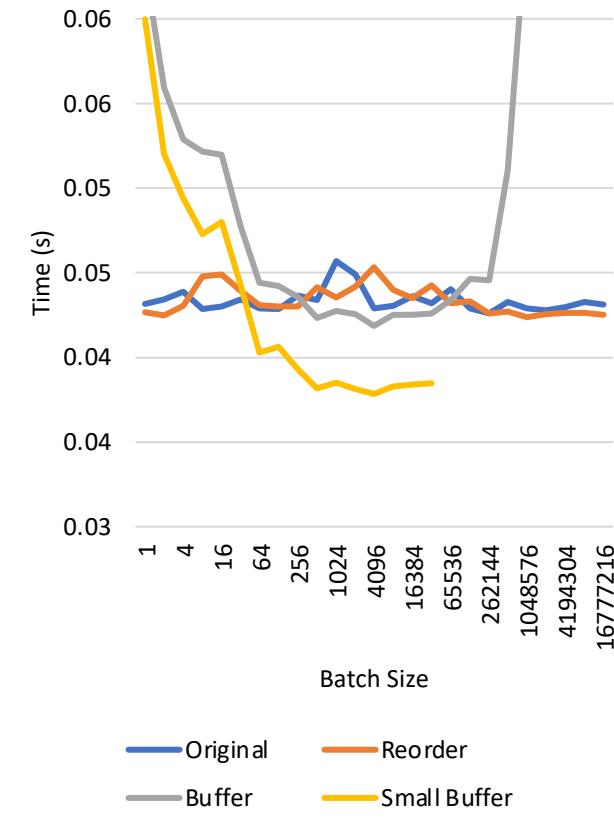
Speedup: 17.6%

Grid Size: 128x128x128



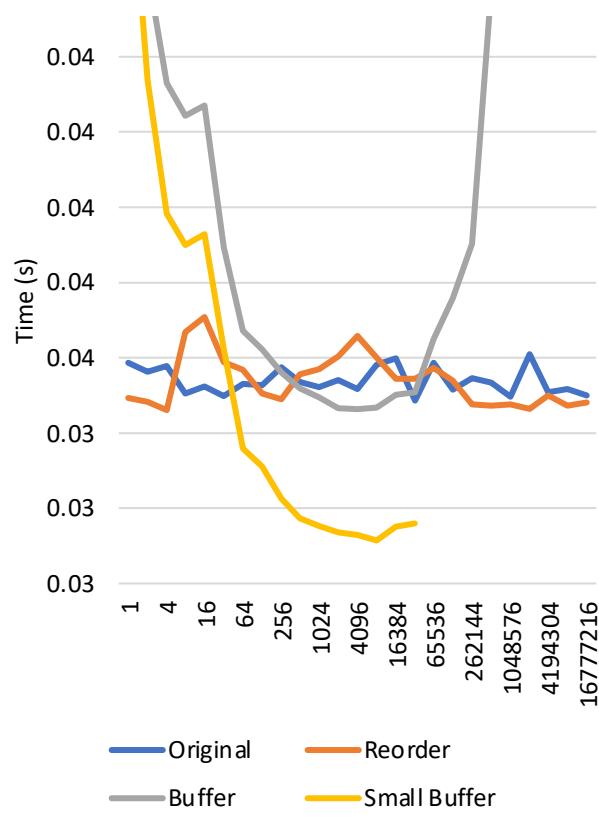
Speedup: 19.3%

Grid Size: 64x64x64



Speedup: 15.7%

Grid Size: 60x60x60



Speedup: 13.0%

Verification: Solving Poisson Equation

$$\nabla^2 \phi = f \quad \frac{\partial^2 \phi}{\partial x^2} + \frac{\partial^2 \phi}{\partial y^2} + \frac{\partial^2 \phi}{\partial z^2} = f$$

Taylor Series Expansion at ϕ_i

$$\phi_{i-1} = \phi_i - \frac{\partial \phi_i}{\partial x} h + \frac{\partial^2 \phi_i}{\partial x^2} \frac{h^2}{2!} - \frac{\partial^3 \phi_i}{\partial x^3} \frac{h^3}{3!} + O(h^4)$$

$$\phi_{i+1} = \phi_i + \frac{\partial \phi_i}{\partial x} h + \frac{\partial^2 \phi_i}{\partial x^2} \frac{h^2}{2!} + \frac{\partial^3 \phi_i}{\partial x^3} \frac{h^3}{3!} + O(h^4)$$

$$\phi_{i-1} + \phi_{i+1} = 2\phi_i + \frac{\partial^2 \phi_i}{\partial x^2} h^2 + O(h^4)$$

$$\frac{\partial^2 \phi}{\partial x^2} = \frac{\phi_{i-1} - 2\phi_i + \phi_{i+1}}{h^2} + O(h^2)$$

$$\frac{\partial^2 \phi}{\partial y^2} = \frac{\phi_{j-1} - 2\phi_j + \phi_{j+1}}{h^2} + O(h^2)$$

$$\frac{\partial^2 \phi}{\partial z^2} = \frac{\phi_{k-1} - 2\phi_k + \phi_{k+1}}{h^2} + O(h^2)$$

3-point Equation:

$$\delta_x^2 \phi = \frac{\phi_{i-1} - 2\phi_i + \phi_{i+1}}{h^2} = f_i + O(h^2)$$

5-point Equation:

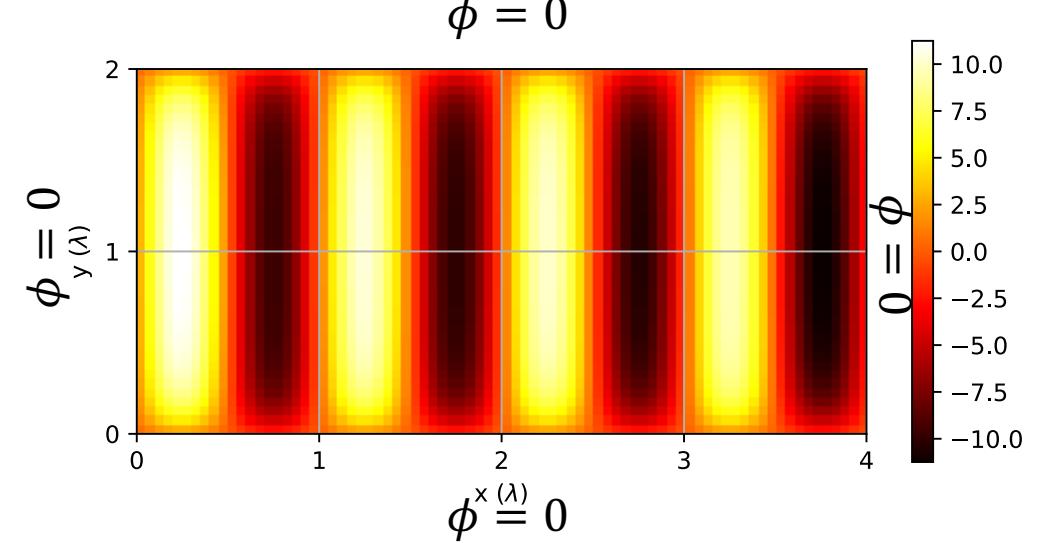
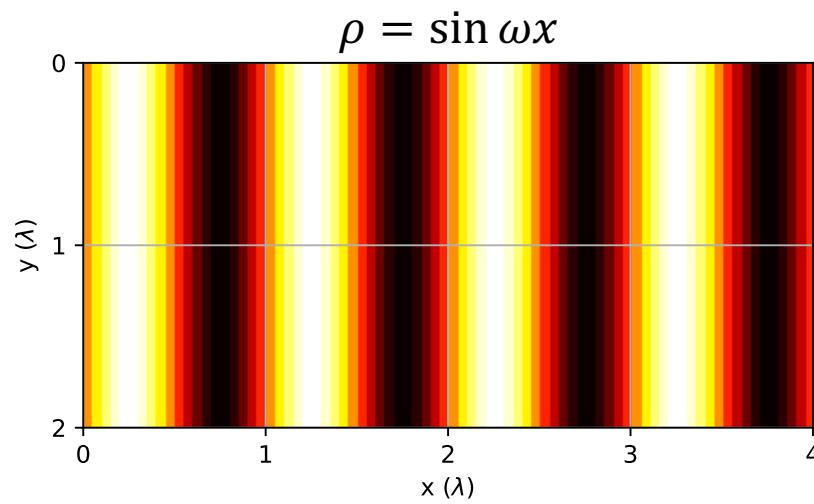
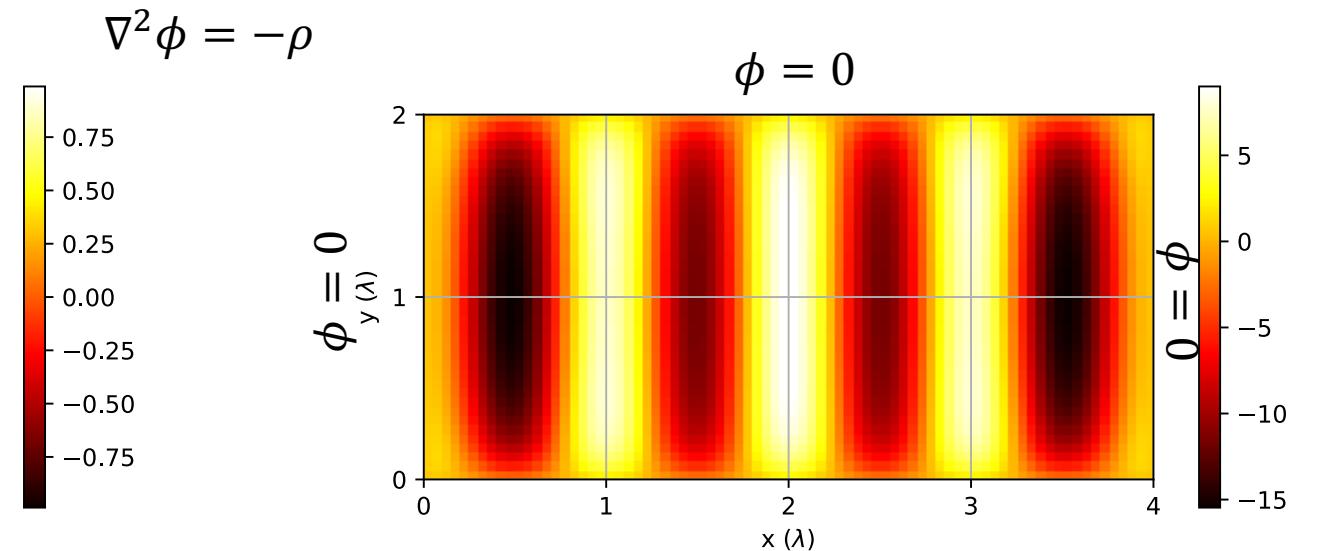
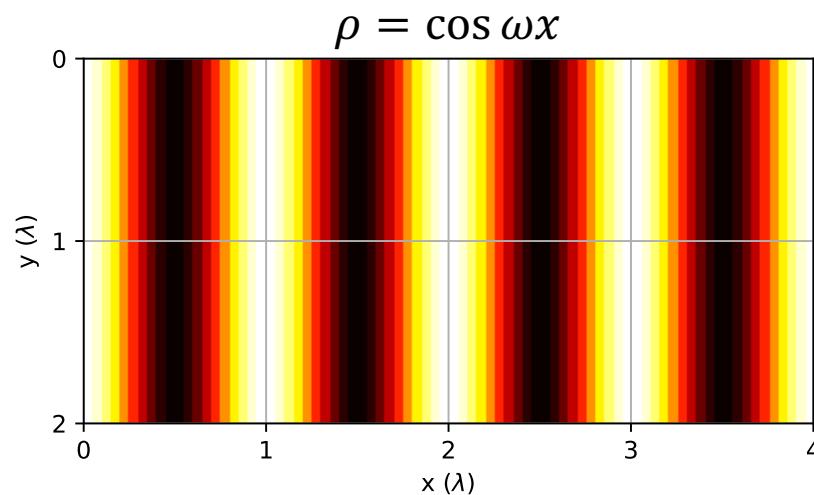
$$[\delta_x^2 + \delta_y^2] \phi = \frac{\phi_{i-1} + \phi_{j-1} - 4\phi_{i,j} + \phi_{i+1} + \phi_{j+1}}{h^2} = f_{i,j} + O(h^2)$$

7-point Equation:

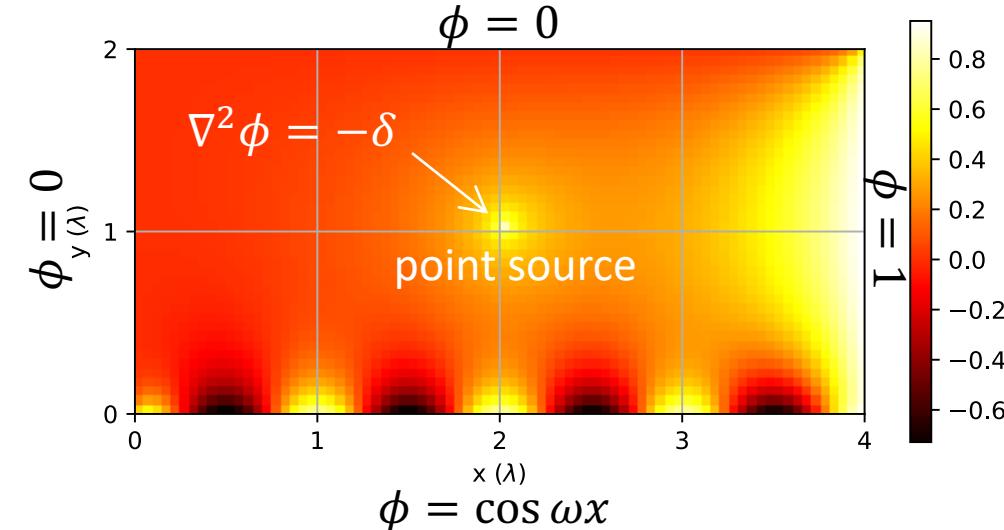
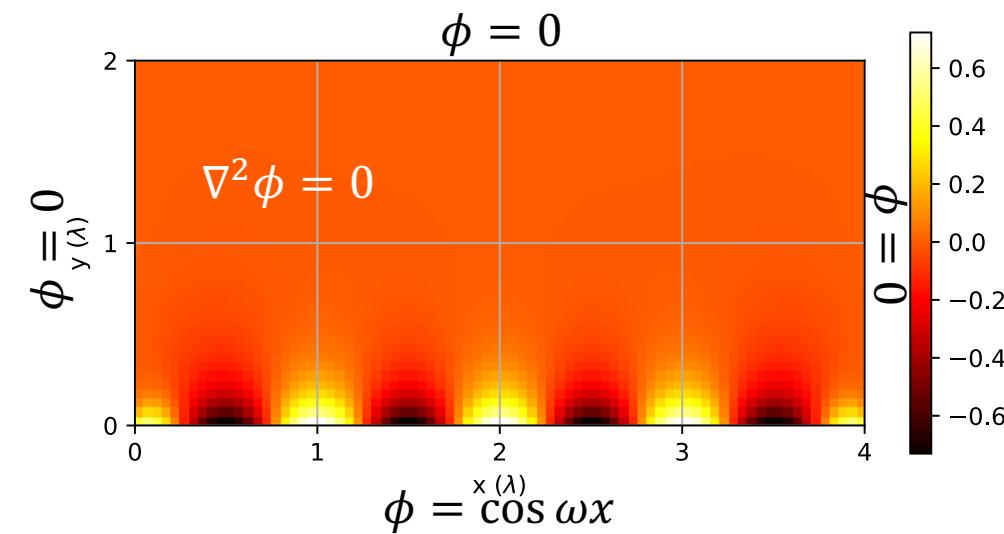
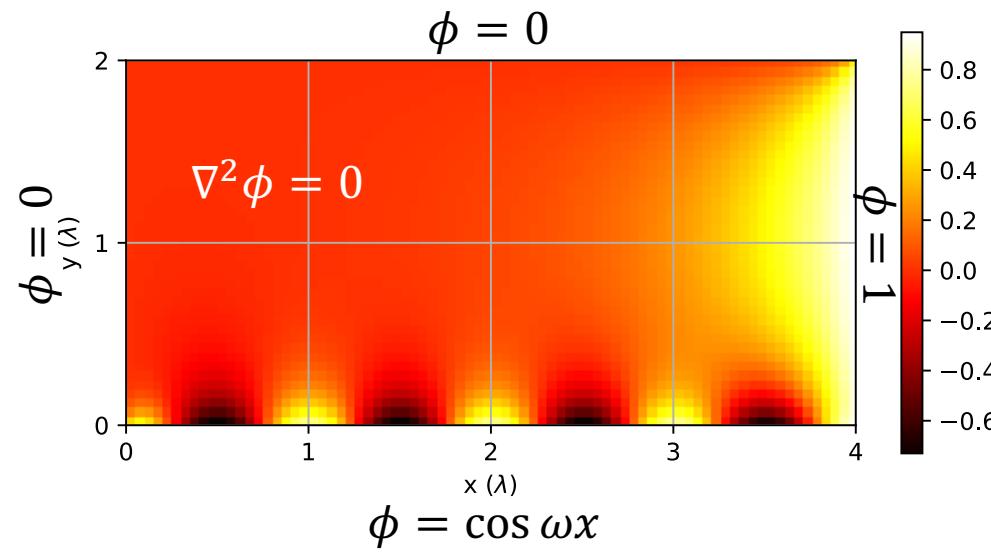
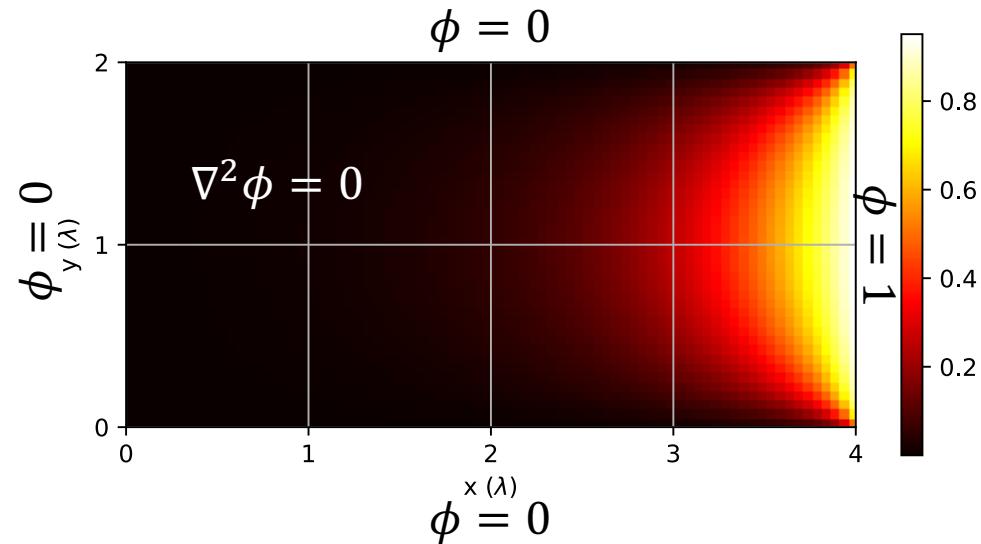
$$[\delta_x^2 + \delta_y^2 + \delta_z^2] \phi = \frac{\phi_{i-1} + \phi_{j-1} + \phi_{k-1} - 6\phi_{i,j,k} + \phi_{i+1} + \phi_{j+1} + \phi_{k+1}}{h^2} = f_{i,j,k} + O(h^2)$$



Verification: Solving Poisson Equation with Dirichlet Boundary Conditions



Verification: Solving Poisson Problem with Dirichlet Boundary Conditions



Verification $\nabla\phi = f$

Function 1:

$$\phi = xyz \quad f = 0$$

7 Points: $\nabla_7^2\phi = [\delta_x^2 + \delta_y^2 + \delta_z^2]\phi = f + \mathcal{O}(h^2)$

19 Points: $\nabla_{19}^2\phi = \left[\nabla_7^2 + \frac{h^2}{6}(\delta_x^2\delta_y^2 + \delta_y^2\delta_z^2 + \delta_x^2\delta_z^2)\right]\phi = f + \frac{h^2}{12}\nabla_7^2f + \mathcal{O}(h^4)$

Function 2:

$$\phi = x^2y^2z^2 \quad f = 2y^2z^2 + 2x^2z^2 + 2x^2y^2$$

27 Points: $\nabla_{27}^2\phi = \left[\nabla_{19}^2 + \frac{h^4}{30}(\delta_x^2 + \delta_y^2 + \delta_z^2)\right]\phi = f + \frac{h^2}{12}\nabla^2f + \frac{h^4}{360}\nabla^4f + \frac{h^4}{180}\left(\frac{\partial^4f}{\partial x^2\partial y^2} + \frac{\partial^4f}{\partial y^2\partial z^2} + \frac{\partial^4f}{\partial x^2\partial z^2}\right) + \mathcal{O}(h^6)$

Function 3:

$$\phi = x^3y^3z^3 \quad f = 6xy^3z^3 + 6x^3yz^3 + 6x^3y^3z$$

Function 4:

$$\phi = \sin(x + y + z)\cos(x + y + z) \quad f = -6\sin(2(x + y + z))$$

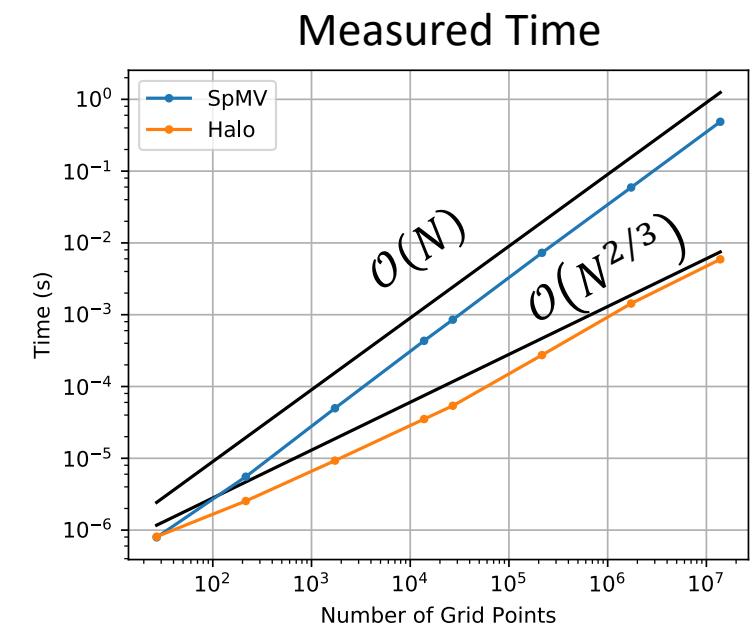
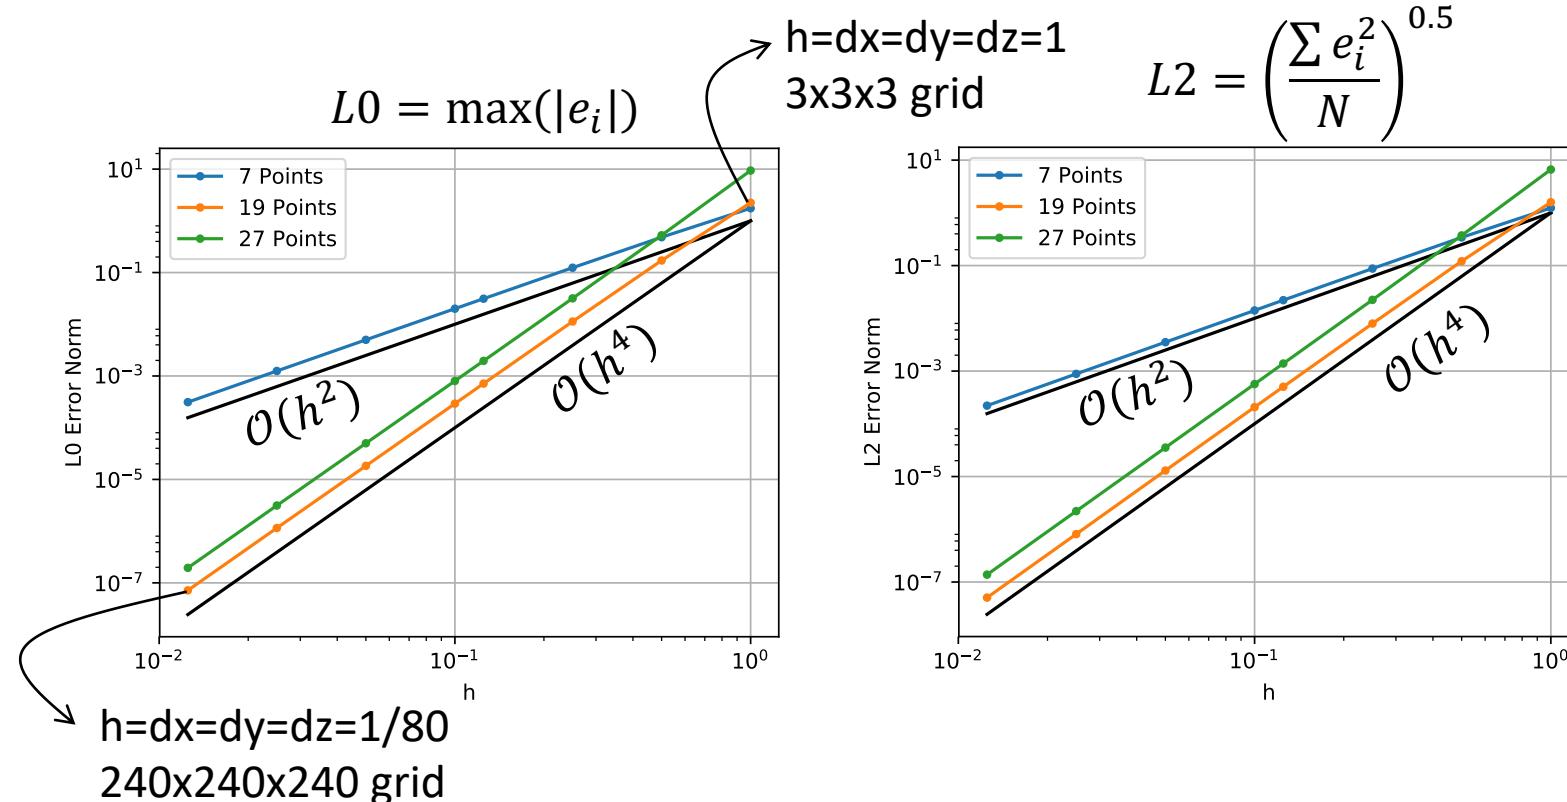
$$\nabla^2f = 72\sin(2(x + y + z))$$

$$\nabla^4f = -864\sin(2(x + y + z))$$

$$\frac{\partial^4f}{\partial x^4} = \frac{\partial^4f}{\partial y^4} = \frac{\partial^4f}{\partial y^4} = \frac{\partial^4f}{\partial x^2y^2} = \frac{\partial^4f}{\partial x^2z^2} = \frac{\partial^4f}{\partial y^2z^2} = -96\sin(2(x + y + z))$$

Verification: Forward Problem

$$\phi = \sin(x + y + z) \cos(x + y + z) \quad \Rightarrow \quad \nabla^2 \phi = -6 \sin(2(x + y + z))$$

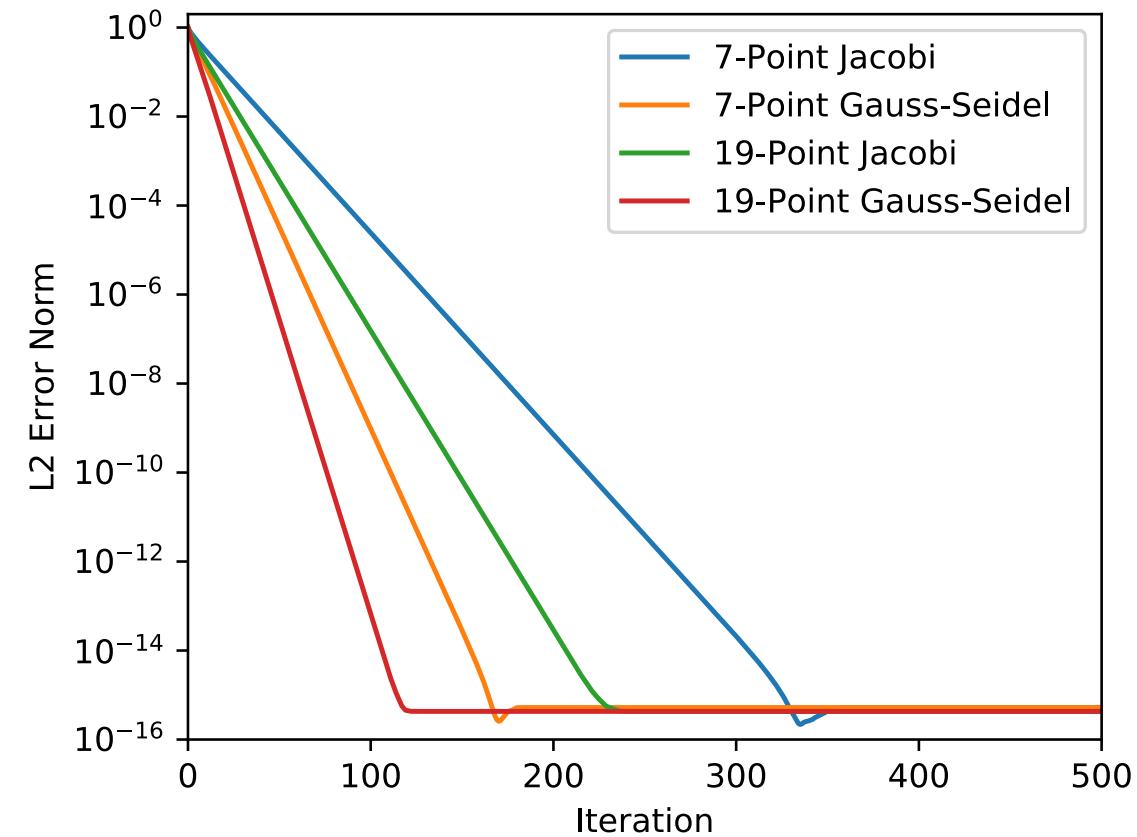
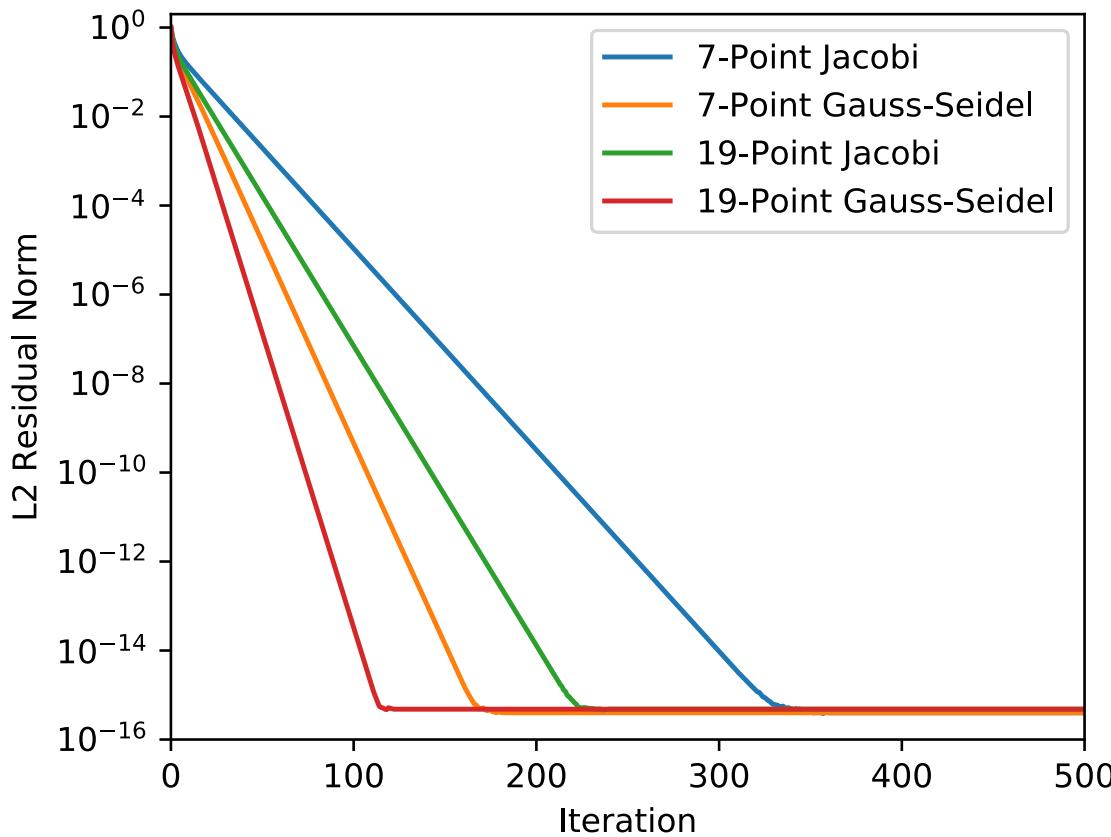


See 7-point, 19-point, and 27-point derivations at: W. F. Spotz and G. F. Carey, "A high-order compact formulation for the 3D Poisson equation," Numerical Methods for Partial Differential Equations, vol. 12, pp. 235-243, 1996.

Verification: Inverse Problem

$$\nabla^2 \phi = -6 \sin(2(x + y + z)) \rightarrow \phi = \sin(x + y + z) \cos(x + y + z)$$

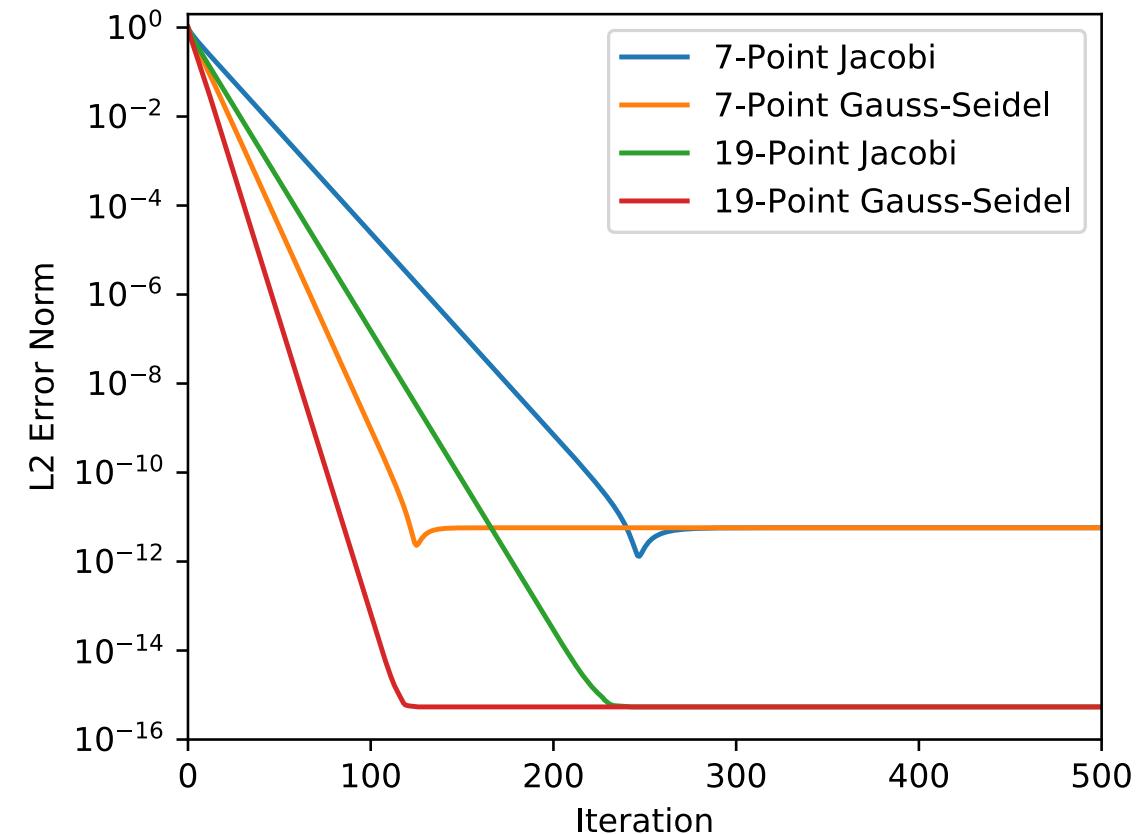
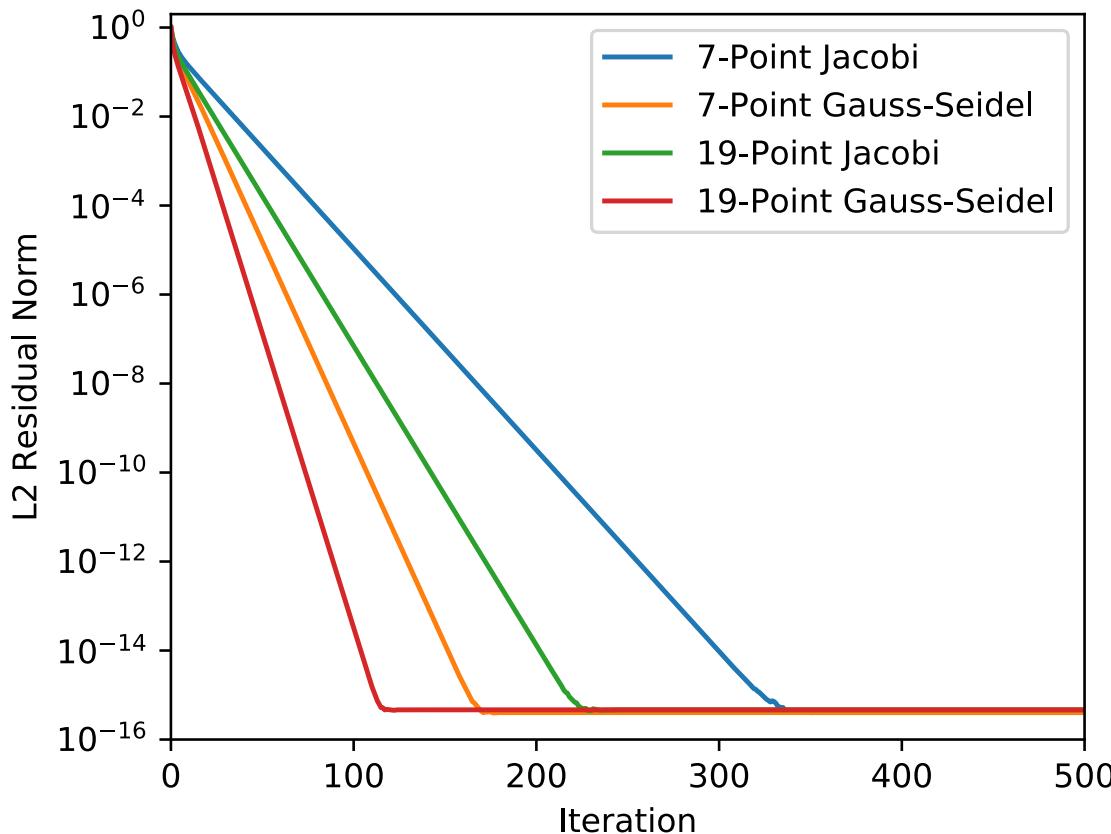
6x6x6
h=0.0001



Verification: Inverse Problem

$$\nabla^2 \phi = -6 \sin(2(x + y + z)) \rightarrow \phi = \sin(x + y + z) \cos(x + y + z)$$

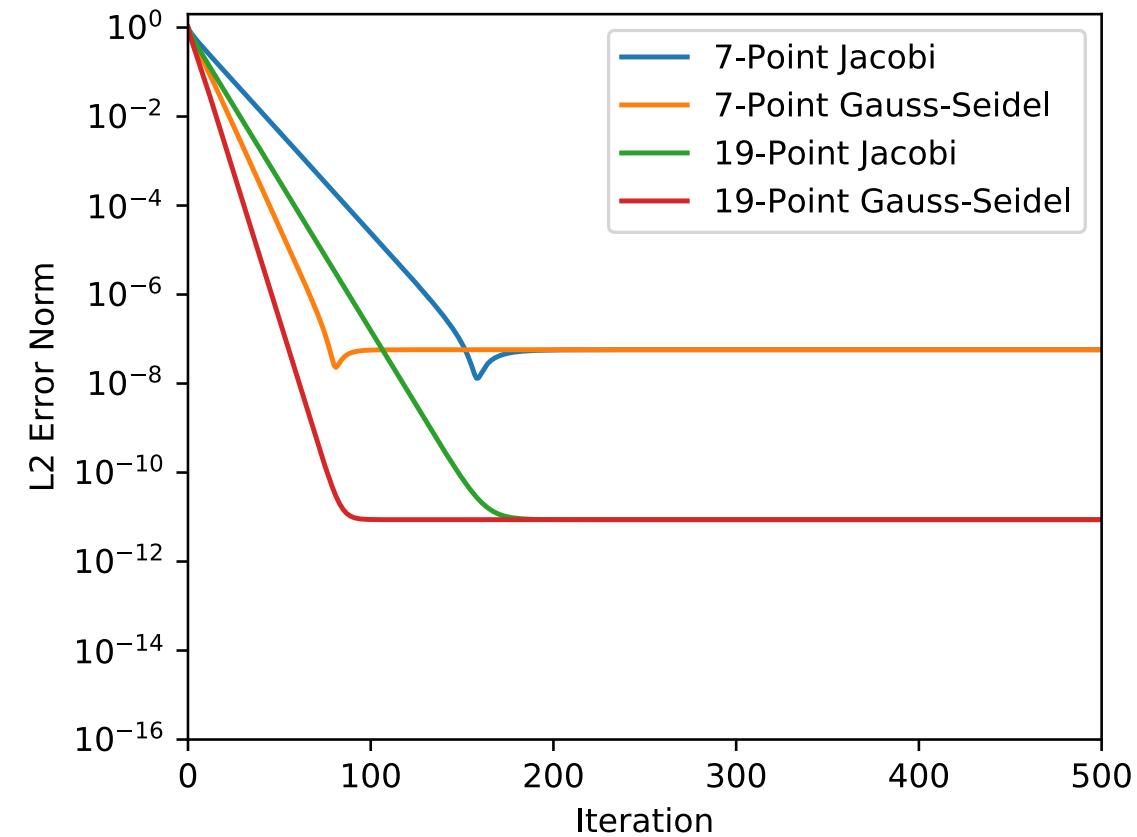
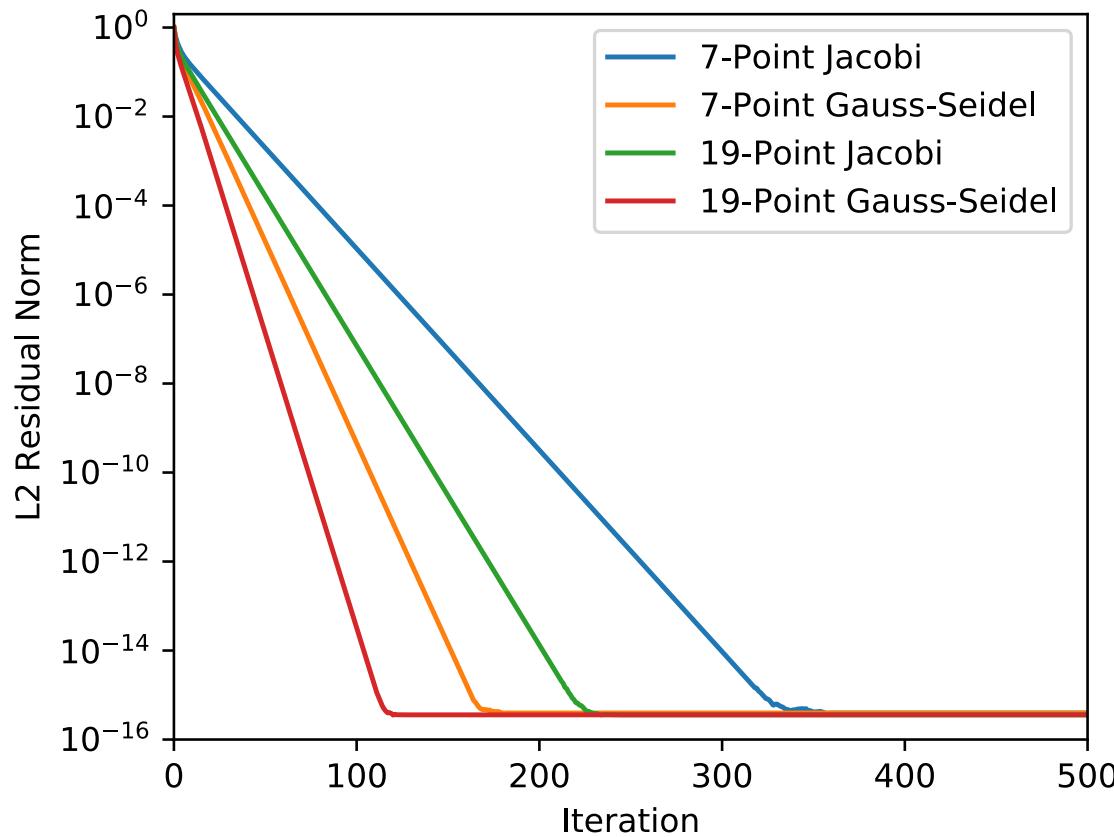
6x6x6
h=0.001



Verification: Inverse Problem

$$\nabla^2 \phi = -6 \sin(2(x + y + z)) \quad \rightarrow \quad \phi = \sin(x + y + z) \cos(x + y + z)$$

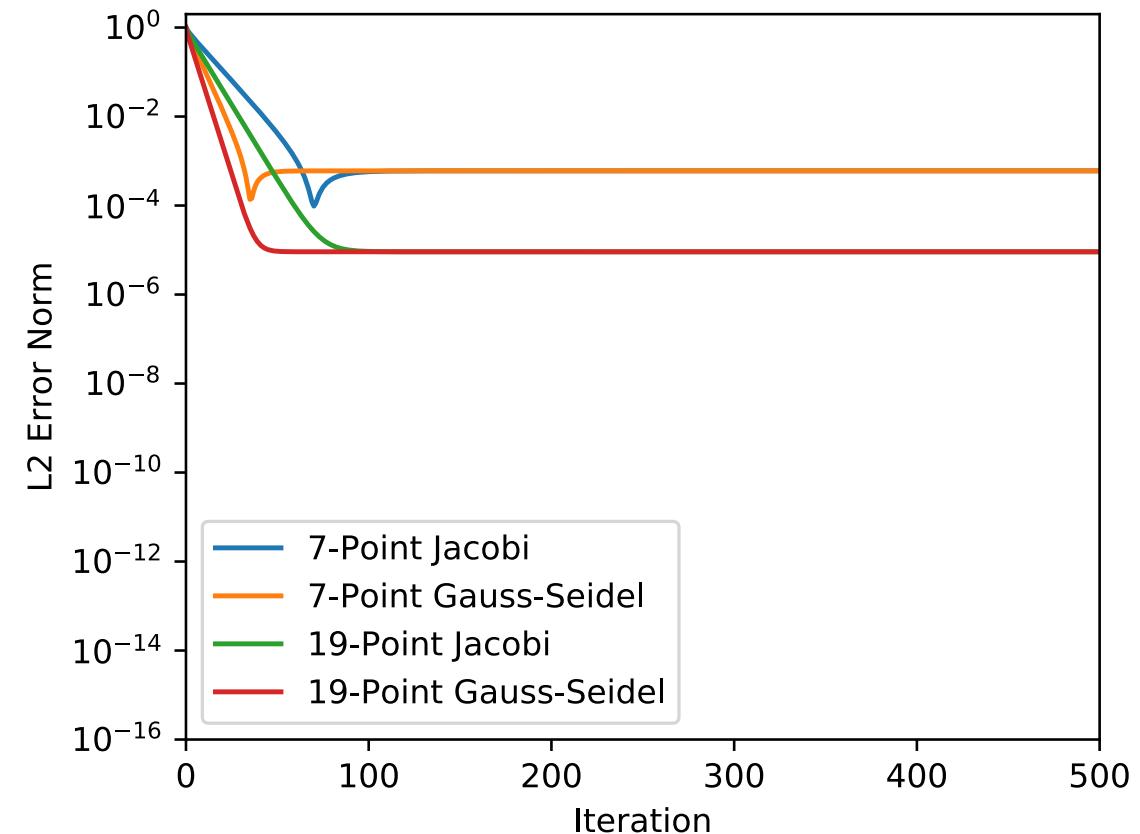
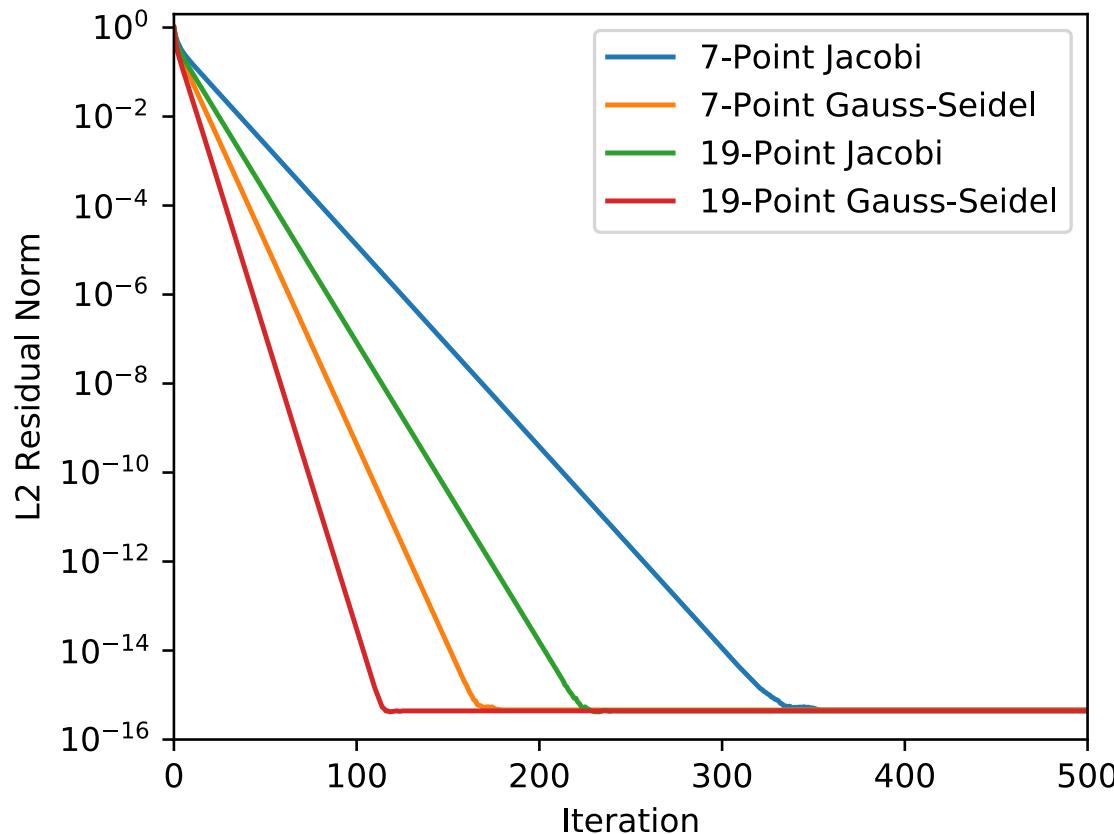
6x6x6
h=0.01



Verification: Inverse Problem

$$\nabla^2 \phi = -6 \sin(2(x + y + z)) \quad \rightarrow \quad \phi = \sin(x + y + z) \cos(x + y + z)$$

6x6x6
h=0.1

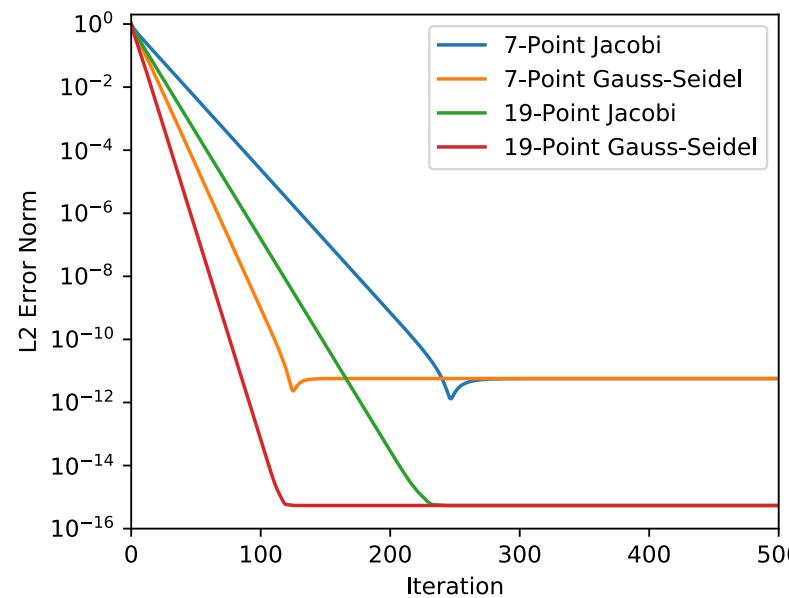


Verification: Batch Reordering & Buffering Algorithms

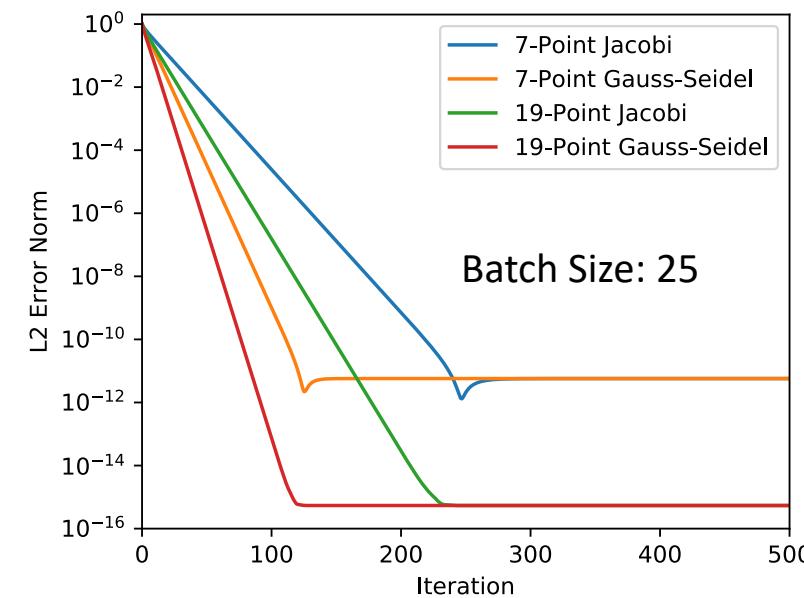
$$\nabla^2 \phi = -6 \sin(2(x + y + z)) \rightarrow \phi = \sin(x + y + z) \cos(x + y + z)$$

6x6x6
h=0.001

Natural Ordering



Batch Reordering



Verification: Forward Problem

Plots Shows $z = 1.55$

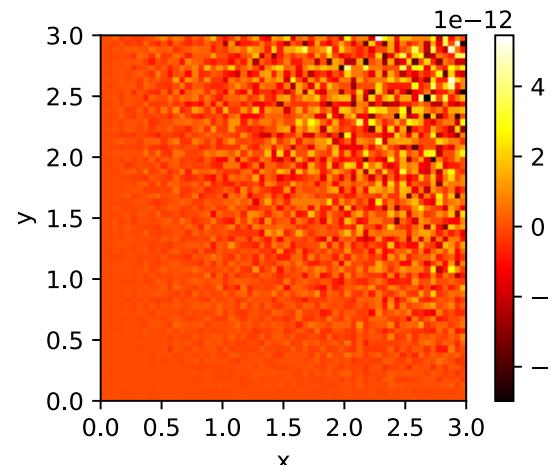
$$\phi = xyz$$

$$\nabla^2 \phi = 0$$

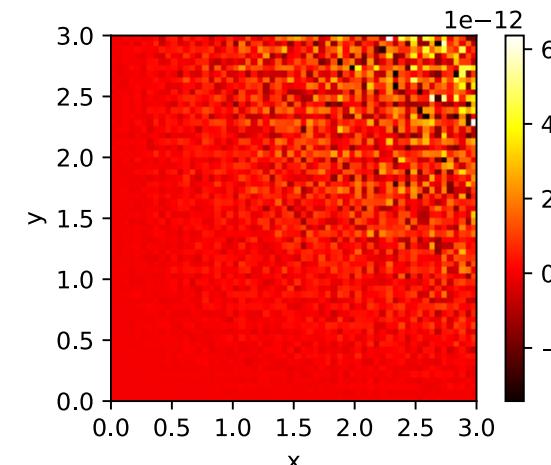
60x60x60, $h = 0.05$, 500 iter.

27-Point Stencil

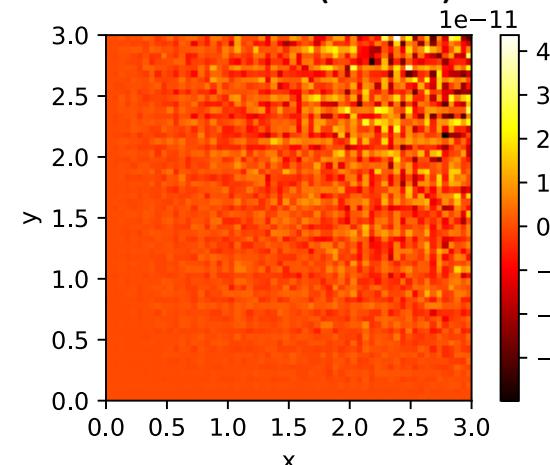
7 Points



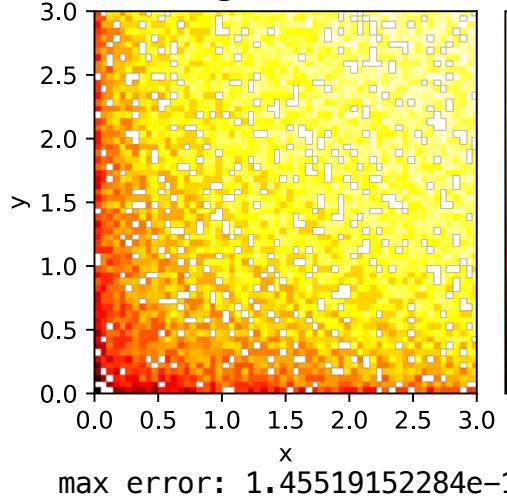
27 Points



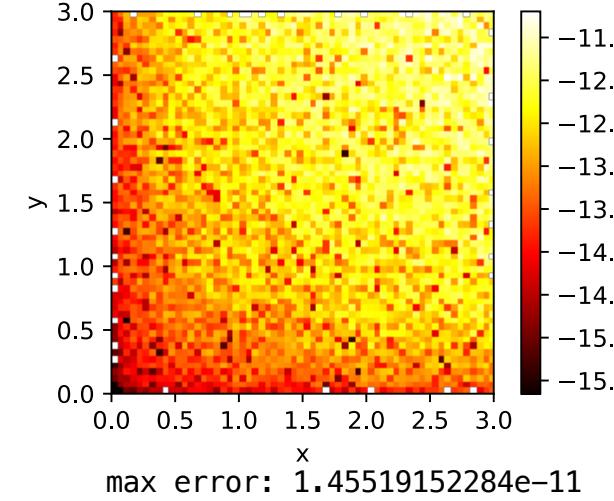
27 Points (HPCG)



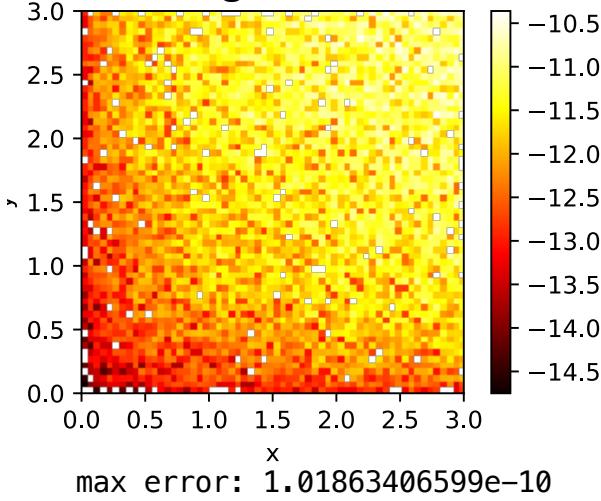
Log10 Error



Log10 Error



Log10 Error



Verification: Forward Problem

Plots Shows $z = 1.55$

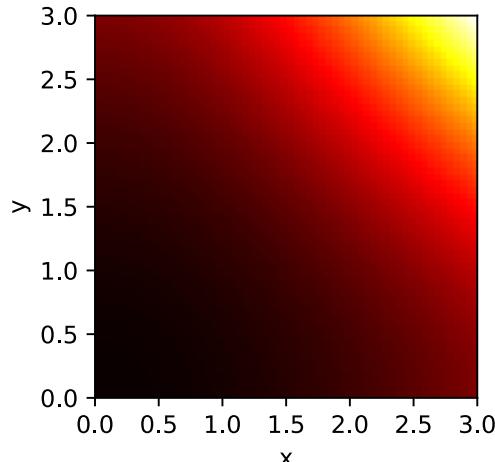
$$\phi = x^2y^2z^2$$

$$\nabla^2\phi = 2y^2z^2 + 2x^2z^2 + 2x^2y^2$$

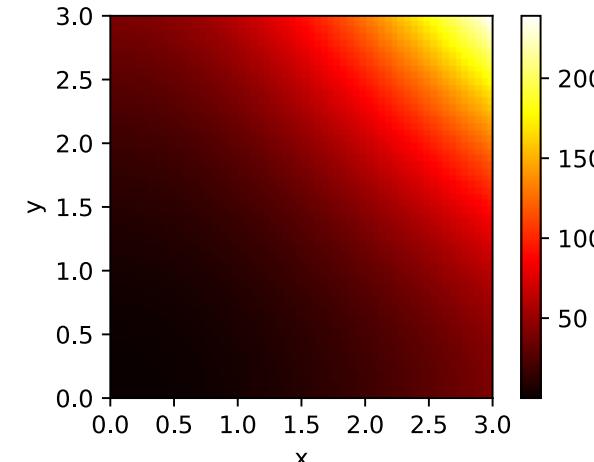
60x60x60, $h = 0.05$, 500 iter.

27-Point Stencil

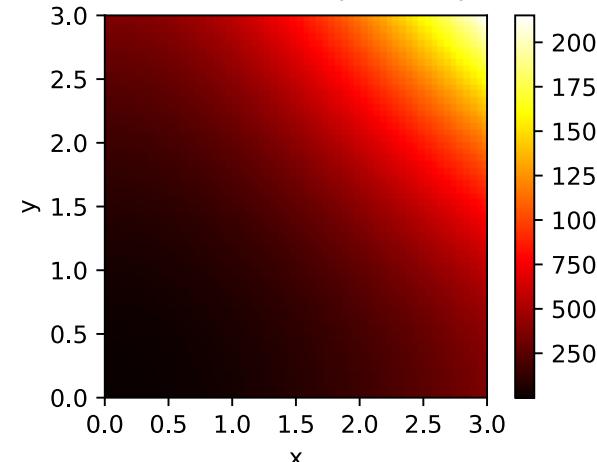
7 Points



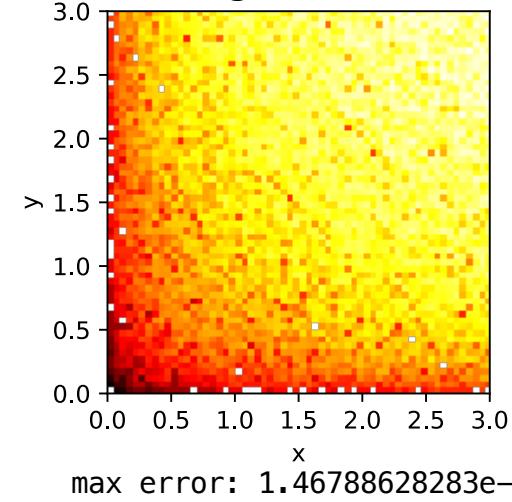
27 Points



27 Points (HPCG)

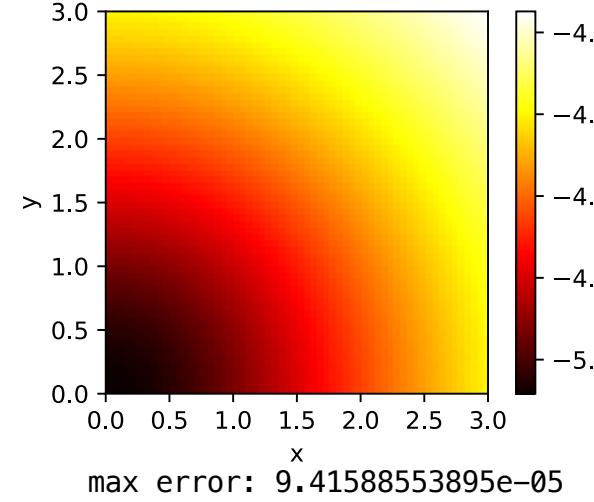


Log10 Error



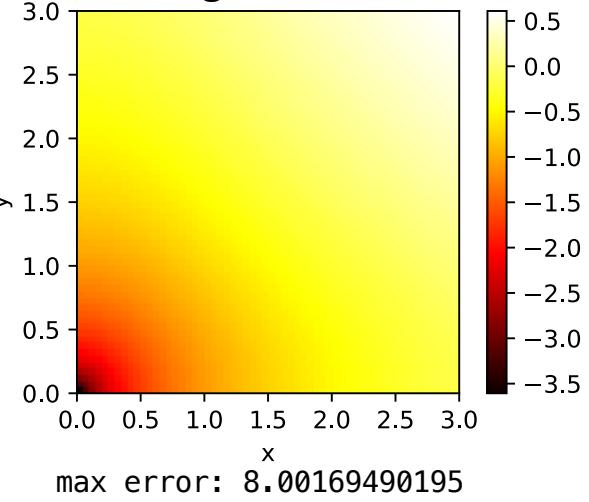
max error: $1.46788628283e-12$

Log10 Error



max error: $9.41588553895e-05$

Log10 Error



max error: 8.00169490195

Verification: Forward Problem

Plots Shows $z = 1.55$

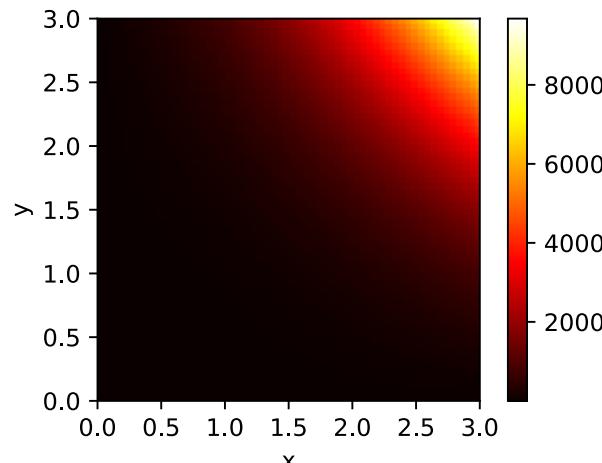
$$\phi = x^3y^3z^3$$

$$\nabla^2\phi = 6xy^3z^3 + 6x^3yz^3 + 6x^3y^3z$$

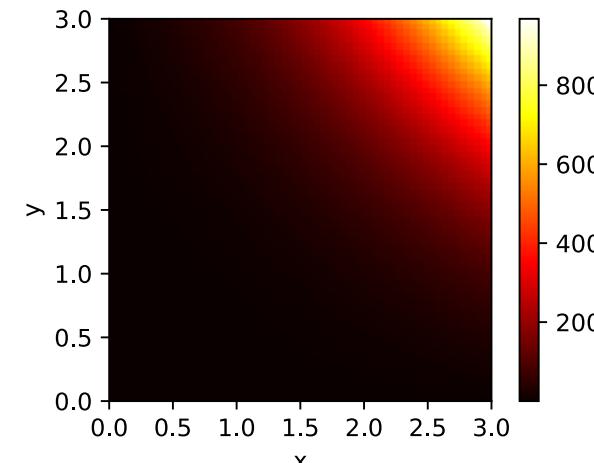
60x60x60, $h = 0.05$, 500 iter.

27-Point Stencil

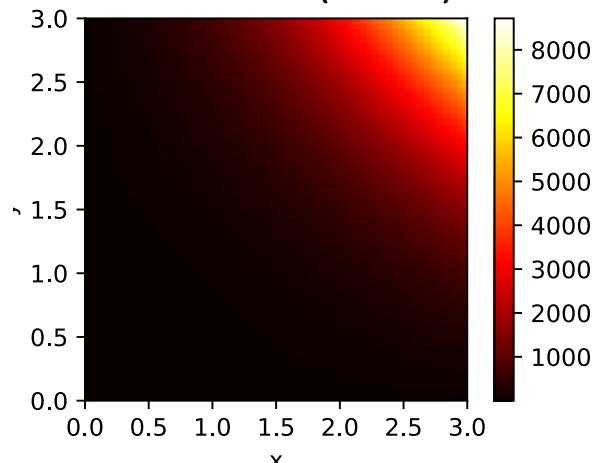
7 Points



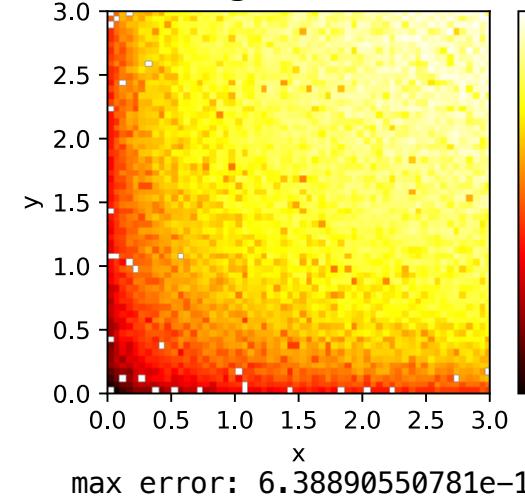
27 Points



27 Points (HPCG)

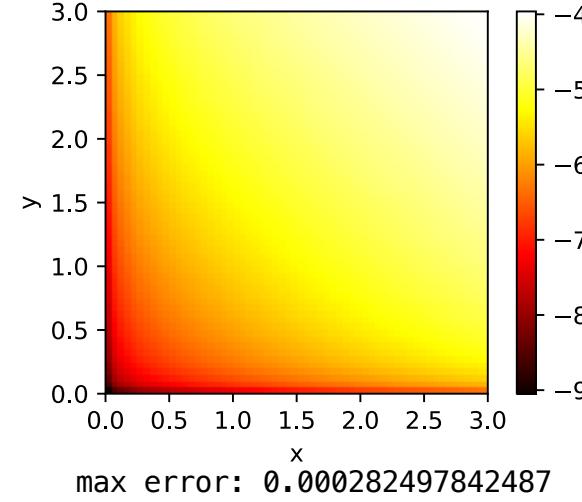


Log10 Error



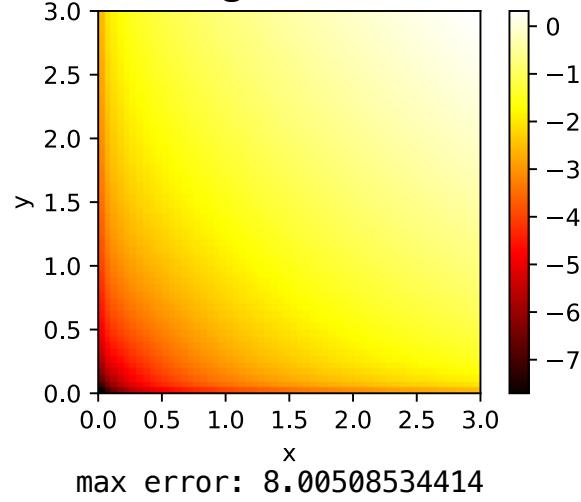
max error: $6.38890550781e-13$

Log10 Error



max error: 0.000282497842487

Log10 Error



max error: 8.00508534414

Verification: Forward Problem

Plots Shows $z = 1.55$

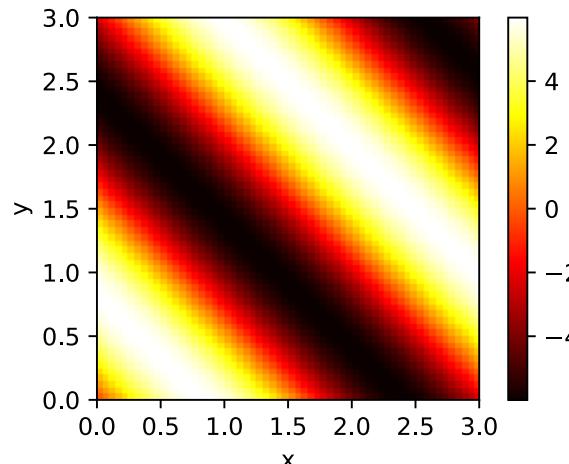
$$\phi = \sin(x + y + z) \cos(x + y + z)$$

$$\nabla^2 \phi = -6 \sin(2(x + y + z))$$

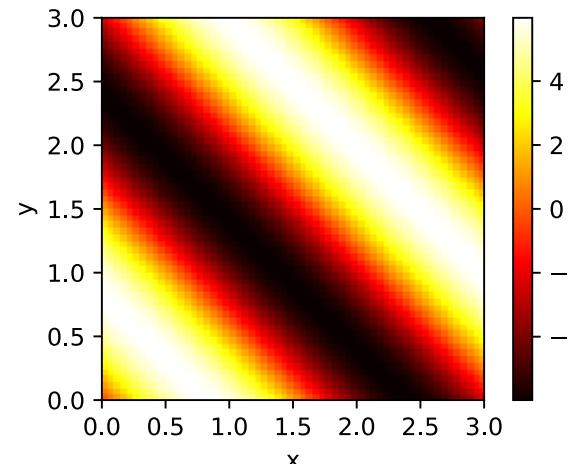
60x60x60, $h = 0.05$, 500 iter.

27-Point Stencil

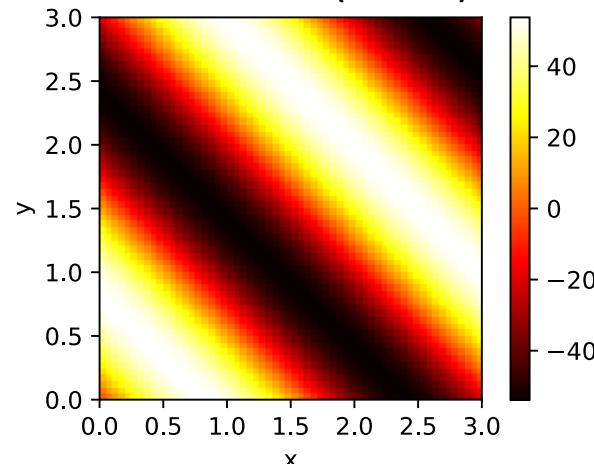
7 Points



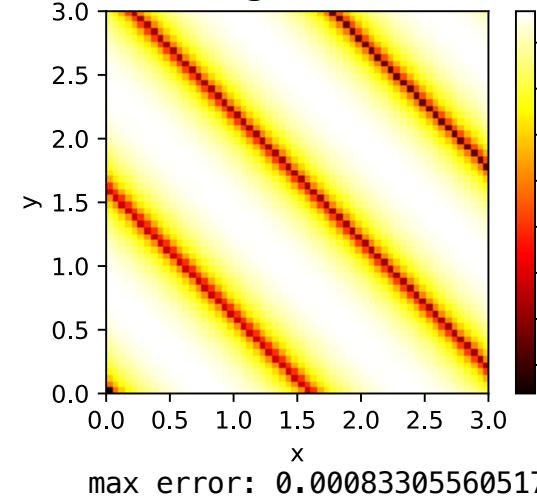
27 Points



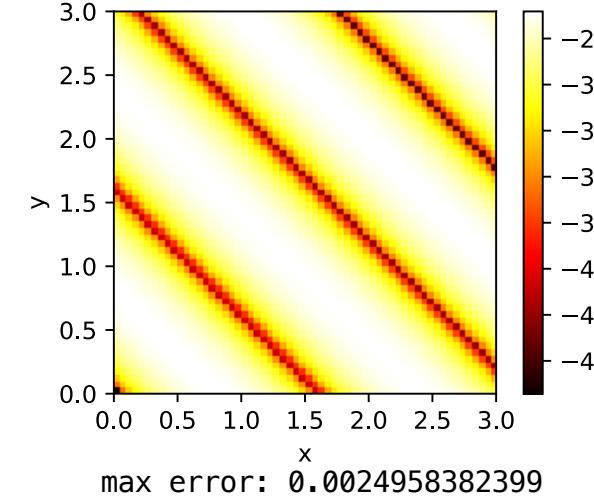
27 Points (HPCG)



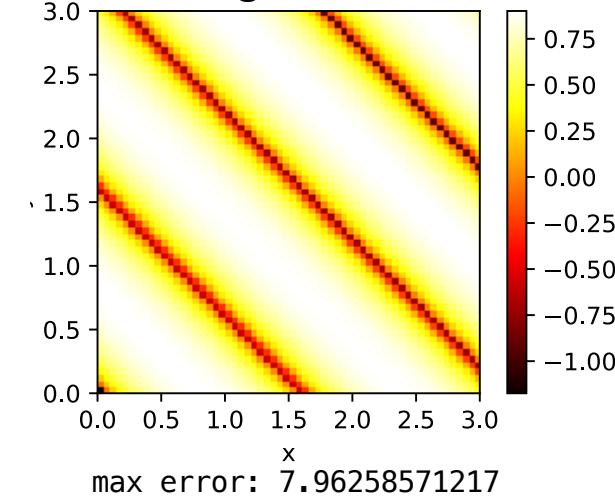
Log10 Error



Log10 Error



Log10 Error



Least Squares: Gradient-Descent Method

$$f(x) = \frac{1}{2}x^T Ax - b^T x + c$$

Finding Gradient:

$$\begin{aligned} f(x + d) &= f(x) + d^T(Ax - b) + \frac{1}{2}d^T Ad \\ &= f(x) + d^T \nabla f(x) + O(d^2) \end{aligned}$$

$$\boxed{\nabla f(x) = (Ax - b) = r}$$

Finding Step Size:

$$f(x + \alpha d) = f(x) + \alpha d^T r + \alpha^2 \frac{1}{2} d^T Ad$$

$$\frac{\partial f}{\partial \alpha} = d^T r + ad^T Ad = 0$$

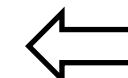
$$\boxed{\alpha = -\frac{d^T r}{d^T Ad}}$$

Gradient-Descent Step:

$$d = \nabla(f) = r$$

$$x \leftarrow x - d \frac{d^T r}{d^T Ad}$$

$$r \leftarrow r - Ad \frac{d^T r}{d^T Ad}$$



Can be implemented
with a single MVM

Sanity Check: step direction has to be orthogonal to the next gradient

$$\langle d, \nabla f(x + \alpha d) \rangle = d^T \left[Ax - Ad \frac{d^T r}{d^T Ad} - b \right] = d^T [Ax - b] - d^T r = 0$$

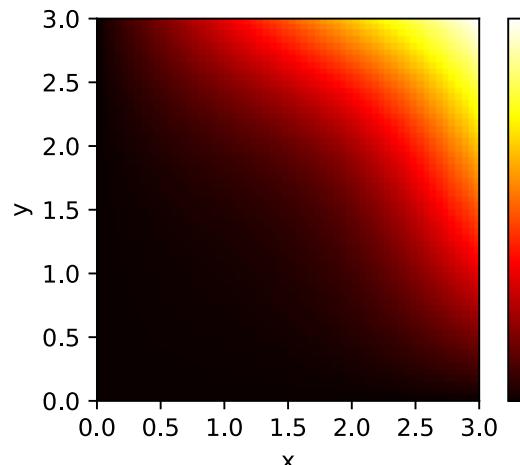
Verification: Inverse Problem

Plots Shows $z = 1.55$

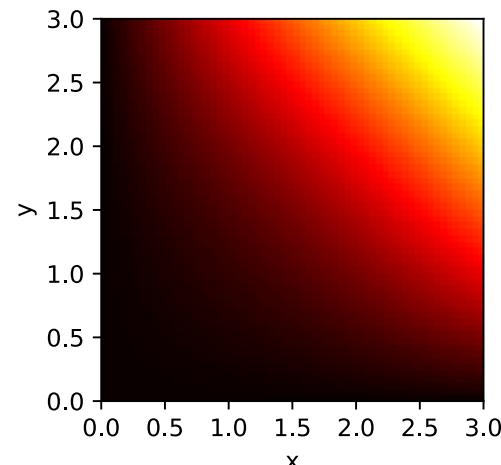
$$\nabla^2 \phi = 0$$
$$\phi = xyz$$

60x60x60, $h = 0.05$, 500 iter.
27-Point Stencil

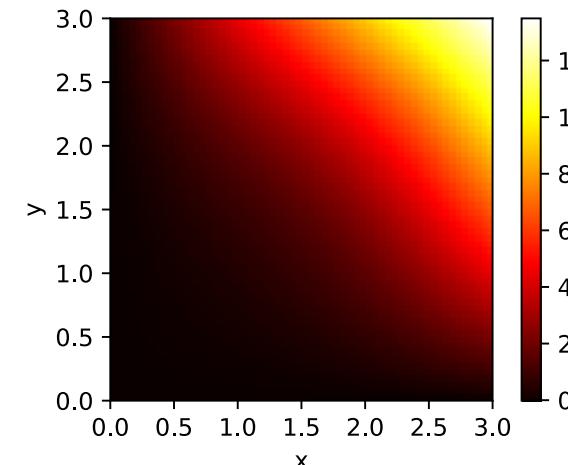
Jacobi



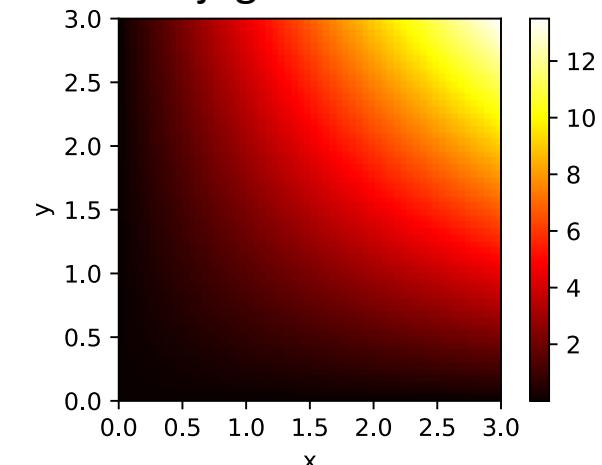
Gauss-Seidel



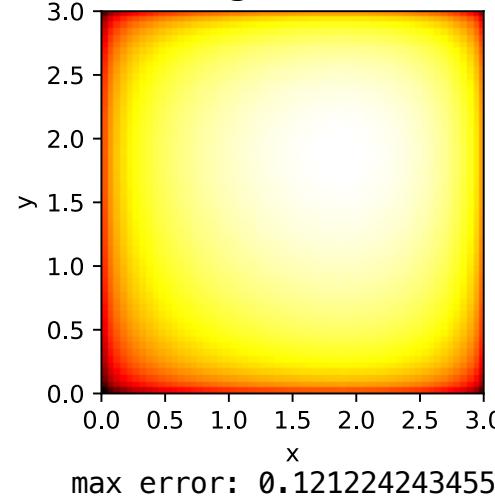
Gradient-Descent



Conjugate-Gradient

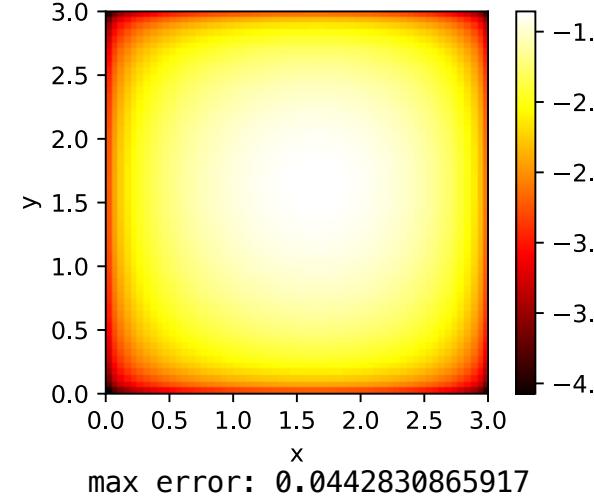


Log10 Error



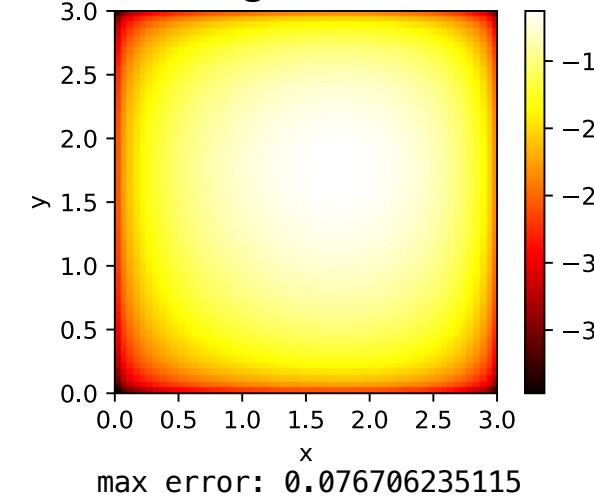
max error: 0.121224243455

Log10 Error



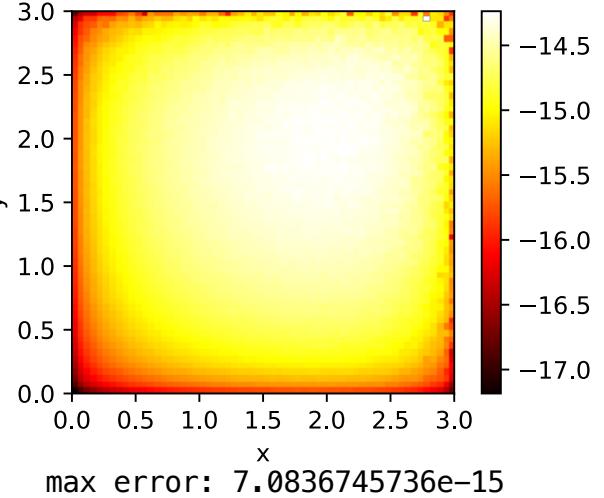
max error: 0.0442830865917

Log10 Error



max error: 0.076706235115

Log10 Error



max error: 7.0836745736e-15

Verification: Inverse Problem

Plots Shows $z = 1.55$

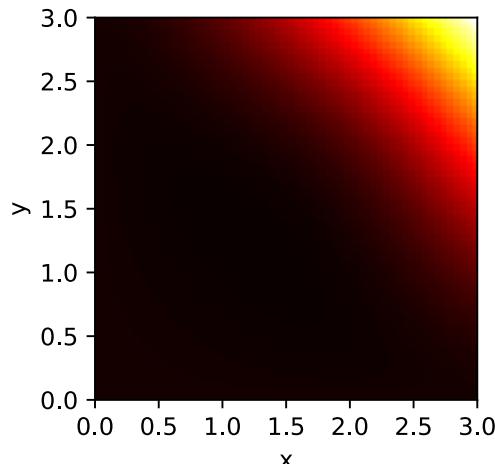
$$\nabla^2\phi = 2y^2z^2 + 2x^2z^2 + 2x^2y^2$$

$$\phi = x^2y^2z^2$$

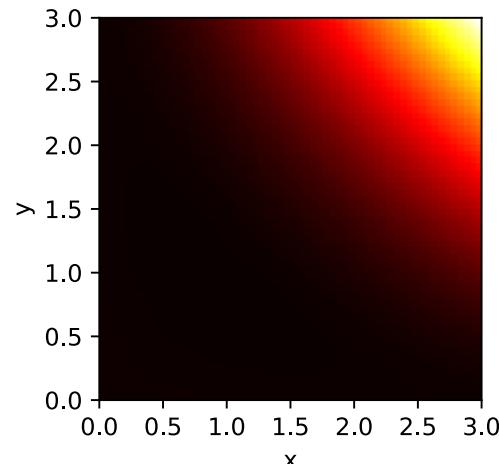
60x60x60, $h = 0.05$, 500 iter.

27-Point Stencil

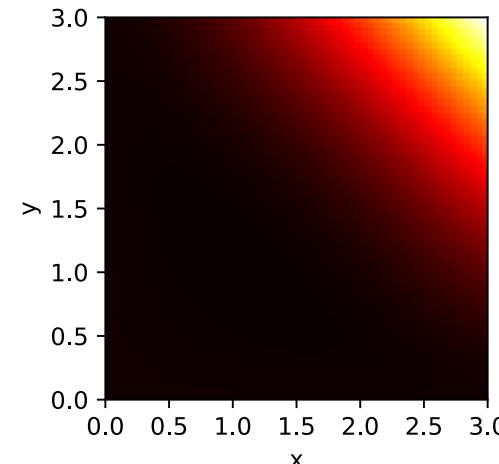
Jacobi



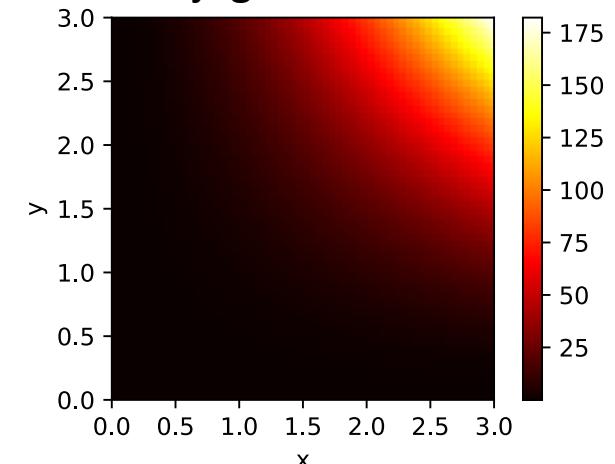
Gauss-Seidel



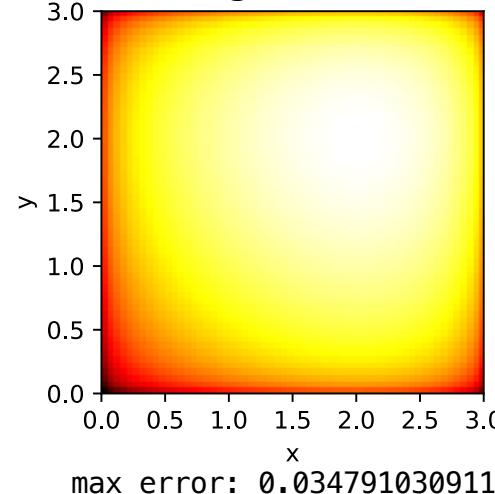
Gradient-Descent



Conjugate-Gradient

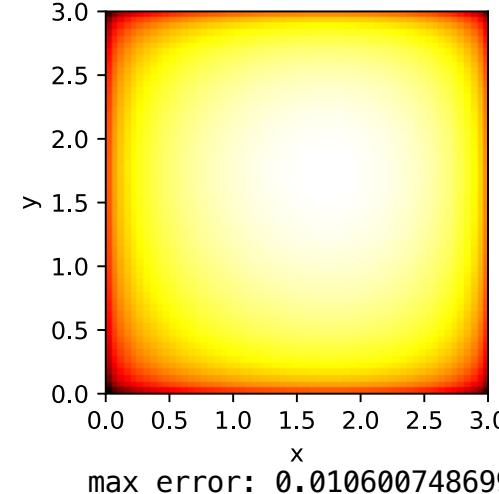


Log10 Error



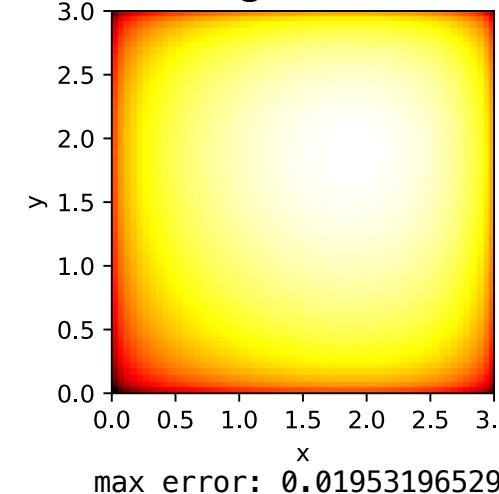
max error: 0.0347910309115

Log10 Error



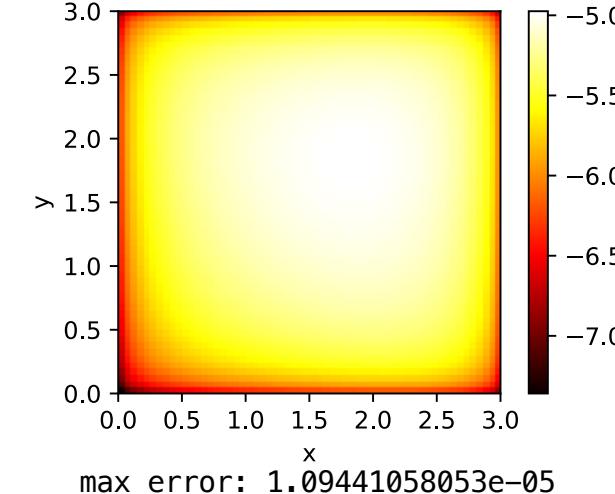
max error: 0.0106007486994

Log10 Error



max error: 0.0195319652925

Log10 Error



max error: 1.09441058053e-05

Verification: Inverse Problem

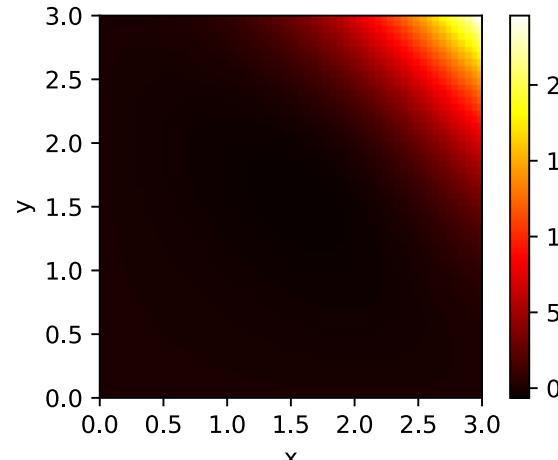
$$\nabla^2 \phi = 6xy^3z^3 + 6x^3yz^3 + 6x^3y^3z \\ \phi = x^3y^3z^3$$

Plots Shows $z = 1.55$

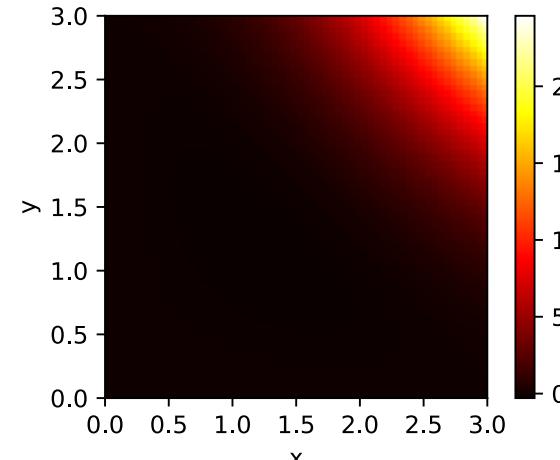
$60 \times 60 \times 60$, $h = 0.05$, 500 iter.

27-Point Stencil

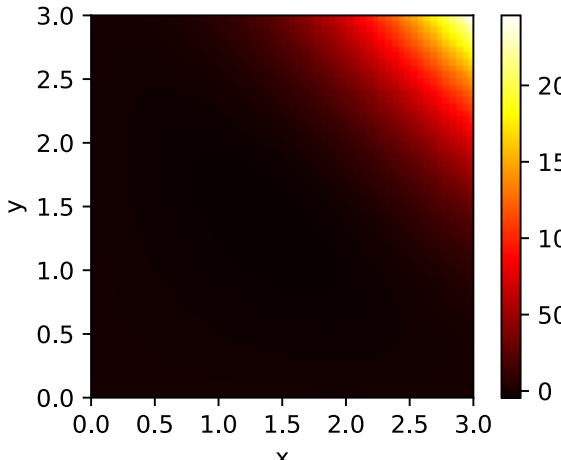
Jacobi



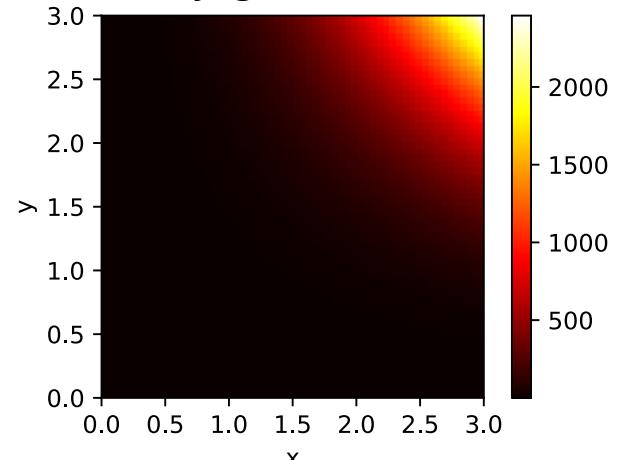
Gauss-Seidel



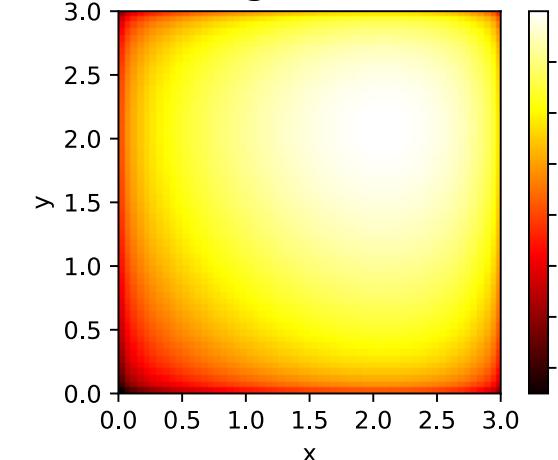
Gradient-Descent



Conjugate-Gradient

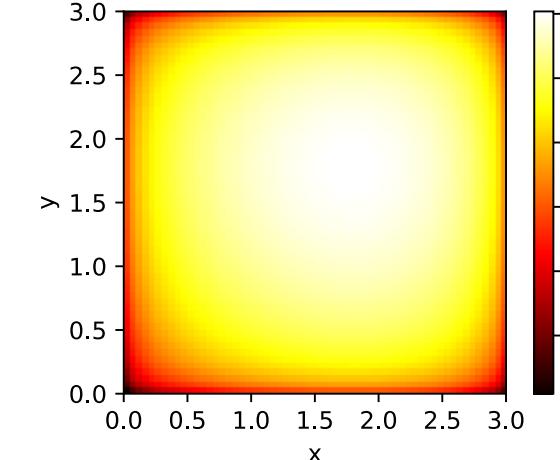


Log10 Error



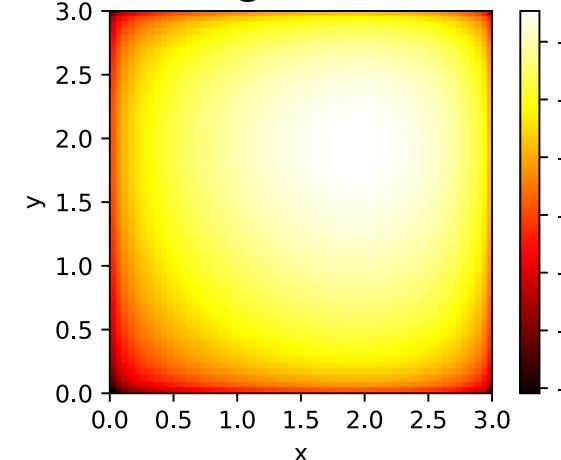
max error: 0.0128737509551

Log10 Error



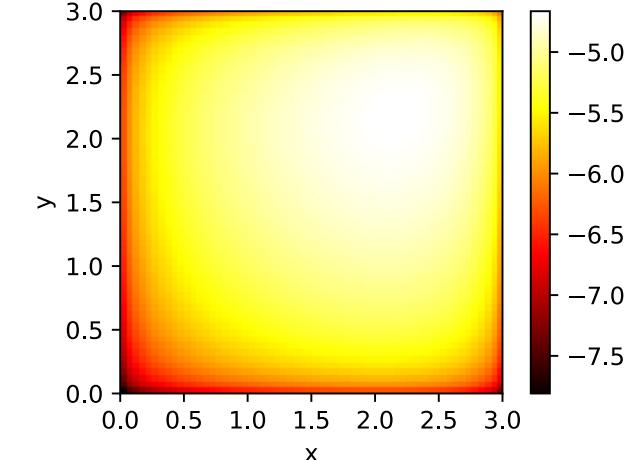
max error: 0.00342426882486

Log10 Error



max error: 0.00661606913659

Log10 Error



max error: 2.86034236171e-05

Verification: Inverse Problem

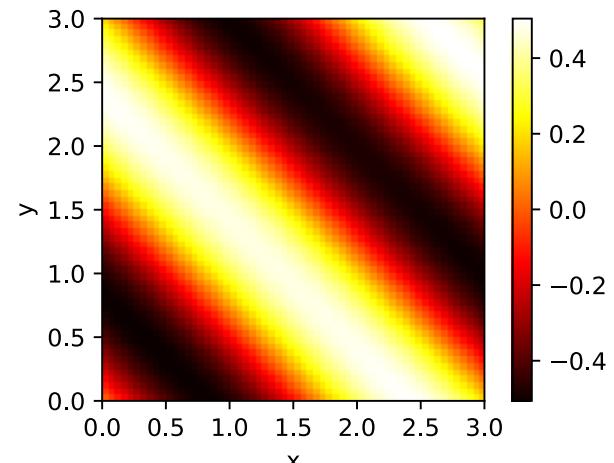
Plots Shows $z = 1.55$

$$\nabla^2 \phi = -6 \sin(2(x + y + z))$$
$$\phi = \sin(x + y + z) \cos(x + y + z)$$

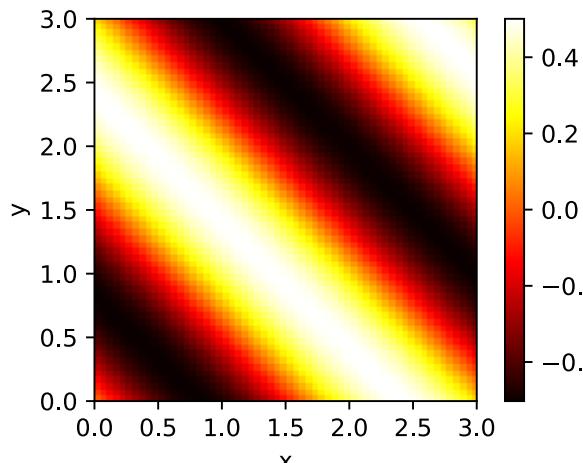
60x60x60, $h = 0.05$, 500 iter.

27-Point Stencil

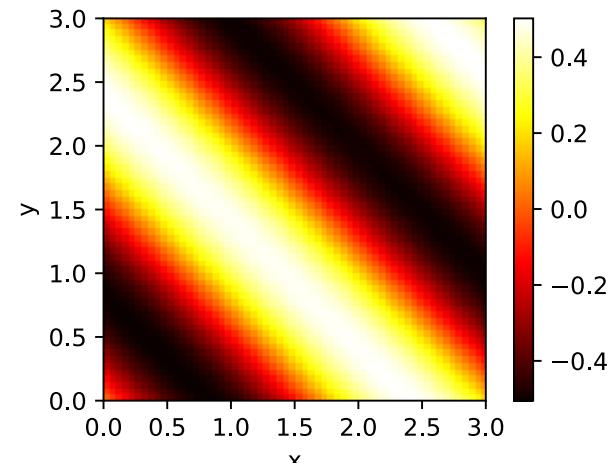
Jacobi



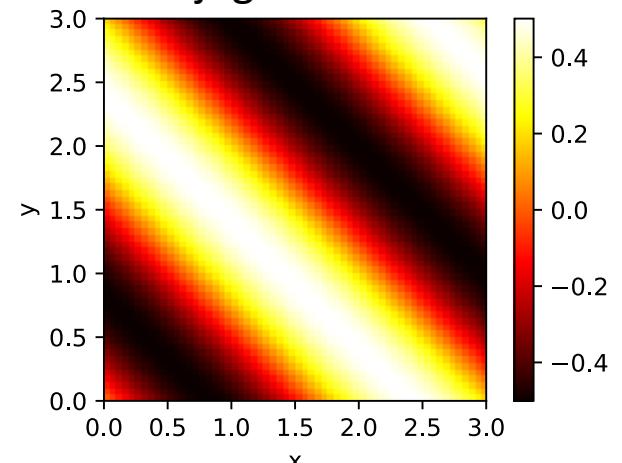
Gauss-Seidel



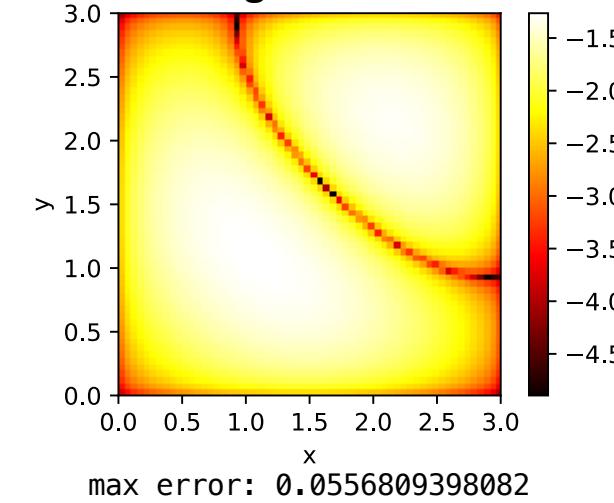
Gradient-Descent



Conjugate-Gradient

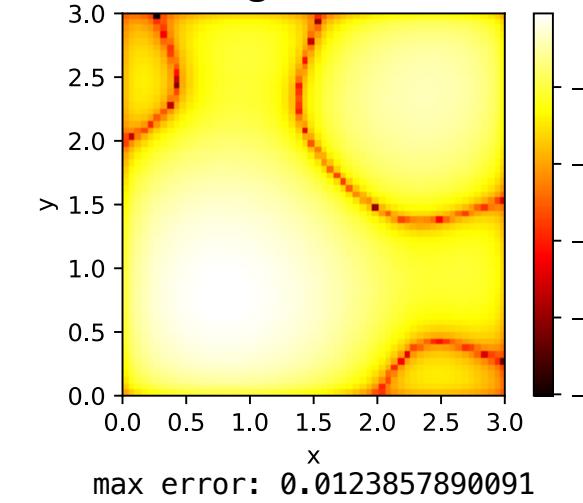


Log10 Error



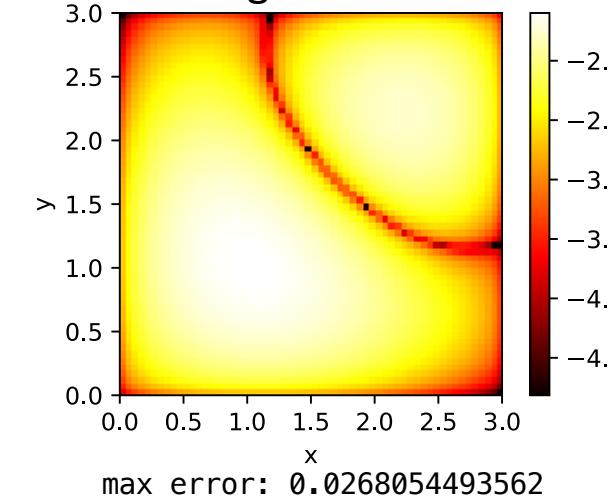
max error: 0.0556809398082

Log10 Error



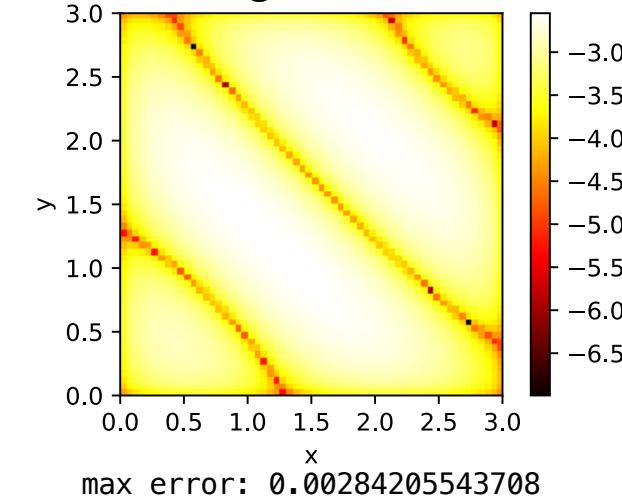
max error: 0.0123857890091

Log10 Error



max error: 0.0268054493562

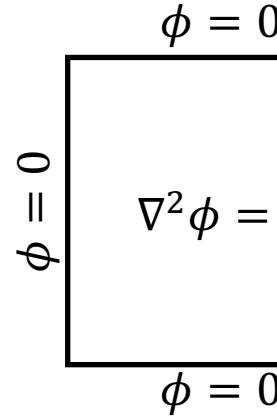
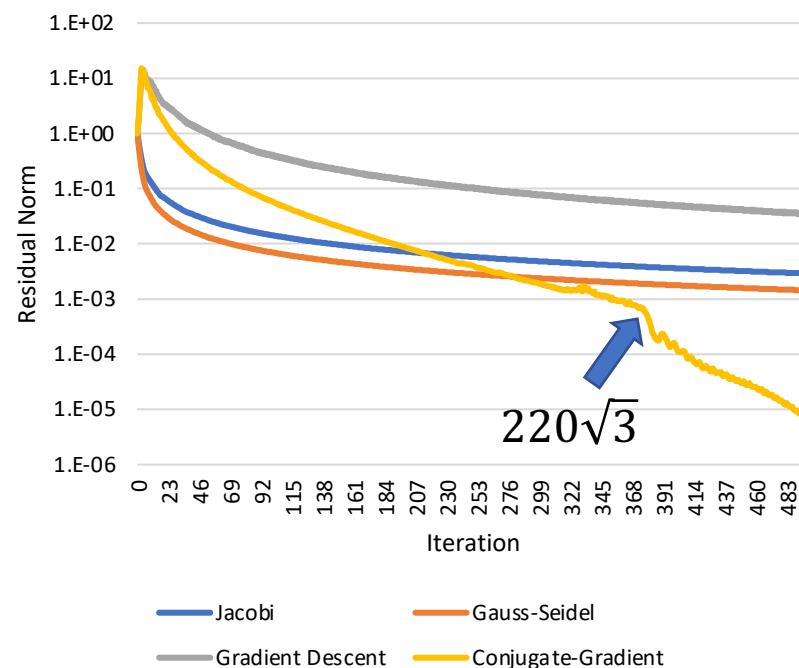
Log10 Error



max error: 0.00284205543708

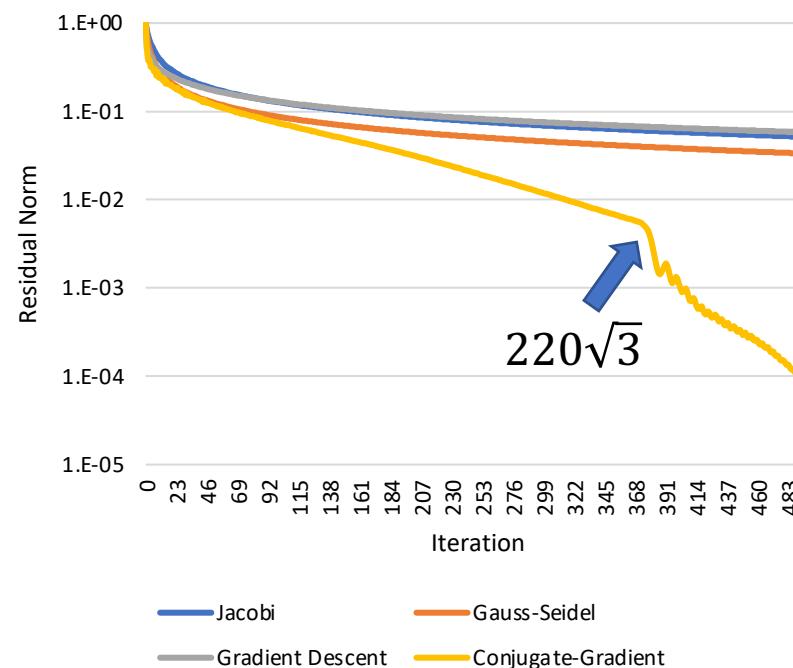
Verification: Convergence

L0 Relative Norm

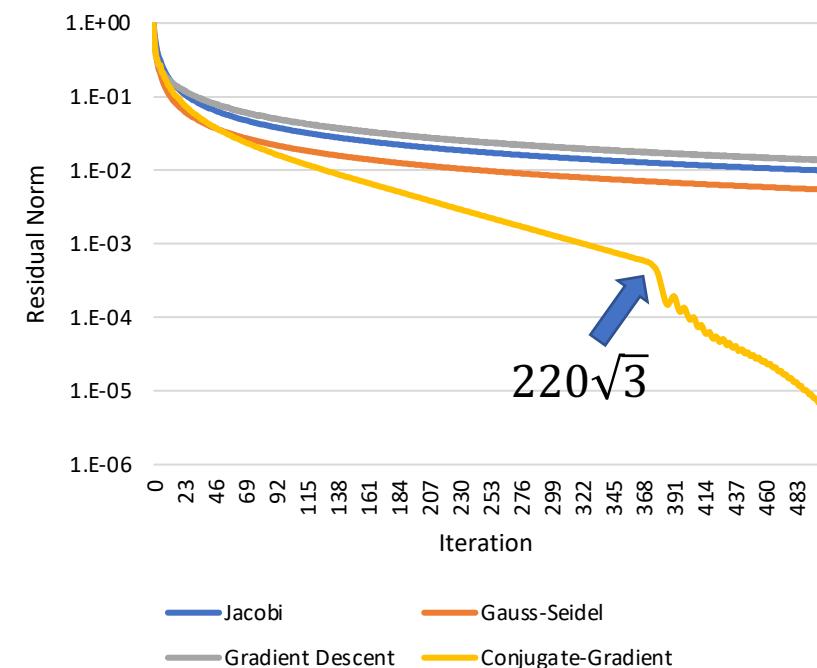


220x220x220
7-point equation

L1 Relative Norm

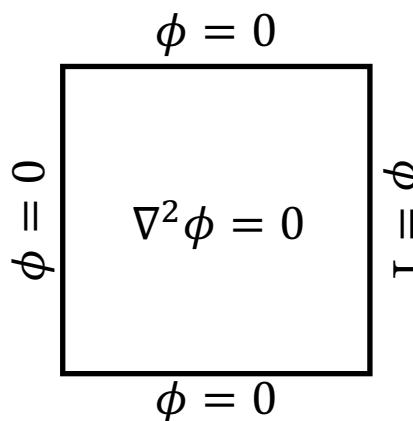


L2 Relative Norm



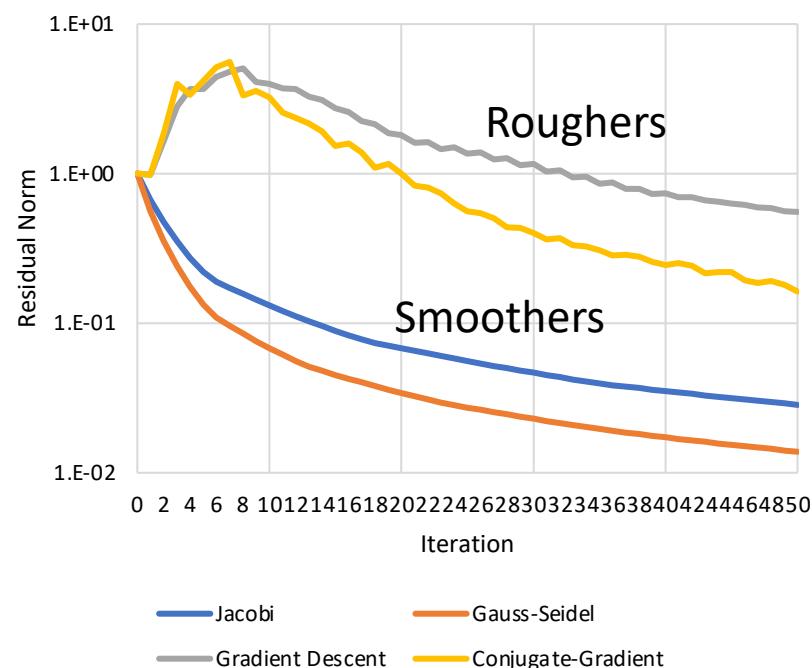
Verification: Convergence

ZOOMED!

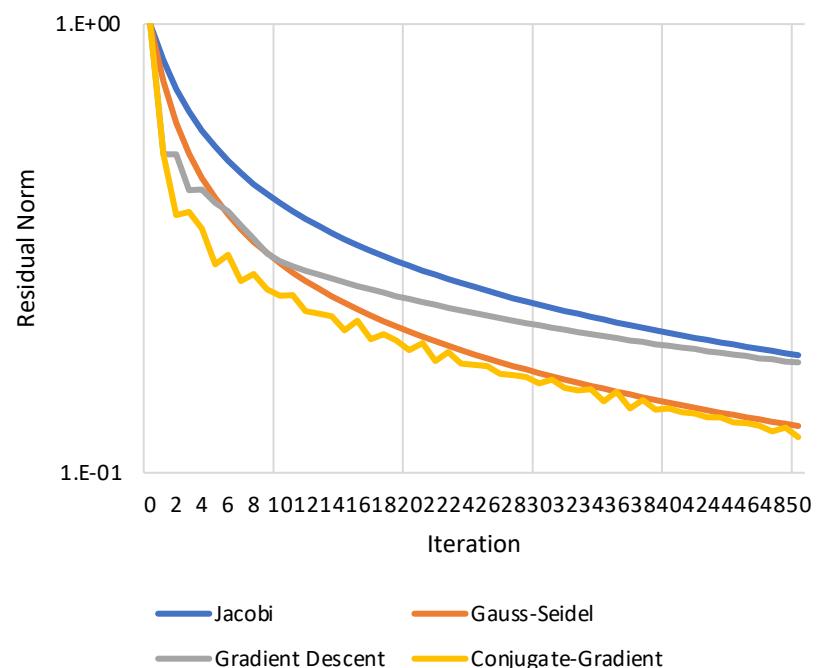


220x220x220
7-point equation

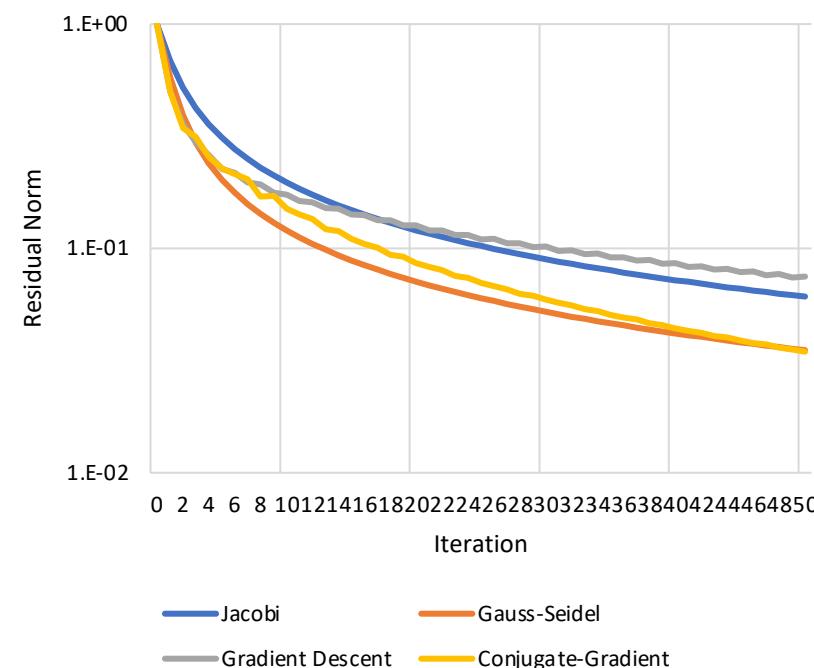
L0 Relative Norm



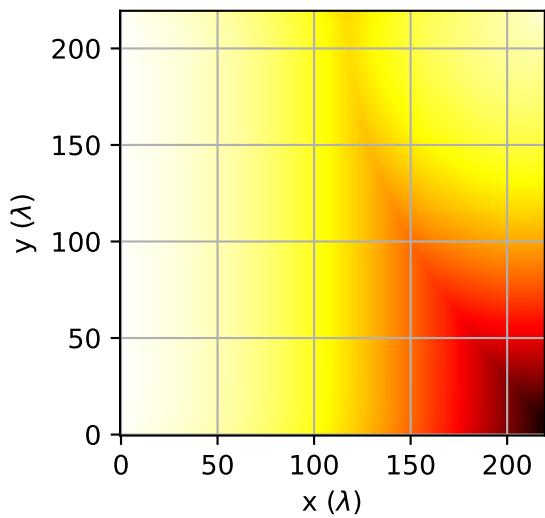
L1 Relative Norm



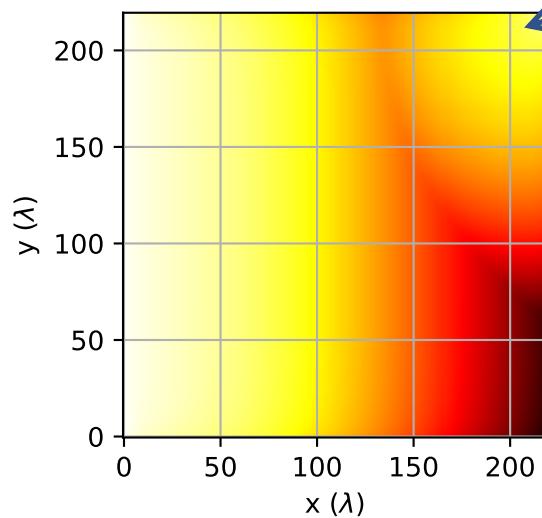
L2 Relative Norm



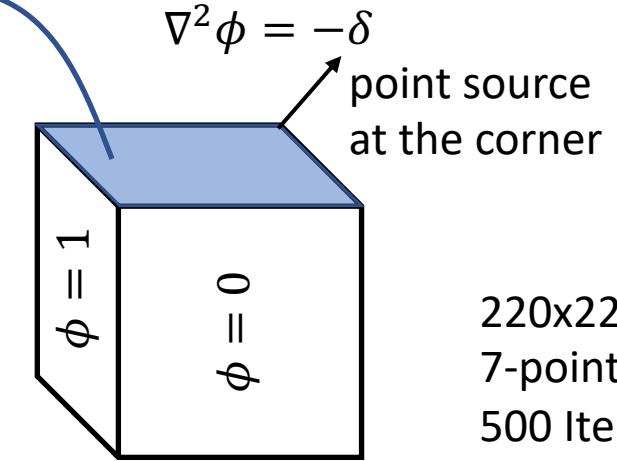
Jacobi



Gauss-Seidel

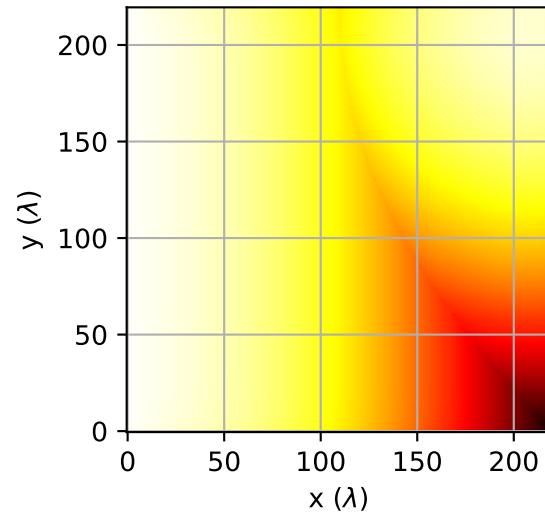


Top Layer

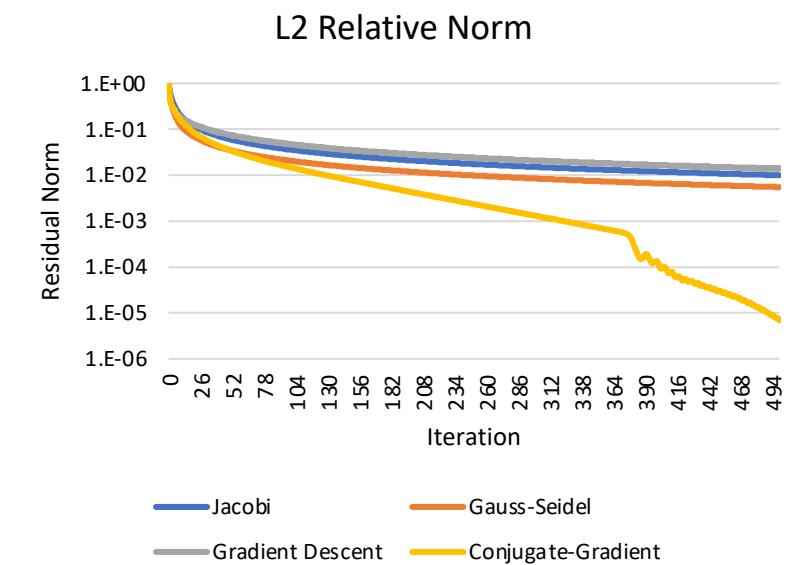
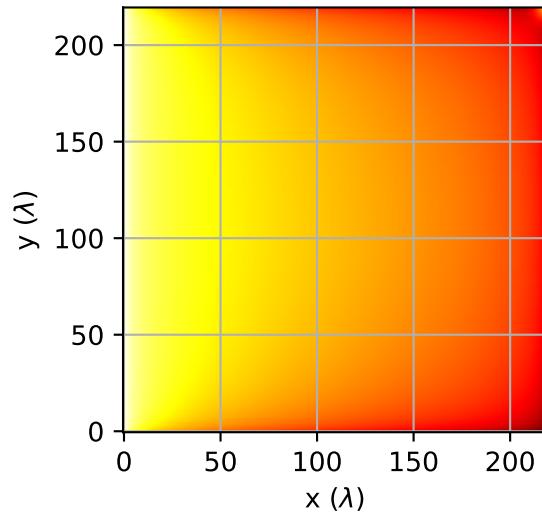


220x220x220
7-point equation
500 Iterations

Gradient Descent

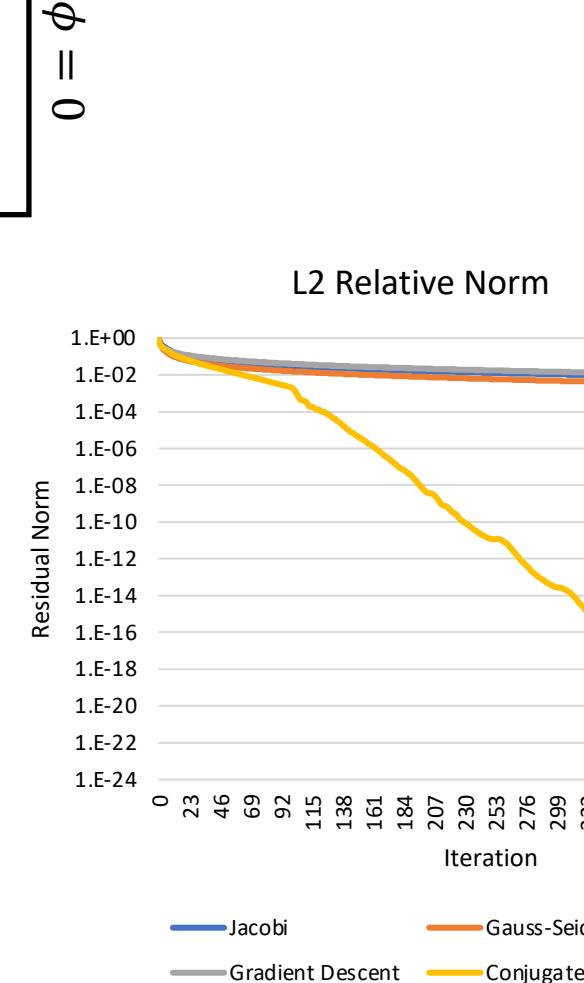
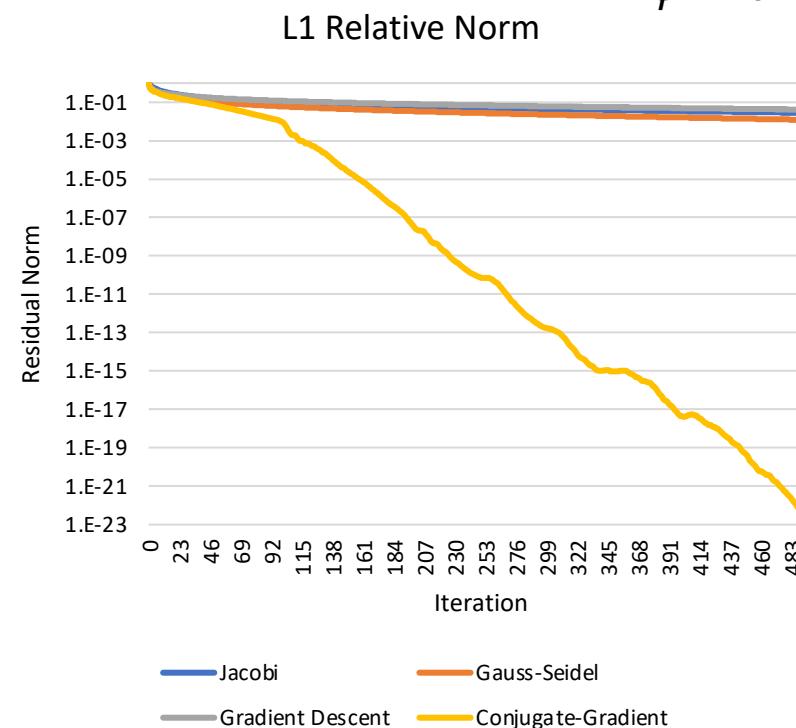
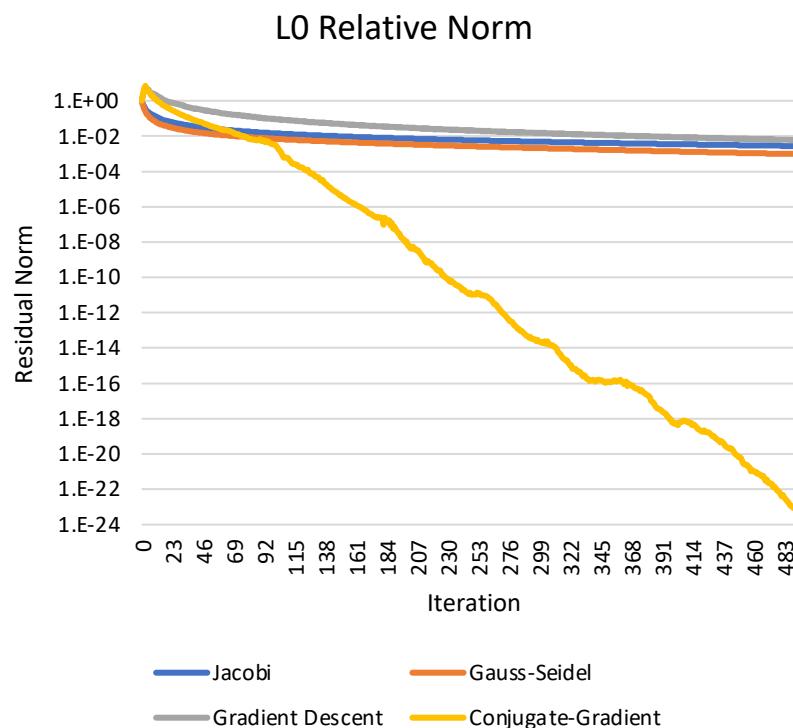


Conjugate Gradient



Verification: Convergence

60x60x60, $h = 0.05$
7-point equation



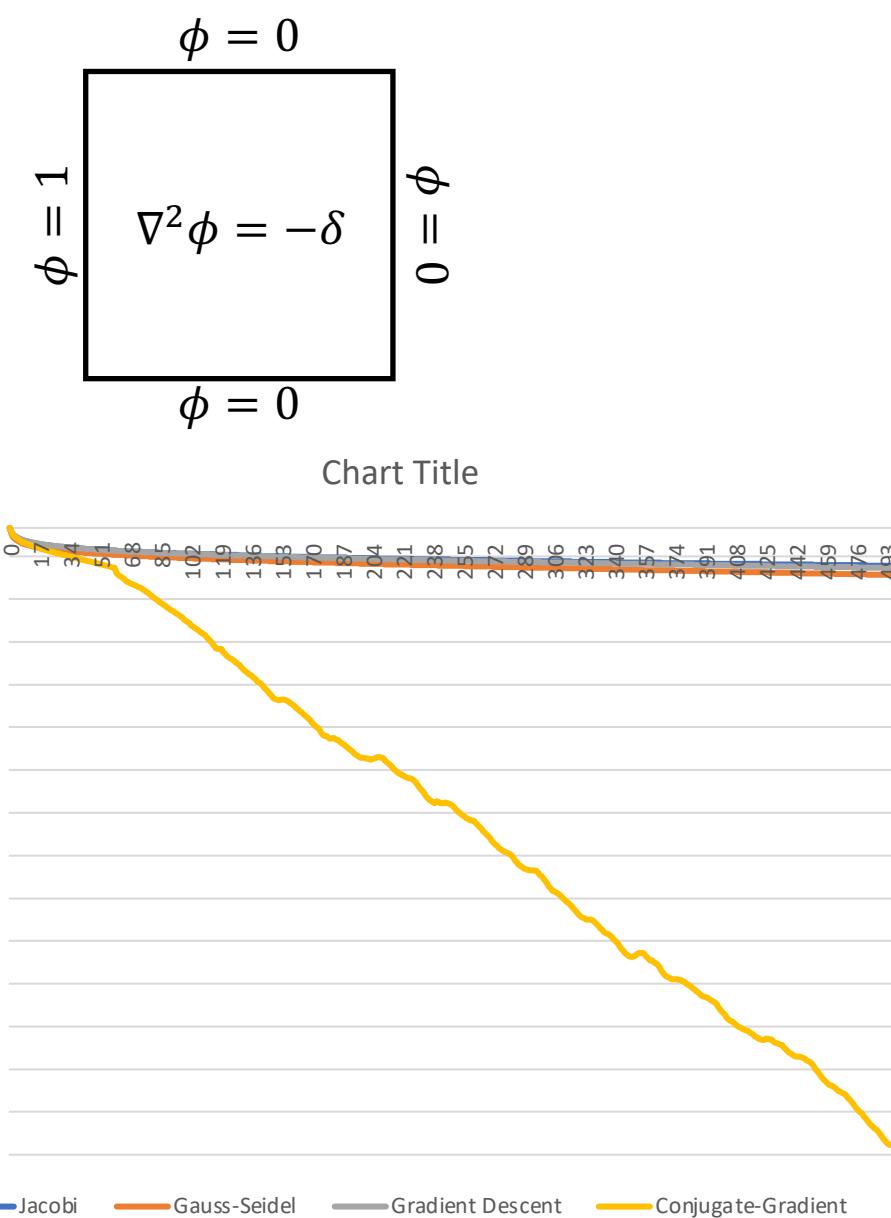
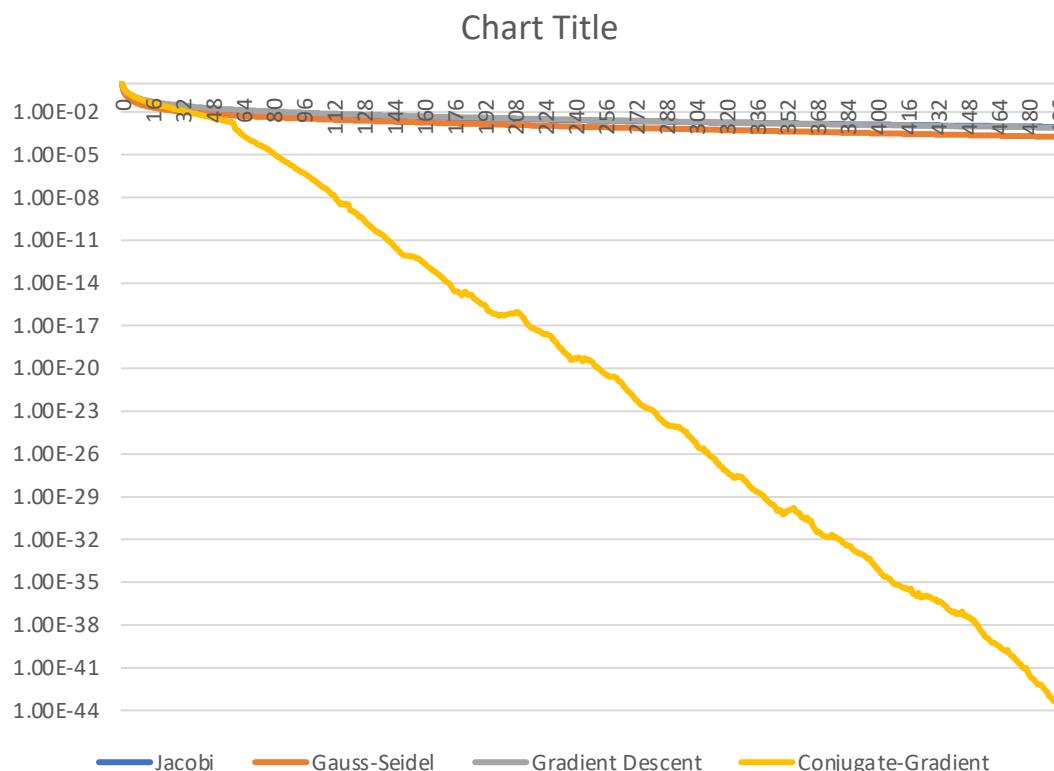
IBM Research



center for
cognitive computing
systems research

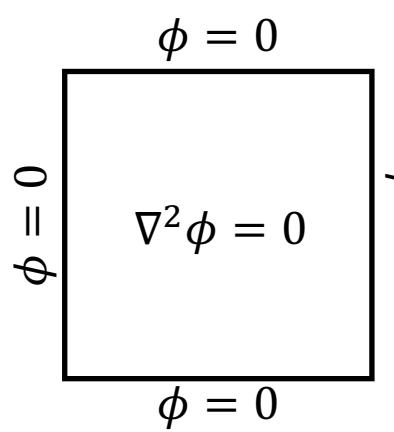
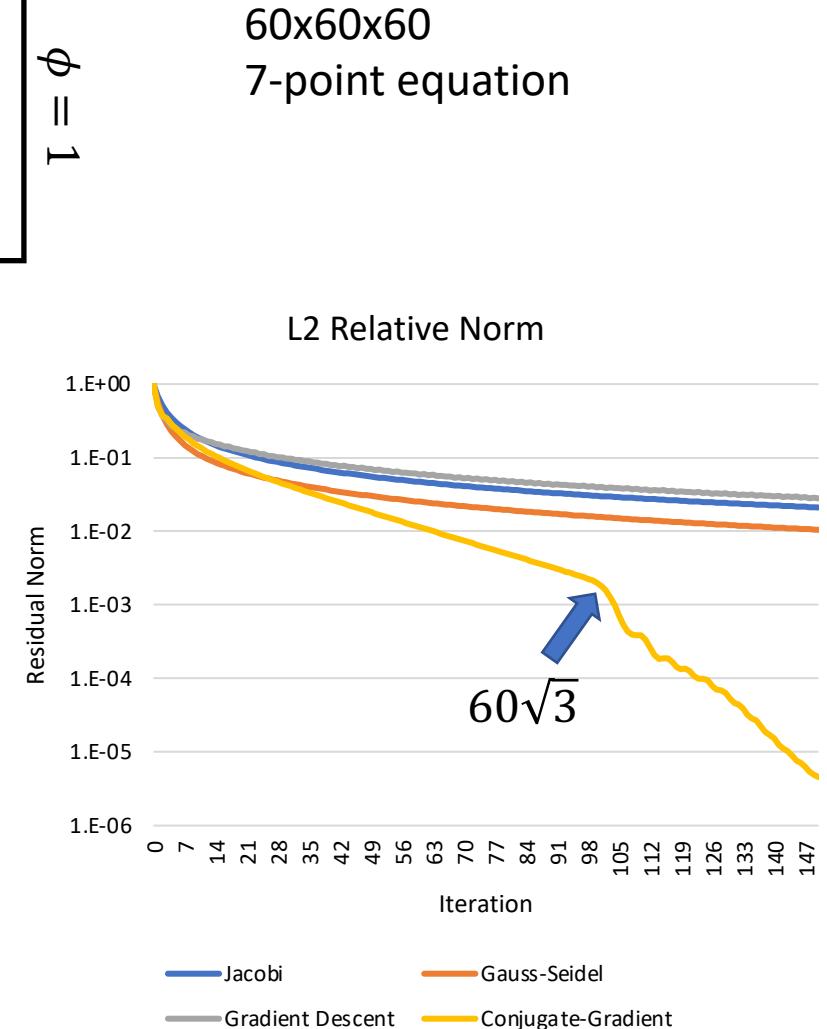
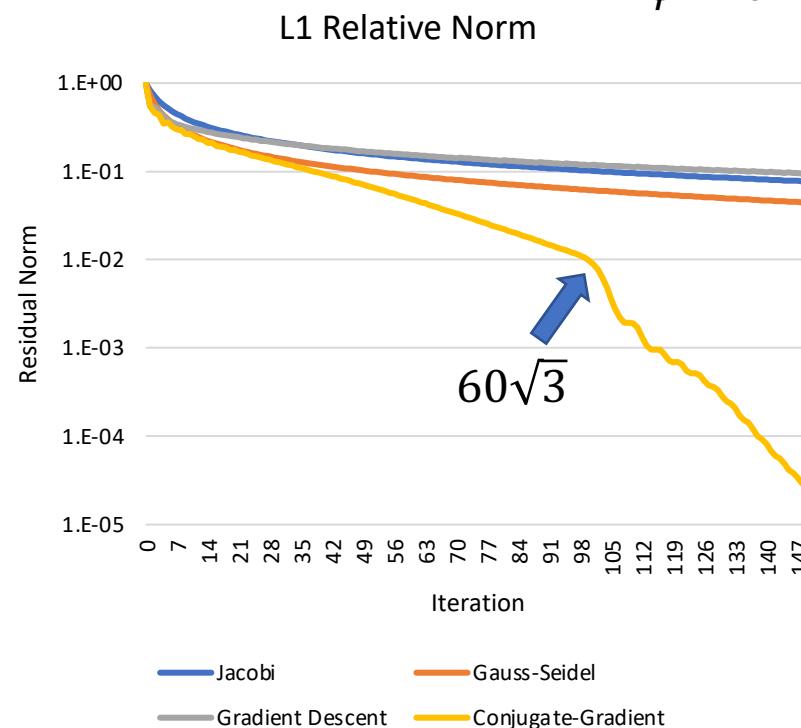
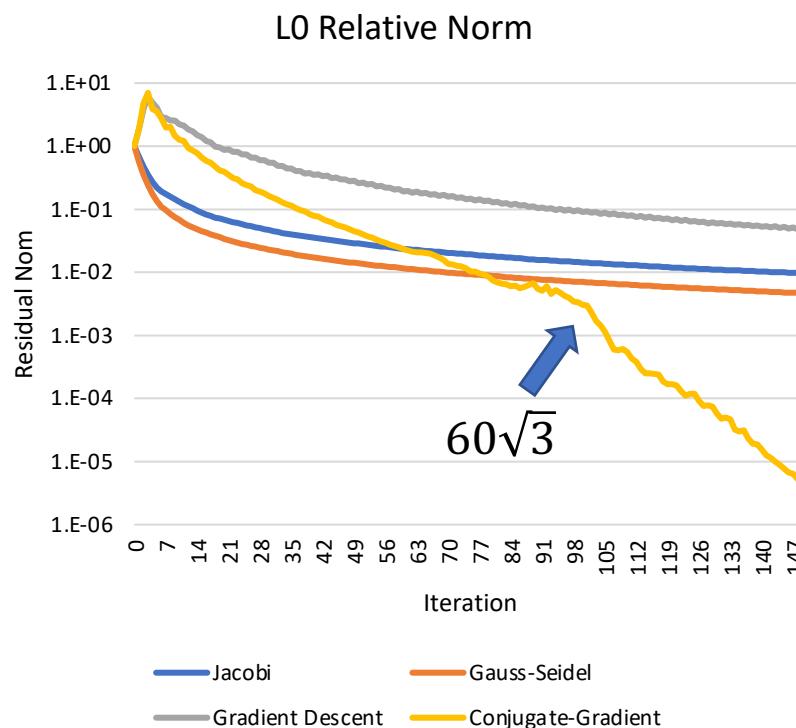
Verification: Convergence

60x60x60, h = 0.05
27-point equation (HPCG)



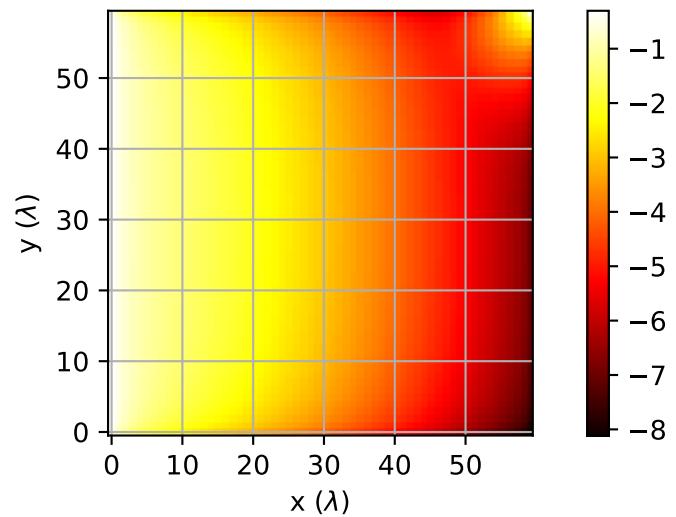
Verification: Convergence

ZOOMED!

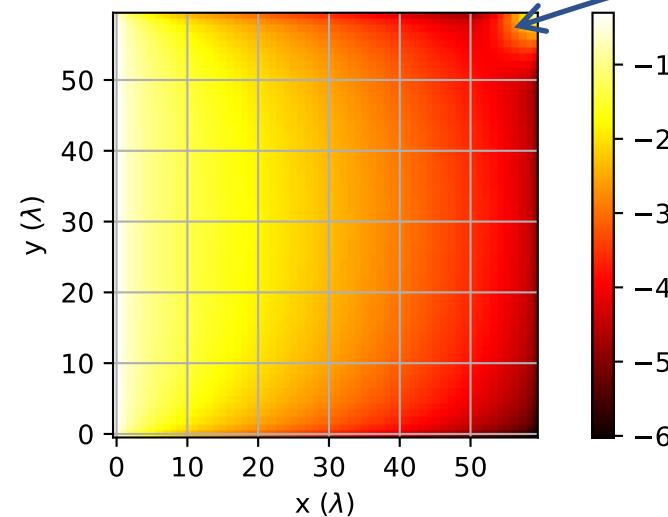


60x60x60
7-point equation

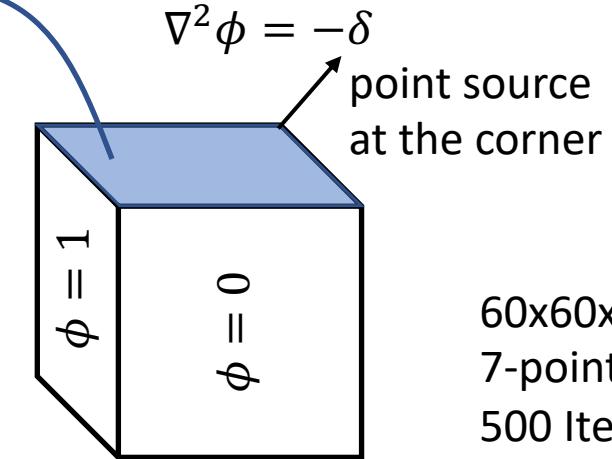
Jacobi



Gauss-Seidel

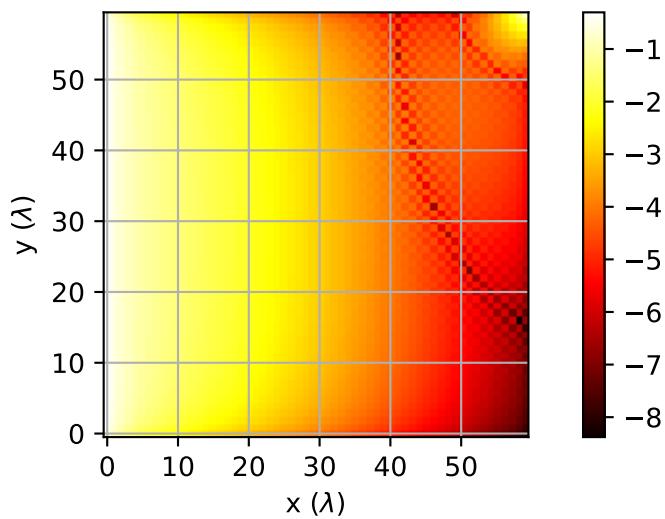


Top Layer

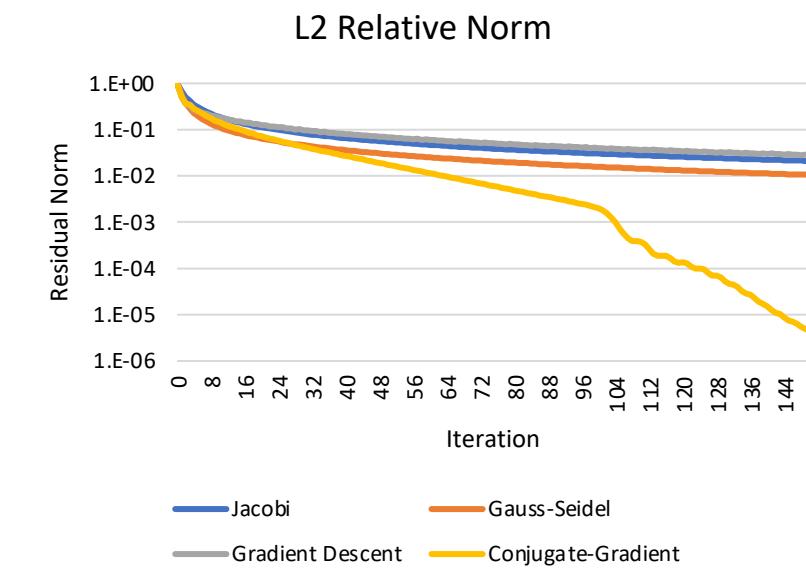
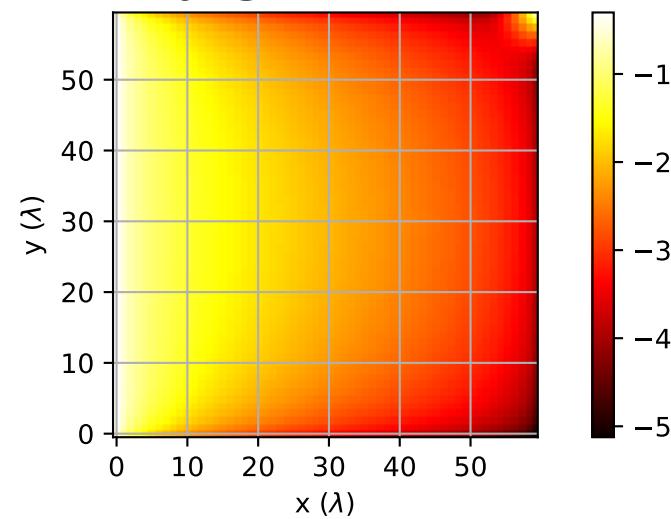


60x60x60
7-point equation
500 Iterations

Gradient Descent



Conjugate Gradient



Computational Kernels in Each Iteration

Other HPCG Operations

```
//GAUSS-SEIDEL FORWARD SWEEP
for(int batch = 0; batch < numbatch; batch++){
    //READ MAPPING DATA
    int start = mapdispl[batch];
    int mapnz = mapdispl[batch+1]-start;
    //ALLOCATE BUFFER
    double buffer[batchsize+mapnz];
    //LOAD INTER-BATCH DATA
    for(int n = 0; n < mapnz; n++)
        buffer[batchsize+n] = x[map[start+n]];
    //LOAD INTRA-BATCH DATA
    start = batch*batchsize;
    for(int n = 0; n < batchsize; n++)
        buffer[n] = x[start+n];
    //PERFORM GAUSS-SEIDEL PER BATCH
    for(int m = start; m < start+batchsize; m++){
        double reduce = b[m];
        for(int n = displ[m]; n < displ[m+1]; n++)
            reduce -= value[n]*buffer[sndex[n]];
        buffer[m-start] += reduce/diag[m];
    }
    //WRITE-BACK INTRA-BATCH DATA
    for(int m = 0; m < batchsize; m++)
        x[start+m] = buffer[m];
}
```

Other HPCG Operations

```
//GAUSS-SEIDEL BACKWARD SWEEP
for(int batch = numbatch-1; batch > -1; batch--){
    //READ MAPPING DATA
    int start = mapdispl[batch];
    int mapnz = mapdispl[batch+1]-start;
    //ALLOCATE BUFFER
    double buffer[batchsize+mapnz];
    //LOAD INTER-BATCH DATA
    for(int n = 0; n < mapnz; n++)
        buffer[batchsize+n] = x[map[start+n]];
    //LOAD INTRA-BATCH DATA
    start = batch*batchsize;
    for(int n = 0; n < batchsize; n++)
        buffer[n] = x[start+n];
    //PERFORM GAUSS-SEIDEL
    for(int m = start+batchsize-1; m > start-1; m--){
        double reduce = b[m];
        for(int n = displ[m+1]-1; n > displ[m]-1; n--)
            reduce -= value[n]*buffer[sndex[n]];
        buffer[m-start] += reduce/diag[m];
    }
    //WRITE-BACK INTRA-BATCH DATA
    for(int m = 0; m < batchsize; m++)
        x[start+m] = buffer[m];
}
```

Other HPCG Operations

```
//SPMV
for(int batch = 0; batch < numbatch; batch++){
    //READ MAPPING DATA
    int start = mapdispl[batch];
    int mapnz = mapdispl[batch+1]-start;
    //ALLOCATE BUFFER
    double buffer[batchsize+mapnz];
    //LOAD INTER-BATCH DATA
    for(int n = 0; n < mapnz; n++)
        buffer[batchsize+n] = x[map[start+n]];
    //LOAD INTRA-BATCH DATA
    start = batch*batchsize;
    for(int n = 0; n < batchsize; n++)
        buffer[n] = x[start+n];
    //PERFORM SPMV
    for(int m = start; m < start+batchsize; m++){
        double reduce = 0;
        for(int n = displ[m]; n < displ[m+1]; n++)
            reduce += value[n]*buffer[sndex[n]];
        y[m] = reduce;
    }
}
```

Computational Kernels in Each Iteration

Other HPCG Operations

Other HPCG Operations

```
//GAUSS-SEIDEL FORWARD SWEEP
for(int batch = 0; batch < numbatch; batch++){
    //READ MAPPING DATA
    int start = mapdispl[batch];
    int mapnz = mapdispl[batch+1]-start;
    //ALLOCATE BUFFER
    double buffer[batchsize+mapnz];
    //LOAD INTER-BATCH DATA
    for(int n = 0; n < mapnz; n++)
        buffer[batchsize+n] = x[map[start+n]];
    //LOAD INTRA-BATCH DATA
    start = batch*batchsize;
    for(int n = 0; n < batchsize; n++)
        buffer[n] = x[start+n];
    //PERFORM GAUSS-SEIDEL PER BATCH
    for(int m = start; m < start+batchsize; m++){
        double reduce = b[m];
        for(int n = displ[m]; n < displ[m+1]; n++)
            reduce -= value[n]*buffer[sndex[n]];
        buffer[m-start] += reduce/diag[m];
    }
    //WRITE-BACK INTRA-BATCH DATA
    for(int m = 0; m < batchsize; m++)
        x[start+m] = buffer[m];
}
```

```
//GAUSS-SEIDEL BACKWARD SWEEP
for(int batch = numbatch-1; batch > -1; batch--){
    //READ MAPPING DATA
    int start = mapdispl[batch];
    int mapnz = mapdispl[batch+1]-start;
    //ALLOCATE BUFFER
    double buffer[batchsize+mapnz];
    //LOAD INTER-BATCH DATA
    for(int n = 0; n < mapnz; n++)
        buffer[batchsize+n] = x[map[start+n]];
    //LOAD INTRA-BATCH DATA
    start = batch*batchsize;
    for(int n = 0; n < batchsize; n++)
        buffer[n] = x[start+n];
    //PERFORM GAUSS-SEIDEL
    for(int m = start+batchsize-1; m > start-1; m--){
        double reduce = b[m];
        for(int n = displ[m+1]-1; n > displ[m]-1; n--)
            reduce -= value[n]*buffer[sndex[n]];
        buffer[m-start] += reduce/diag[m];
    }
    //WRITE-BACK INTRA-BATCH DATA
    for(int m = 0; m < batchsize; m++)
        x[start+m] = buffer[m];
}
```

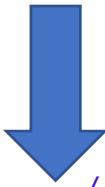


IDEA: G-S can be fused with SpMV to reuse matrix data from cache

```
//SPMV
for(int batch = 0; batch < numbatch; batch++){
    //READ MAPPING DATA
    int start = mapdispl[batch];
    int mapnz = mapdispl[batch+1]-start;
    //ALLOCATE BUFFER
    double buffer[batchsize+mapnz];
    //LOAD INTER-BATCH DATA
    for(int n = 0; n < mapnz; n++)
        buffer[batchsize+n] = x[map[start+n]];
    //LOAD INTRA-BATCH DATA
    start = batch*batchsize;
    for(int n = 0; n < batchsize; n++)
        buffer[n] = x[start+n];
    //PERFORM SPMV
    for(int m = start; m < start+batchsize; m++){
        double reduce = 0;
        for(int n = displ[m]; n < displ[m+1]; n++)
            reduce += value[n]*buffer[sndex[n]];
        y[m] = reduce;
    }
}
```

Computational Kernels in Each Iteration

Other HPCG Operations



```
//GAUSS-SEIDEL FORWARD SWEEP
for(int batch = 0; batch < numbatch; batch++){
    //READ MAPPING DATA
    int start = mapdispl[batch];
    int mapnz = mapdispl[batch+1]-start;
    //ALLOCATE BUFFER
    double buffer[batchsize+mapnz];
    //LOAD INTER-BATCH DATA
    for(int n = 0; n < mapnz; n++)
        buffer[batchsize+n] = x[map[start+n]];
    //LOAD INTRA-BATCH DATA
    start = batch*batchsize;
    for(int n = 0; n < batchsize; n++)
        buffer[n] = x[start+n];
    //PERFORM GAUSS-SEIDEL PER BATCH
    for(int m = start; m < start+batchsize; m++){
        double reduce = b[m];
        for(int n = displ[m]; n < displ[m+1]; n++)
            reduce -= value[n]*buffer[sndex[n]];
        buffer[m-start] += reduce/diag[m];
    }
    //WRITE-BACK INTRA-BATCH DATA
    for(int m = 0; m < batchsize; m++)
        x[start+m] = buffer[m];
}
```

Other HPCG Operations



```
//FUSED GAUSS-SEIDEL BACKWARD SWEEP & SPMV
for(int batch = numbatch-1; batch > -1; batch--){
    //READ MAPPING DATA
    int start = mapdispl[batch];
    int mapnz = mapdispl[batch+1]-start;
    //ALLOCATE BUFFER
    double buffer[batchsize+mapnz];
    //LOAD INTER-BATCH DATA
    for(int n = 0; n < mapnz; n++)
        buffer[batchsize+n] = x[map[start+n]];
    //LOAD INTRA-BATCH DATA
    start = batch*batchsize;
    for(int n = 0; n < batchsize; n++)
        buffer[n] = x[start+n];
    //PERFORM GAUSS-SEIDEL
    for(int m = start+batchsize-1; m > start-1; m--){
        double reduce = b[m];
        for(int n = displ[m+1]-1; n > displ[m]-1; n--)
            reduce -= value[n]*buffer[sndex[n]];
        buffer[m-start] += reduce/diag[m];
    }
    //PERFORM SPMV
    for(int mn = spmvdispl[batch]; mn < spmvdispl[batch+1]; mn++){
        int m = spmvmap[mn];
        double reduce = 0;
        for(int n = displ[m]; n < displ[m+1]; n++)
            reduce += value[n]*buffer[sndex[n]];
        y[m] = reduce;
    }
    //WRITE-BACK INTRA-BATCH DATA
    for(int m = 0; m < batchsize; m++)
        x[start+m] = buffer[m];
}
```

Other HPCG Operations



To-Do List:

- Validate correctness by solving a 3-D Laplace equation
- Investigate convergence vs. speedup trade-off
- Incorporate kernels into HPCG and validate again
- Scale on Summit ???