

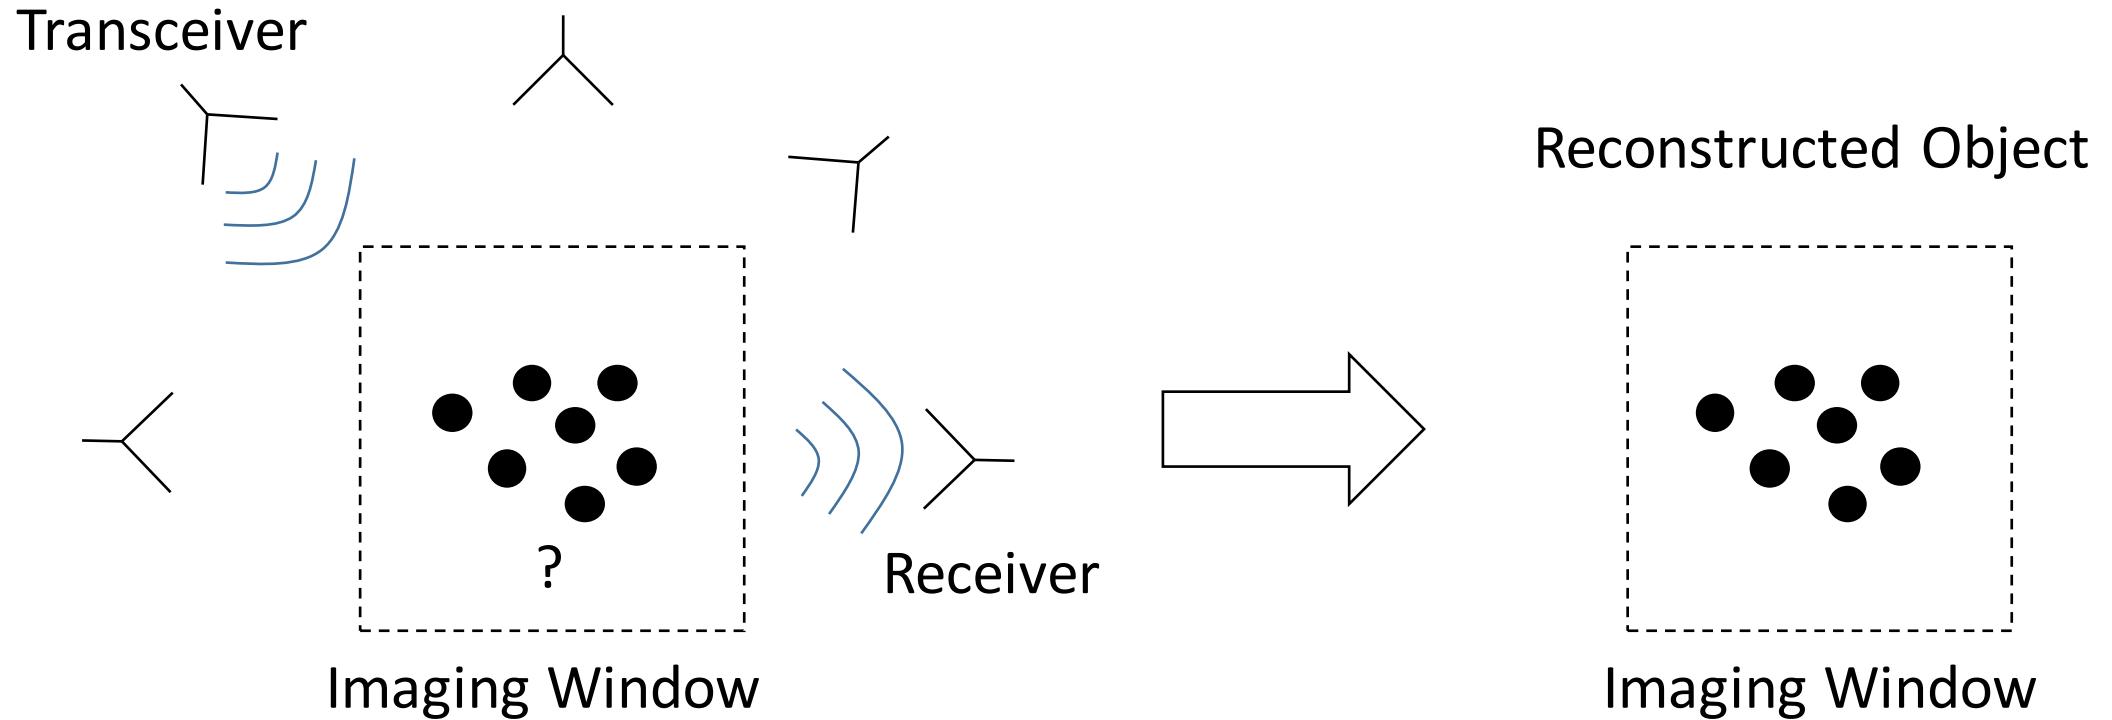
# Large-Scale Inverse Multiple-Scattering Imaging on GPU Supercomputers with Real Data

Mert Hidayetoğlu  
Stanford University

In courtesy of Weng Cho Chew, Michael Oelze, and Wen-mei Hwu  
University of Illinois at Urbana-Champaign

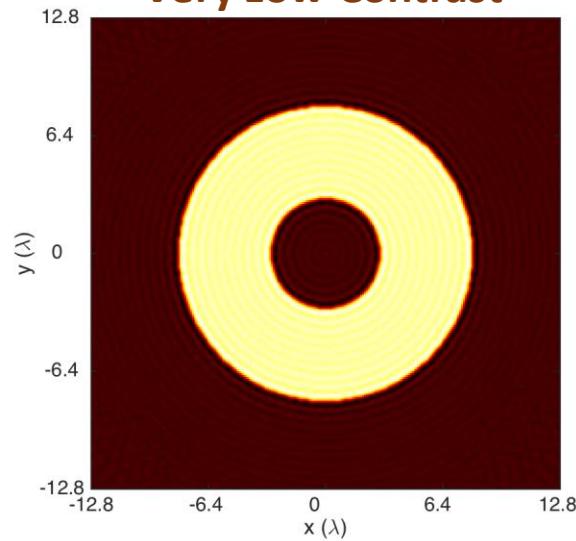
[merth@stanford.edu](mailto:merth@stanford.edu)

# Imaging with Inverse-Scattering

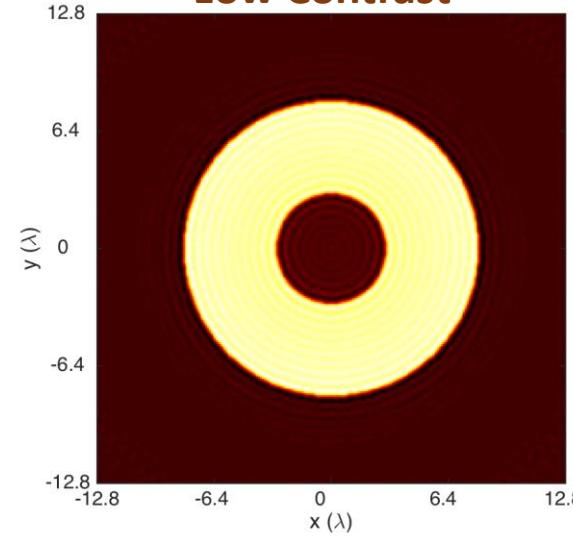


# DT Breaks Down with Multiple Scattering

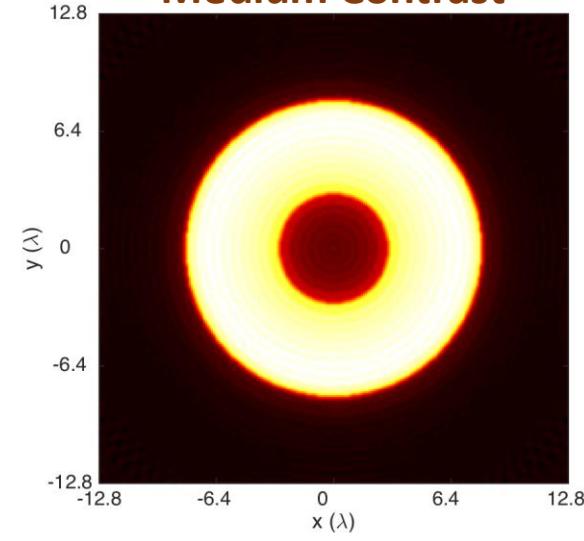
Very Low Contrast



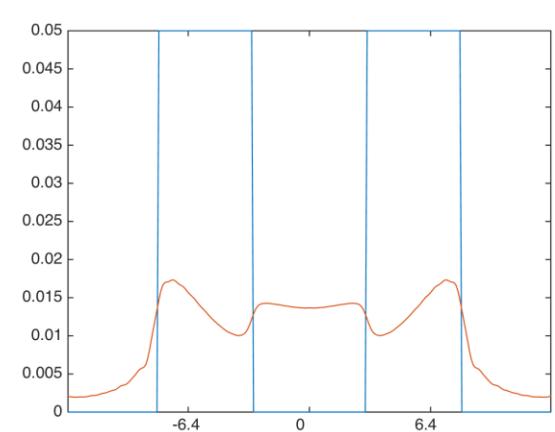
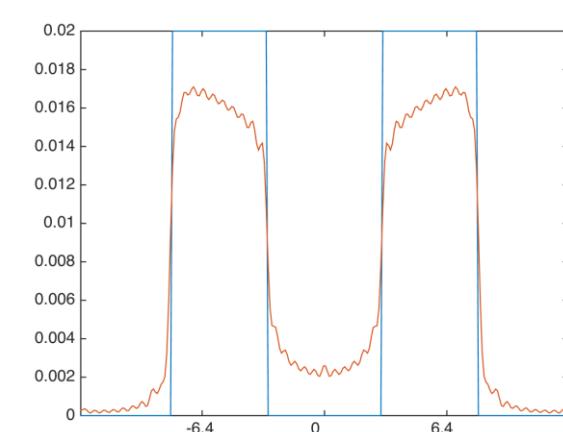
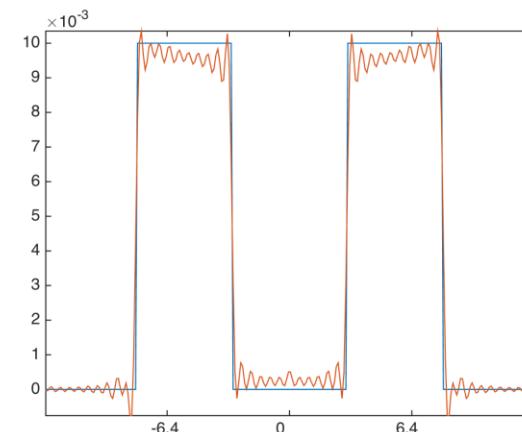
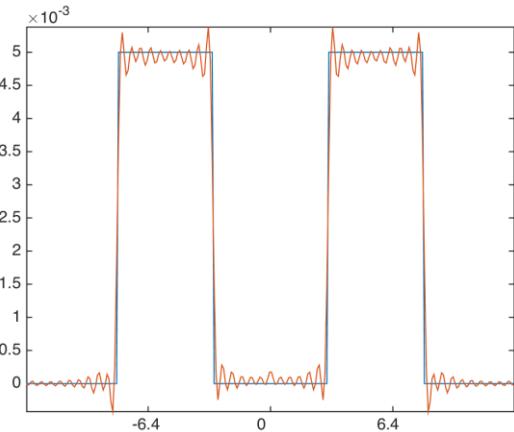
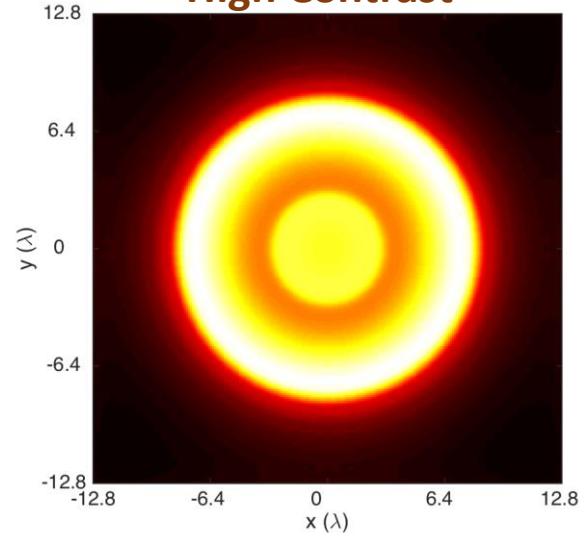
Low Contrast



Medium Contrast



High Contrast

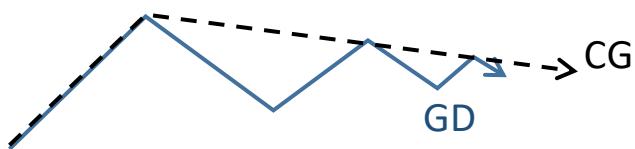


# Gradient Search for Nonlinear Optimization

- Gradient-Descent

$$-\nabla \Phi = -\bar{F}^H r$$

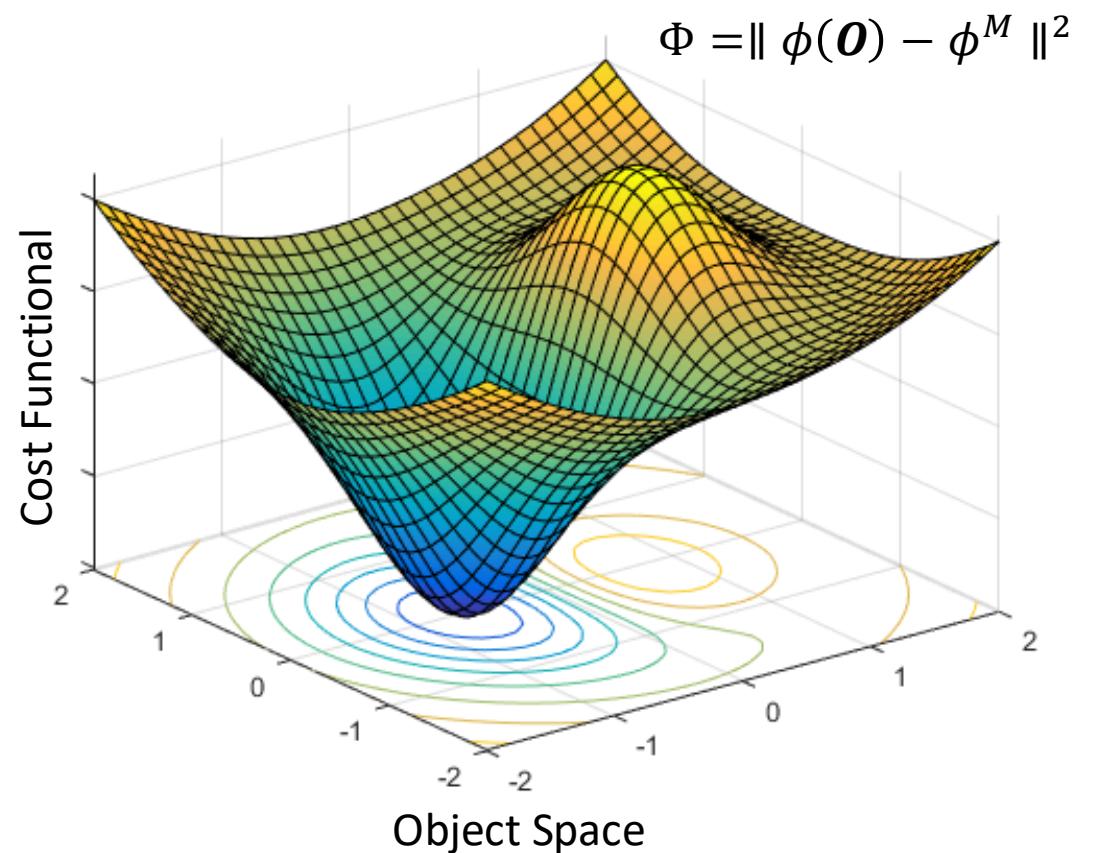
- Conjugate-Gradient



- Newton-Type Methods

$$[\mu^2 \bar{I} + \bar{F}^H \bar{F}] \delta \theta = \nabla \Phi$$

$\bar{F}$ : Functional Derivative Operator



# Finding Derivative with Distorted-Born Approximation

Scattering Equation:

$$\mathcal{G} = \mathcal{G}_0 + \mathcal{G}_0 \mathcal{O} \mathcal{G}$$

Object Perturbation:

$$\mathcal{O} = \mathcal{O}_b + \delta\mathcal{O}$$

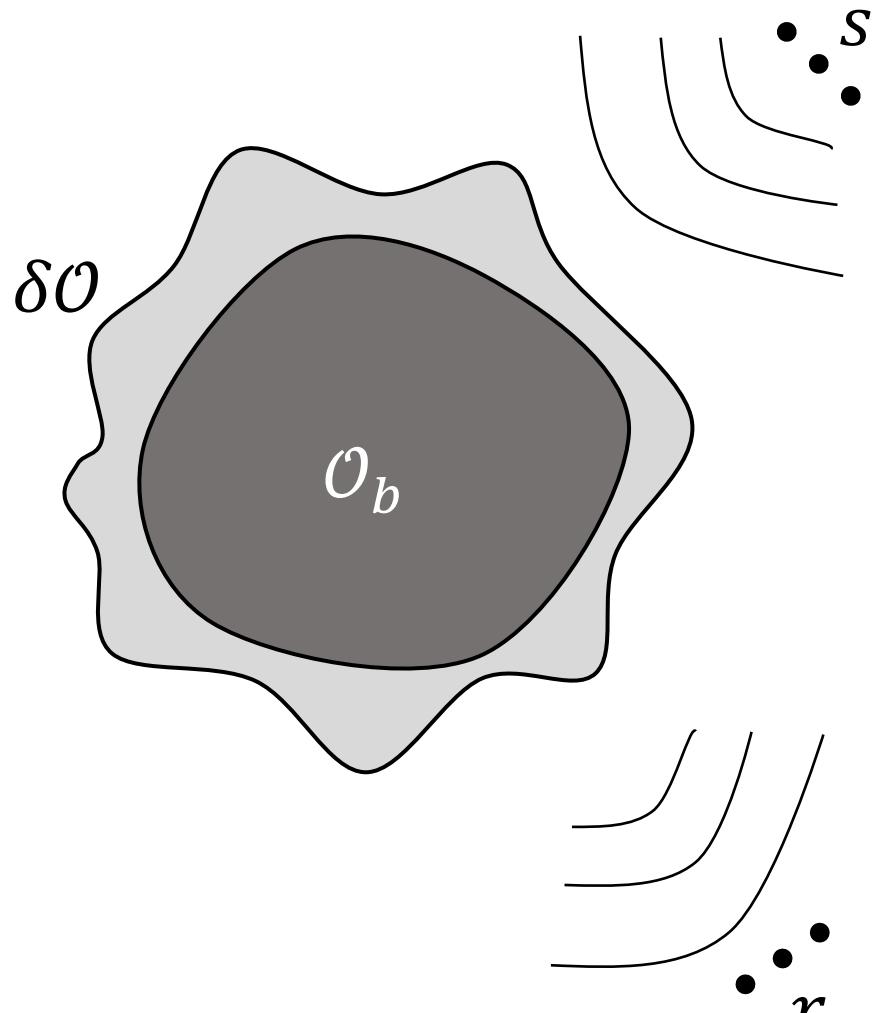
$$\delta\mathcal{G} = \mathcal{G}_b \delta\mathcal{O} \mathcal{G}_b + \underbrace{\mathcal{G}_b \delta\mathcal{O} \delta\mathcal{G}}_{\text{Higher-order var.}}$$

$$\delta\mathcal{G} \approx \delta\mathcal{G}^{(1)} = \mathcal{G}_b \delta\mathcal{O} \mathcal{G}_b \quad \xleftarrow{\text{Distorted Born Approximation}}$$

When operates on a source  $s$ :

$$\delta\phi \approx \mathcal{G}_b \delta\mathcal{O} \phi_b = \mathcal{G}_b \phi_b \delta\mathcal{O}$$

$$\delta\phi^{(1)} = \mathcal{F} \delta\mathcal{O}$$



## Distorted-Born Approximation Requires two Scattering Solutions

$$-\nabla\Phi = -\mathcal{F}^\dagger r$$

MoM

$\mathcal{F} = \mathcal{G}_b \phi_b$

$\overline{\mathcal{F}} = \overline{\mathcal{G}}_R \left\{ \overline{\mathbf{I}} + \overline{\mathbf{O}}_b \underbrace{\left[ \overline{\mathbf{I}} - \overline{\mathcal{G}}_0 \overline{\mathbf{O}}_b \right]^{-1} \overline{\mathcal{G}}_0}_{\text{MLFMA}} \right\} \text{diag} \left\{ \underbrace{\left[ \overline{\mathbf{I}} - \overline{\mathcal{G}}_0 \overline{\mathbf{O}}_b \right]^{-1} \overline{\mathcal{G}}_T s}_{\text{MLFMA}} \right\}$

$\overline{\mathcal{G}}_0$ : Dense,  $N \times N$

$\overline{\mathcal{G}}_T$ : Dense,  $N \times T$

$\overline{\mathcal{G}}_R$ : Dense,  $R \times N$

1 Million Pixels ( $100\lambda$ ): 16TB  
16 Million Pixels ( $400\lambda$ ): 4PB (!)

$\overline{\mathbf{O}}_b$ : Diagonal,  $N \times N$

$\overline{\mathbf{I}}$ : Diagonal,  $N \times N$

# Algorithmic Speedup with MLFMA

Direct Methods:  $\mathcal{O}(N^3)$

Iterative Methods:  $\mathcal{O}(N^2)$

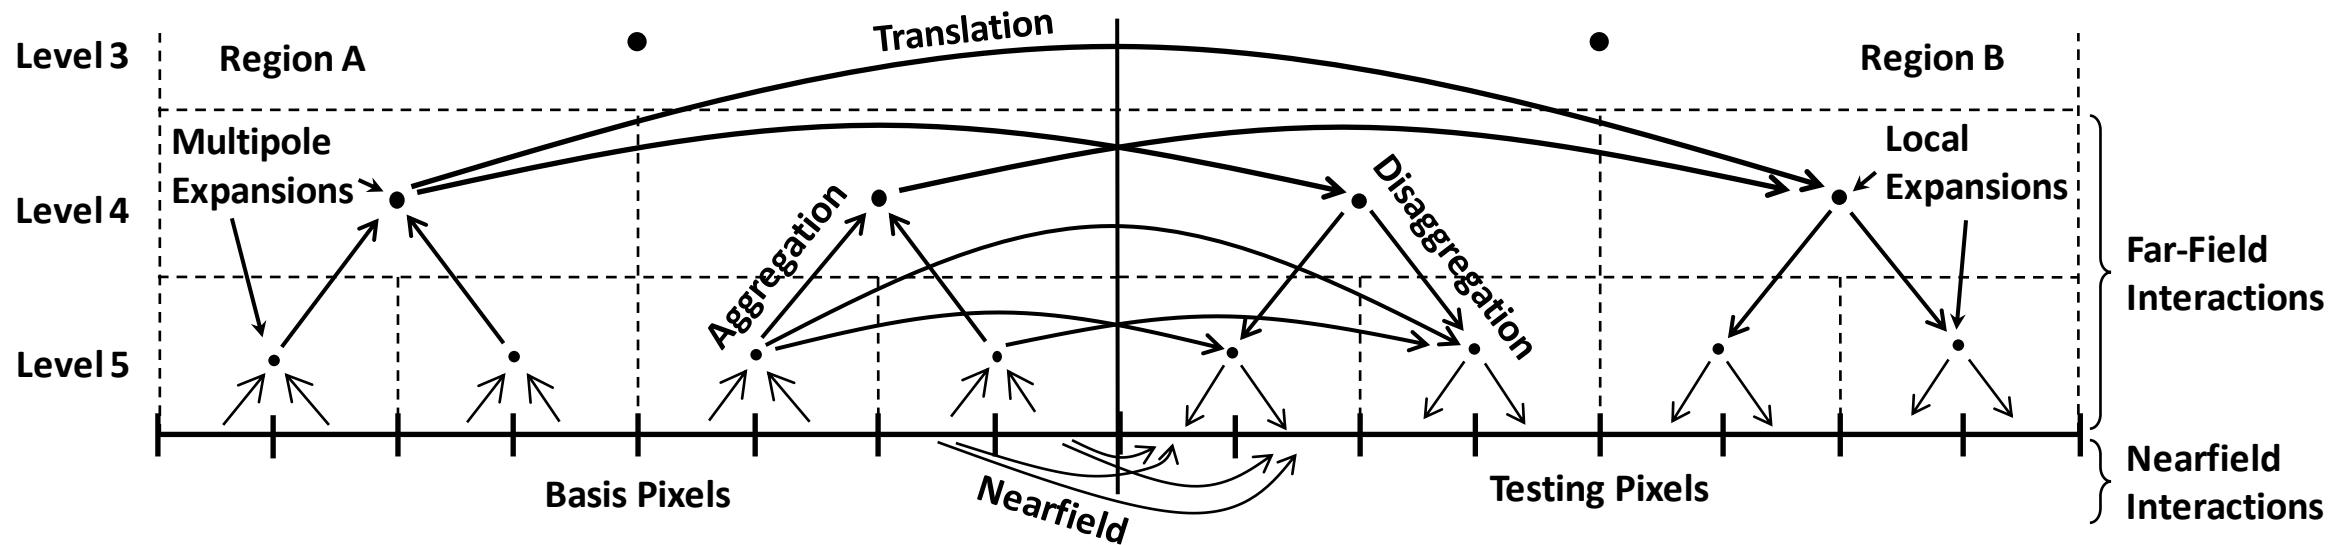
Fast Multipole Method:  $\mathcal{O}(N^{1.4}) - \mathcal{O}(N^{1.5})$

Multilevel Fast Multipole Algorithm:  $\mathcal{O}(N) - \mathcal{O}(N \log N)$

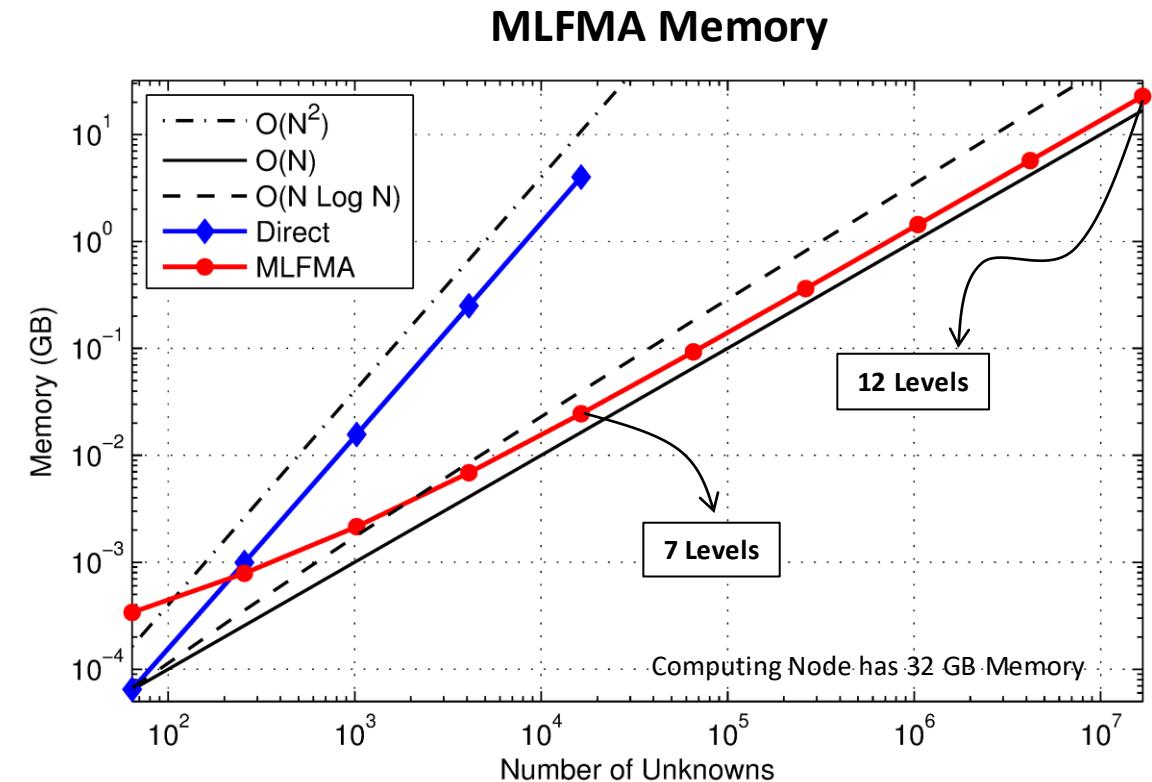
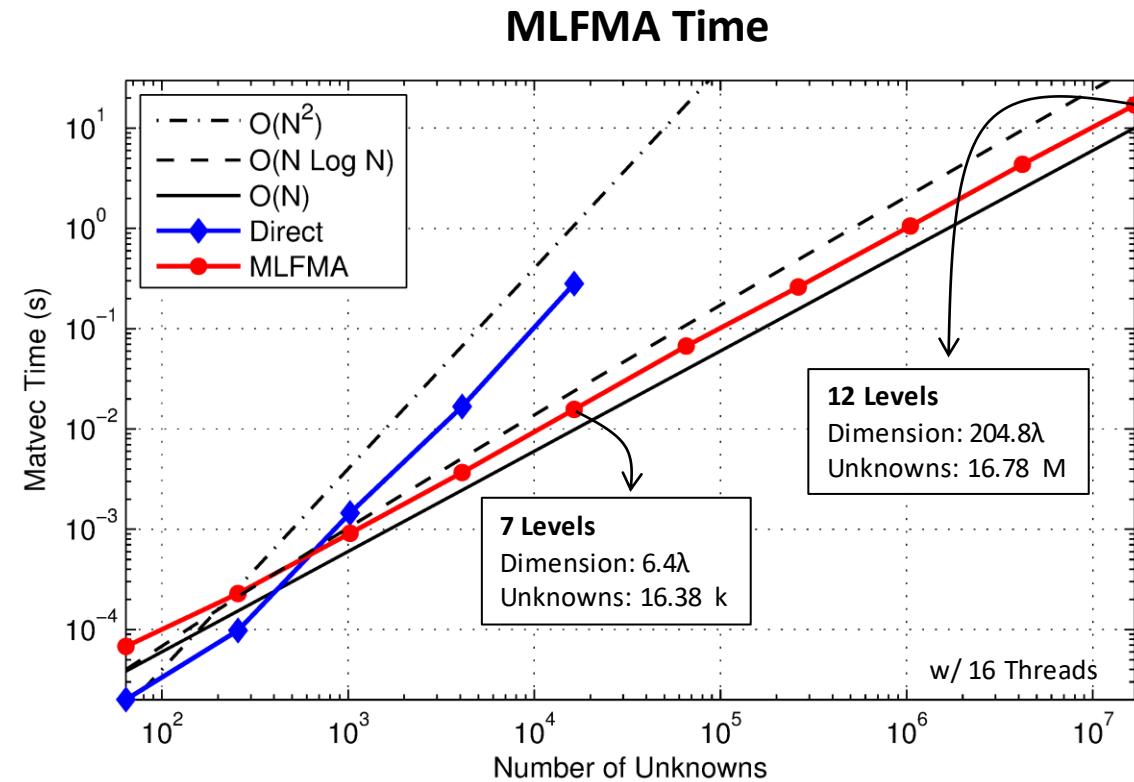
$$\phi = [\bar{I} - \bar{G}_0 \bar{O}]^{-1} \phi_0$$

↓  
4 PB

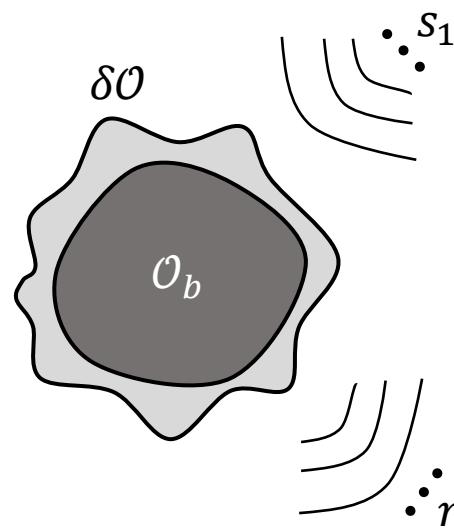
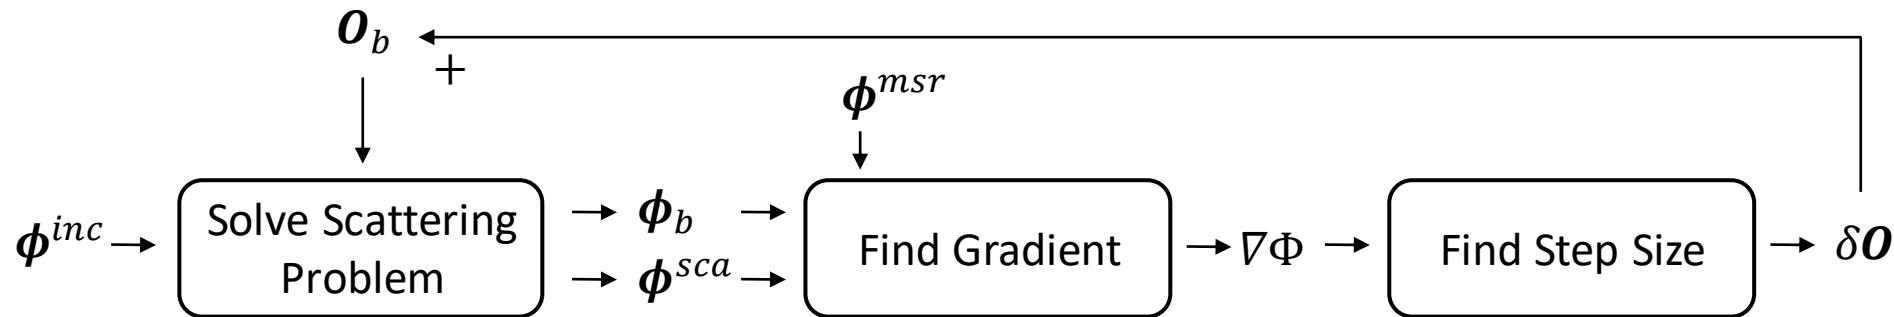
## MLFMA Schematic



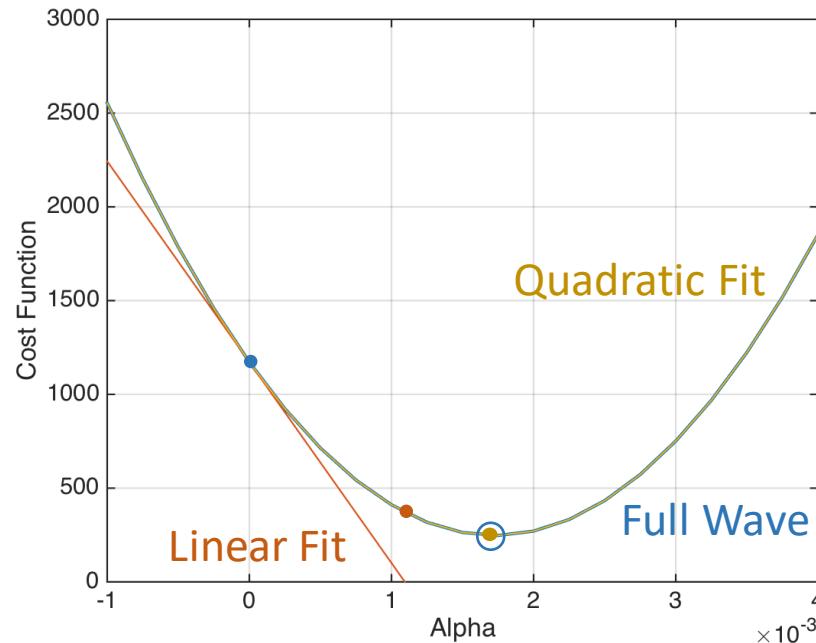
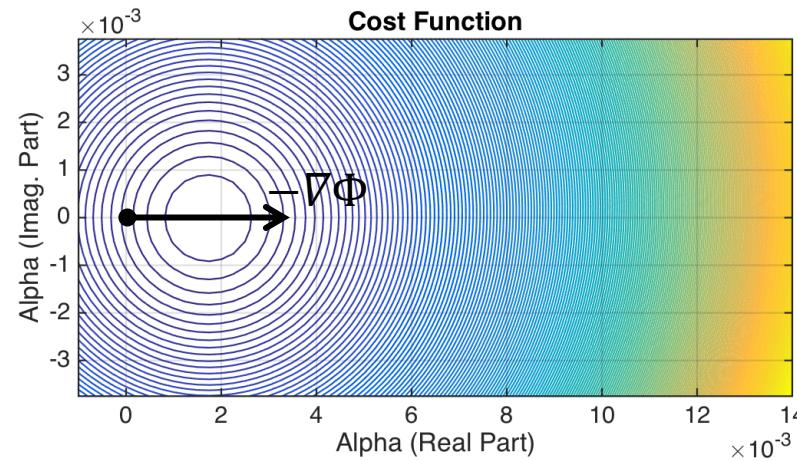
# Algorithmic Speedup with MLFMA



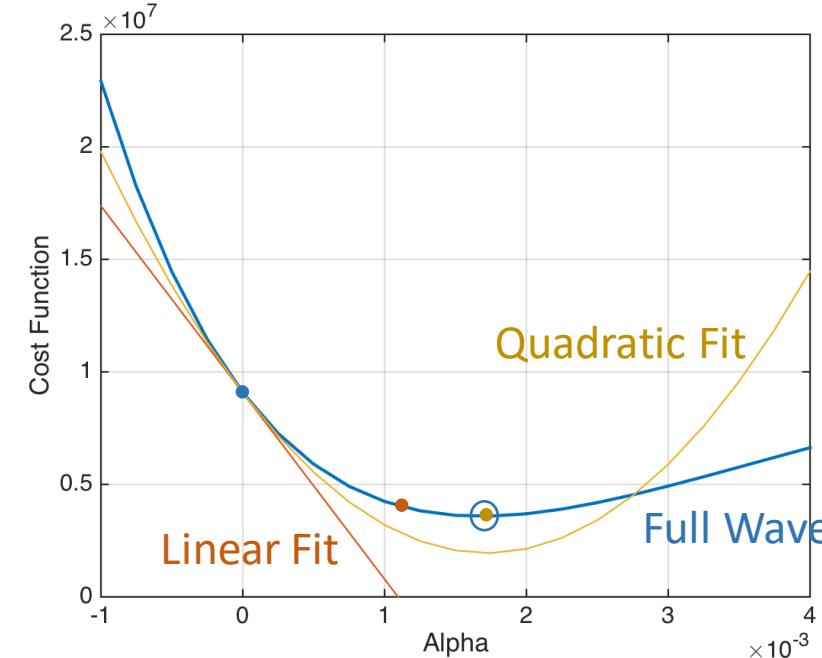
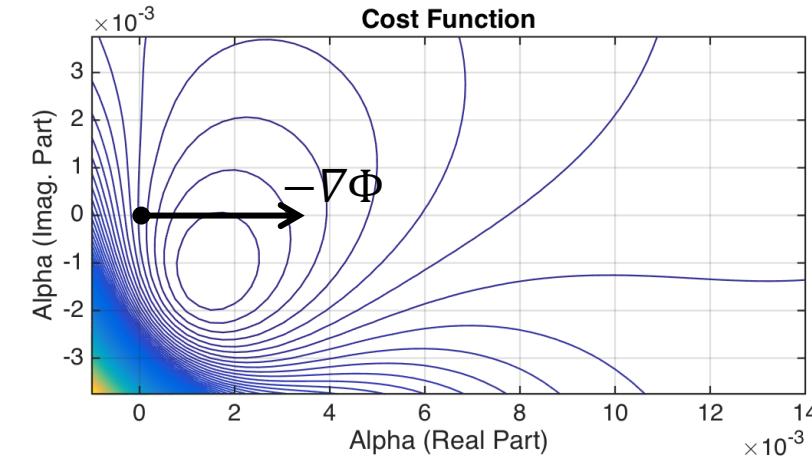
## Three MLFMA Solutions per Iteration



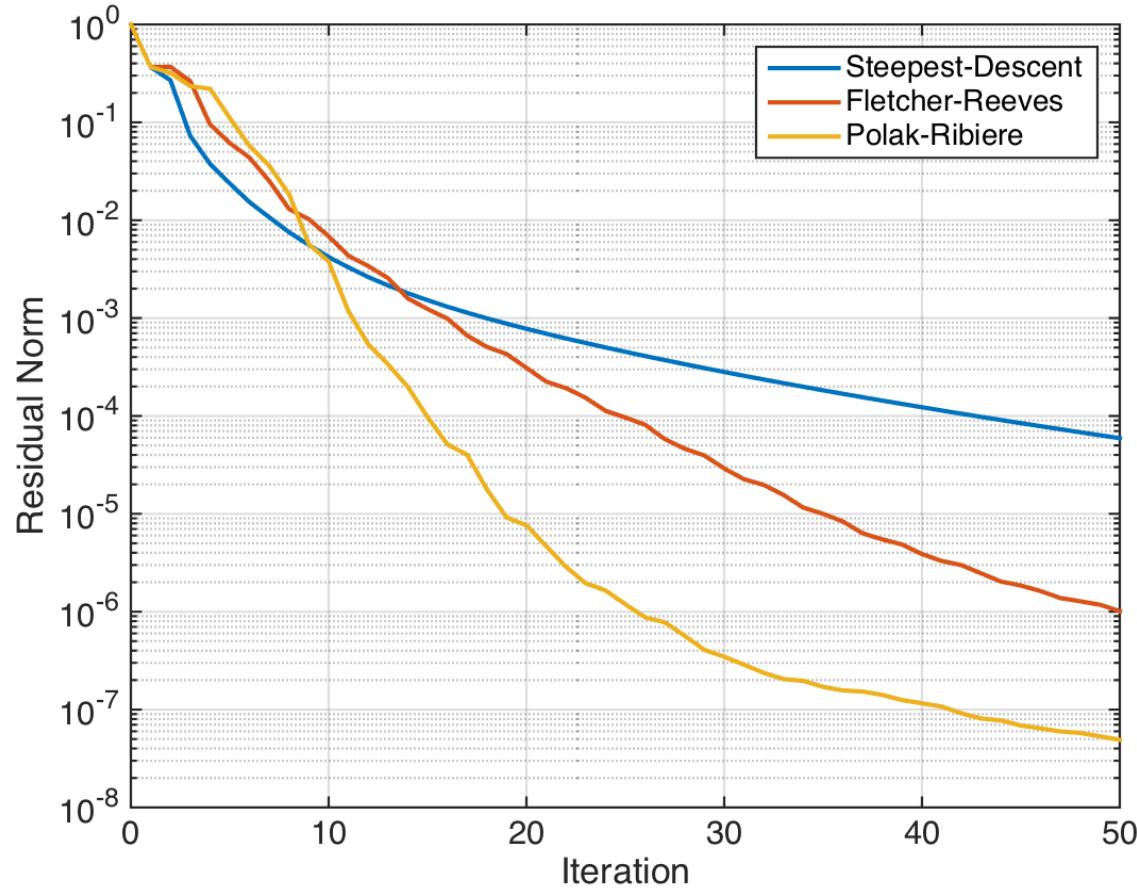
## Weak Scatterer



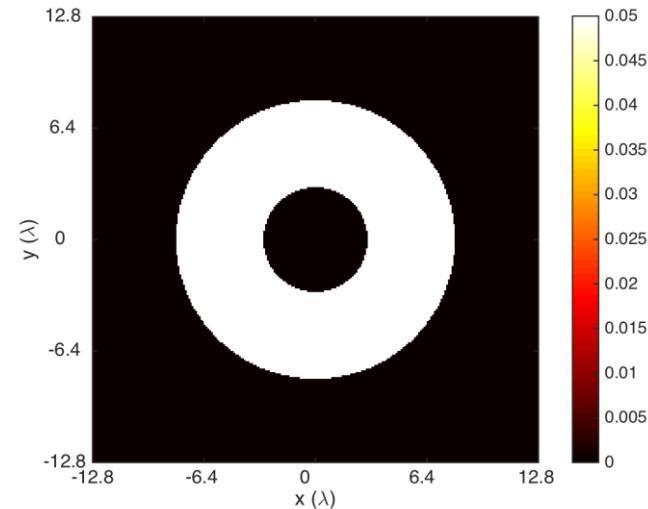
## Medium Scatterer



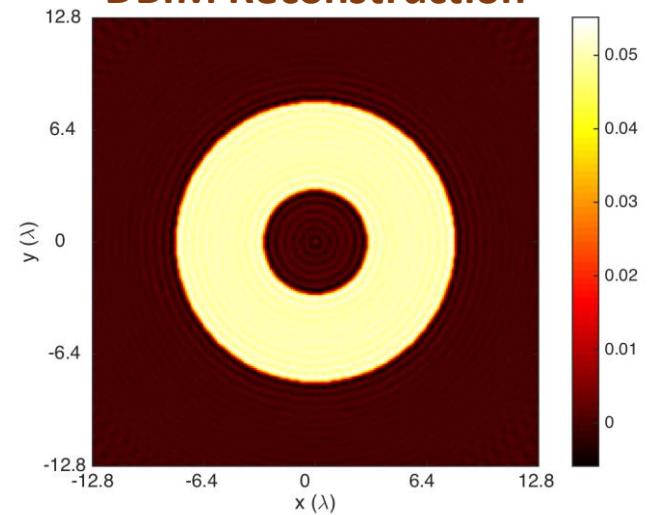
# Conjugate-Gradient Convergence



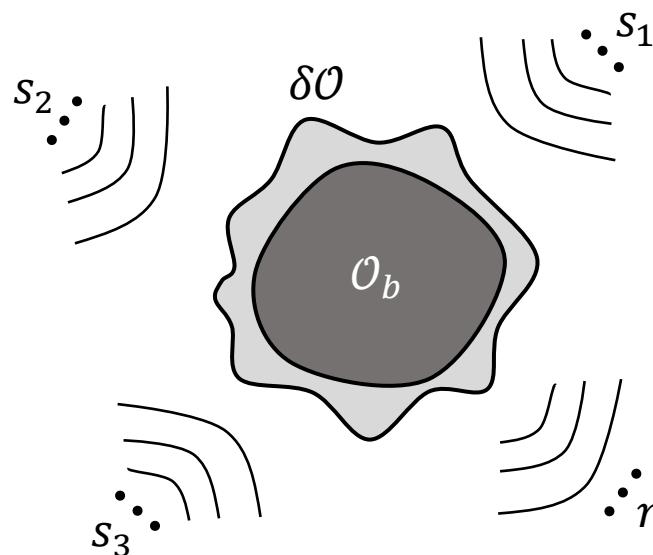
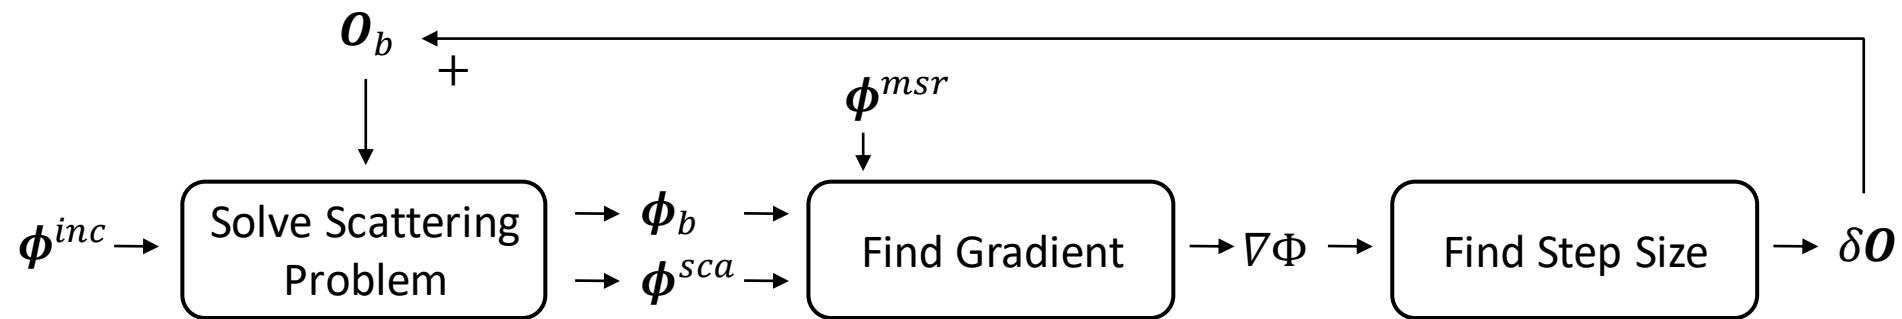
Numerical Phantom



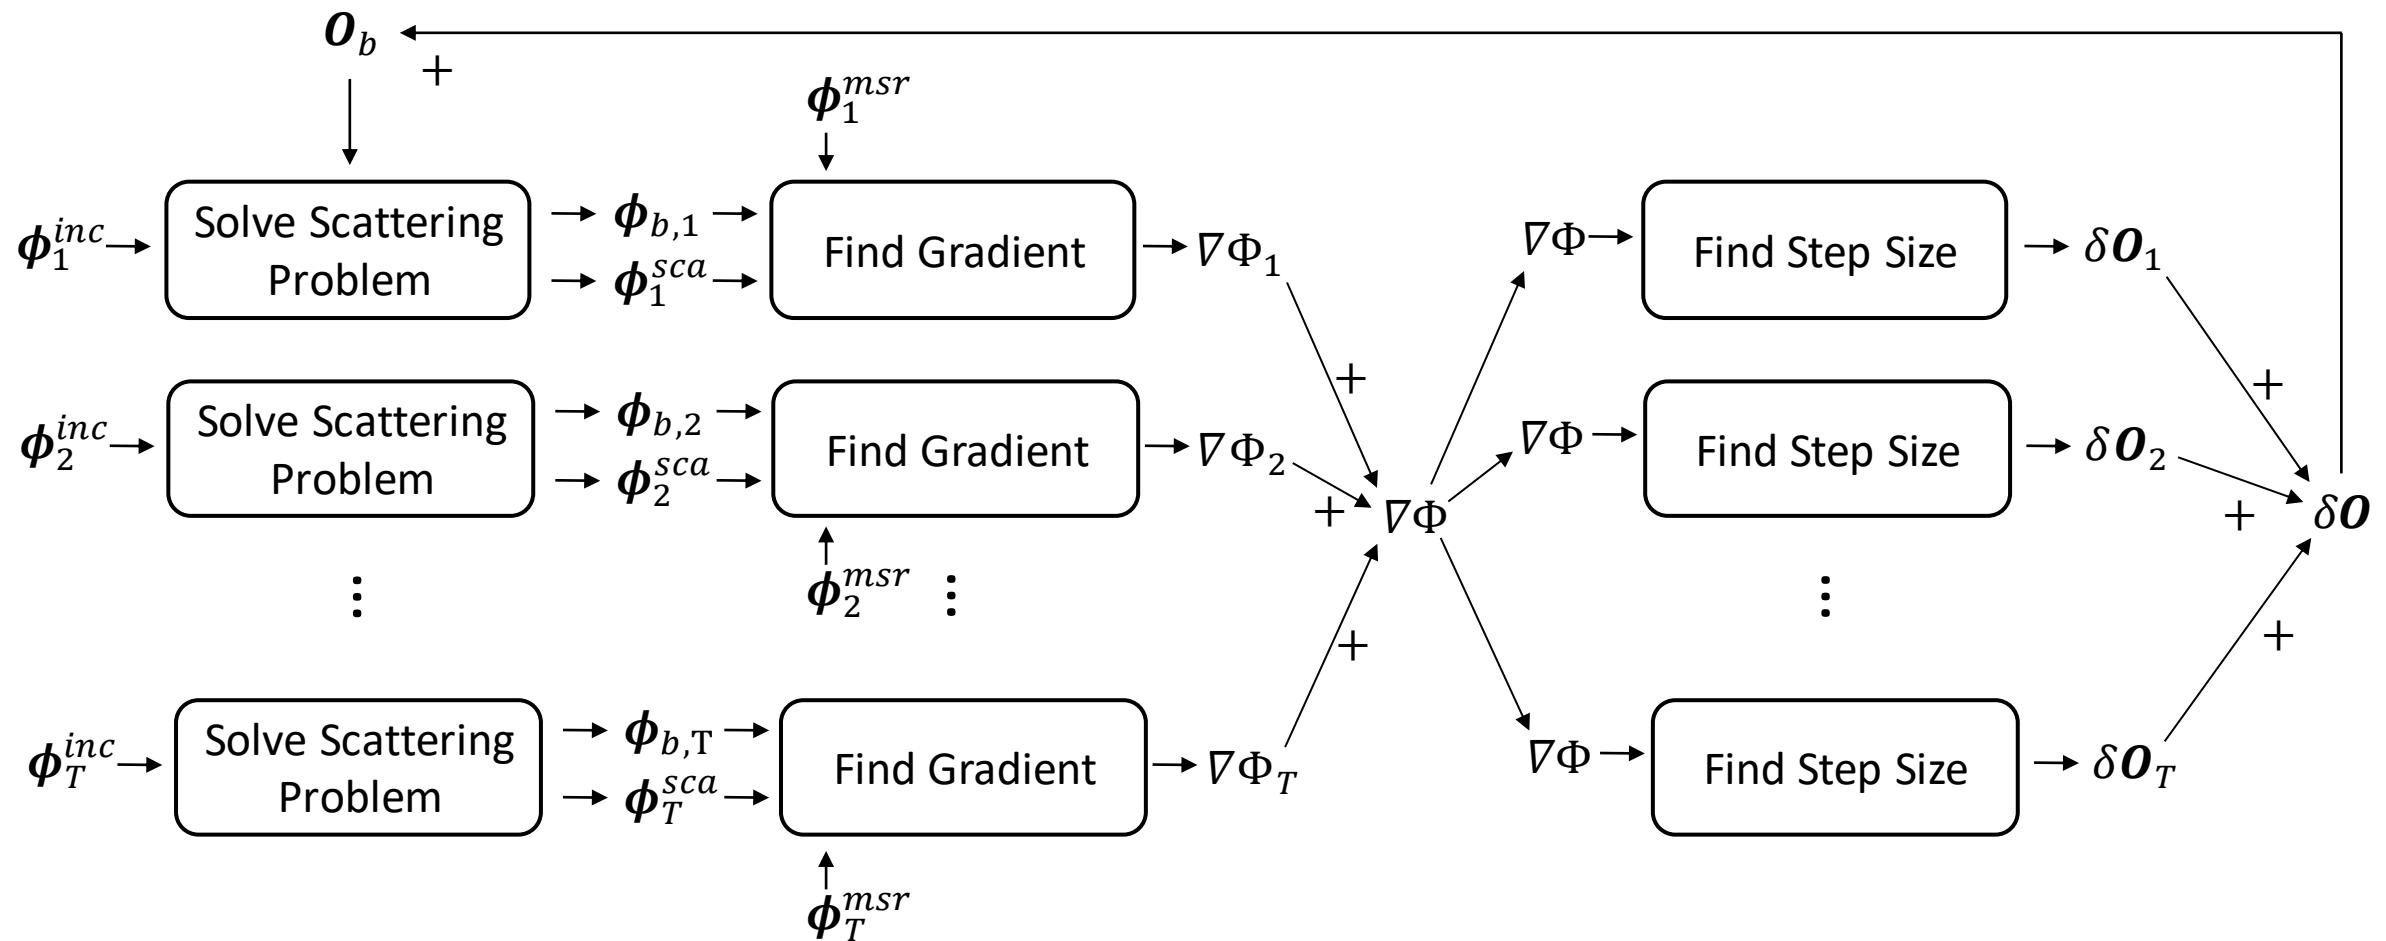
DBIM Reconstruction



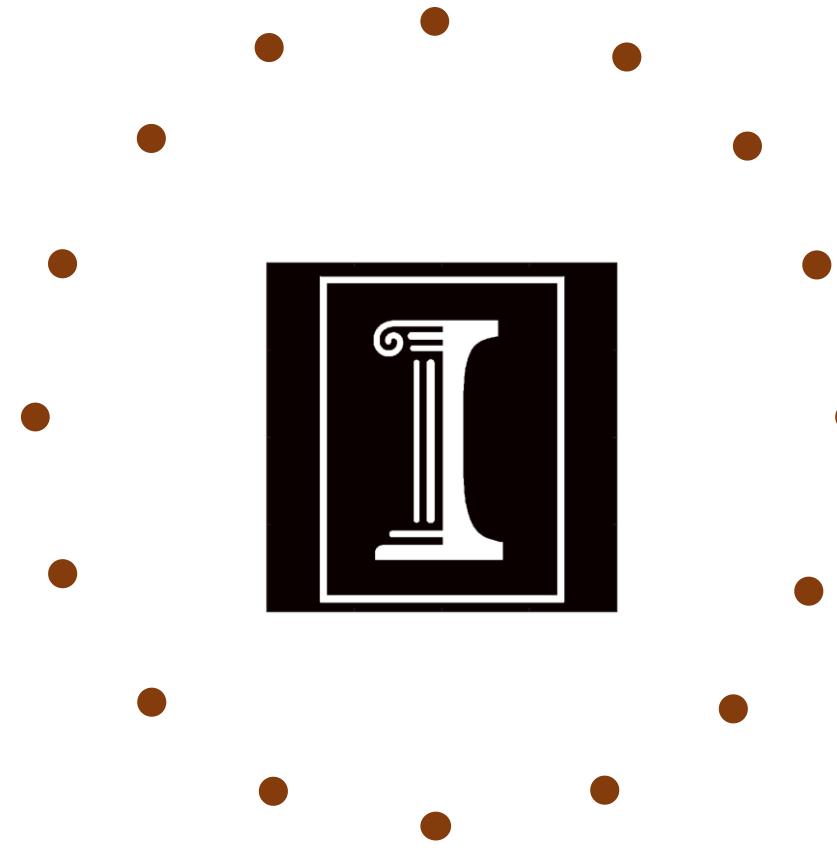
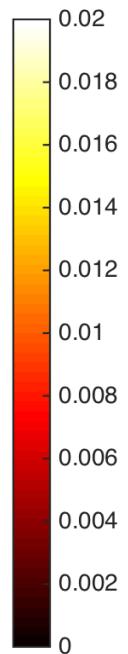
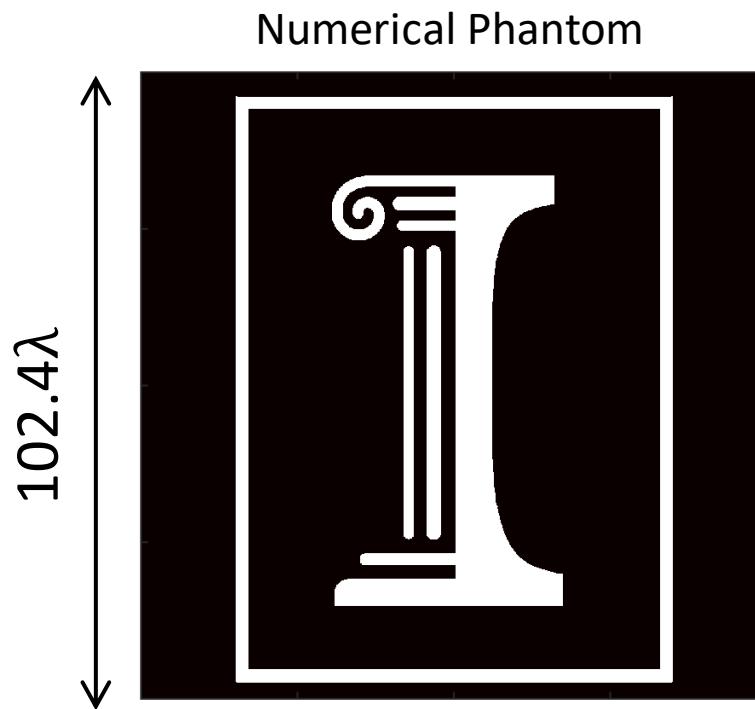
# Multiple Illuminations Provide Diversity



## Many MLFMA Solutions per Iteration

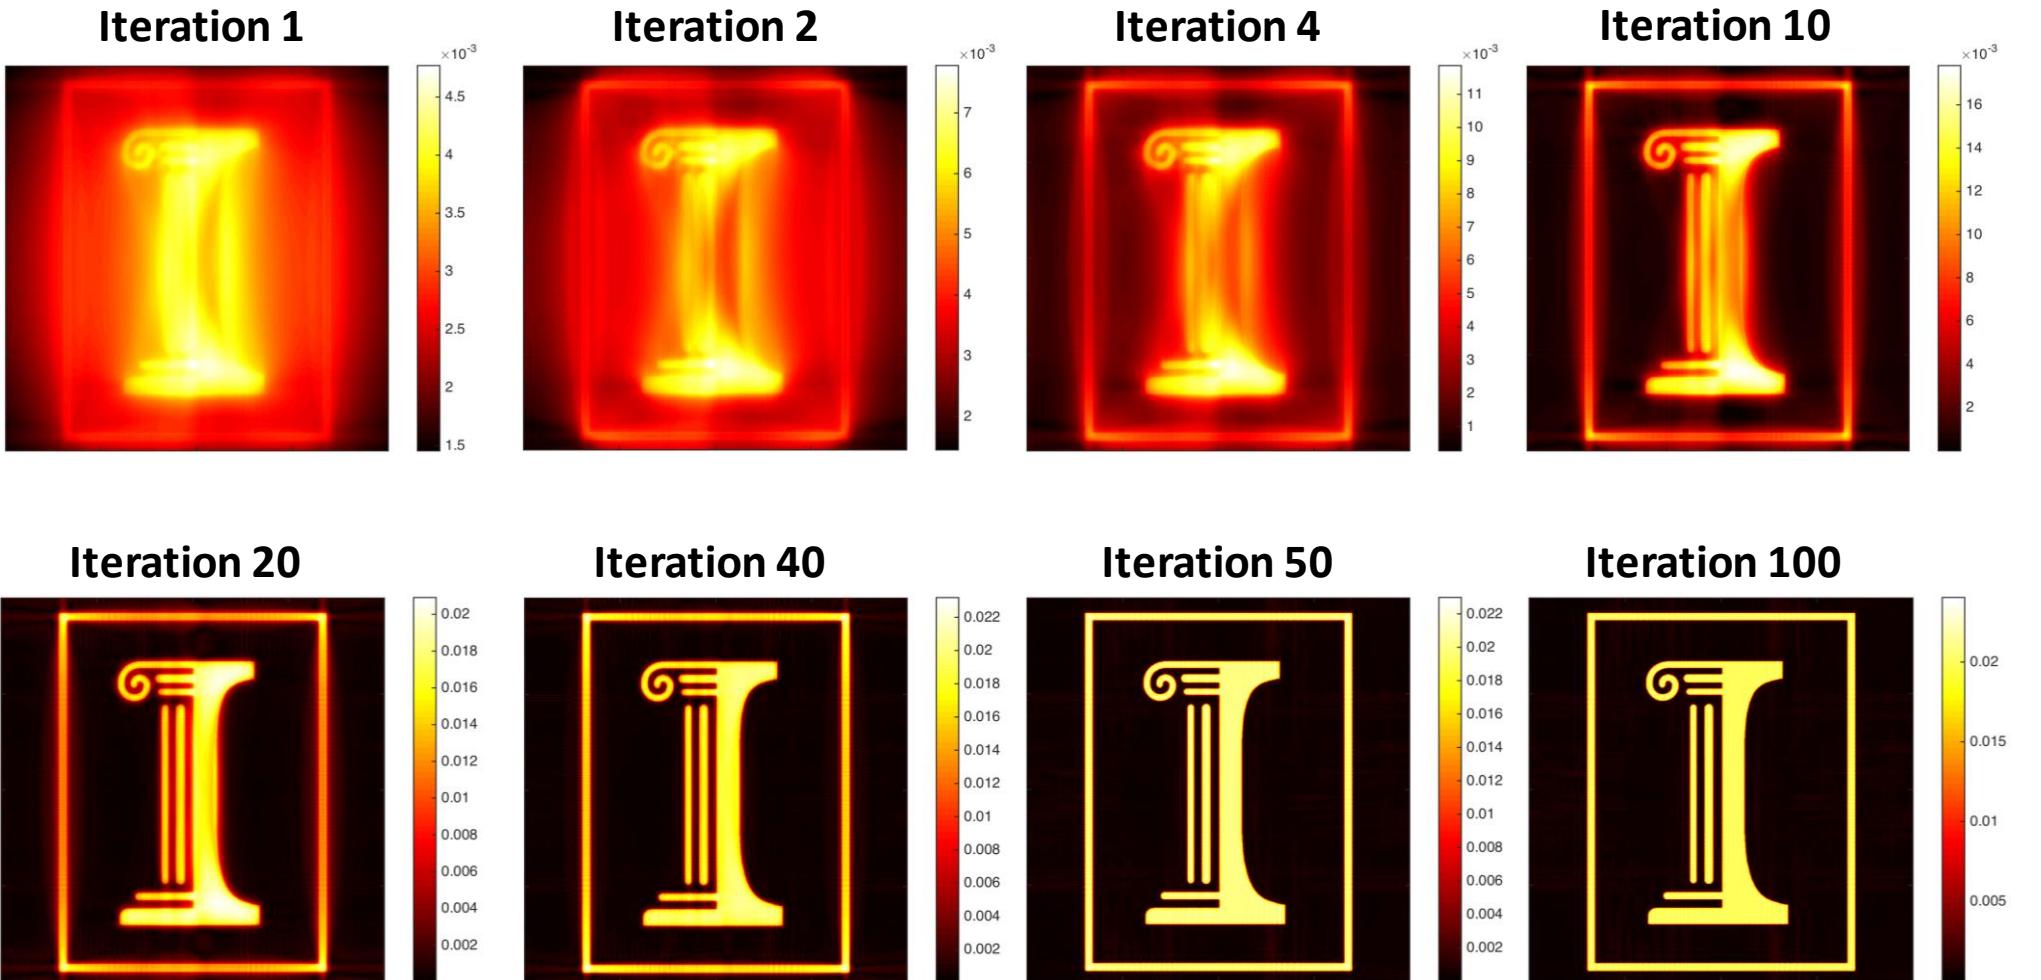


# Numerical Setup for Monochromatic Reconstructions

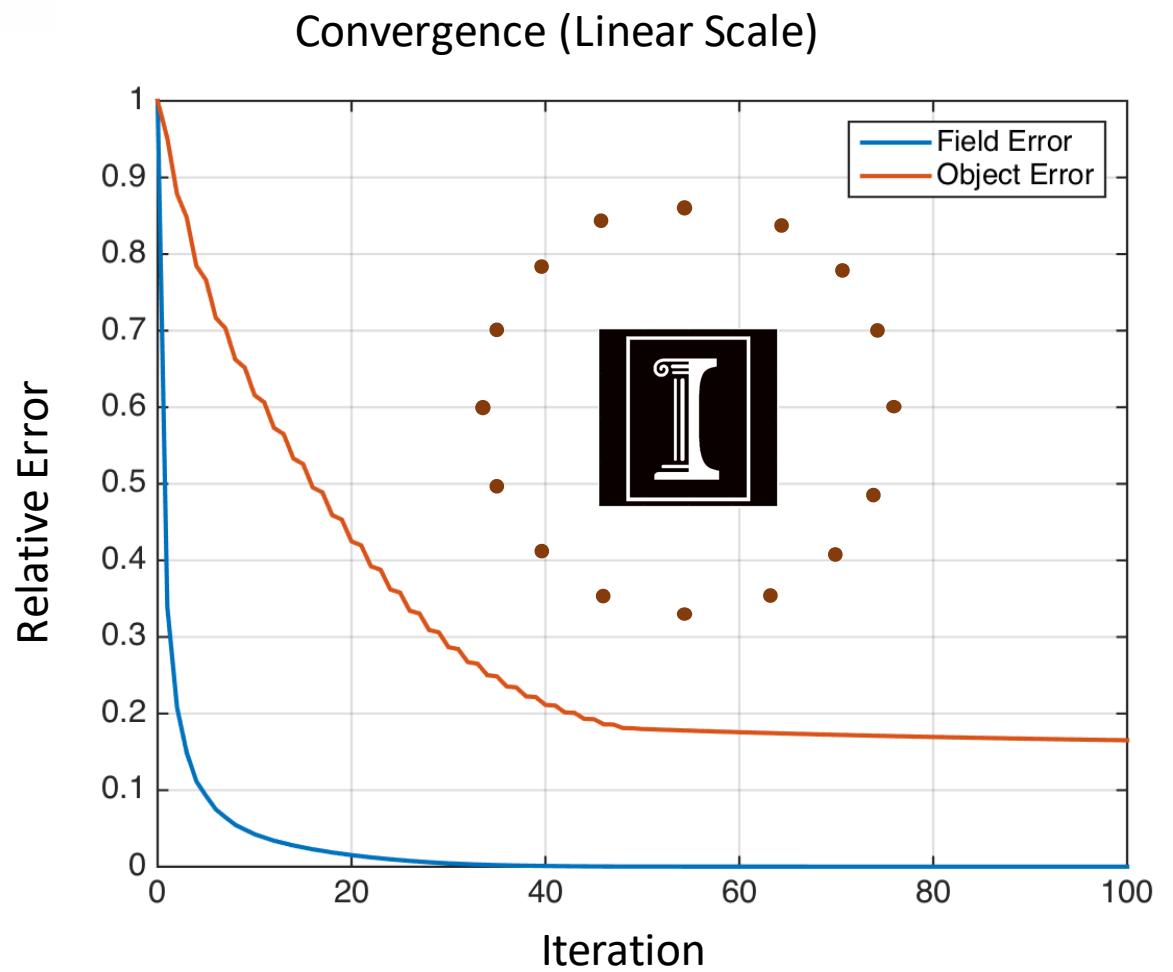
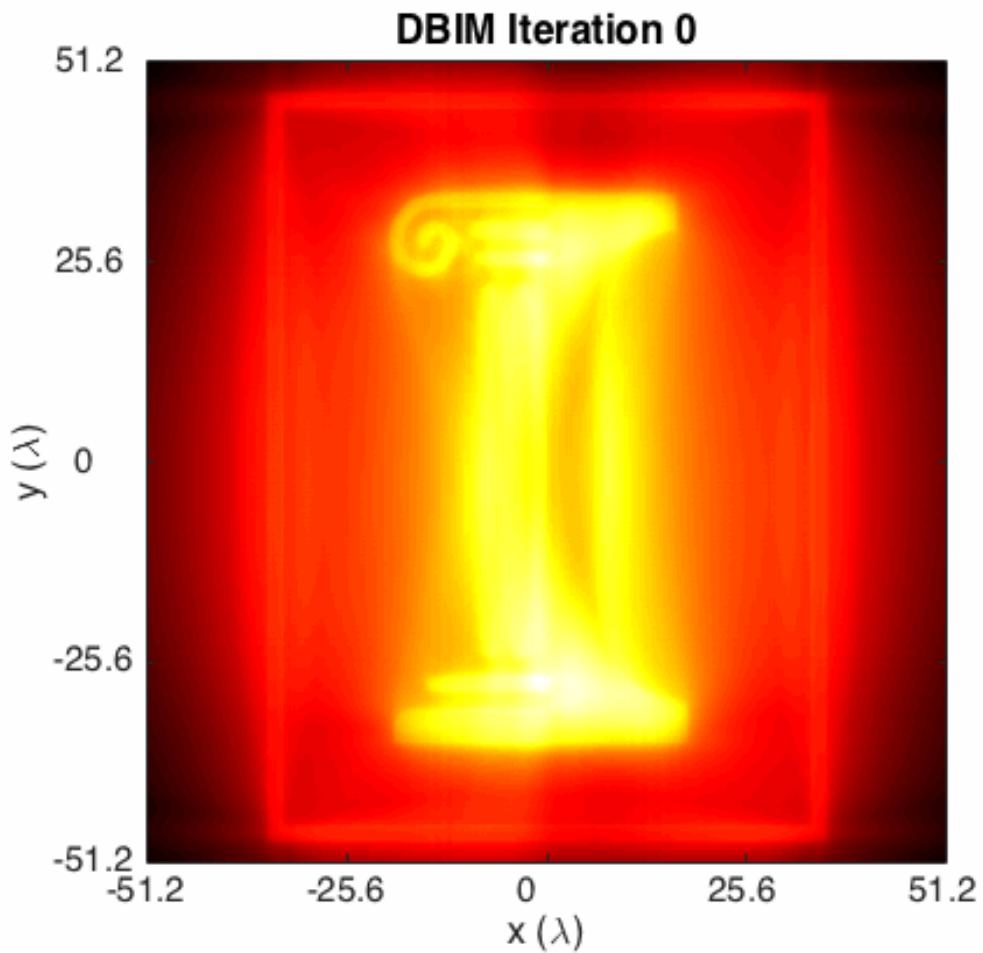


- 256 Transmitters
- 1,024 Receivers
- Monochromatic

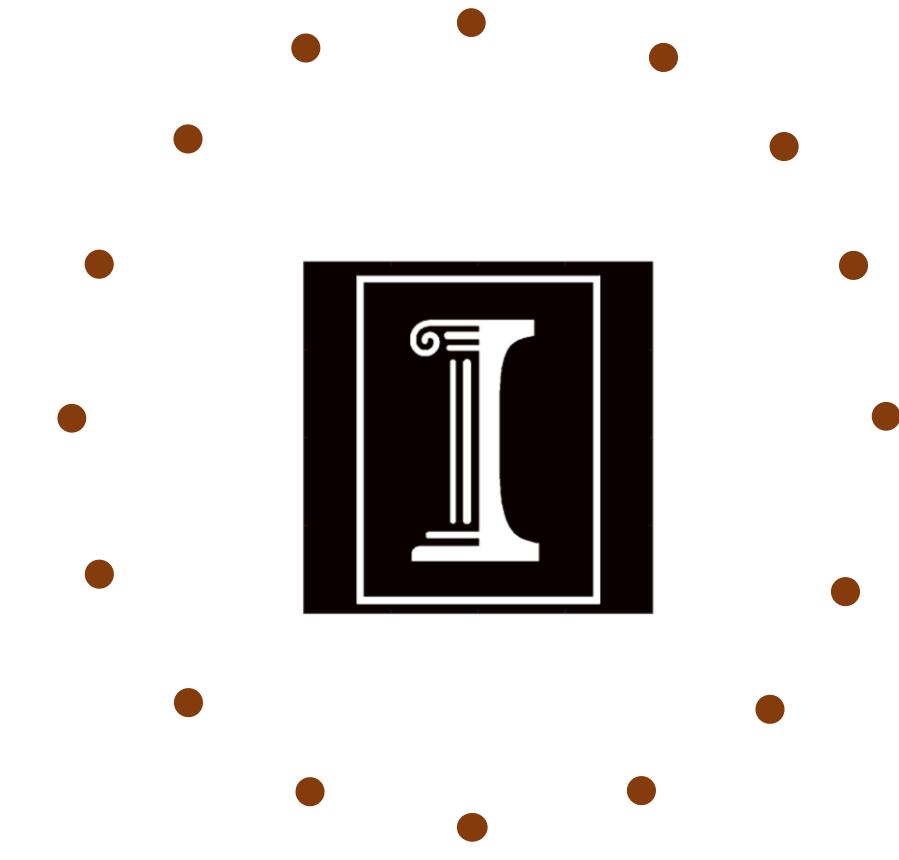
# Large-Scale Reconstruction—Almost Perfect!



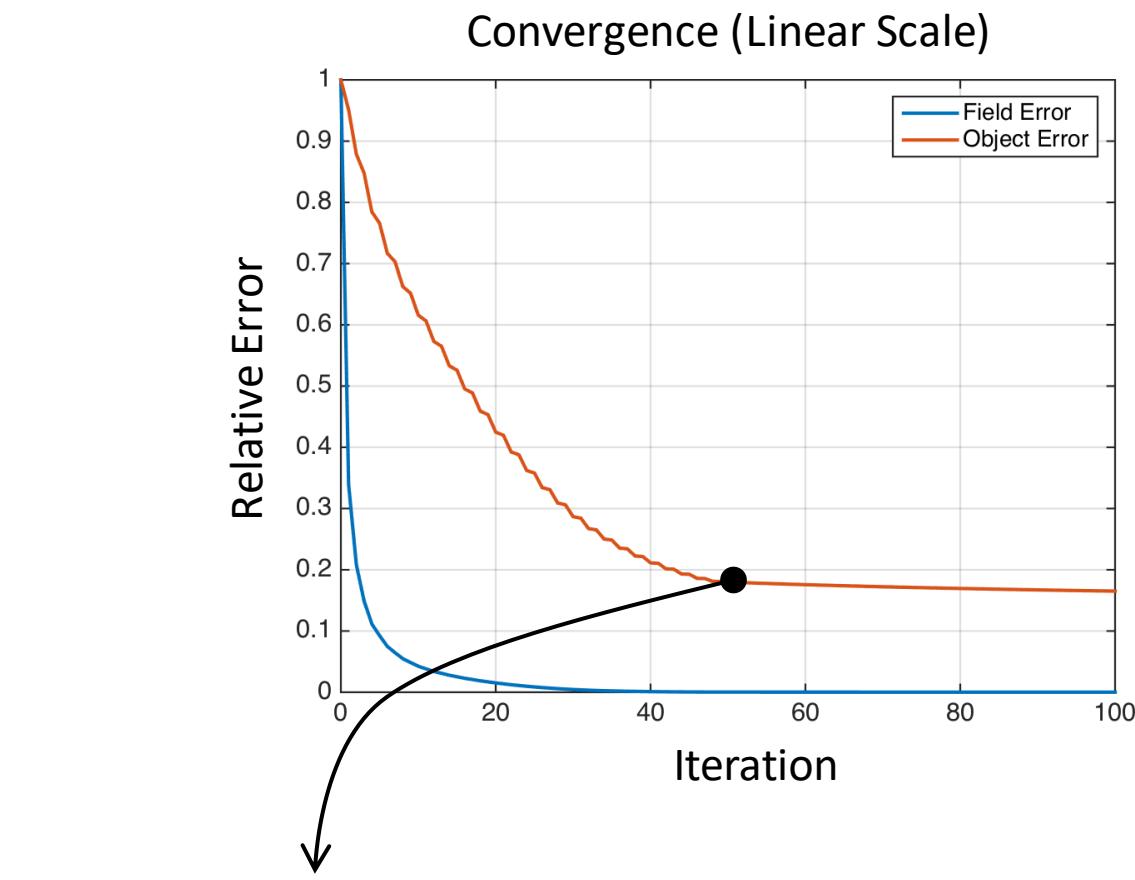
# Large-Scale Reconstruction—Almost Perfect!



# Large-Scale Reconstruction—Almost Perfect!

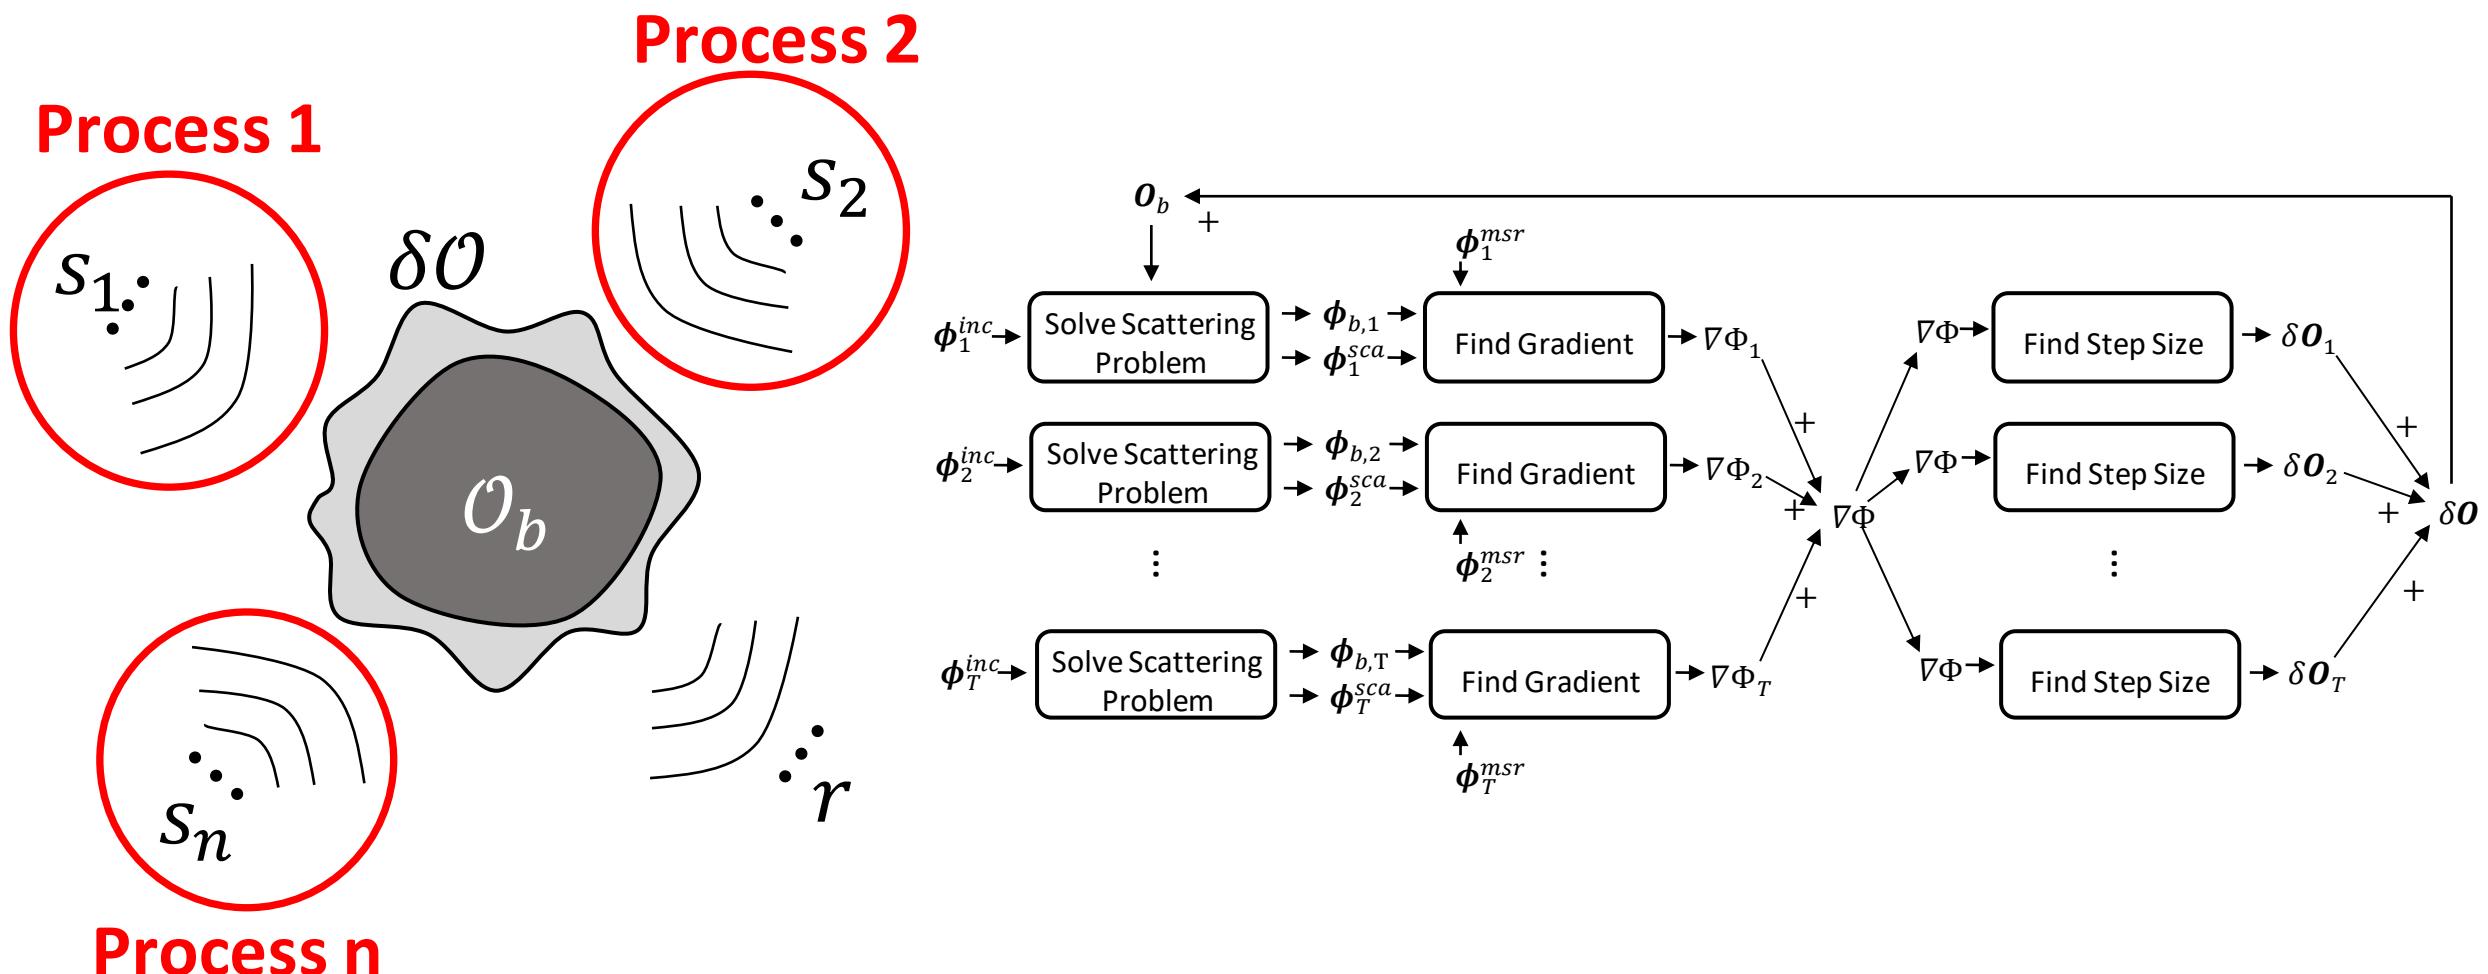


- 256 Transmitters
- 1,024 Receivers
- Monochromatic



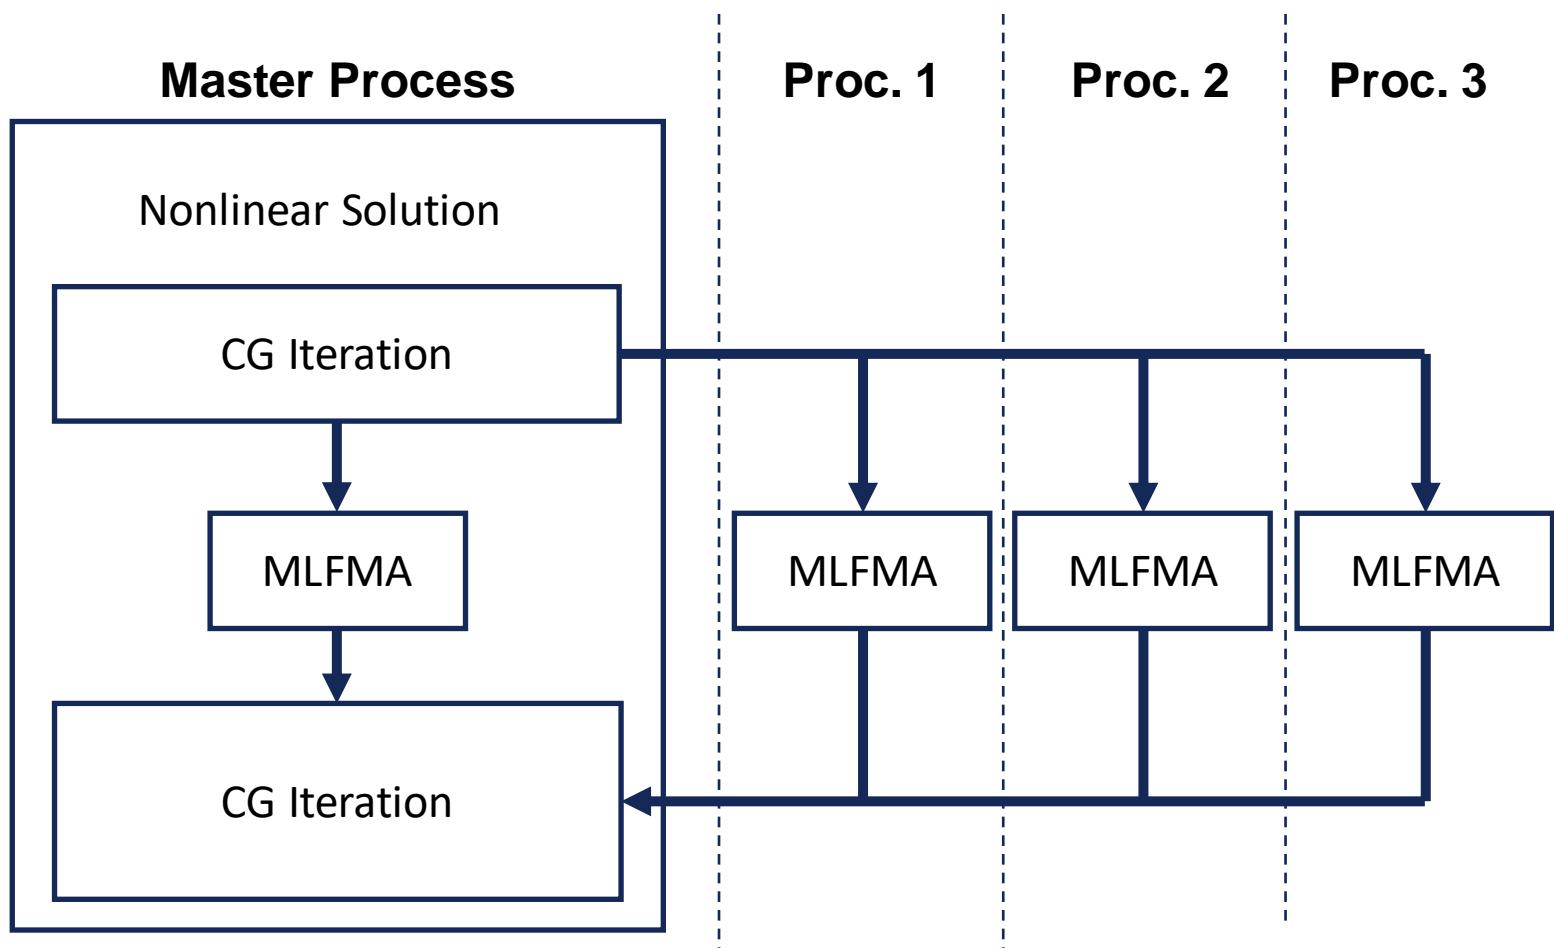
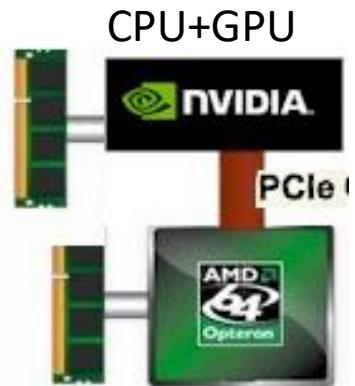
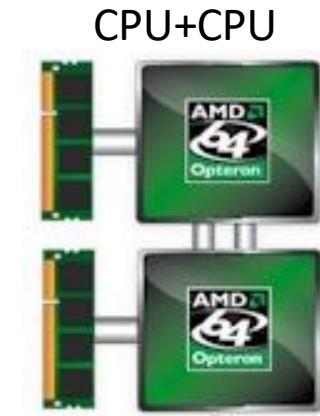
256 Trans. x 50 Iter. x 3 = 38,400 Scattering Solutions

# Multiple Illuminations Provide Diversity

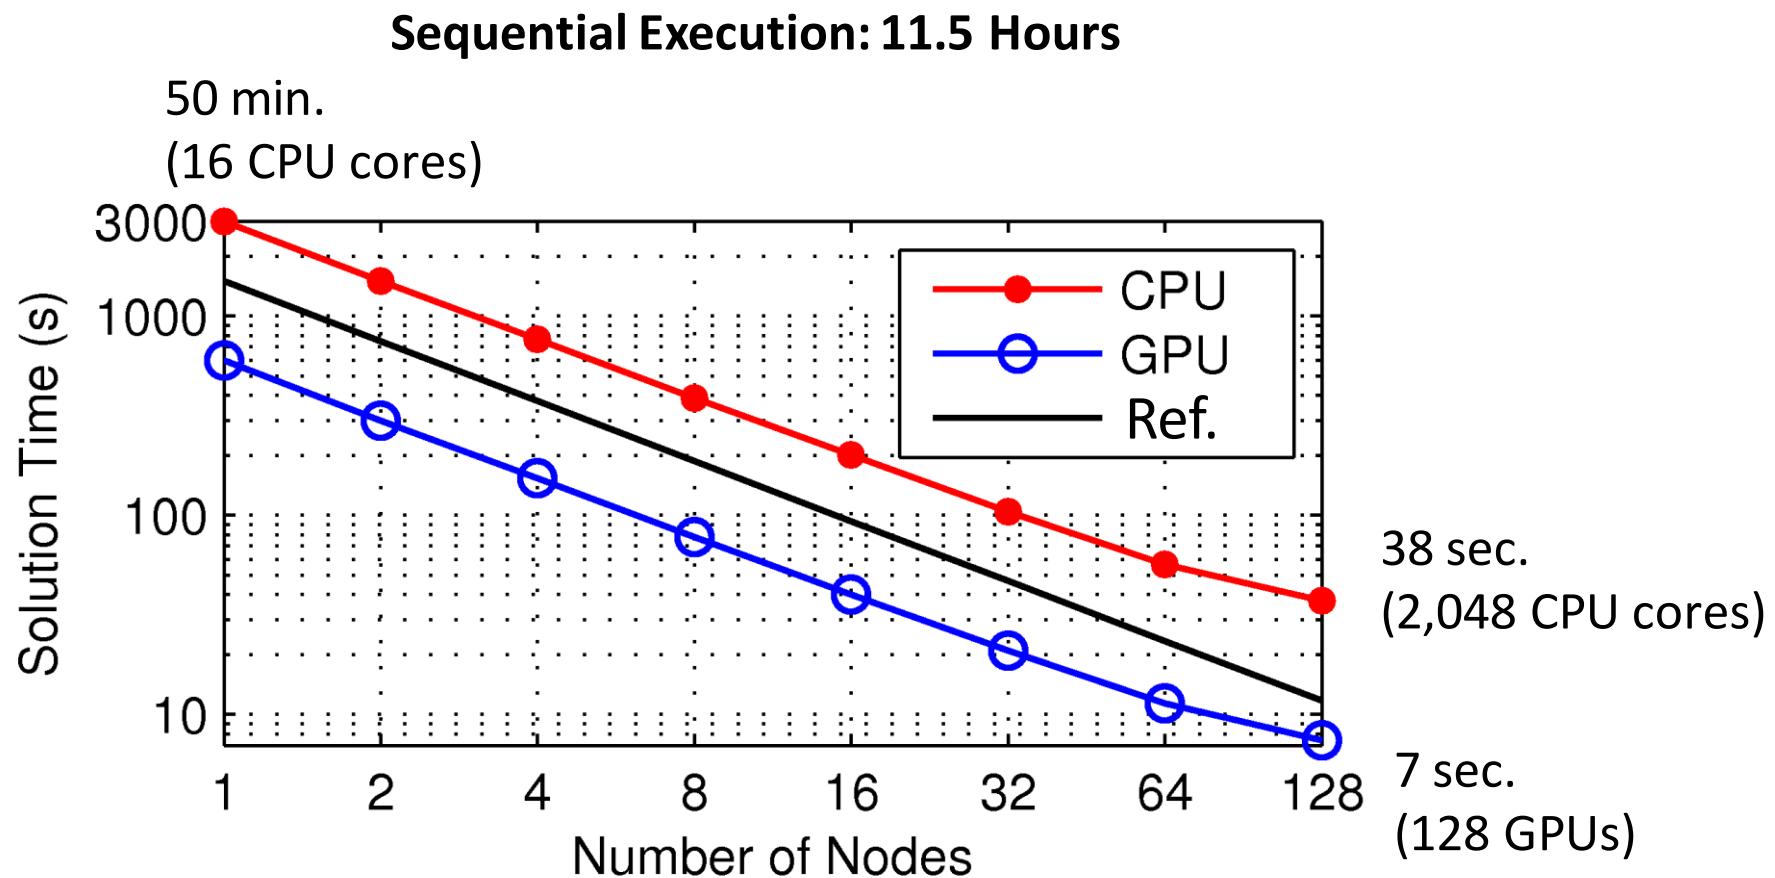


# BLUE WATERS

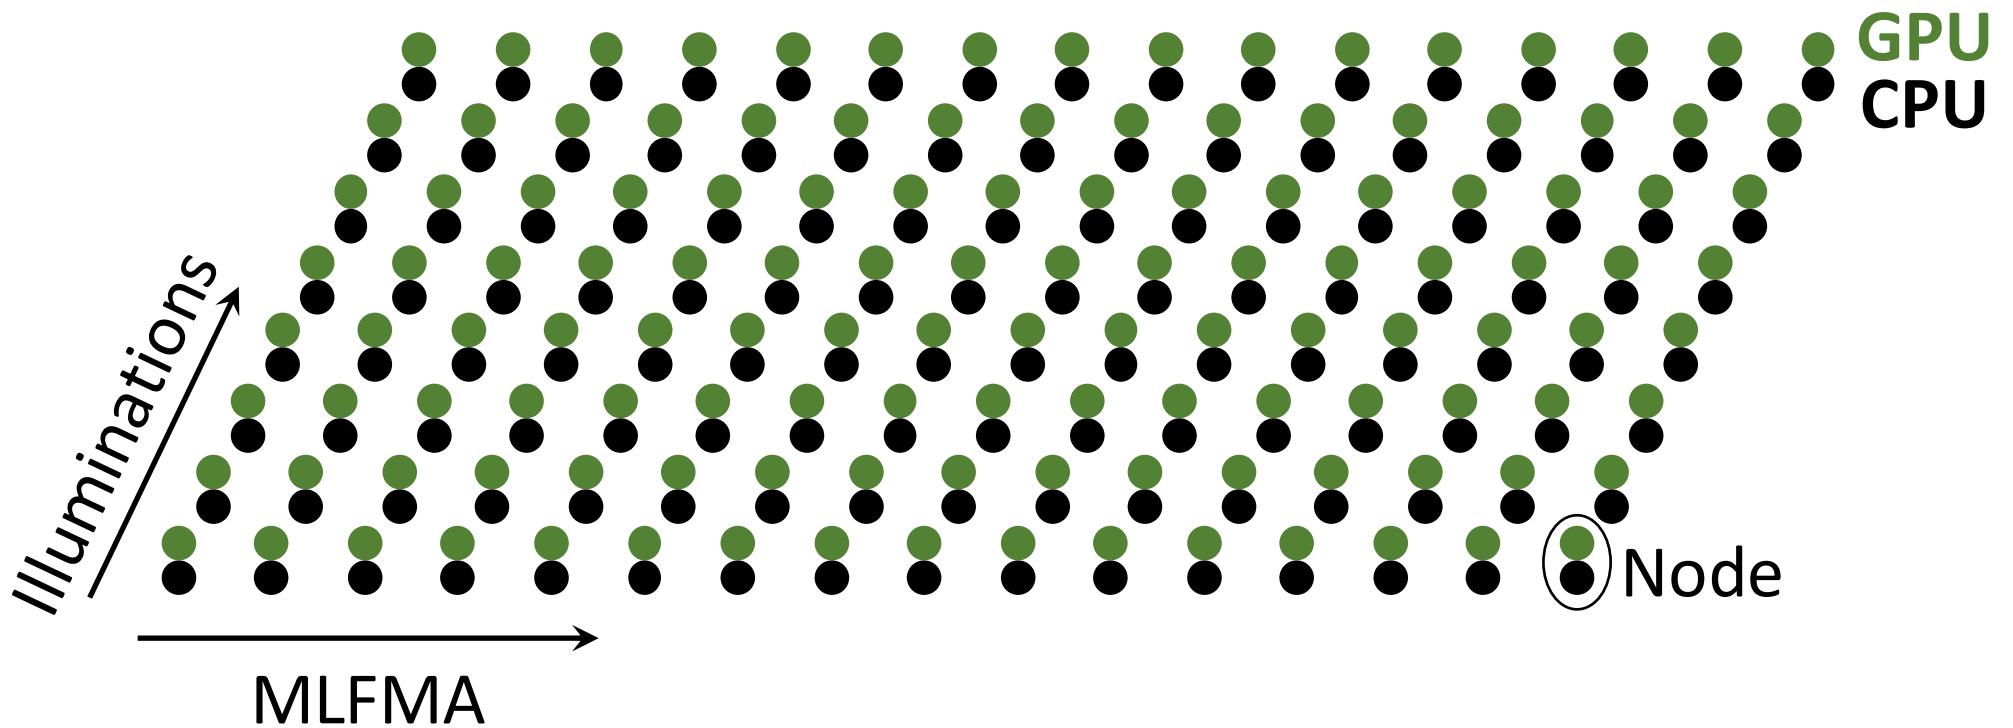
SUSTAINED PETASCALE COMPUTING



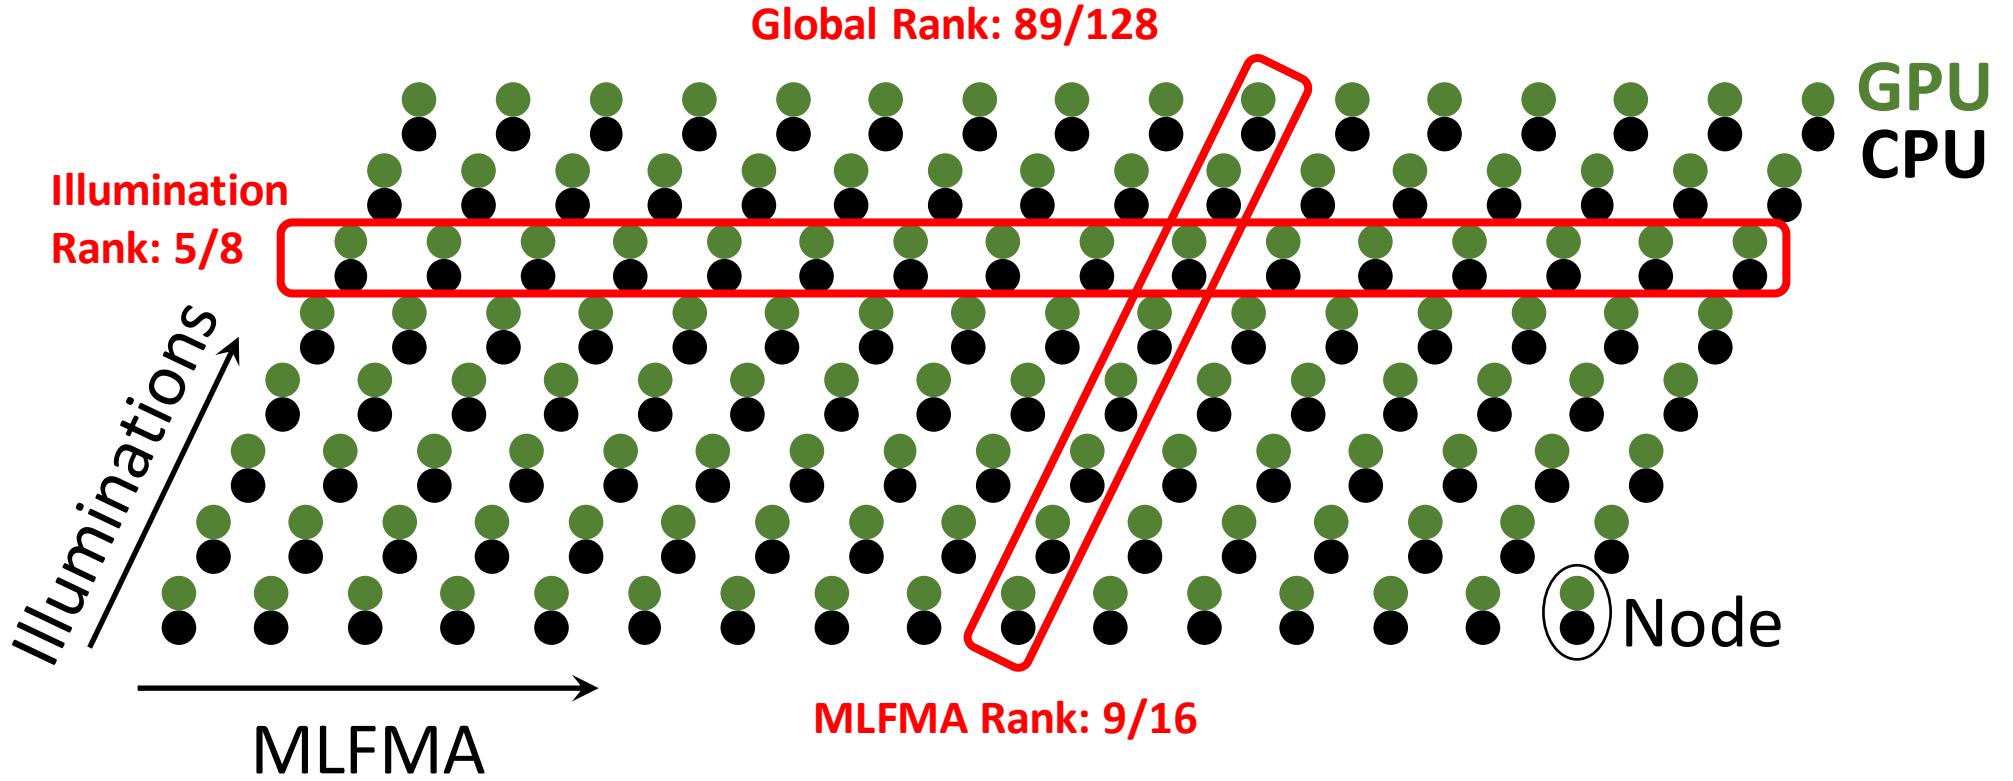
# CPU Nodes vs. GPU Nodes



# Massively-Parallel Inverse Solution Scheme



# Massively-Parallel Inverse Solution Scheme



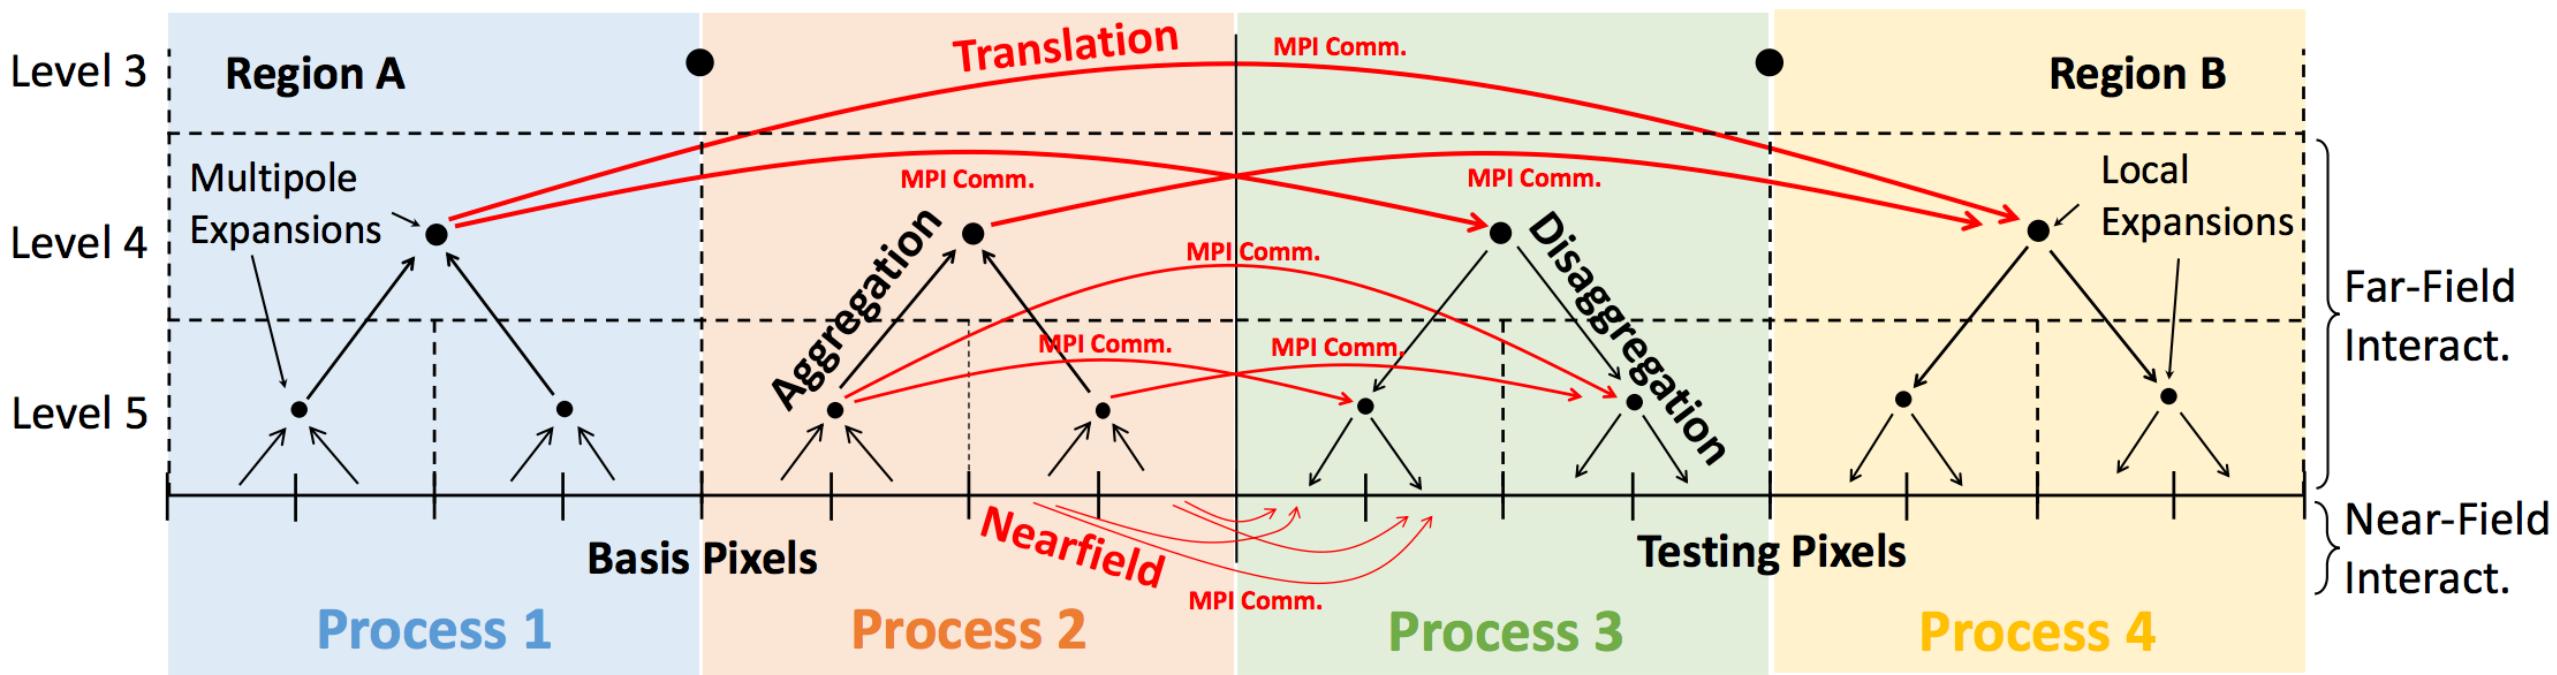
```
MPI_Comm_split( MPI_Comm comm, int color, int key, MPI_Comm* newcomm)
```

# Parallelization of MLFMA

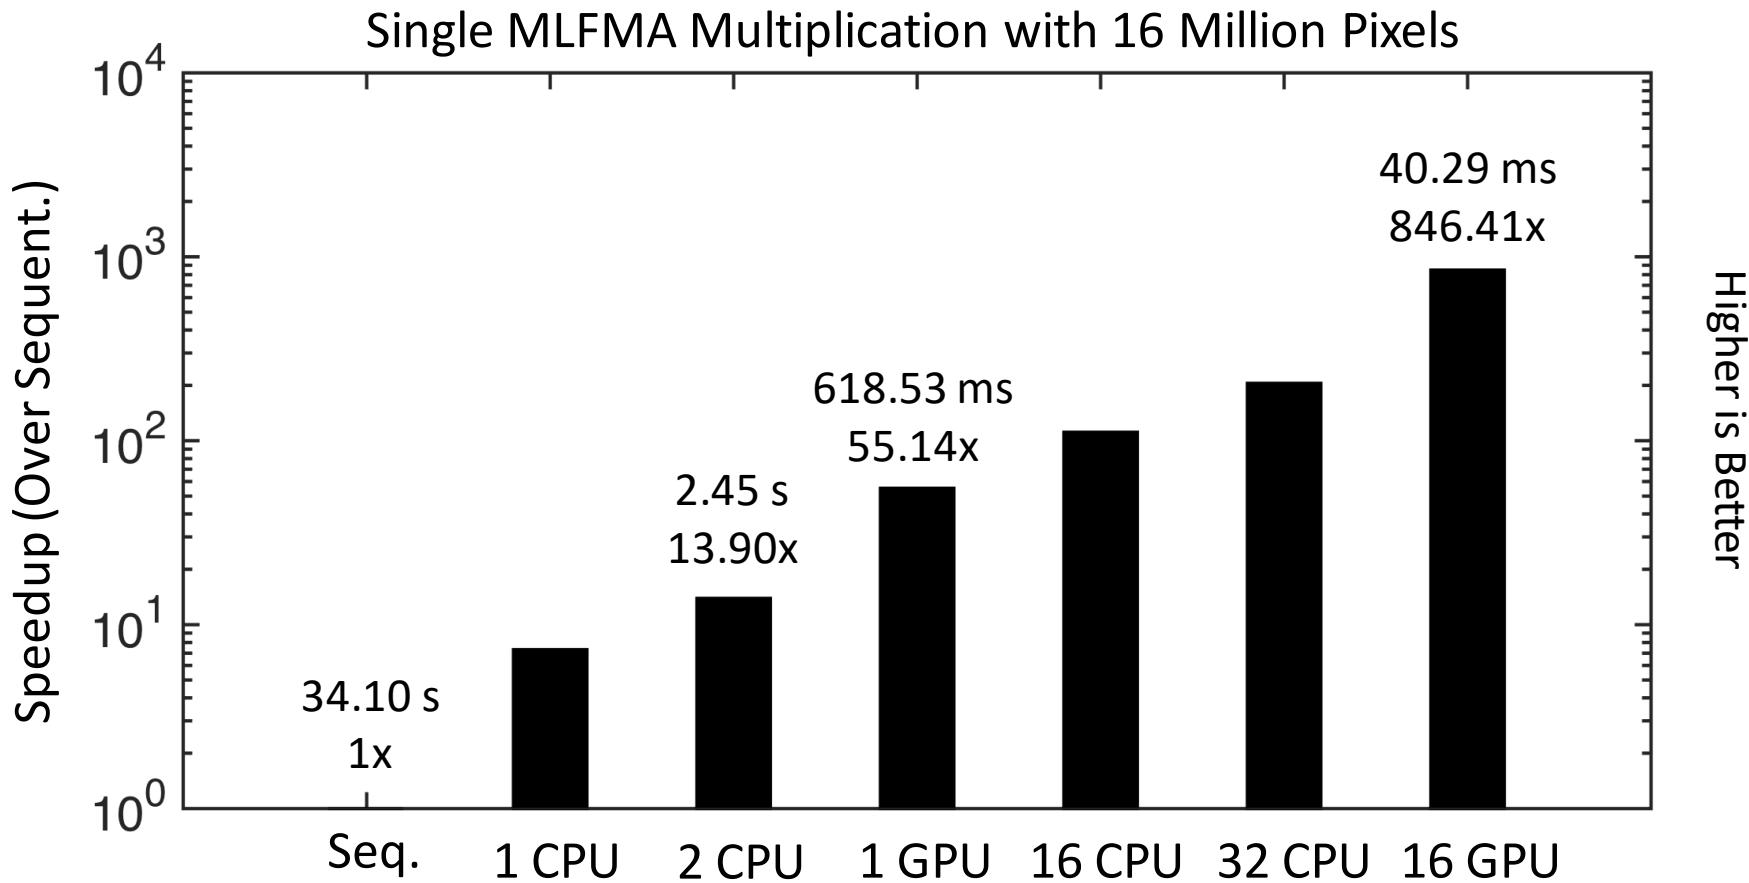
$$\phi = [\bar{I} - \bar{G}_0 \bar{O}]^{-1} \phi_0$$

↓  
4 PB

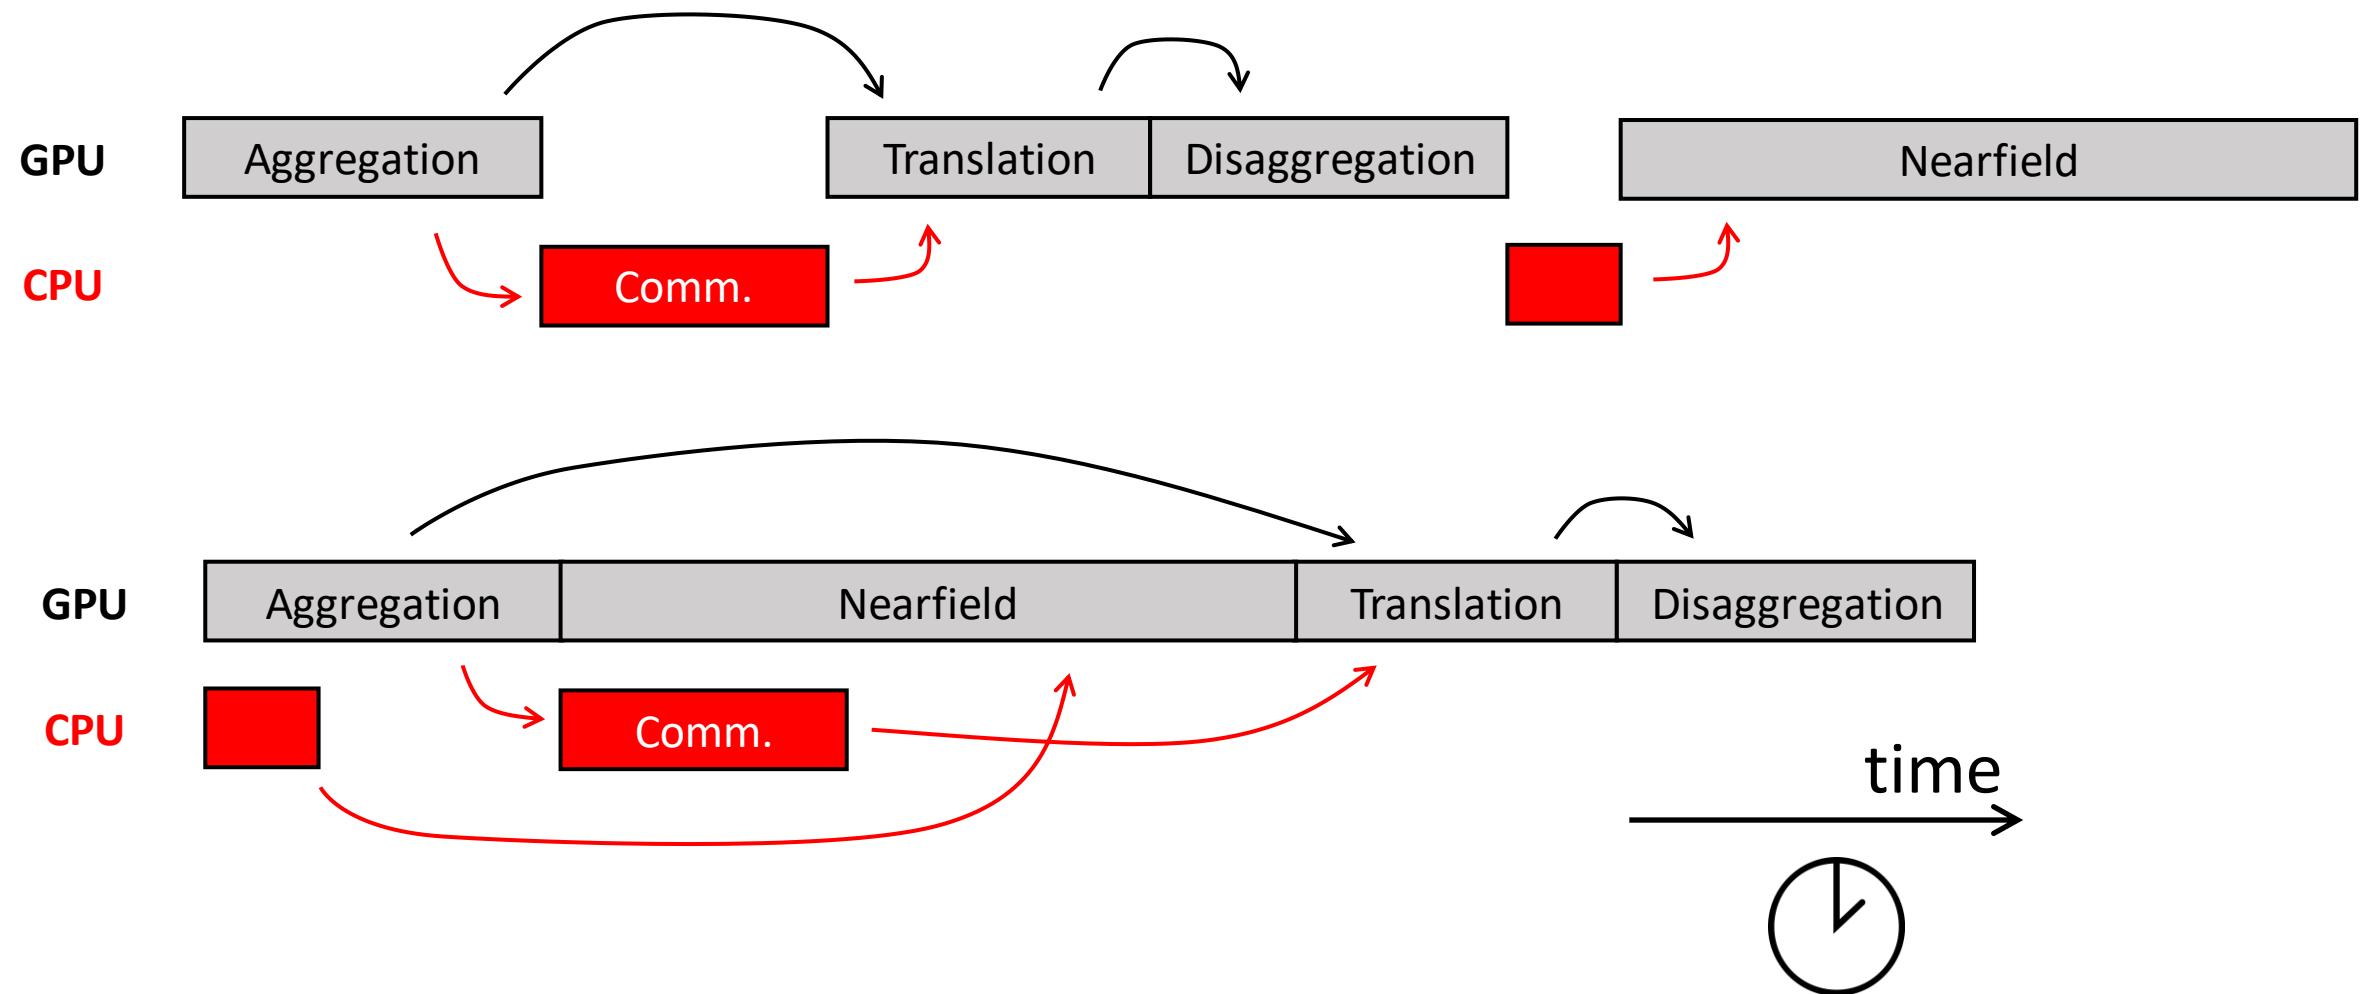
MLFMA Schematic



# GPU Provides About 55x Per-Node Speedup

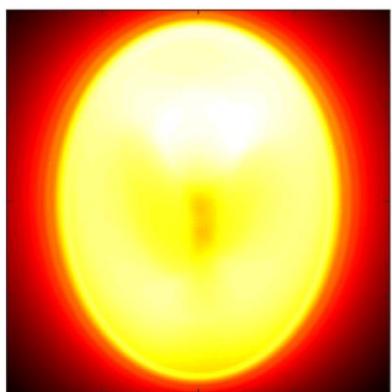


# MPI Communications can be Completely Hidden

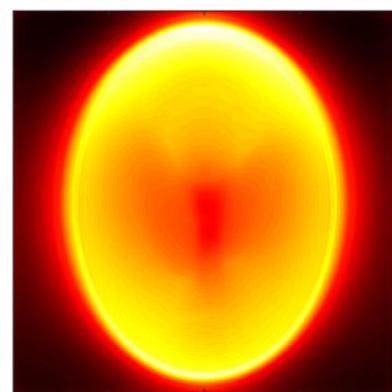


# Large-Scale Reconstruction—Almost Perfect!

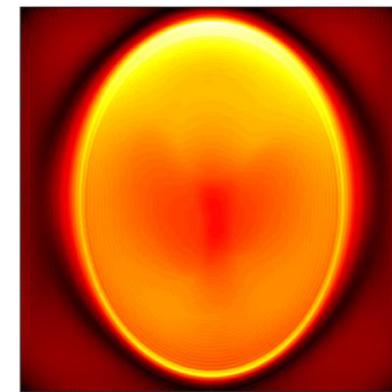
Iteration 1



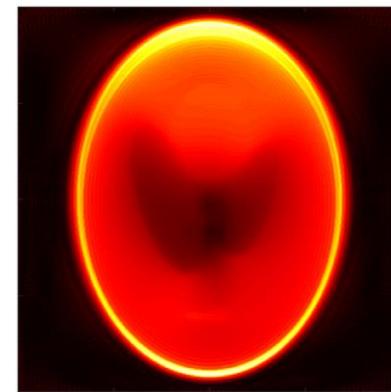
Iteration 2



Iteration 4



Iteration 10



Iteration 15



Iteration 20

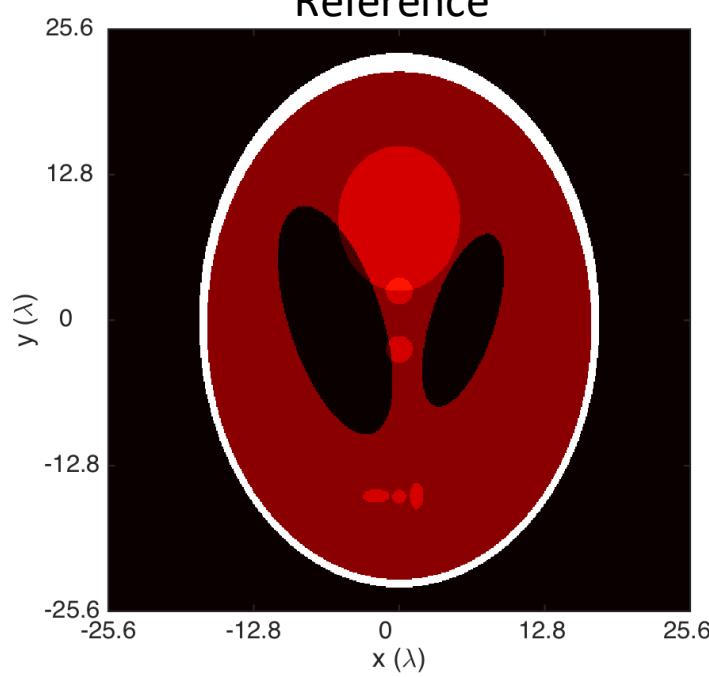
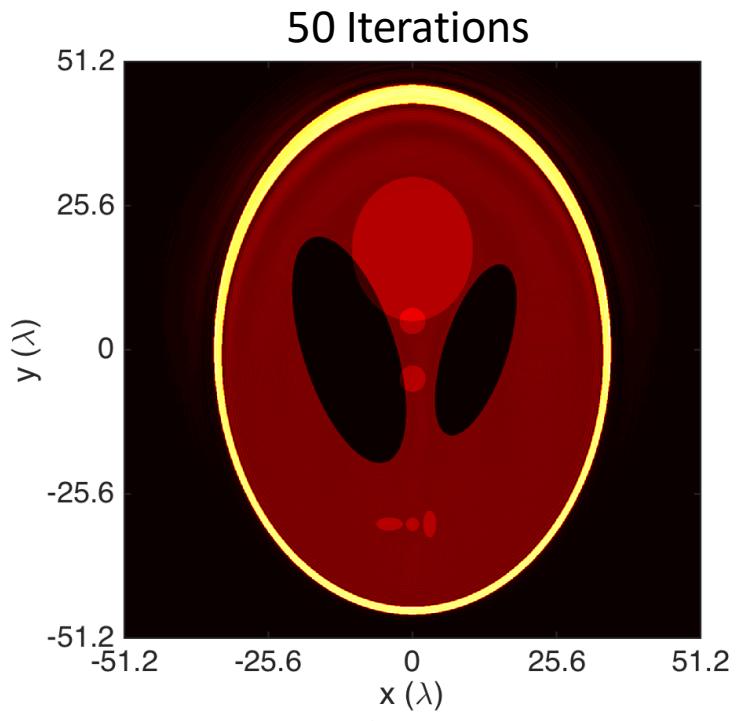


Iteration 60

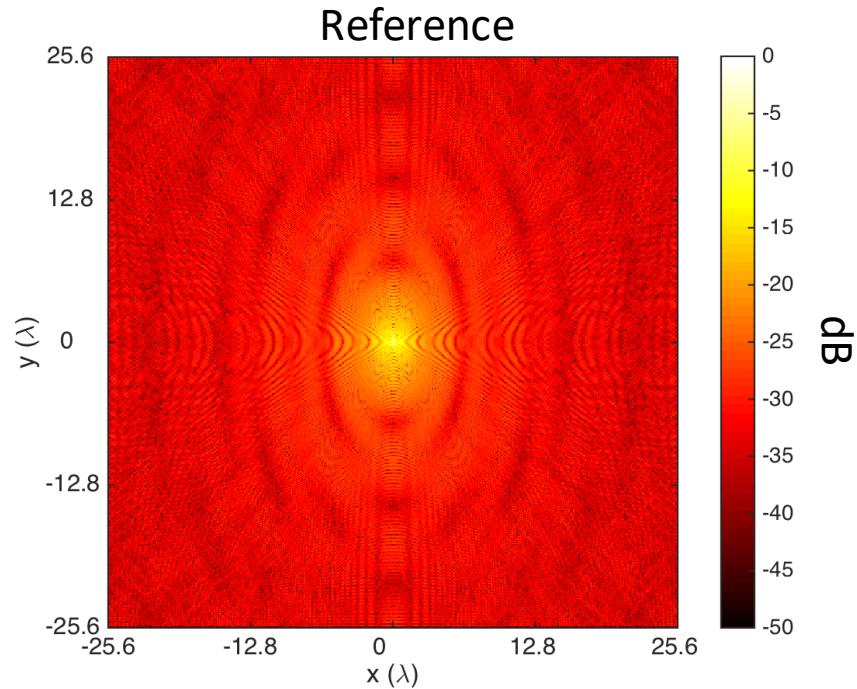
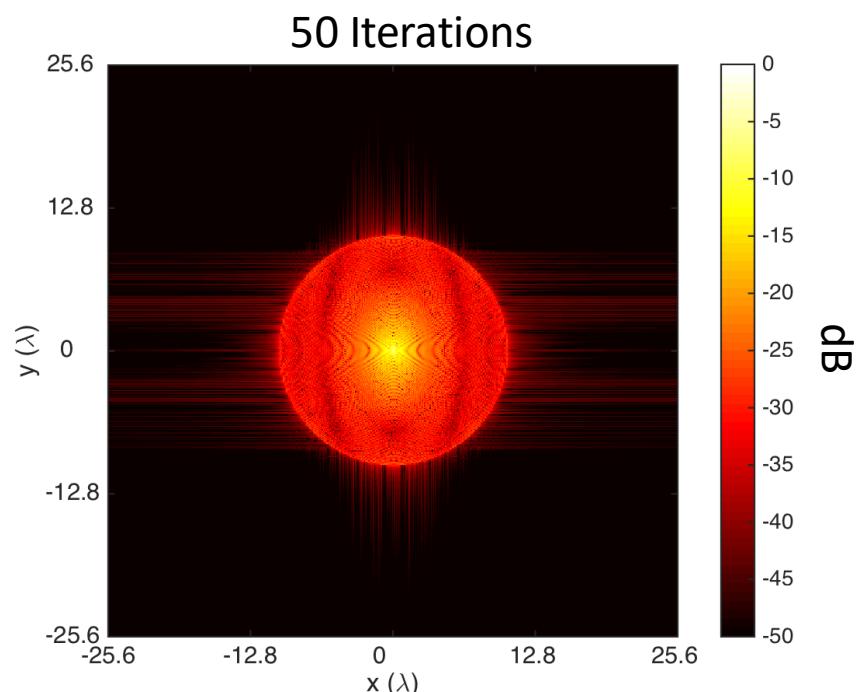


Iteration 100



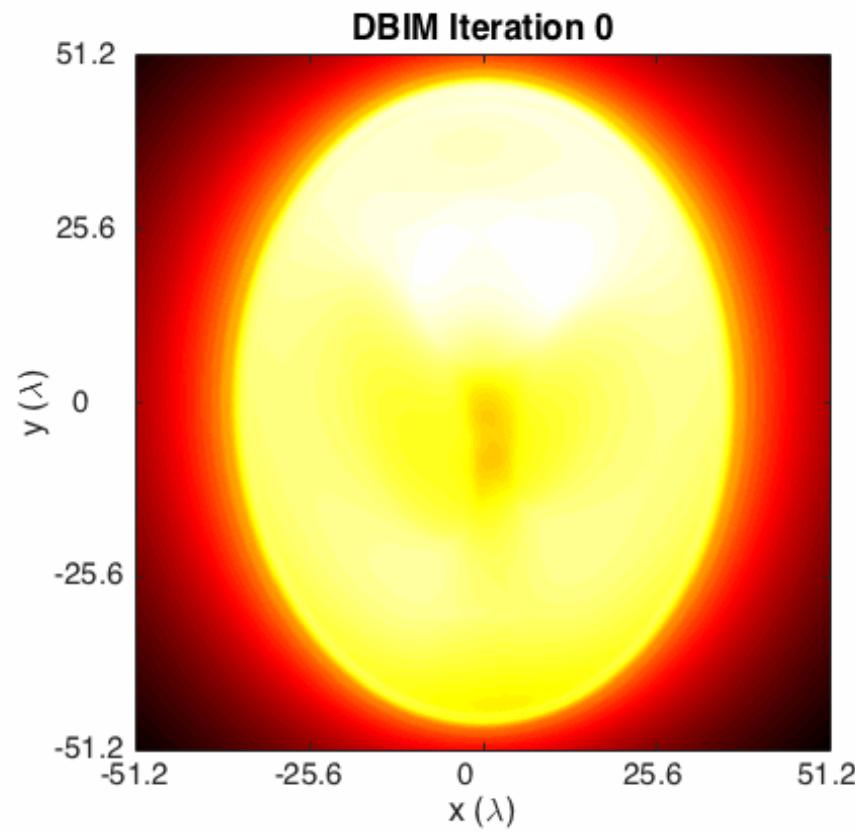


Spatial Domain

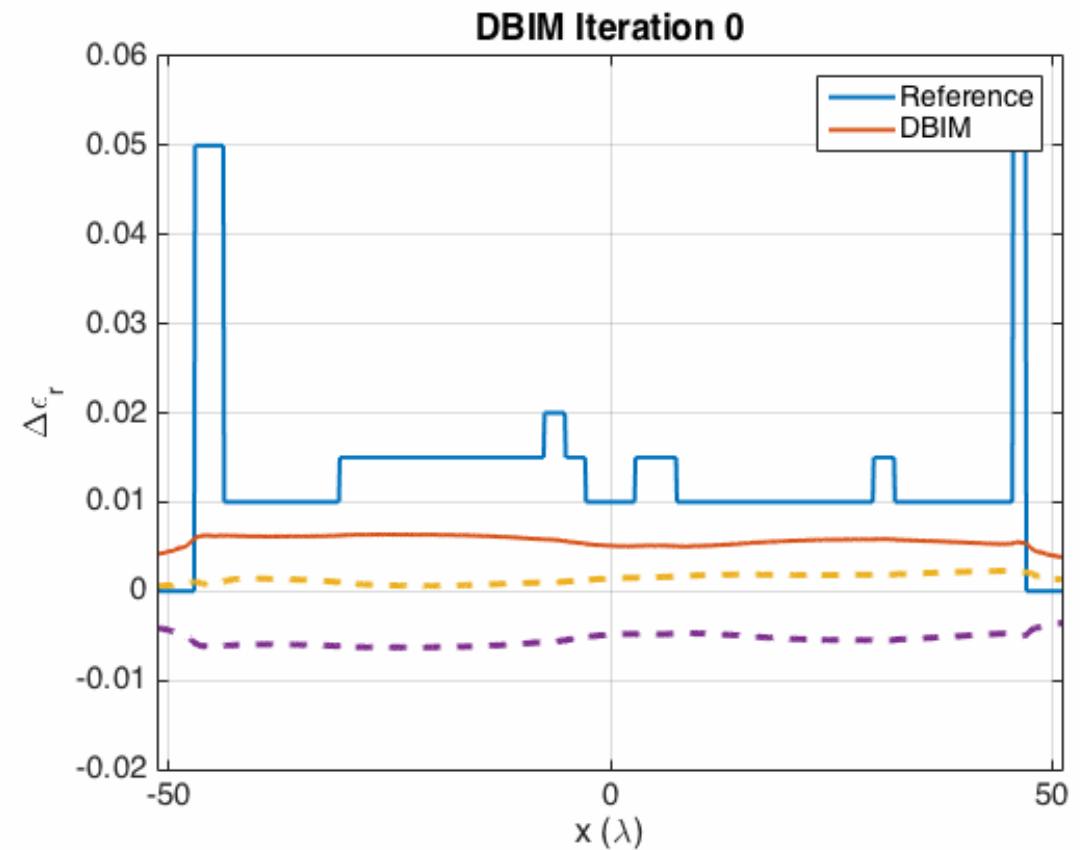


Spectral Domain

# Large-Scale Reconstruction—Almost Perfect!

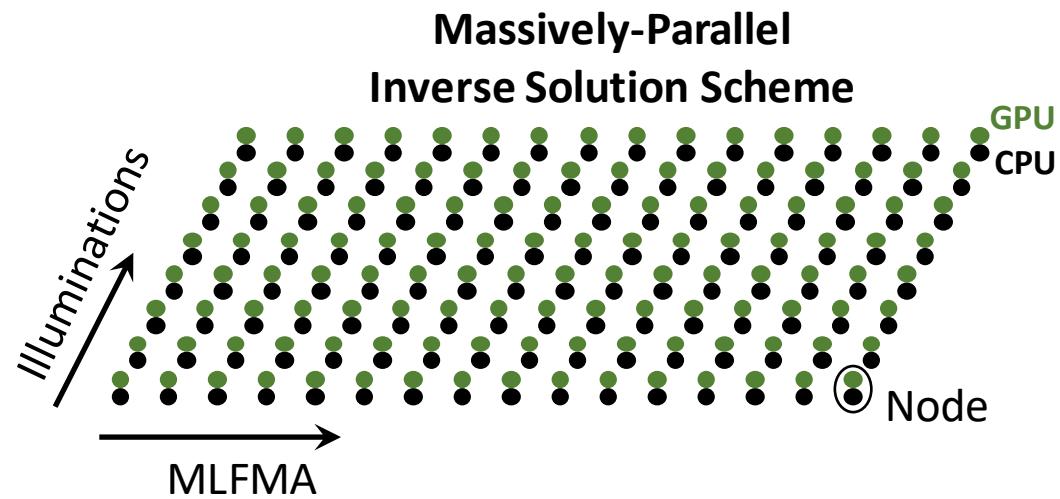
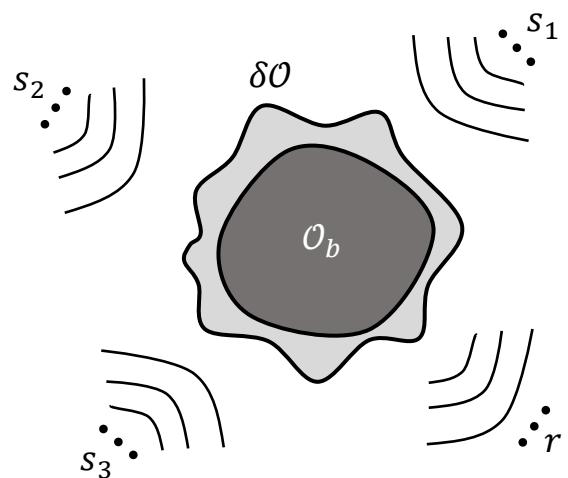


**1024 Transmitters**  
**1024 Receivers**  
**Monochromatic**

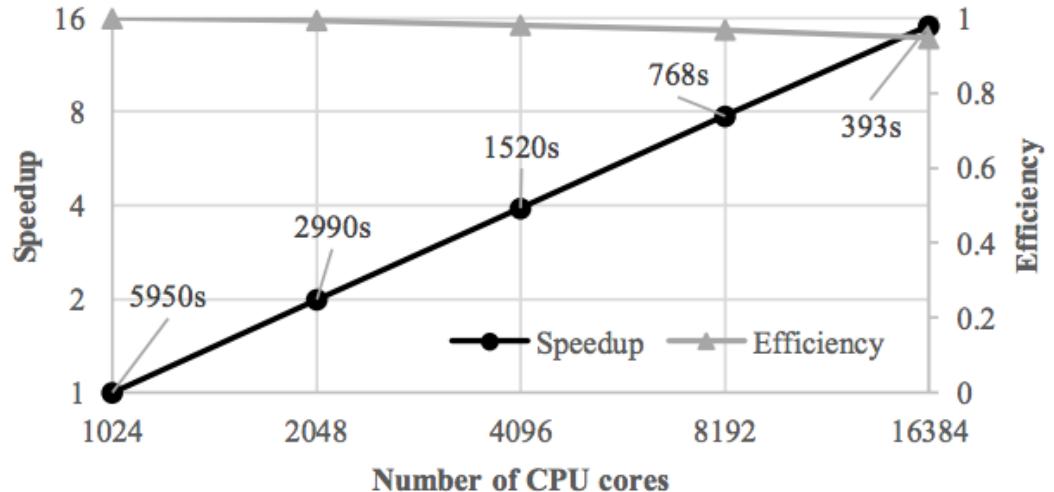


**1024 CPU Nodes: 558 sec.**  
**1024 GPU Nodes: 142 sec.**

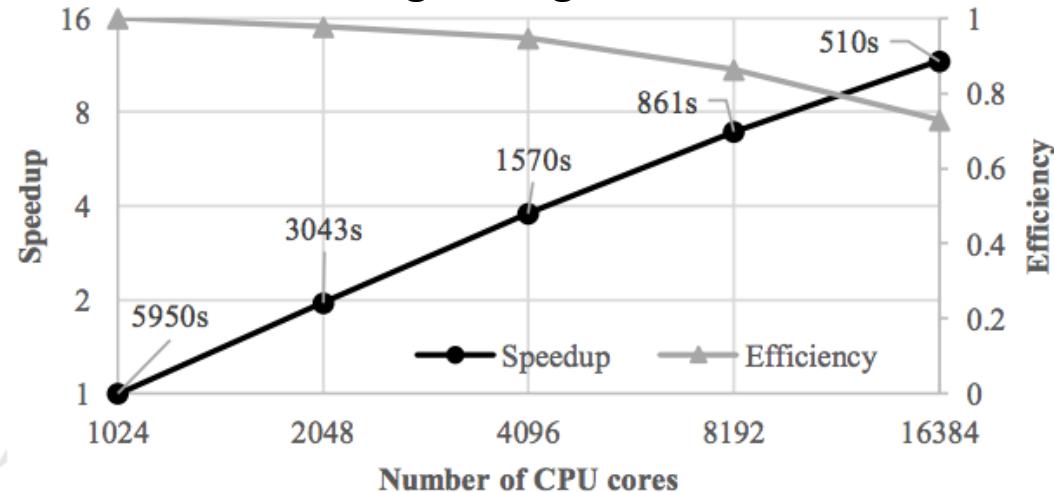
# Scalable Parallel Inverse Solutions



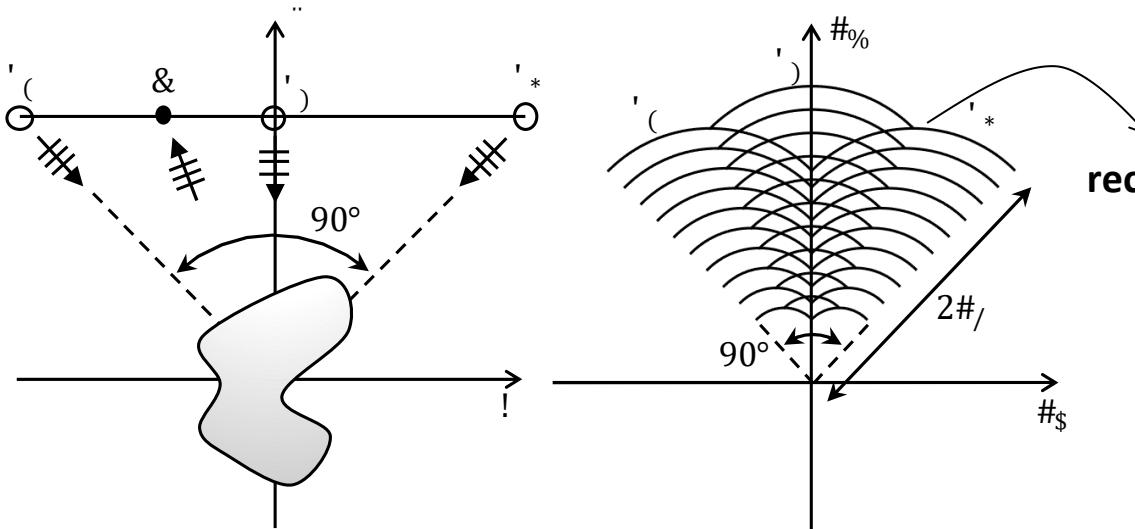
Strong Scaling of Illumination



Strong Scaling of MLFMA

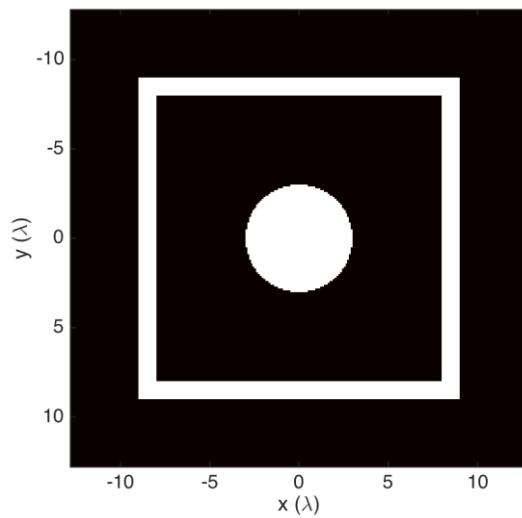


# Seeing the “Invisible”

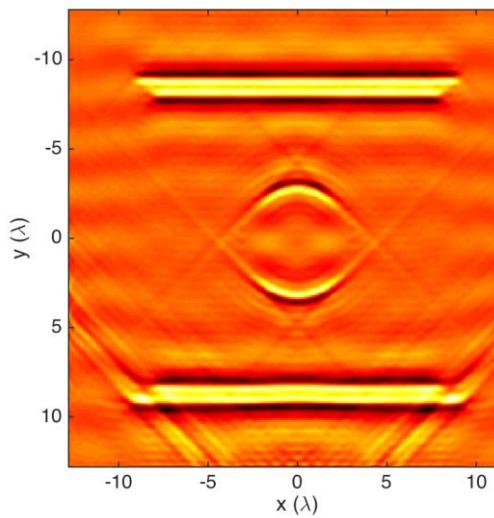


recoverable spectral area with  $90^\circ$  angle sector of illumination/observation

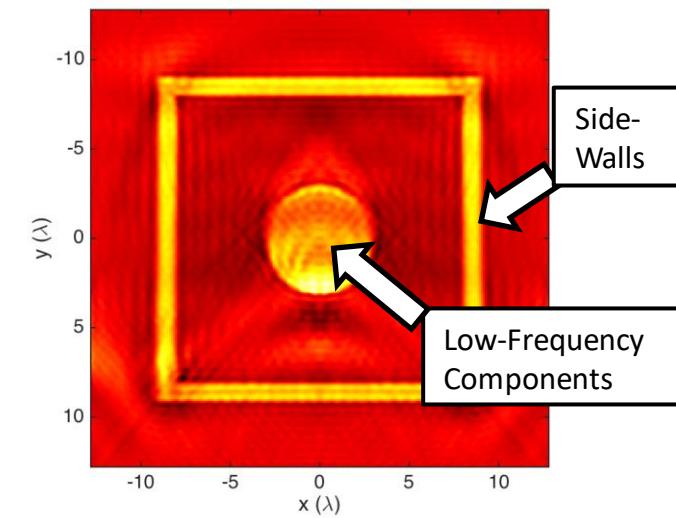
Numerical Phantom



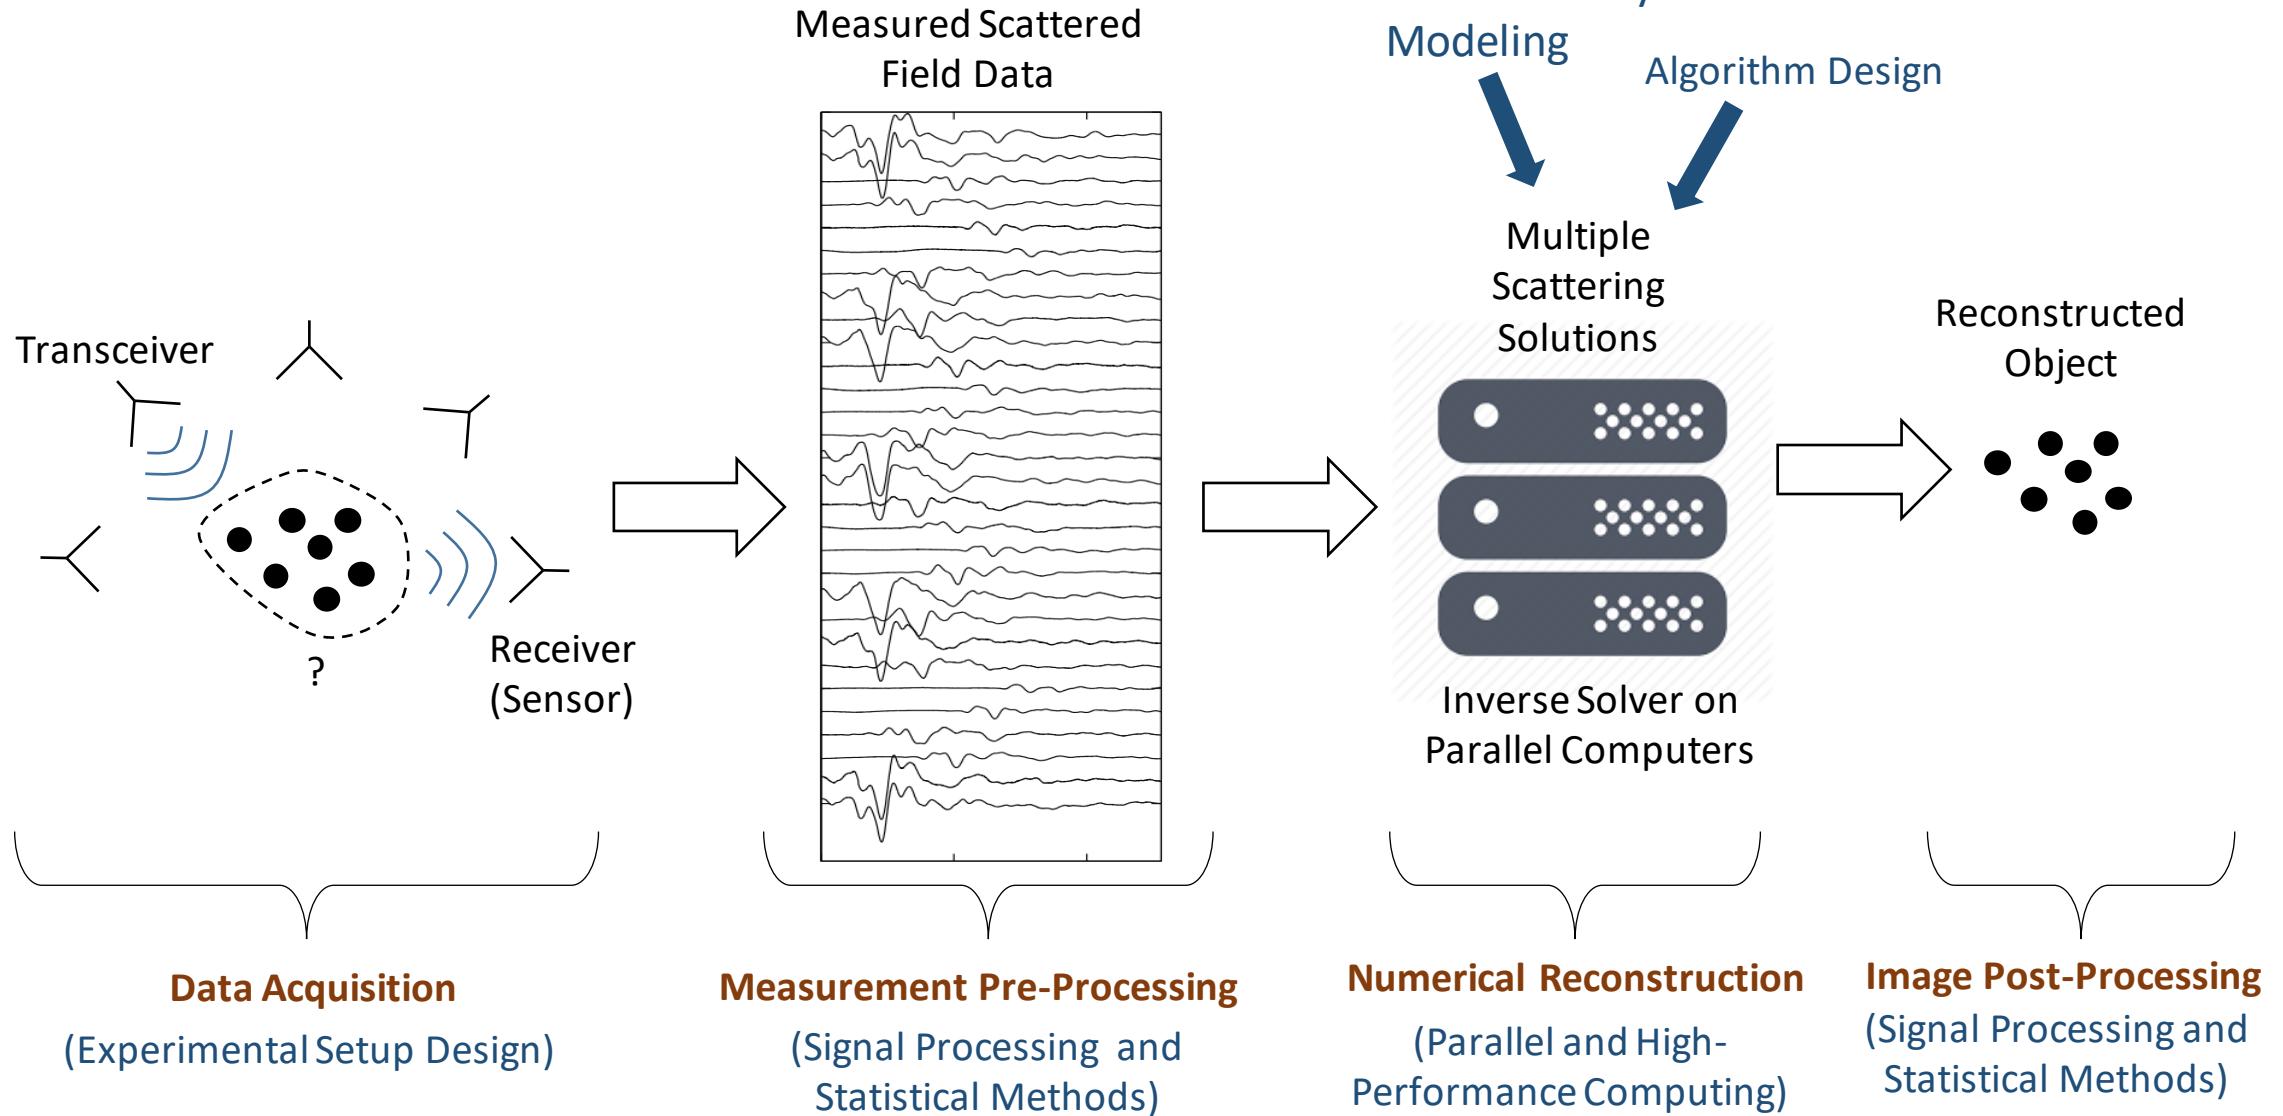
Diffraction-Tomography Imaging



Multiple-Scattering Imaging



# The Big Picture



ECE ILLINOIS

Department of Electrical  
and Computer Engineering

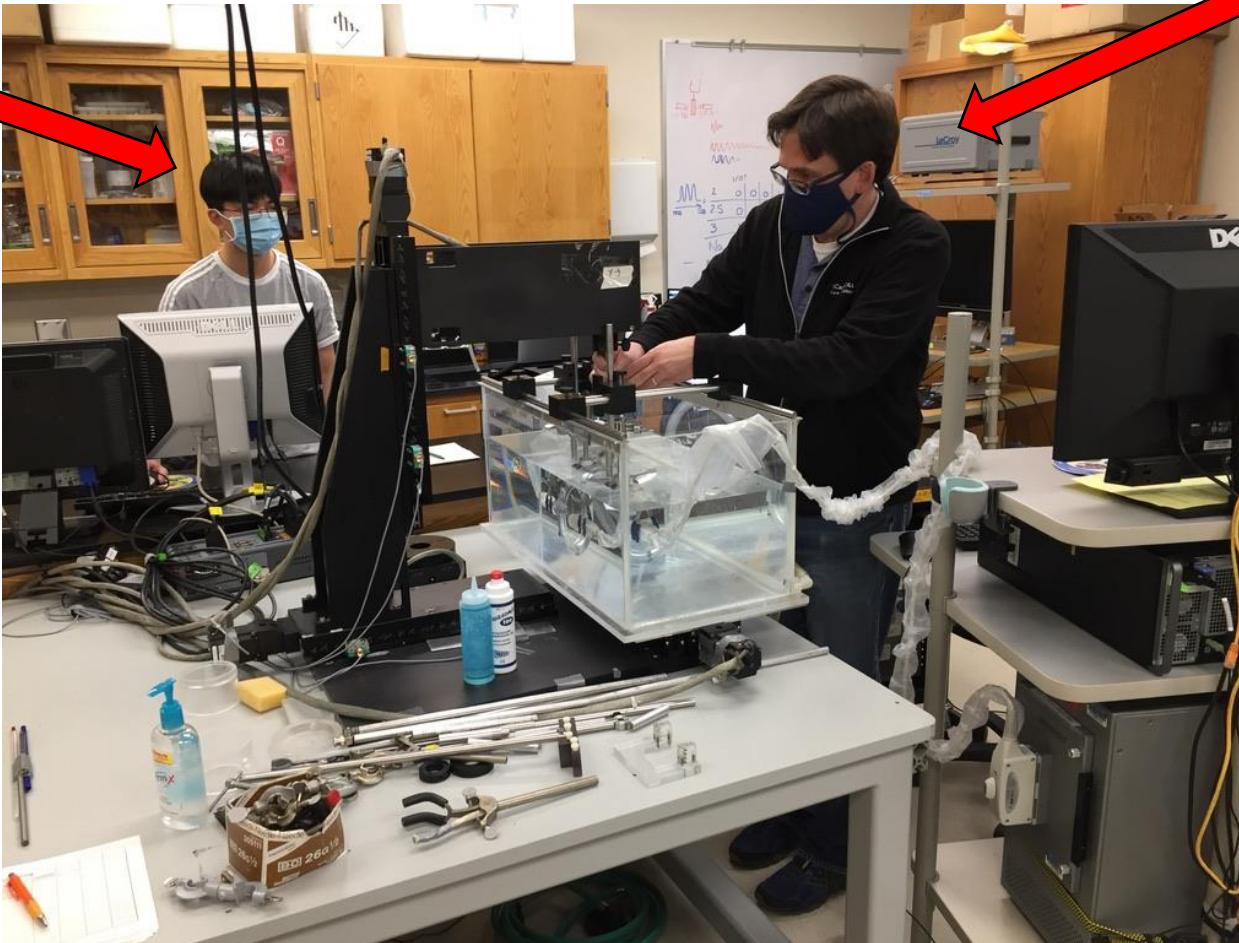
**I** Computational  
Science and  
Engineering

The Dark Side of the Moon:  
**Calibration**

# Calibration Step 1: Go down the lab!

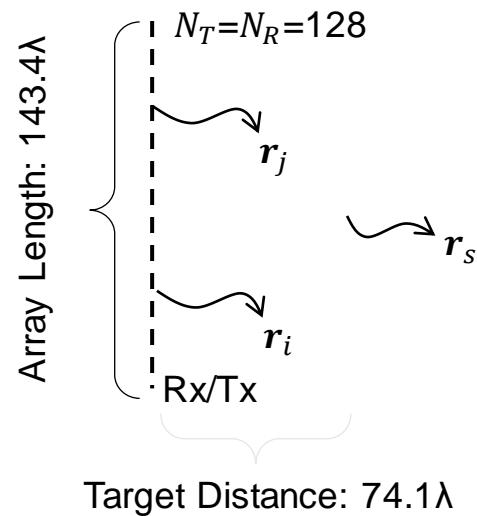
PhD Student

Dr. Oelze

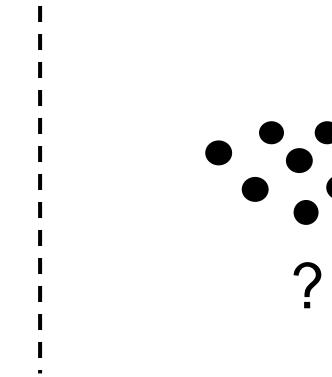


## Calibration Step 2: Take Reference Measurements

Reference Experiment

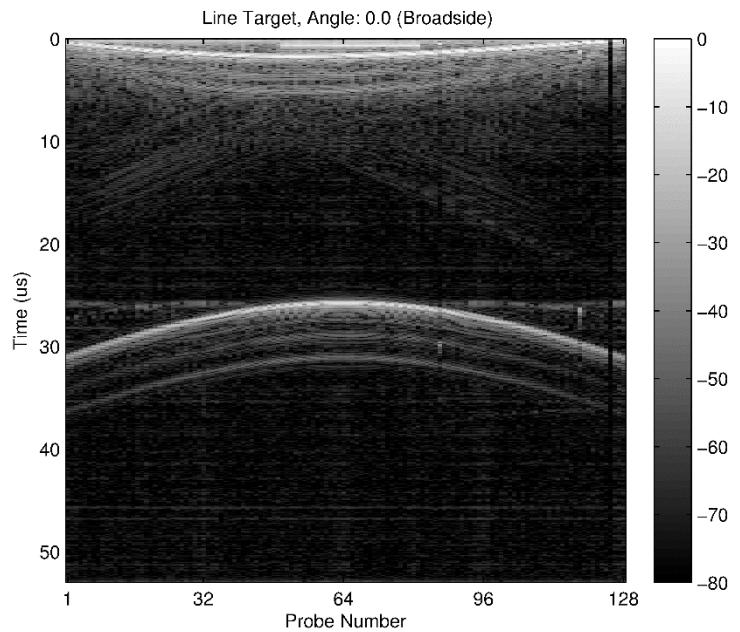


Actual Experiment

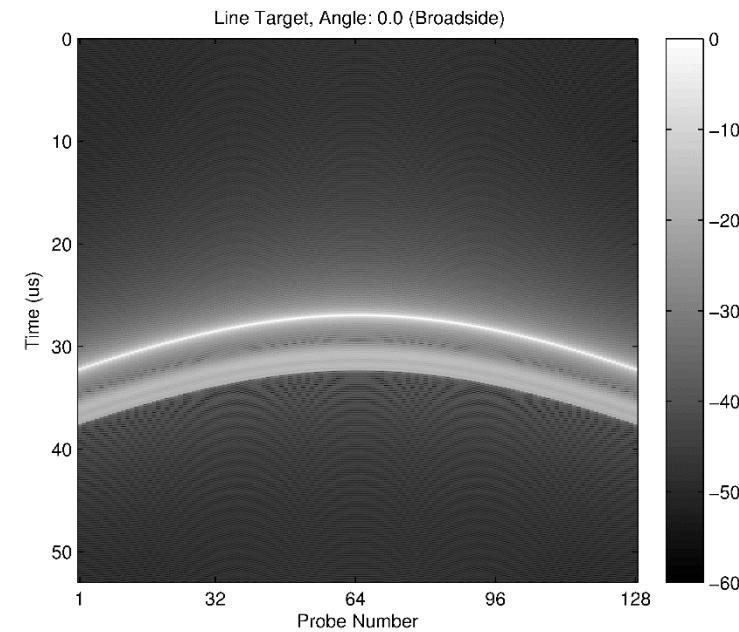


# Calibration Step 2: Take Reference Measurements

**Measured (For Calibration)**

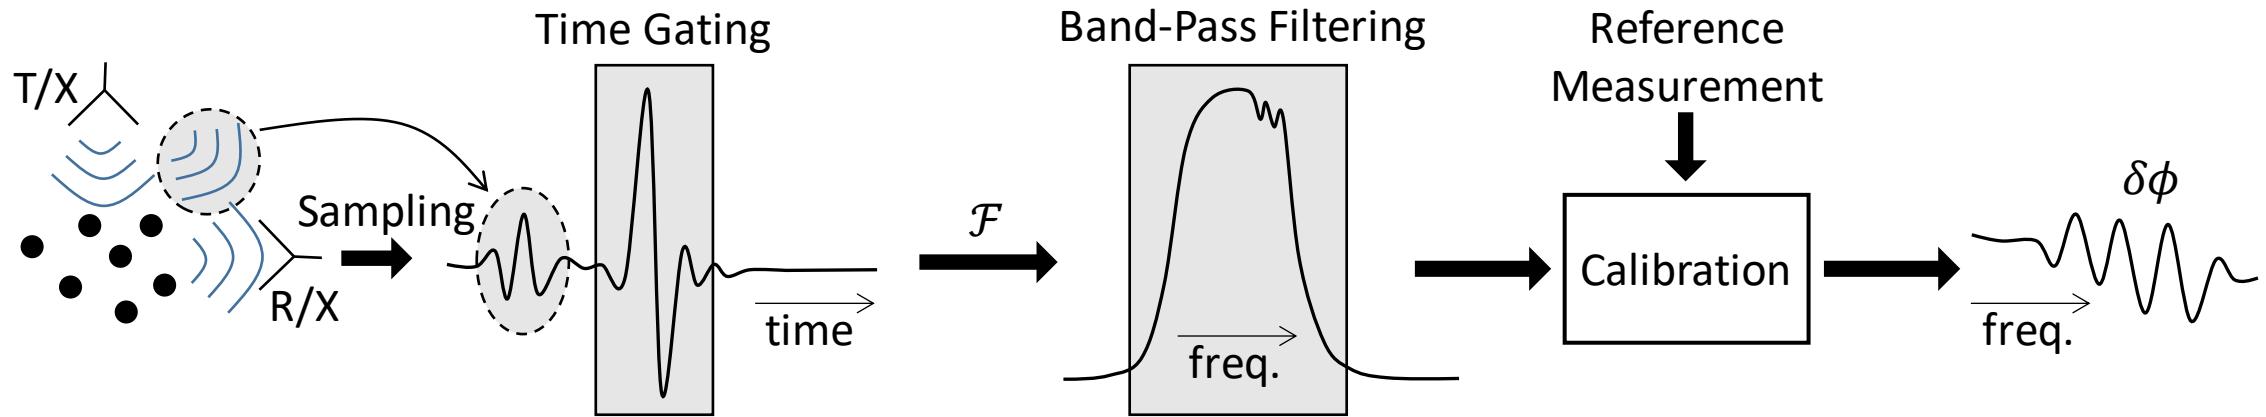


**Theoretical (Ideal)**



In Time Domain

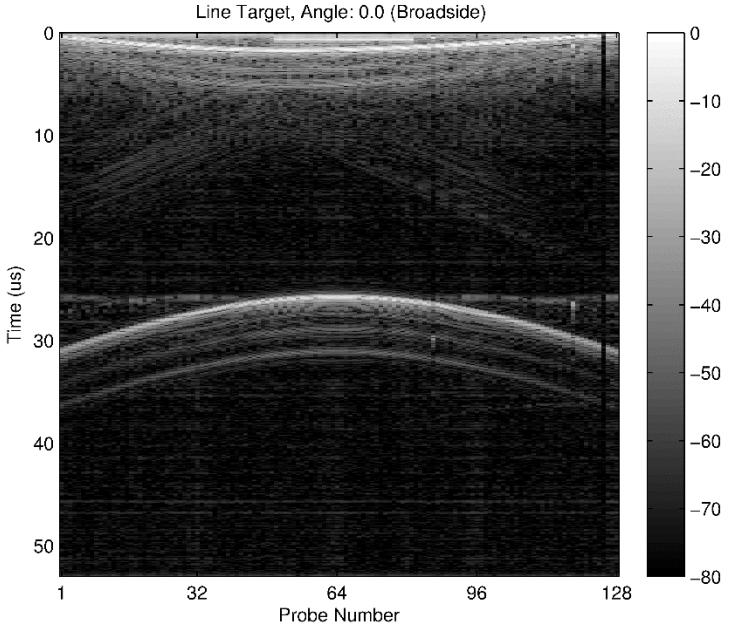
# Calibration Step 3: Find Calibration Parameters



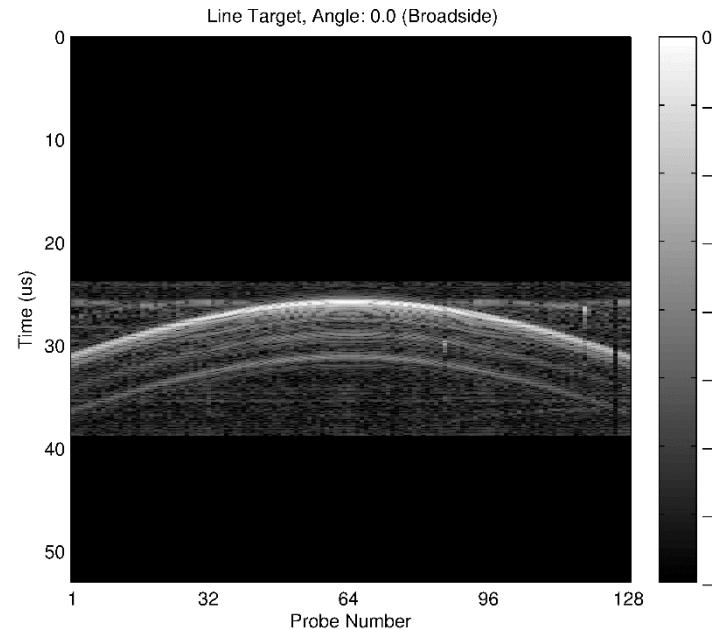
Distorted Born Approximation

$$\delta\phi \approx \mathcal{G}_b \phi_b \delta\mathcal{O}$$

**Measured (For Calibration)**

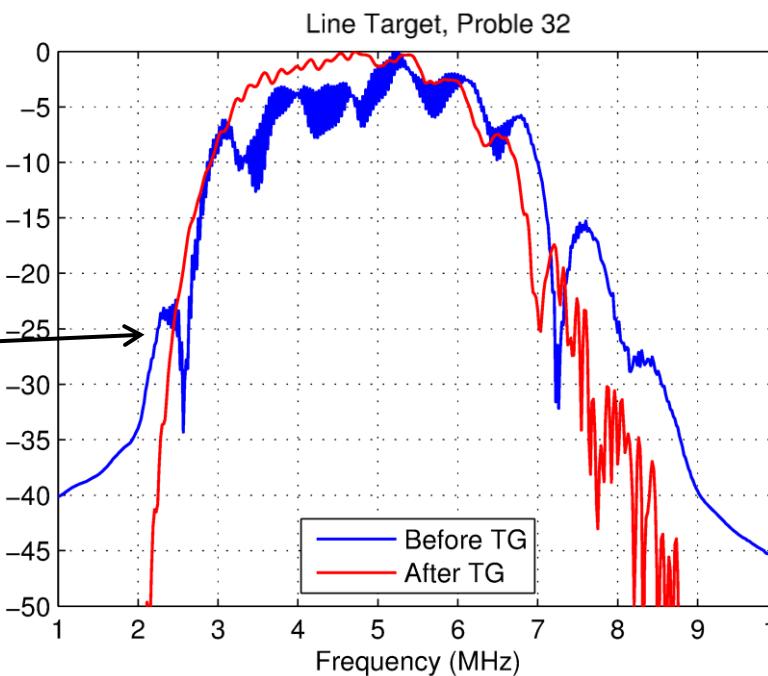


**Measured (For Calibration)**

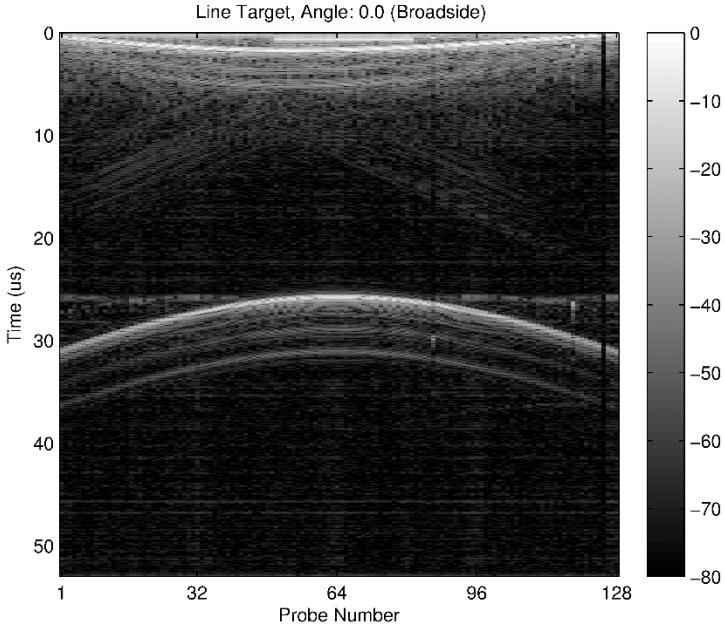


Time  
Gating  
(Simple)

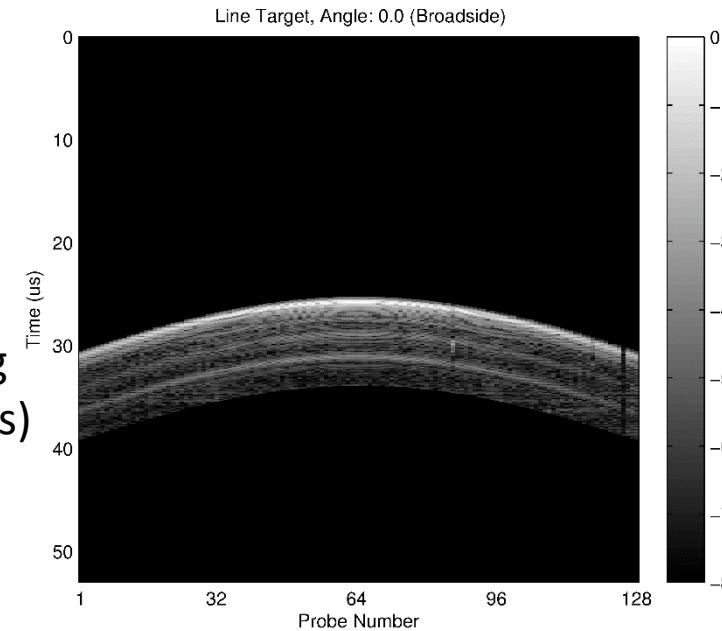
Frequency Band:  
2.5 – 7.5 MHz



**Measured (For Calibration)**



**Measured (For Calibration)**



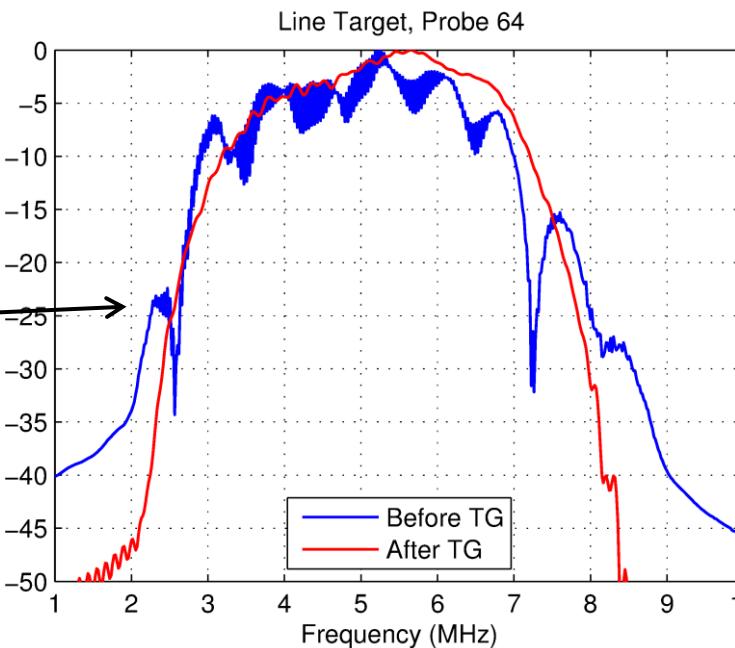
(Clever)

Time

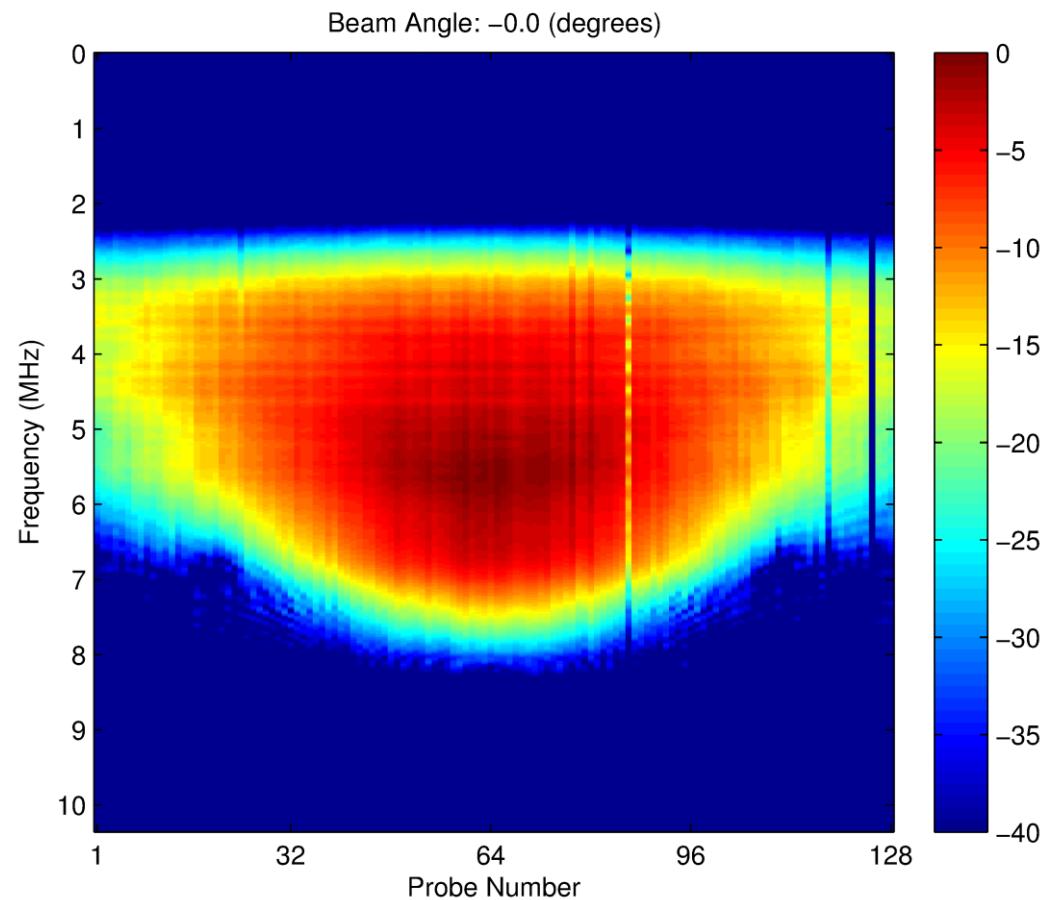
Gating

→  
(By calculating  
round-trip times)

Frequency Band:  
2.4 – 7.7 MHz



## Reference Measurement



**Maximum Bandwidth: 2.5 – 8 MHz (in middle probes)**  
**Minimum Bandwidth: 2.5 – 6.5 MHz (in side-probes)**

$$M_{Rj,T}(\omega) = r_j(\omega) G_0(\mathbf{r}_j, \mathbf{r}_s) \overbrace{S(\mathbf{r}_s, \mathbf{r}_s)}^1 \sum_{i=1}^{N_T} G_0(\mathbf{r}_s, \mathbf{r}_i) t_i(\omega)$$

$$M_{Rj,T}(\omega) = r_j(\omega) G_0(\mathbf{r}_j, \mathbf{r}_s) t(\omega) \sum_{i=1}^{N_T} G_0(\mathbf{r}_s, \mathbf{r}_i)$$

$$M_{Rj,T}(\omega) = r_j(\omega) t(\omega) G_0(\mathbf{r}_j, \mathbf{r}_s) \sum_{i=1}^{N_T} G_0(\mathbf{r}_s, \mathbf{r}_i)$$

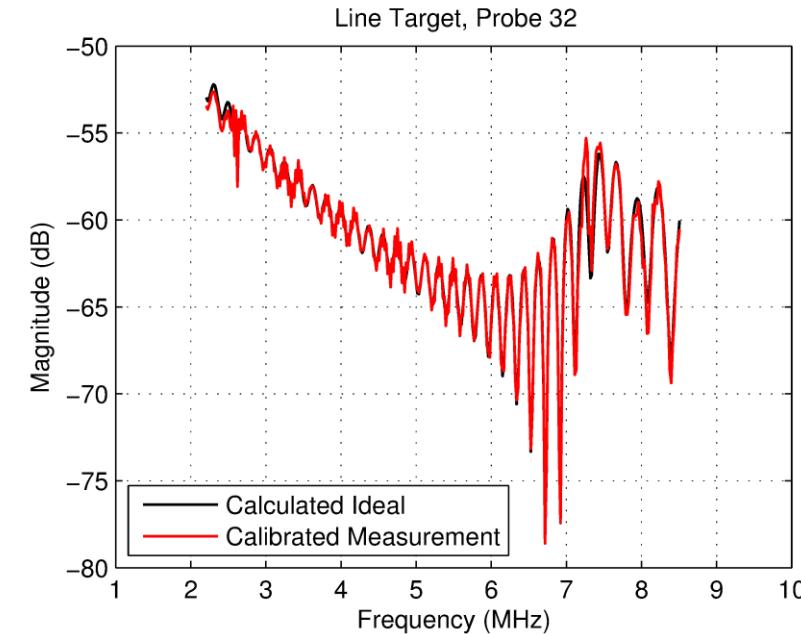
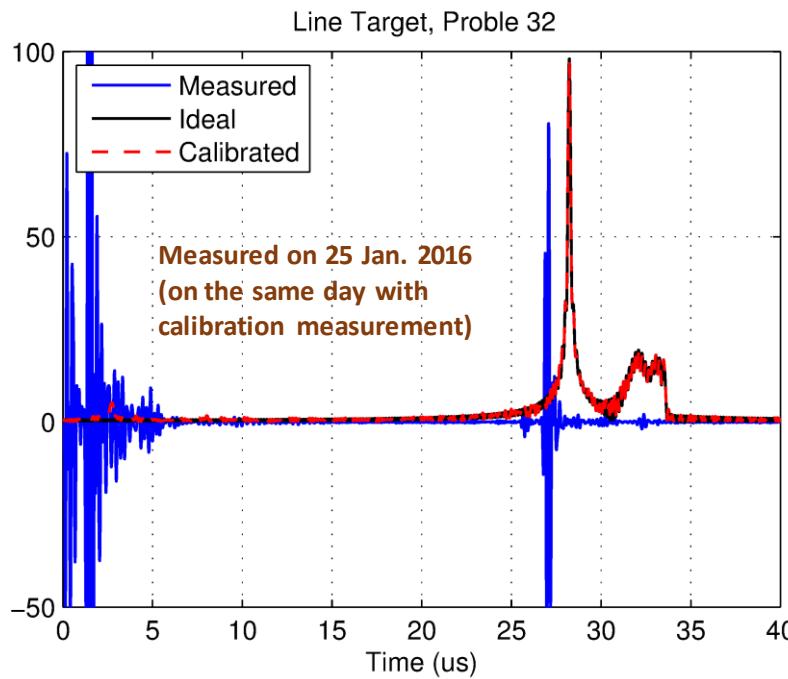
**Assumption!**

$$t_i(\omega) \approx t(\omega)$$

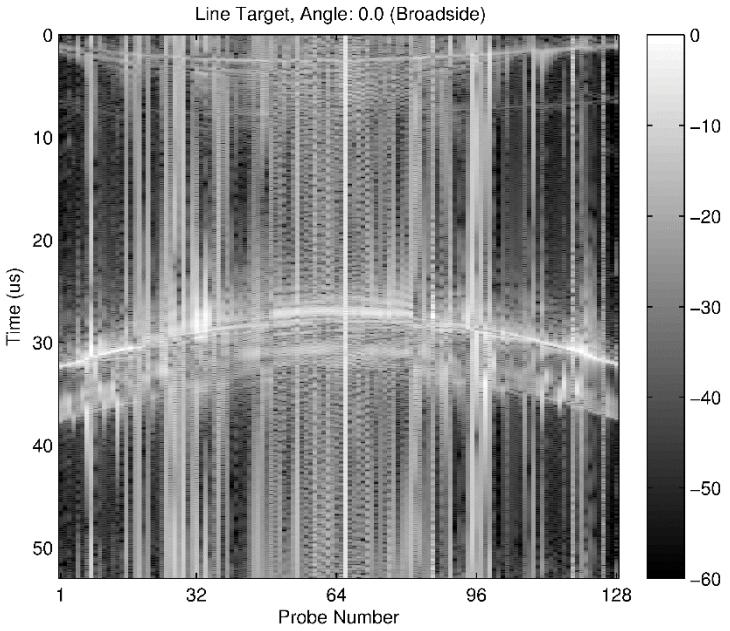
$$\underbrace{r_j(\omega) t(\omega)}_{\text{Weights}} = \frac{M_{Rj,T}(\omega)}{G_0(\mathbf{r}_j, \mathbf{r}_s) \sum_{i=1}^{N_T} G_0(\mathbf{r}_s, \mathbf{r}_i)}$$

Measured

Calculated (Ideal)

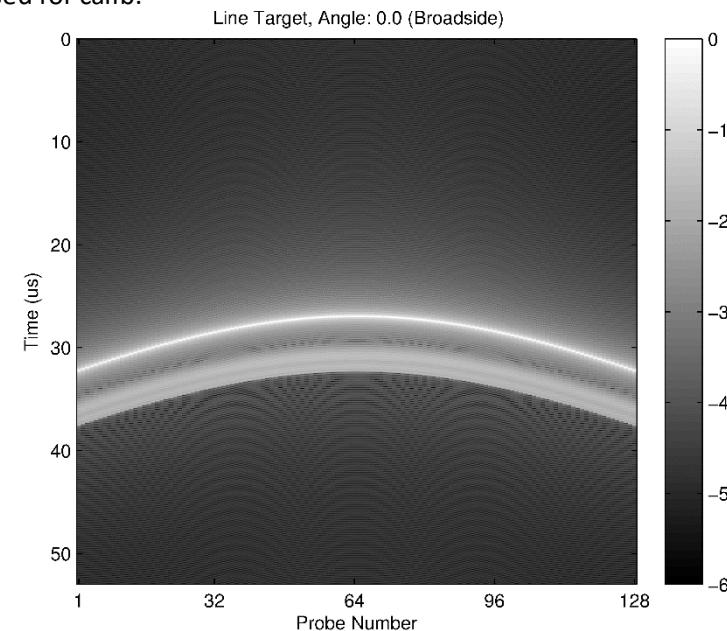


## Measured (Calibrated) ←

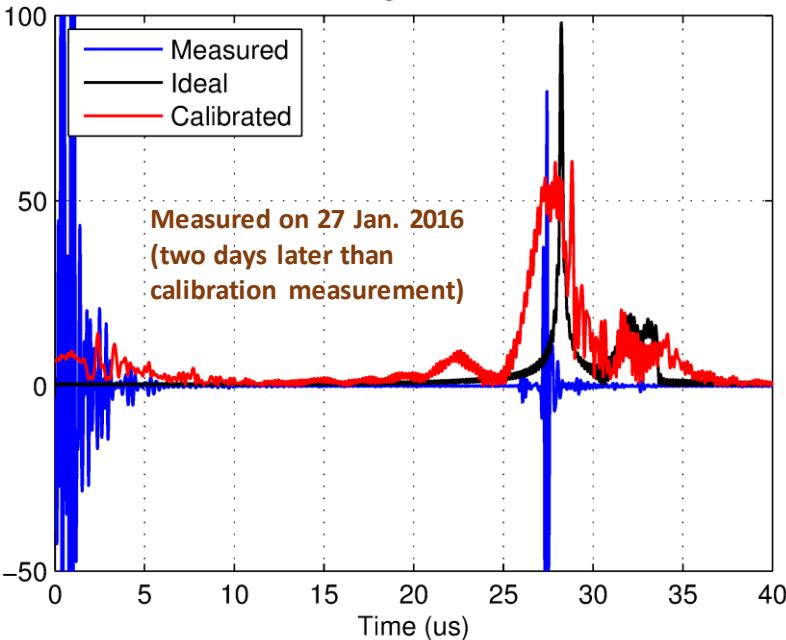


Another meas. than  
what we used for calib.

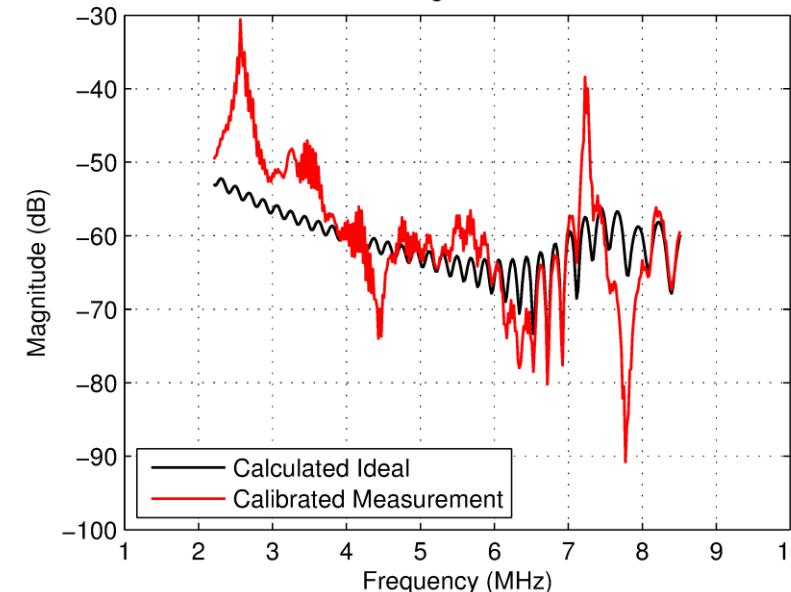
## Theoretical



Line Target, Probe 32



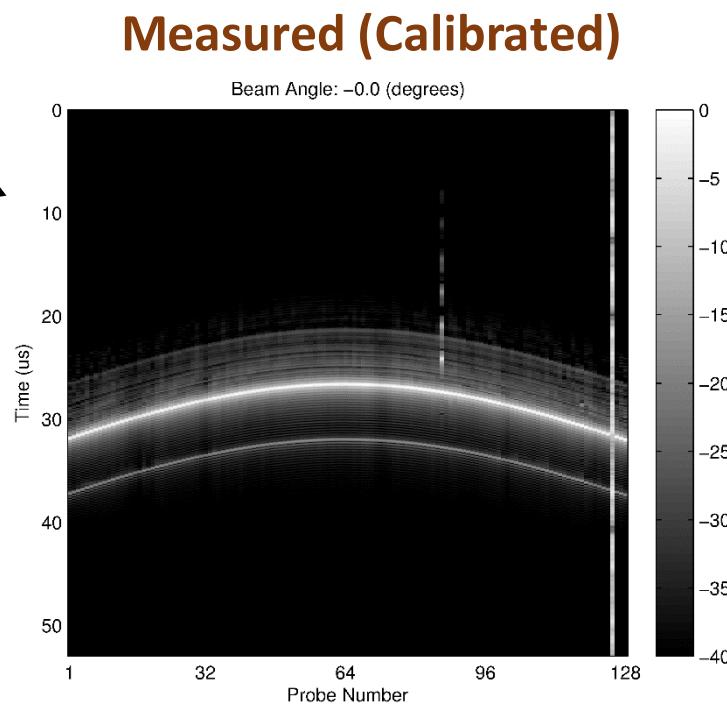
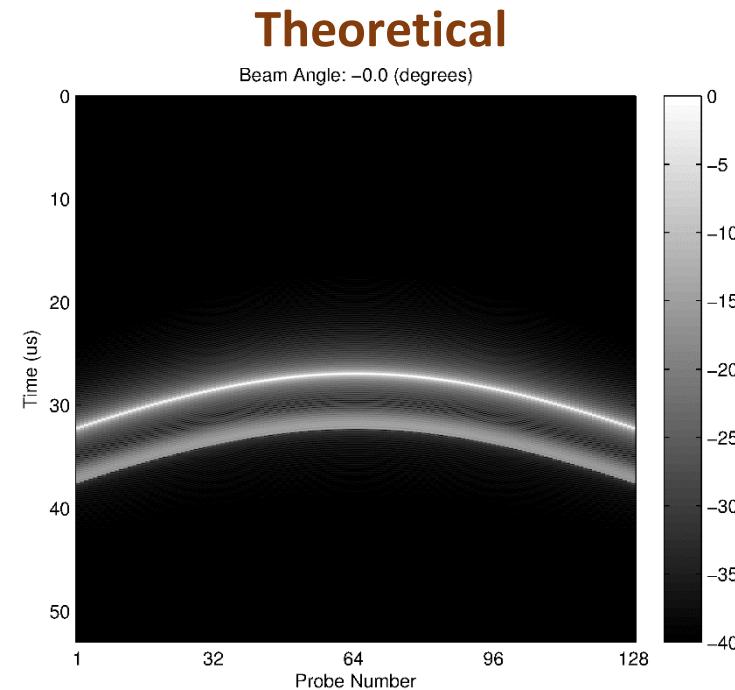
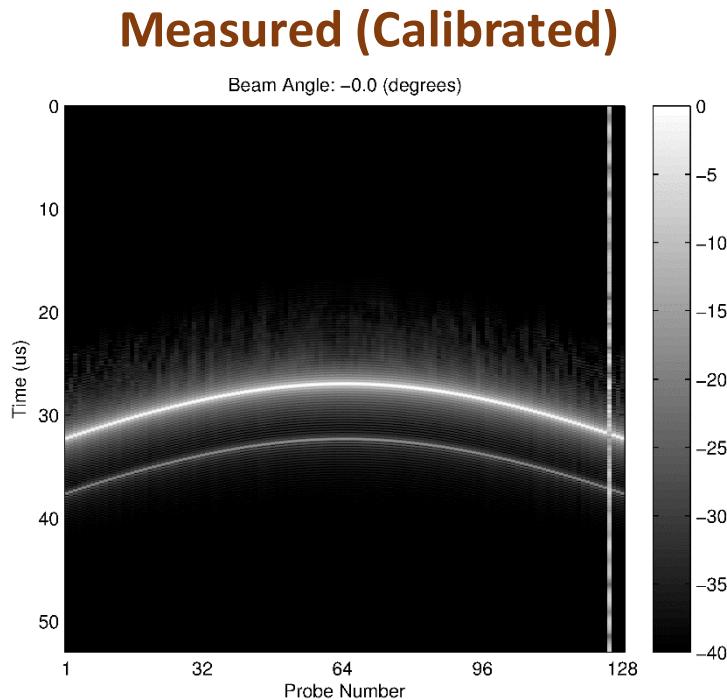
Line Target, Probe 32



# Bandwidth 2.5 – 6.5 MHz

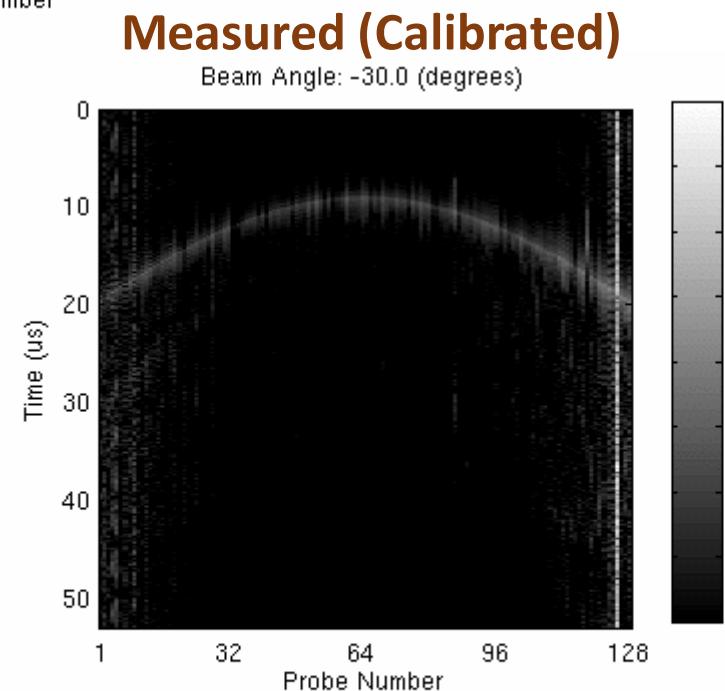
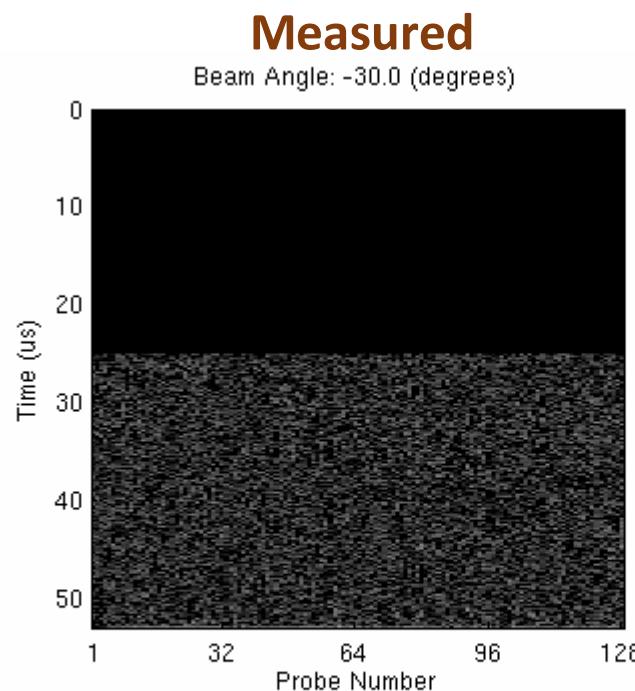
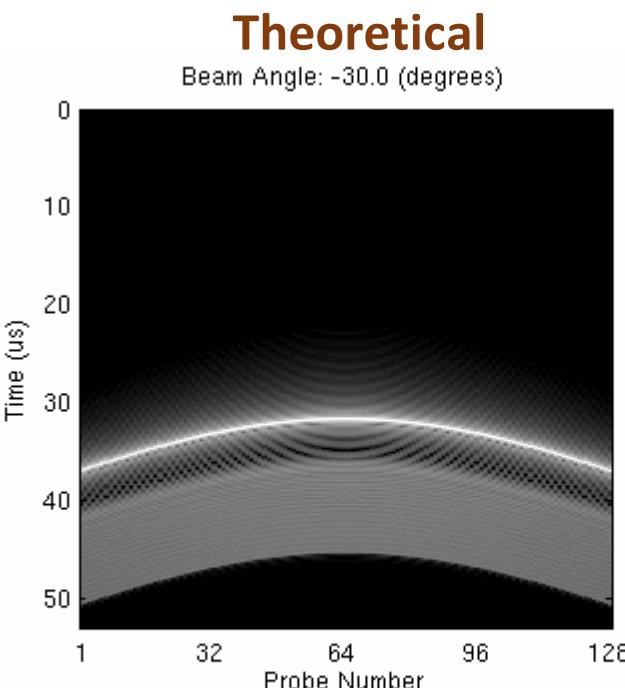
Measured on the  
same day with the  
ref. measurement

Measured two days  
after the ref.  
measurement



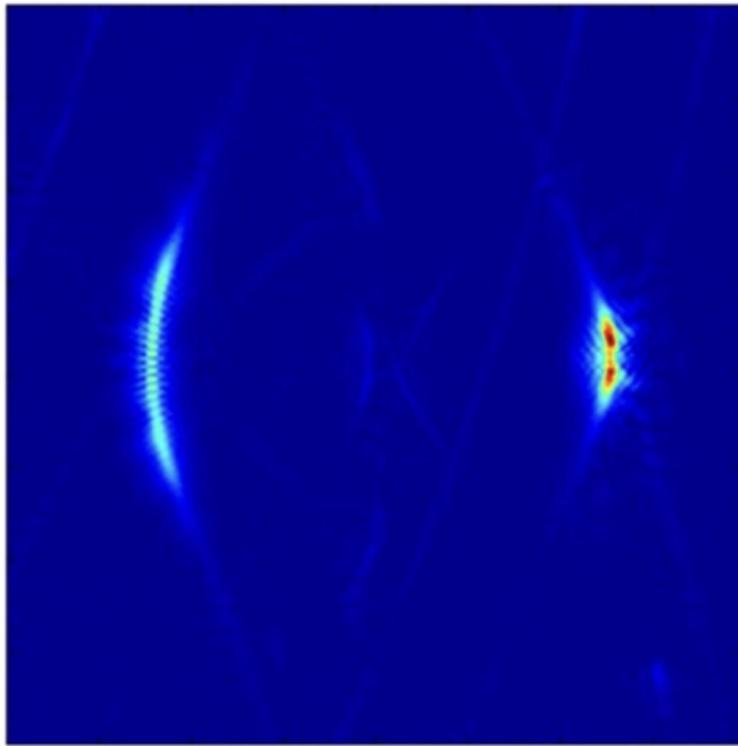
**Bandwidth 2.5 – 6.5 MHz**

Complex Mag. →

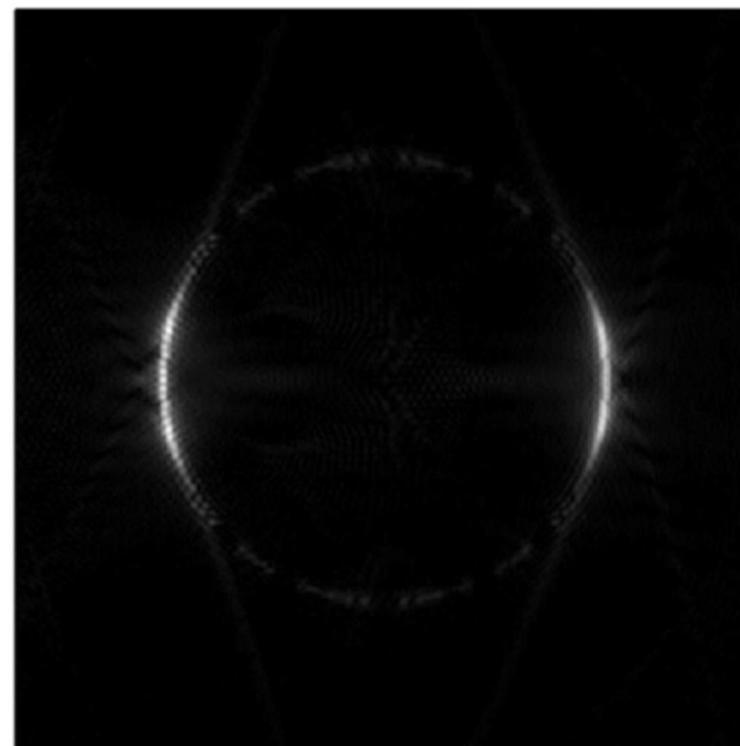


# Calibration Step 4: Run the Numerical Solver!

Real Scatterer



Numerical Scatterer



# Large-Scale Inverse Multiple-Scattering Imaging on GPU Supercomputers with Real Data

Mert Hidayetoğlu  
Stanford University

In courtesy of Weng Cho Chew, Michael Oelze,  
and Wen-mei Hwu

[merth@stanford.edu](mailto:merth@stanford.edu)

